

Solving the Quantum Many-Body Problem with Neural-Network Quantum States



MICHIGAN STATE
UNIVERSITY

Jane Kim
Michigan State University
Facility for Rare Isotope Beams
PhD Dissertation Defense
20 July 2023



Acknowledgements

- **Research Advisor:** Morten Hjorth-Jensen
- **Guidance Committee:** Filomena Nunes
Artemis Spyrou
Stuart Tessmer
Michael Murillo
Huey-Wen Lin
(Matthew Hirn)
- **Collaborators:** Alessandro Lovato (ANL)
Bryce Fore (ANL)
Stefano Gandolfi (LANL)
Giuseppe Carleo (EPFL)
Gabriel Pescia (EPFL)
Jannes Nys (EPFL)
Even Nordhagen (UiO)
- **Funding:** US NSF Grants No. PHY-1404159 and PHY-2013047

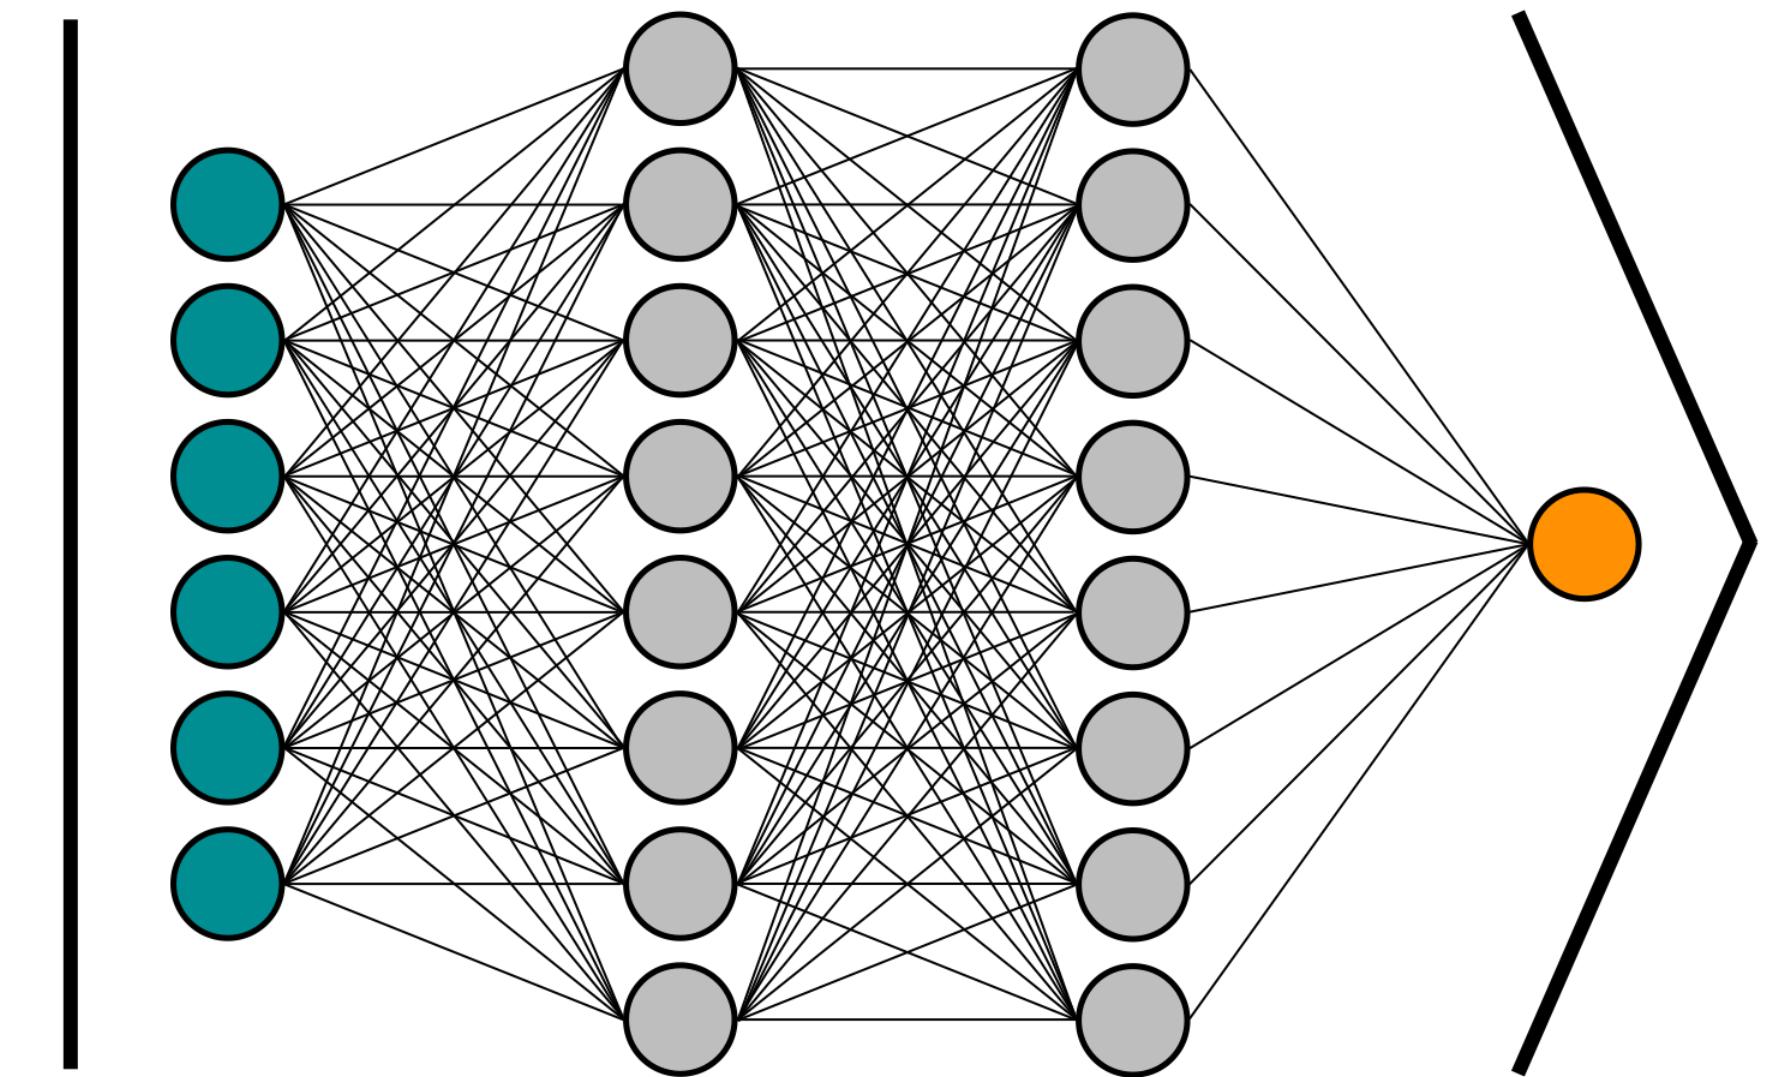


MICHIGAN STATE
UNIVERSITY



Overview

- The quantum many-body problem
- Quantum Monte Carlo
- Machine learning and artificial neural networks
- Implementations in C++/MPI and Python/JAX
- Applications:
 - Calogero-Sutherland model
 - Dilute neutron matter
 - Homogeneous electron gas
 - Ultra-cold Fermi gases



The Quantum Many-Body Problem

Assumptions

- Time-independent Hamiltonian \implies Solve time-independent Schrödinger equation

$$\hat{H}|\Psi\rangle = E|\Psi\rangle$$

- Indistinguishable particles $\implies |\Psi\rangle$ must be symmetric or antisymmetric
- Work in continuous position (discrete spin, isospin) basis $\implies |\Psi\rangle = \int d\mathbf{X} \Psi(\mathbf{X}) |\mathbf{X}\rangle$
- Simulate N particles in d dimensions $\implies |\mathbf{X}\rangle = |\mathbf{x}_1\rangle \otimes |\mathbf{x}_2\rangle \otimes \dots \otimes |\mathbf{x}_N\rangle$
 $|\mathbf{x}_i\rangle = |\mathbf{r}_i\rangle \otimes |\mathbf{s}_i\rangle, \quad \mathbf{r}_i \in \mathbb{R}^d$

Goals

- Find the ground state wave function $\Psi_0(\mathbf{X})$ of a system of interacting particles
- Develop a flexible, compact, and low-bias ansatz represented by neural networks (neural-network quantum states (NQS) first applied by Carleo and Troyer, Science, 2017)
- Achieve state-of-the-art results (benchmark with diffusion Monte Carlo)
- **Unique nuclear theory perspective:** spin and isospin exchange, strong and short-range interaction, and a possible three-body term motivates ansätze with even more generality and flexibility than those developed from a condensed matter or quantum chemistry perspective

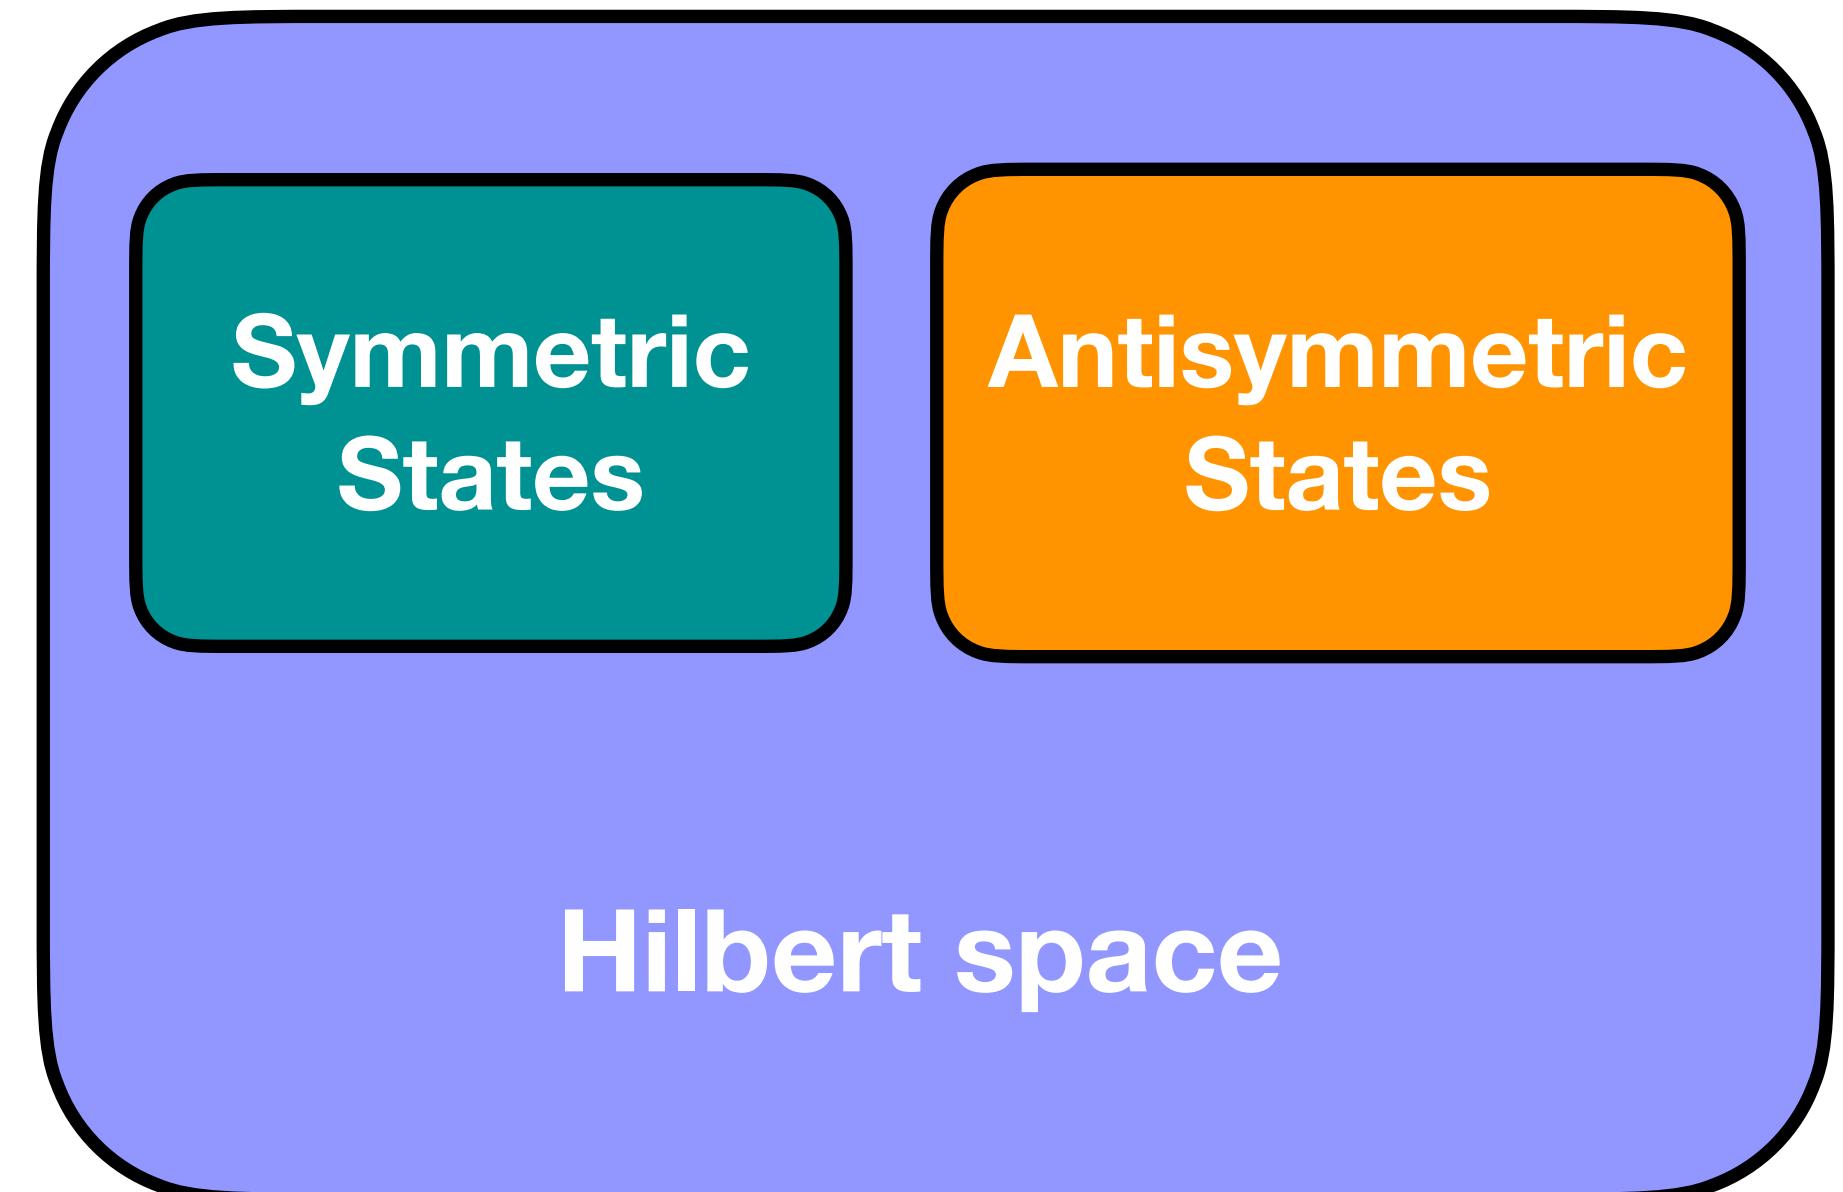
Computational Challenges

- Storing particle degrees of freedom in memory treats particles as distinguishable
- Hilbert space dimension grows exponentially with system size N

$$\mathcal{H} = \mathcal{H}_1 \otimes \mathcal{H}_2 \otimes \cdots \otimes \mathcal{H}_N$$

- But not all states within Hilbert space are physical!
- We need to enforce particle indistinguishability ourselves to restrict our search to the relevant areas of Hilbert space

$$\mathcal{H}_S \oplus \mathcal{H}_A \subseteq \mathcal{H}$$



Computational Challenges

- Enforcing other symmetries can further help reduce the volume of Hilbert space we need to search
- Continuous-space problems \implies infinite basis
- Quantum many-body methods aim to mitigate the problem of exponential scaling of the Hilbert space dimension, for example:
 - Configuration Interaction
 - Hartree-Fock theory
 - Many-Body Perturbation Theory
 - Coupled Cluster
 - In-Medium Similarity Renormalization Group
 - And many more...

Quantum Monte Carlo Methods

Variational Principle

- The expectation value of a Hamiltonian \hat{H} computed with any trial wave function $|\Psi\rangle$ is an upper bound to the true ground state energy E_0

$$E[\Psi] \equiv \frac{\langle \Psi | \hat{H} | \Psi \rangle}{\langle \Psi | \Psi \rangle} \geq E_0$$

- The variance of the energy is zero only if $|\Psi\rangle$ is an eigenstate of \hat{H}
- Caveat: $|\Psi\rangle$ must satisfy the fundamental symmetries of the system described by \hat{H}

Variational Monte Carlo

- Direct application of the variational principle:

$$\min_{\theta} E[\Psi_{\theta}] \geq E_0$$

Variational parameters of trial wave function

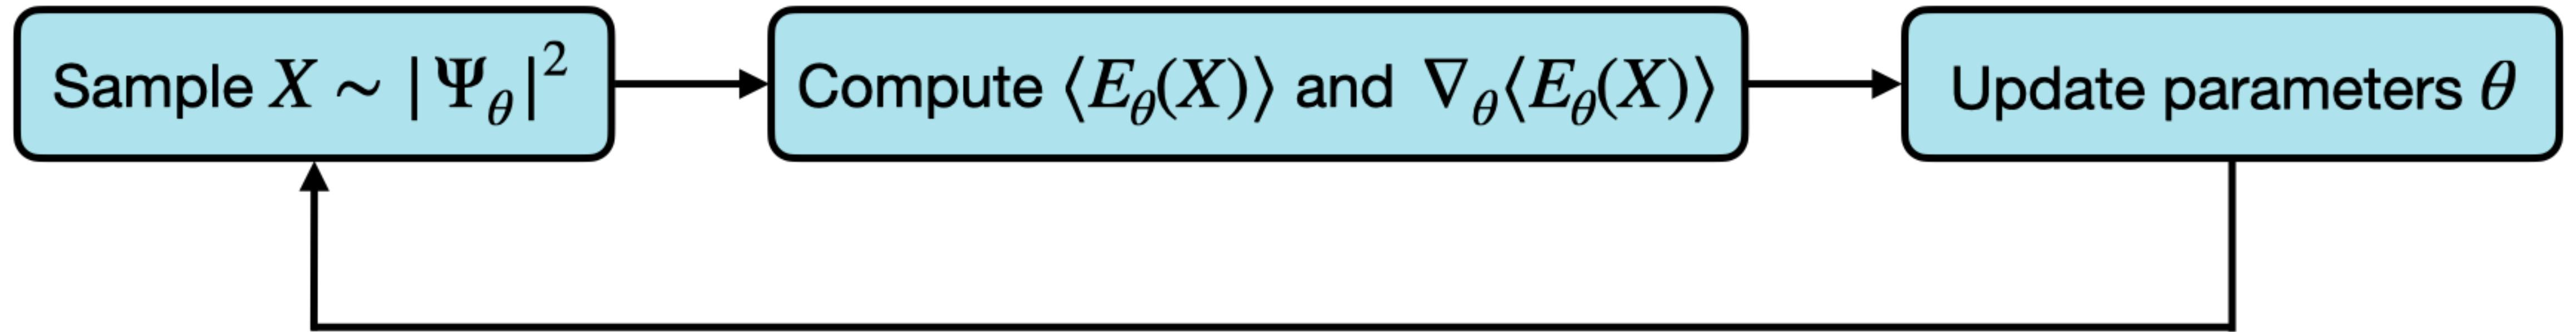
- Efficiently compute the variational energy via Monte Carlo integration:

$$E[\Psi_{\theta}] \equiv \frac{\langle \Psi_{\theta} | \hat{H} | \Psi_{\theta} \rangle}{\langle \Psi_{\theta} | \Psi_{\theta} \rangle} = \frac{\int d\mathbf{X} |\Psi_{\theta}(\mathbf{X})|^2 E_{\theta}(\mathbf{X})}{\int d\mathbf{X} |\Psi_{\theta}(\mathbf{X})|^2} \approx \frac{1}{N_s} \sum_{i=1}^{N_s} E_{\theta}(\mathbf{X}_i) \equiv \langle E_{\theta}(\mathbf{X}) \rangle$$

$$\text{Local energy } E_{\theta}(\mathbf{X}) \equiv \frac{\langle \mathbf{X} | \hat{H} | \Psi_{\theta} \rangle}{\langle \mathbf{X} | \Psi_{\theta} \rangle} = \frac{\hat{H}\Psi_{\theta}(\mathbf{X})}{\Psi_{\theta}(\mathbf{X})}$$

Configurations sampled from $|\Psi_{\theta}(\mathbf{X})|^2$

Optimization



- Variational parameters θ are optimized using gradient descent methods

$$\theta_{t+1} = \theta_t - \eta \nabla_\theta \langle E_\theta(\mathbf{X}) \rangle$$

- Compute gradient stochastically:

$$\nabla_\theta \langle E_\theta(\mathbf{X}) \rangle = 2 \left(\langle E_\theta(\mathbf{X}) \mathbf{O}_\theta(\mathbf{X}) \rangle - \langle E_\theta(\mathbf{X}) \rangle \langle \mathbf{O}_\theta(\mathbf{X}) \rangle \right) \text{ with } \mathbf{O}_\theta(\mathbf{X}) \equiv \frac{\nabla_\theta \Psi_\theta(\mathbf{X})}{\Psi_\theta(\mathbf{X})}$$

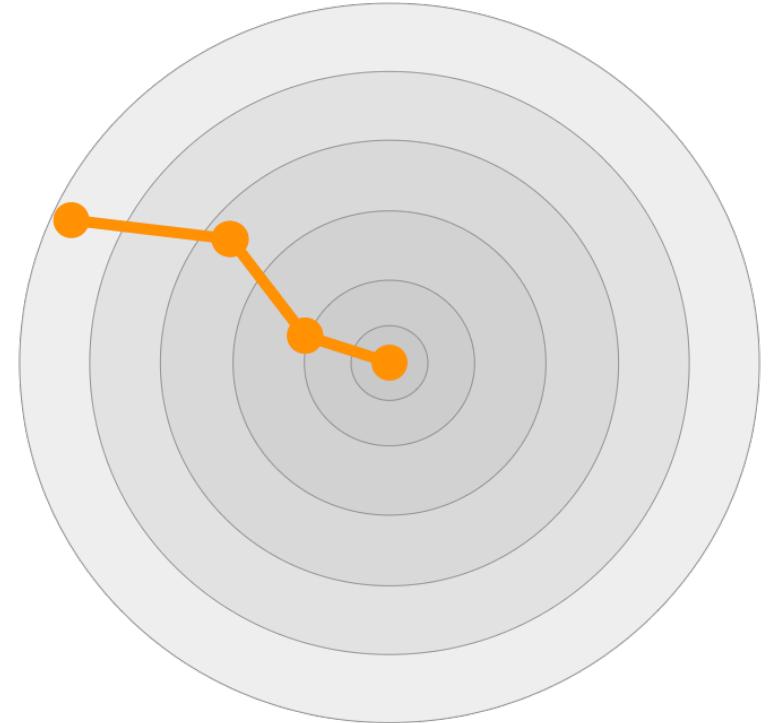
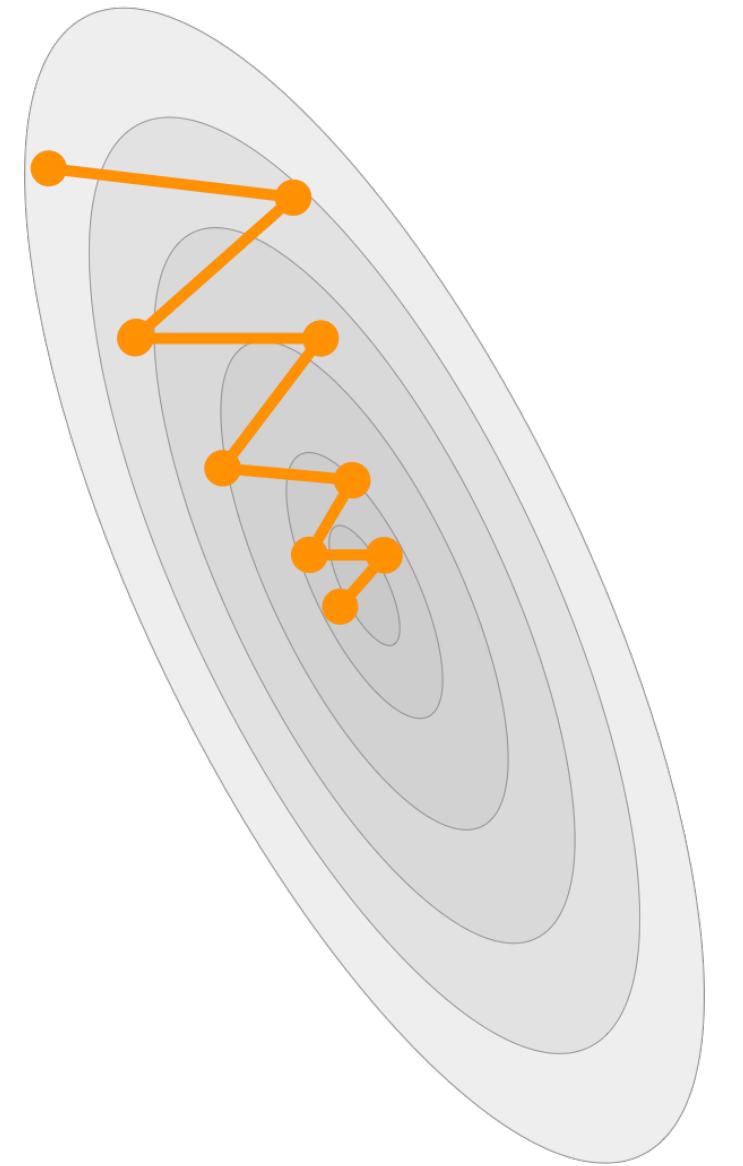
Stochastic Reconfiguration

- Second-order optimization method that reshapes the energy landscape
- Derived by matching the variational state to a state obtained by applying the imaginary-time propagator
- Parameters are updated as

$$\theta_{t+1} = \theta_t - \eta S_\theta^{-1} \nabla_\theta \langle E_\theta(\mathbf{X}) \rangle$$

where S_θ is the quantum geometric tensor / Fisher information matrix

$$[S_\theta]_{ij} = \frac{\langle \partial_i \Psi_\theta | \partial_j \Psi_\theta \rangle}{\langle \Psi_\theta | \Psi_\theta \rangle} - \frac{\langle \partial_i \Psi_\theta | \Psi_\theta \rangle \langle \Psi_\theta | \partial_j \Psi_\theta \rangle}{\langle \Psi_\theta | \Psi_\theta \rangle}$$



Stochastic Reconfiguration

- It is possible for S_θ to be singular if some variational parameters are redundant
- Typically apply a small diagonal shift to stabilize the inverse $S_\theta^{-1} \mapsto [S_\theta + \varepsilon I]^{-1}$
- But constant diagonal shifts ignore the relative magnitudes of the parameters
- Better stabilization developed by Lovato et al. in Phys. Rev. Research **4**, 043178 (Dec 2022), inspired by the RMSprop algorithm

$$\begin{aligned}\mathbf{v}_{t+1} &= \beta \mathbf{v}_t + (1 - \beta) (\nabla_\theta \langle E_\theta(\mathbf{X}) \rangle)^2 \\ \theta_{t+1} &= \theta_t - \eta [S_\theta + \text{diag}(\sqrt{\mathbf{v}_{t+1}} + \varepsilon)]^{-1} \nabla_\theta \langle E_\theta(\mathbf{X}) \rangle\end{aligned}$$

Trial Wave Functions

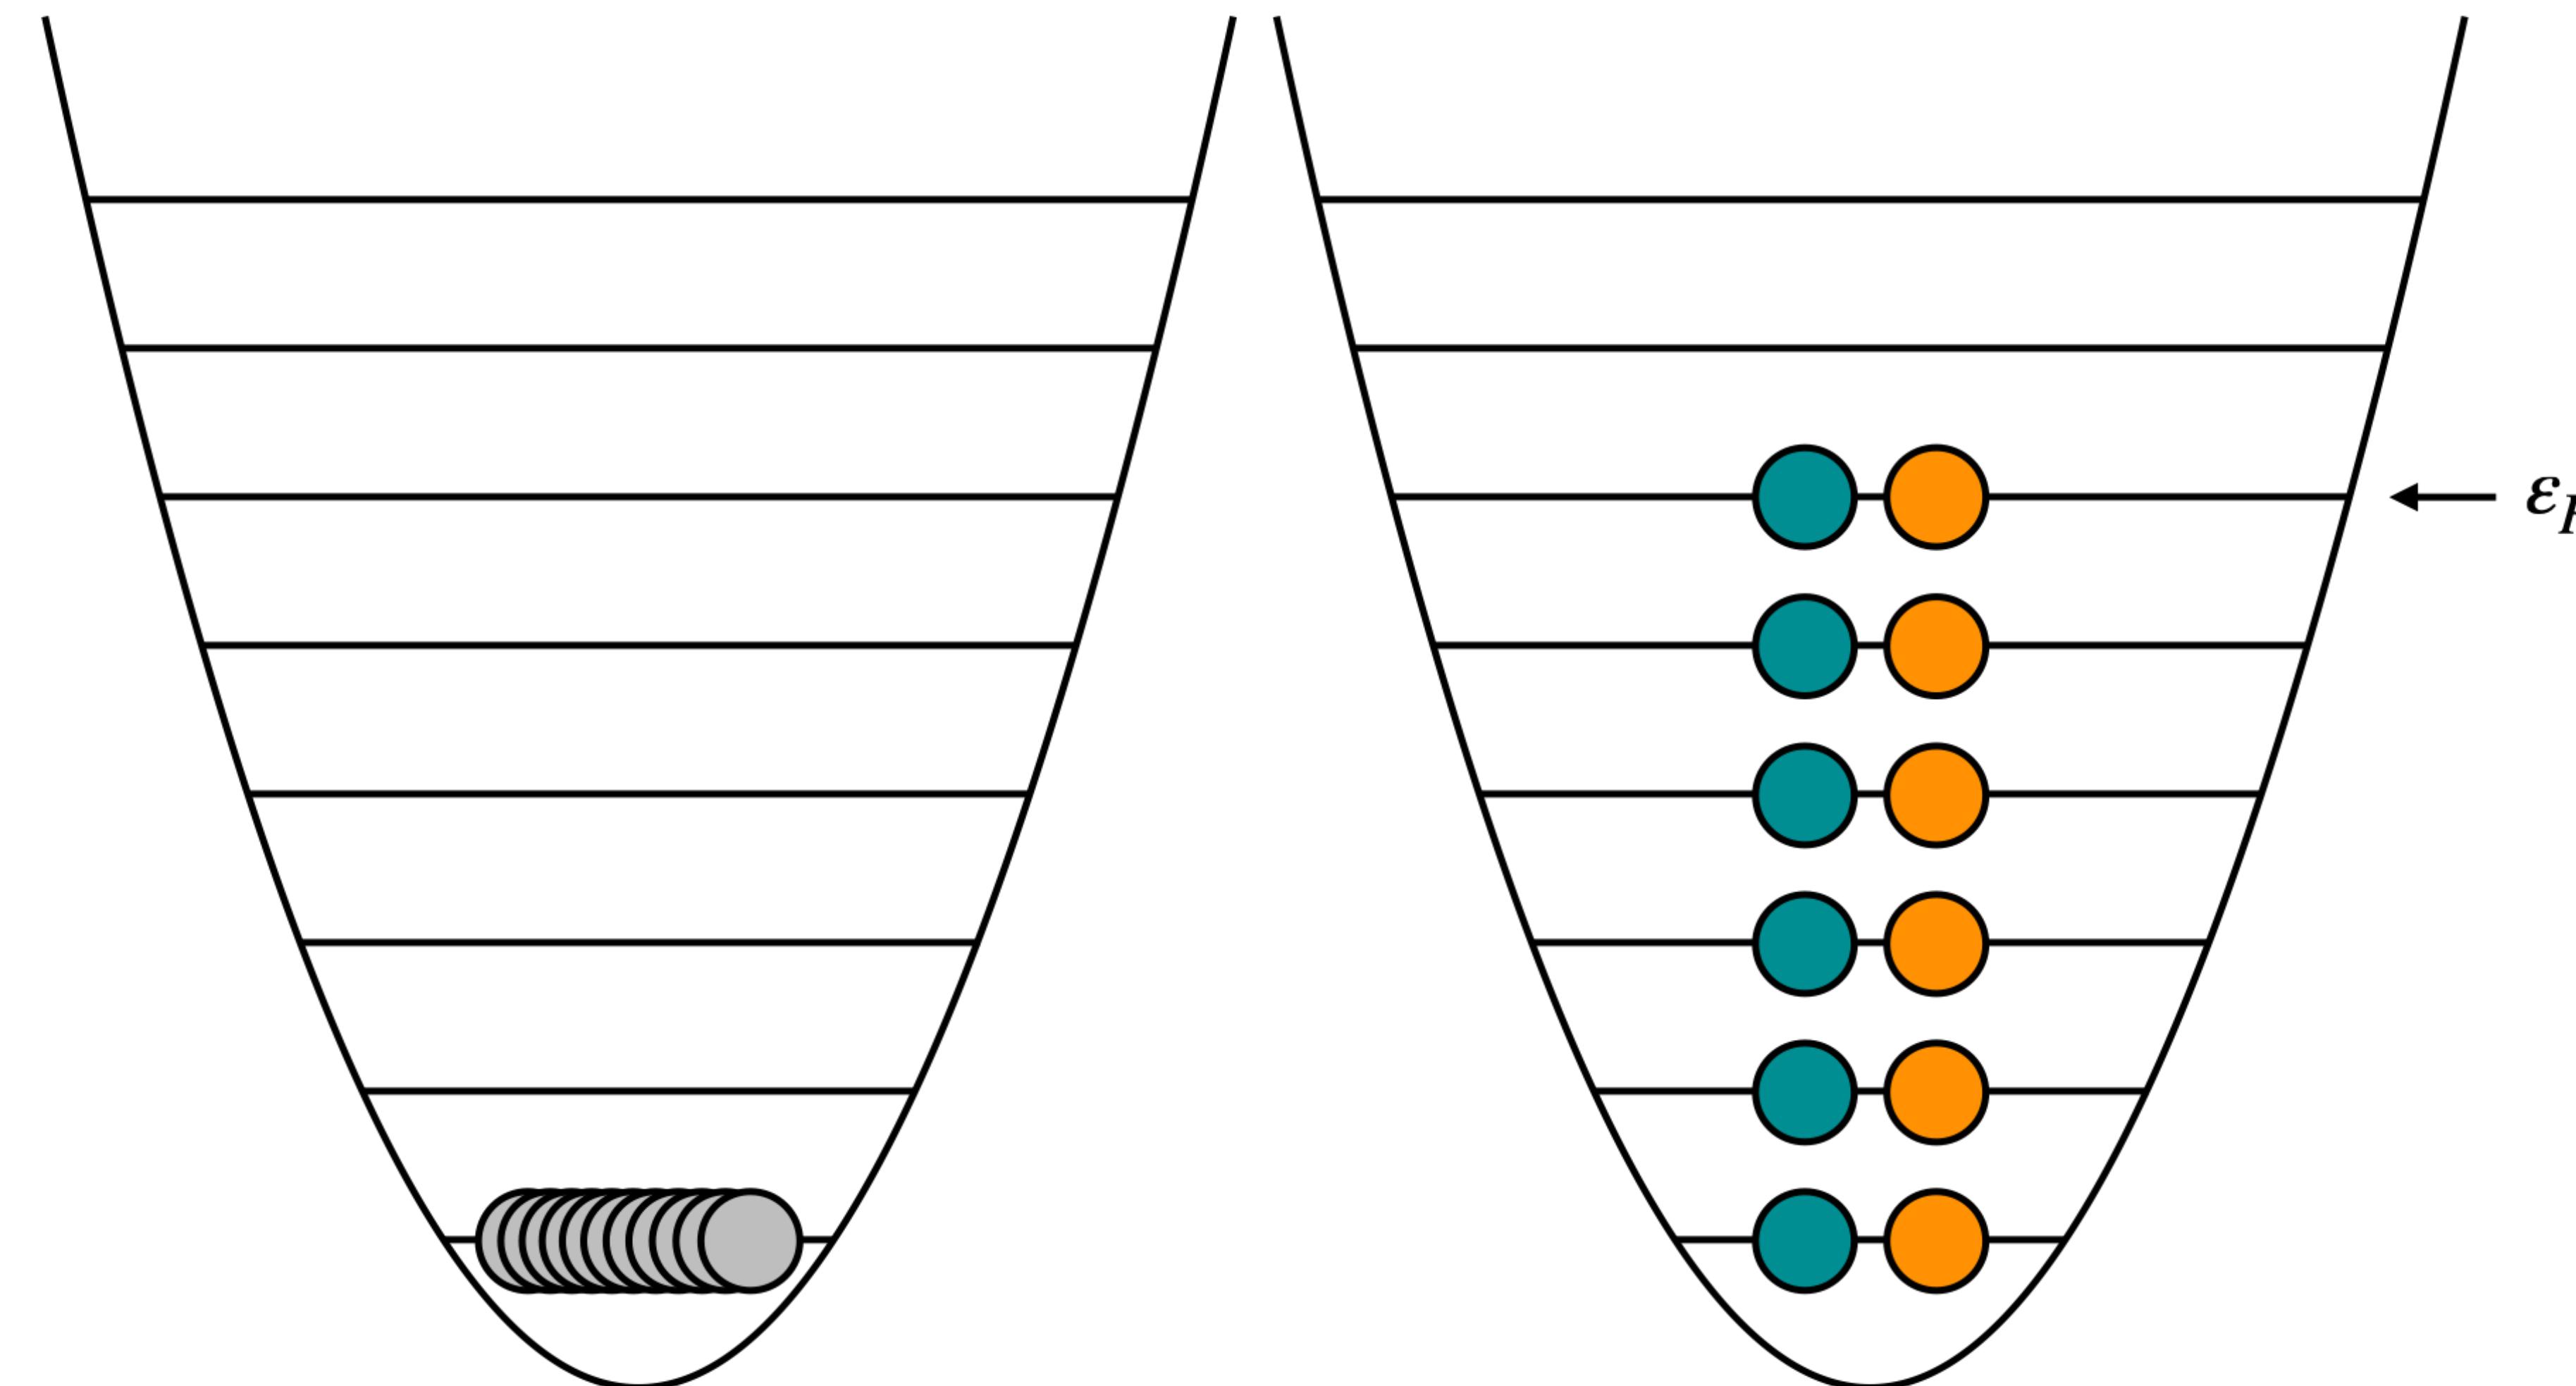
- In standard variational Monte Carlo calculations, the trial wave function is carefully designed for the specific system of interest
- Generally, the trial wave function has the form

$$\Psi_{\theta}(\mathbf{X}) = \Phi_0(\mathbf{X})e^{J(\mathbf{X})}$$

- The noninteracting ground state constrains permutation symmetry/antisymmetry
- Jastrow correlator is designed to include correlations and satisfy Kato's cusp condition:

$$\lim_{r_{ij} \rightarrow 0} E[\Psi_0] < \infty \text{ for all pairs } i, j$$

Bosons vs. Fermions



Slater Determinants

- The simplest way to write down the antisymmetrized product of single-particle states:

$$\Phi_0(\mathbf{X}) = \det \begin{bmatrix} \varphi_1(\mathbf{x}_1) & \varphi_1(\mathbf{x}_2) & \cdots & \varphi_1(\mathbf{x}_N) \\ \varphi_2(\mathbf{x}_1) & \varphi_2(\mathbf{x}_2) & \cdots & \varphi_2(\mathbf{x}_N) \\ \vdots & \vdots & \ddots & \vdots \\ \varphi_N(\mathbf{x}_1) & \varphi_N(\mathbf{x}_2) & \cdots & \varphi_N(\mathbf{x}_N) \end{bmatrix}$$

φ_α = single-particle spin-orbitals

\mathbf{x}_i = single-particle degrees of freedom

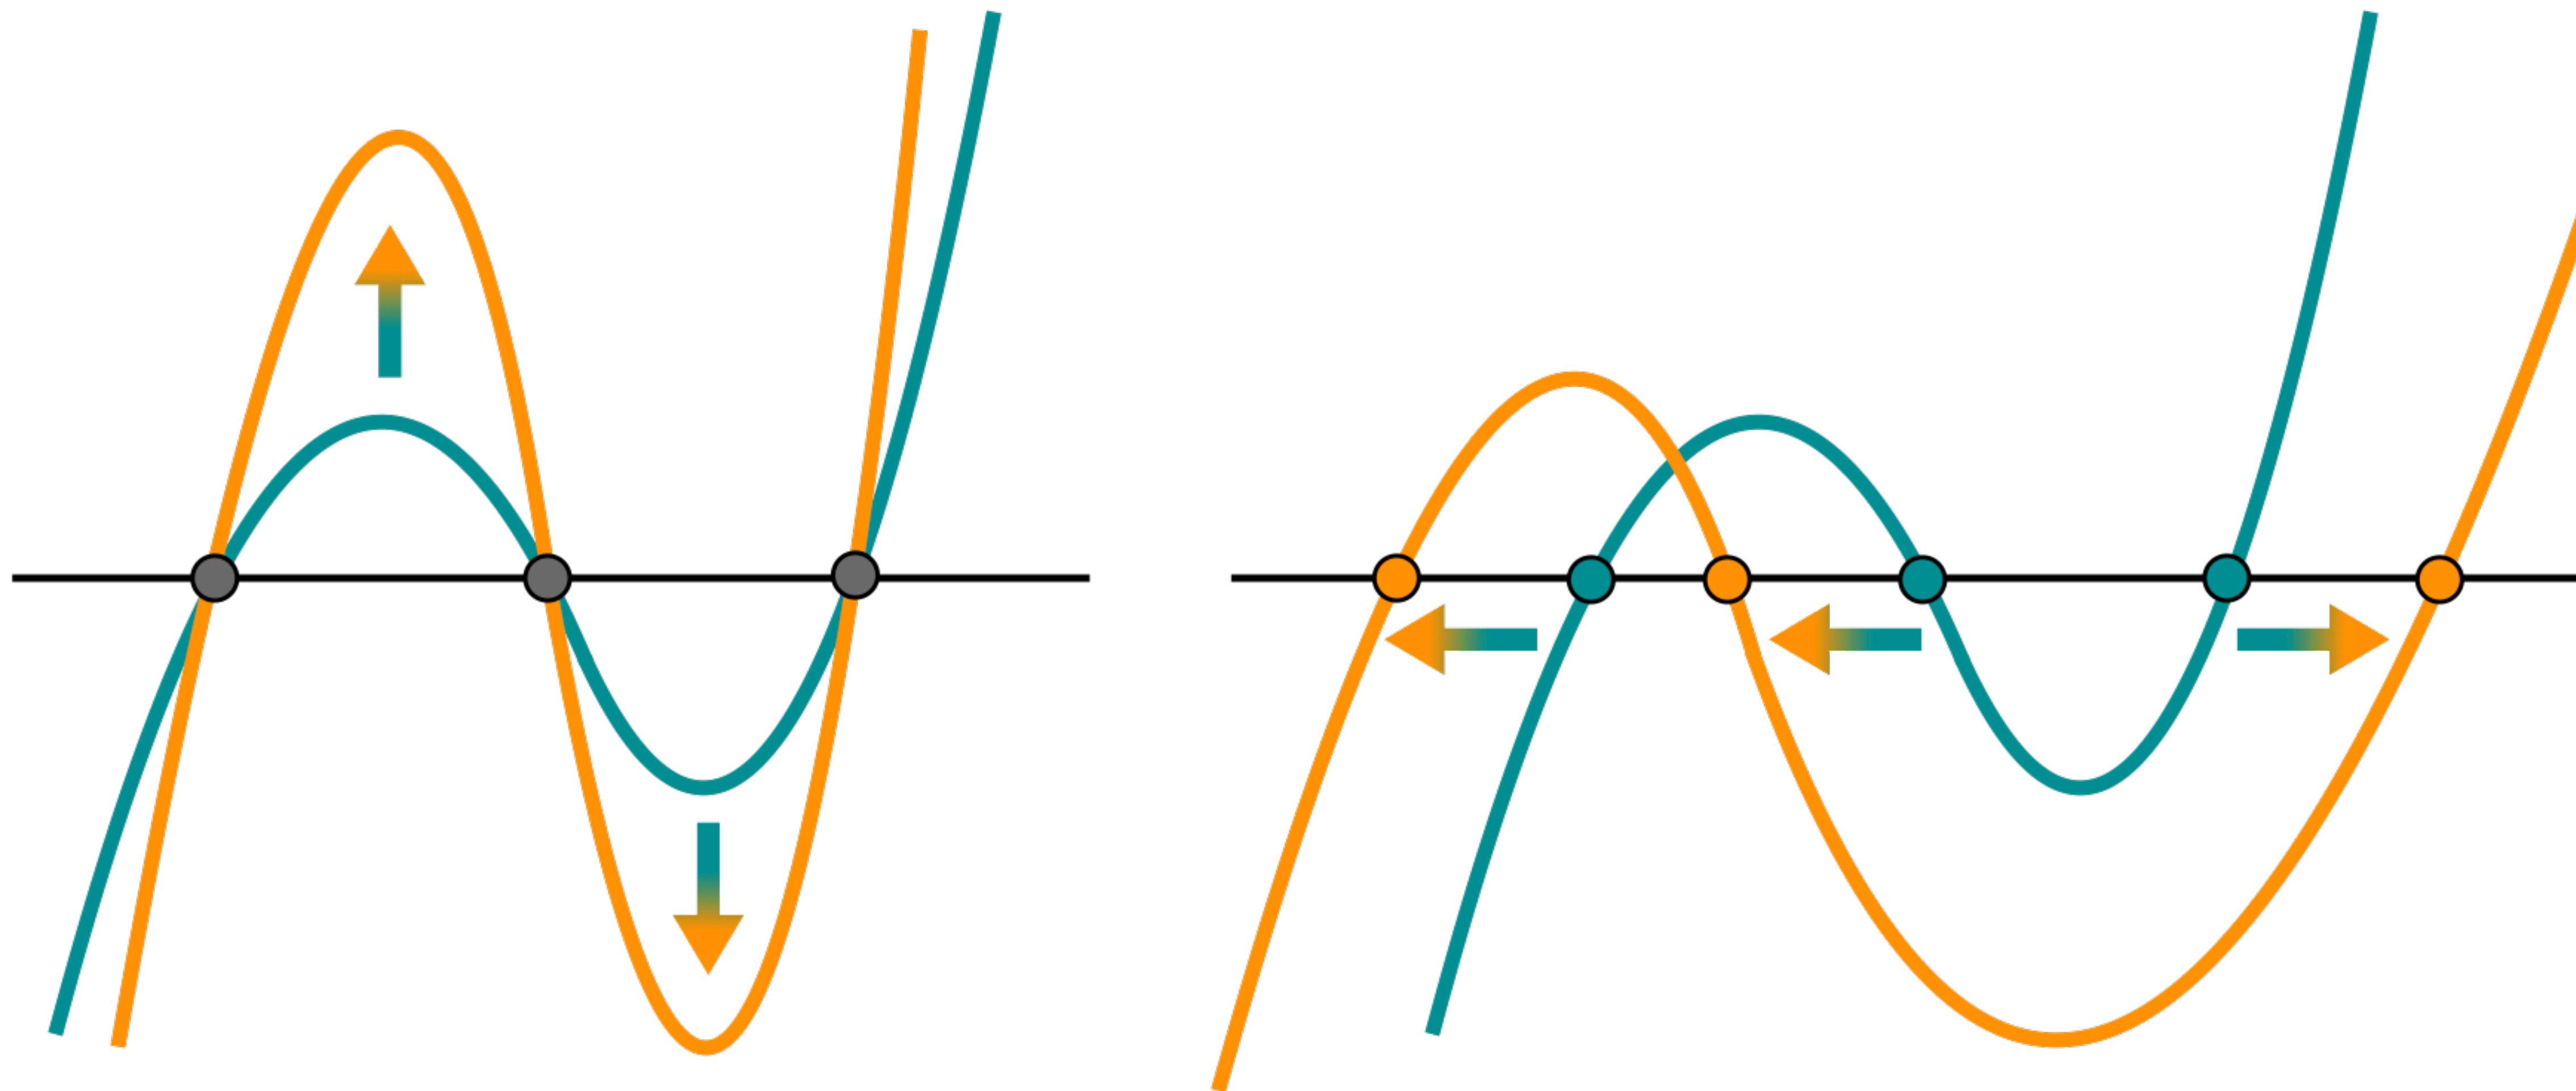
Backflow Transformations

- Backflow correlations refer to the influence of all surrounding particles on the state of a given particle
- Jastrow factors encode some backflow correlations *but* they do not change the nodes of the fermionic wave function
- Backflow transformations typically alter the single-particle positions as

$$\mathbf{r}_i \mapsto \mathbf{r}_i + \sum_{j \neq i}^N \eta(r_{ij})(\mathbf{r}_i - \mathbf{r}_j)$$

- The new positions are used as input to the Slater determinant, allowing the nodes to change

Jastrow Factors vs. Backflow Transformations



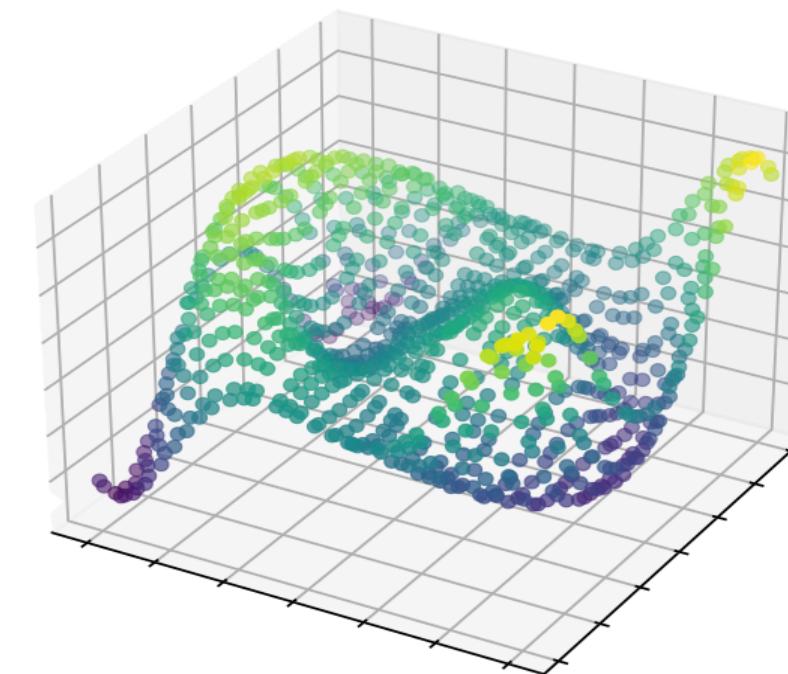
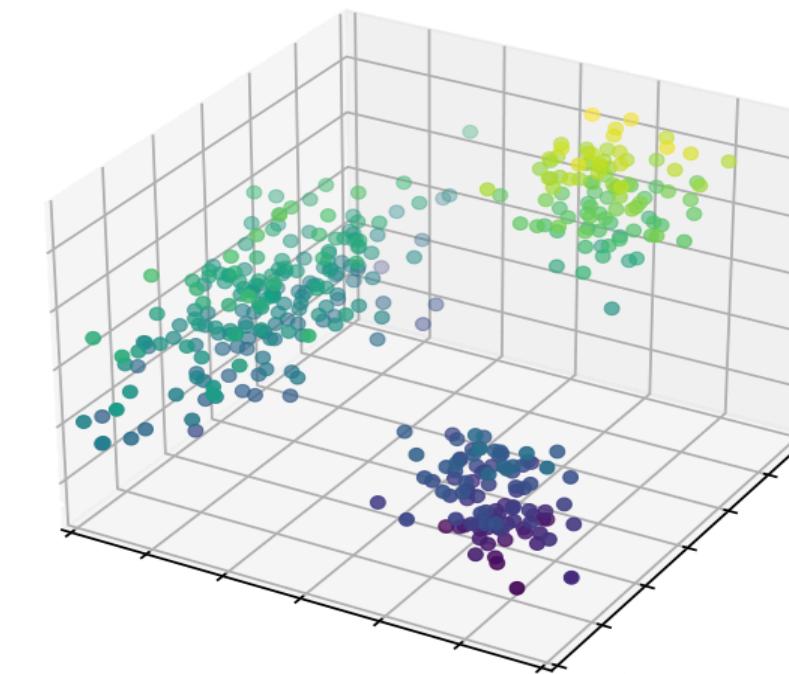
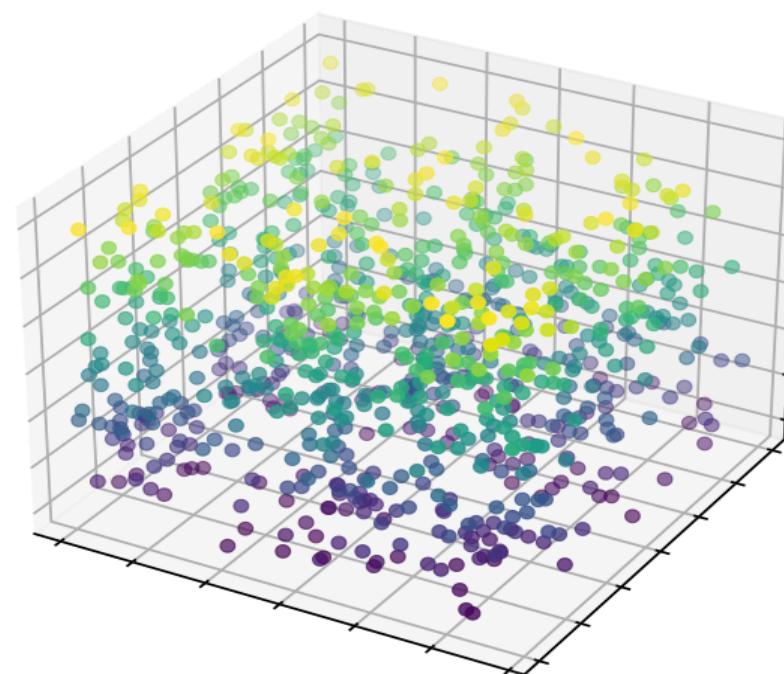
VMC vs. DMC

- Variational Monte Carlo is a direct application of the variational principle and Monte Carlo integration
- Relies on a “good” functional form for the trial wave function
- Higher bias than DMC, but intuitive meaning of the variational parameters
- No fermion sign problem
- Lower computational cost than DMC in standard applications
- Diffusion Monte Carlo (DMC) is a state-of-the-art method based on imaginary-time propagation
- Wave function is determined by the distribution of Monte Carlo walkers (no functional form)
- Relies on a good initial state, usually obtained through a VMC calculation
- Suffers from the sign problem, nodes must be fixed → residual dependence on VMC calculation

Machine Learning

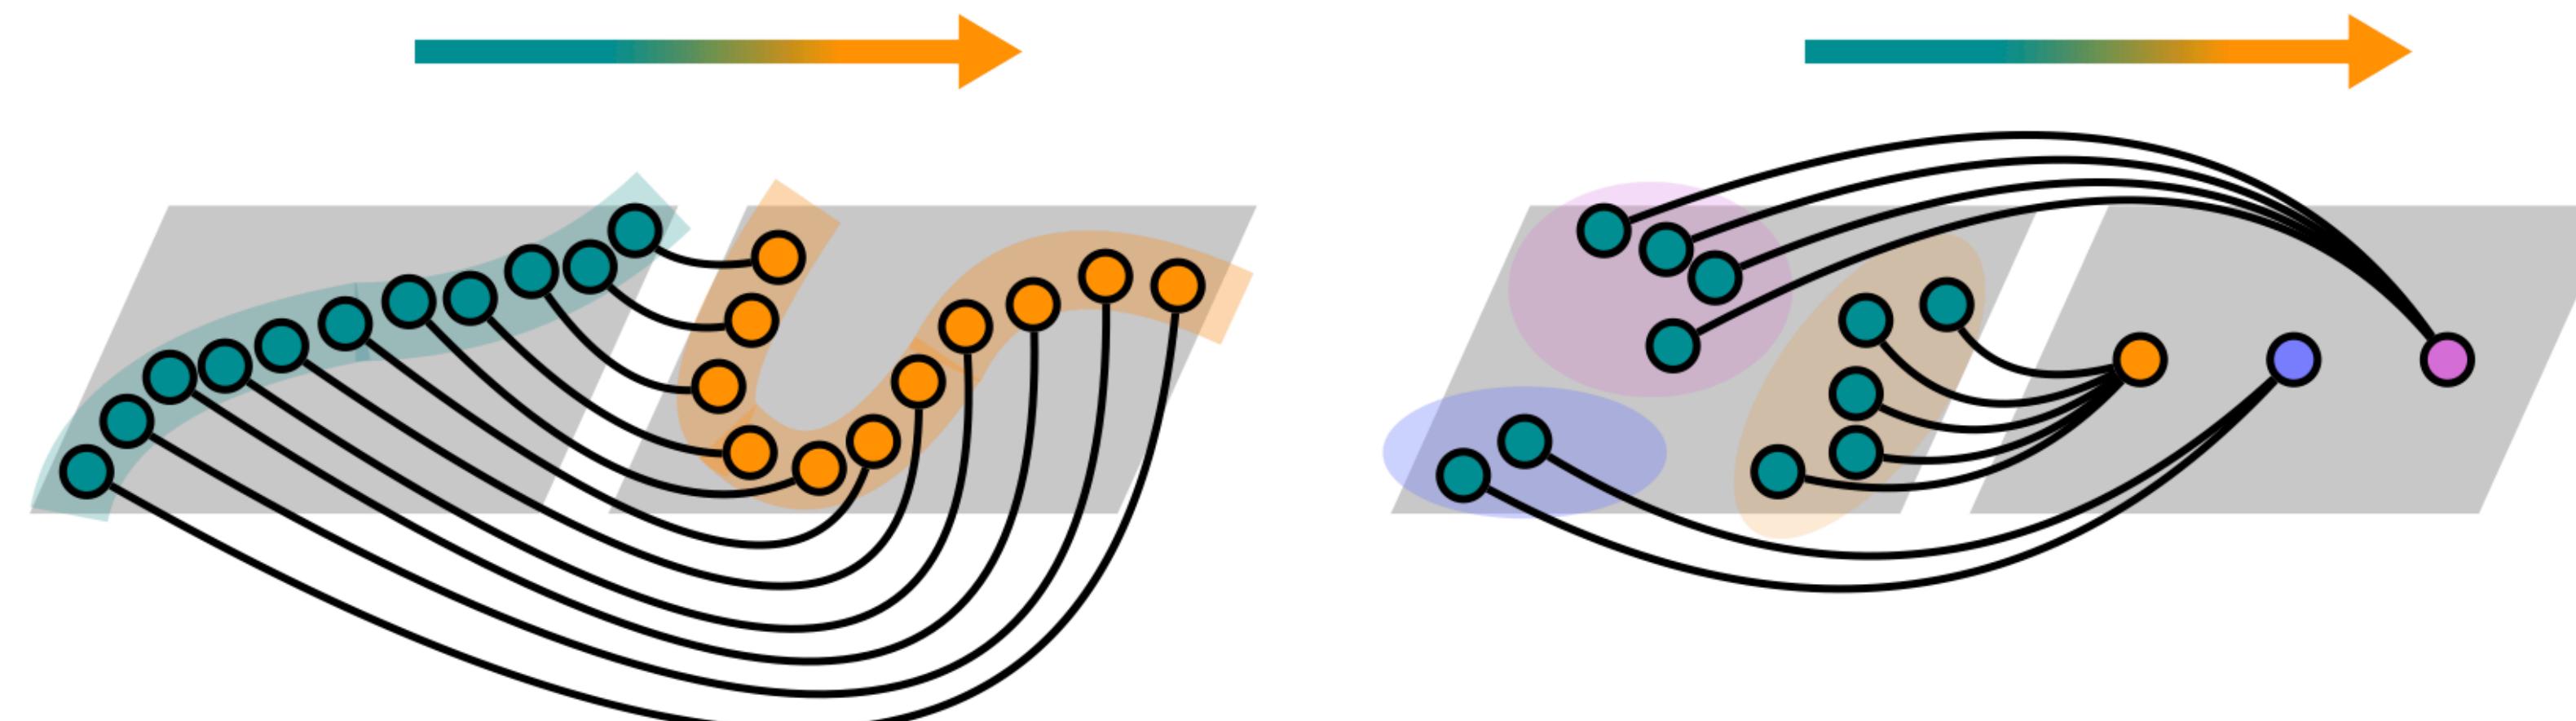
Machine Learning

- Machine learning is a subset of artificial intelligence, a field that aims to endow computers with human-like cognition
- Encompasses a wide range of computational models, including neural networks
- Proven to be successful despite the **curse of dimensionality**: the density of a fixed number of data points decreases exponentially with the number of dimensions
- Why? The information content in a data set implies inherent structure



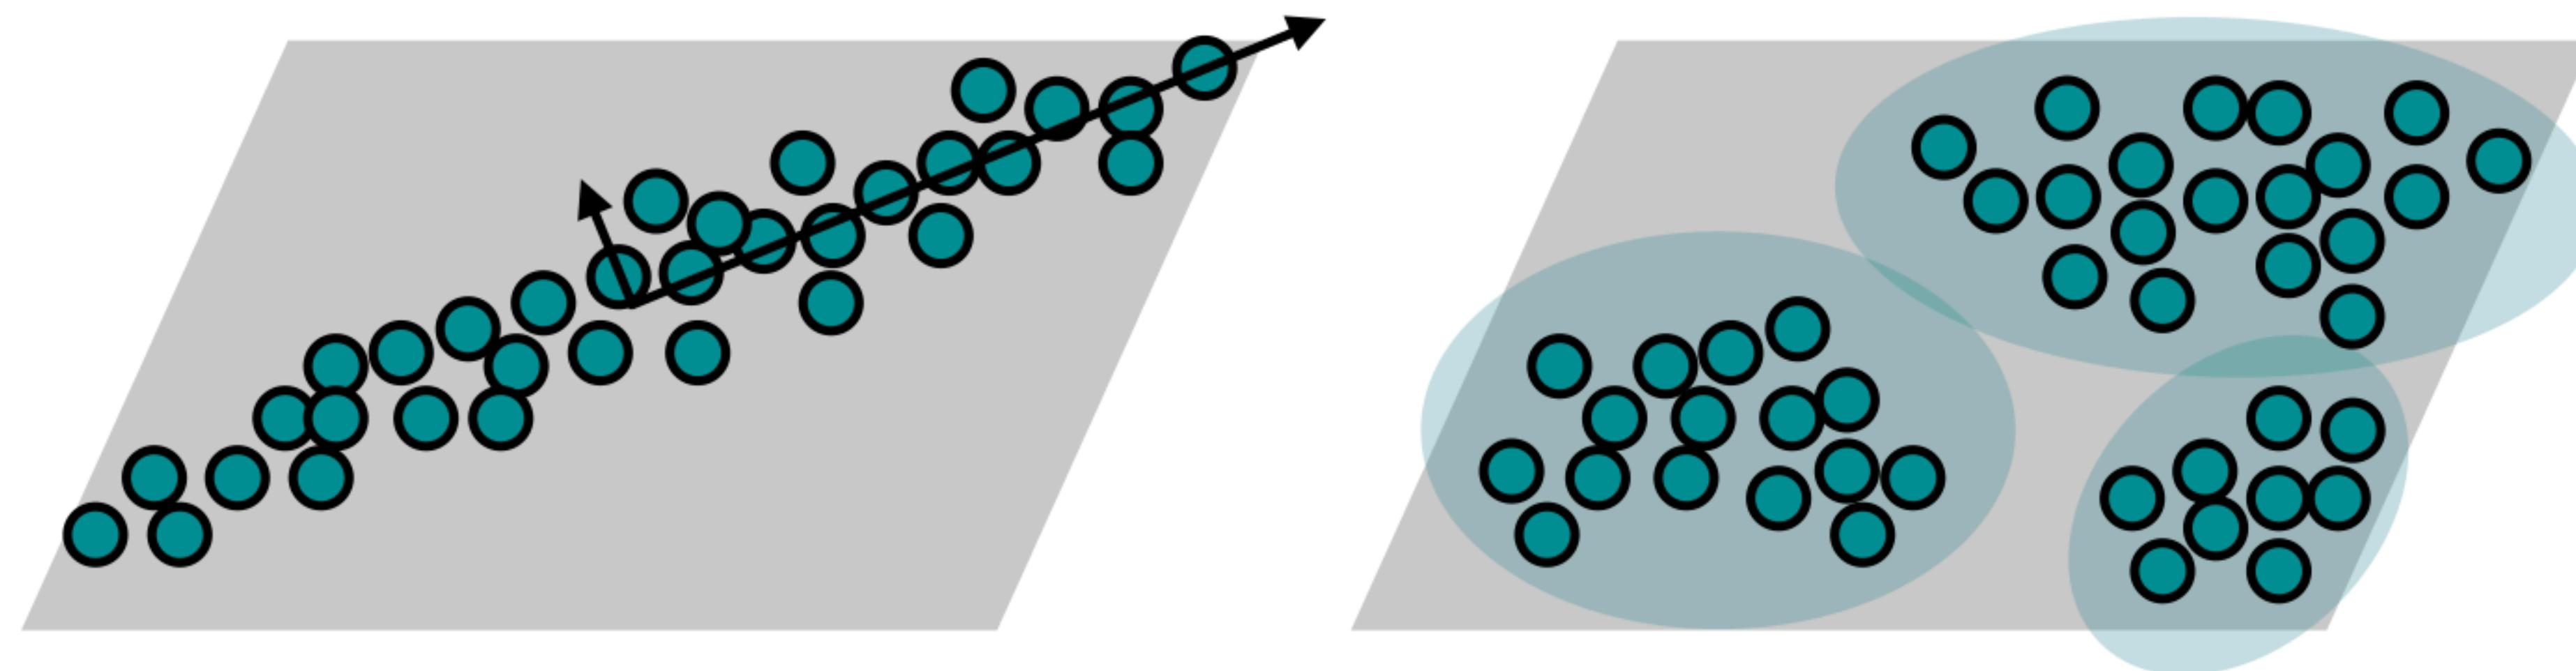
Supervised Learning

- Aims to find a mapping from the input data \mathbf{x} to corresponding output values \mathbf{y}
- The cost function has the generic form $C(\mathbf{y}(\mathbf{x}), \hat{\mathbf{y}}(\mathbf{x}))$, where $\hat{\mathbf{y}}$ is the output of the supervised learning model
- Main categories: regression and classification



Unsupervised Learning

- Aims to learn the underlying structure of the data set \mathbf{x}
- Cost function only depends on inputs $C(\mathbf{x})$
- Main categories: clustering, dimensionality reduction, and anomaly detection



Reinforcement Learning

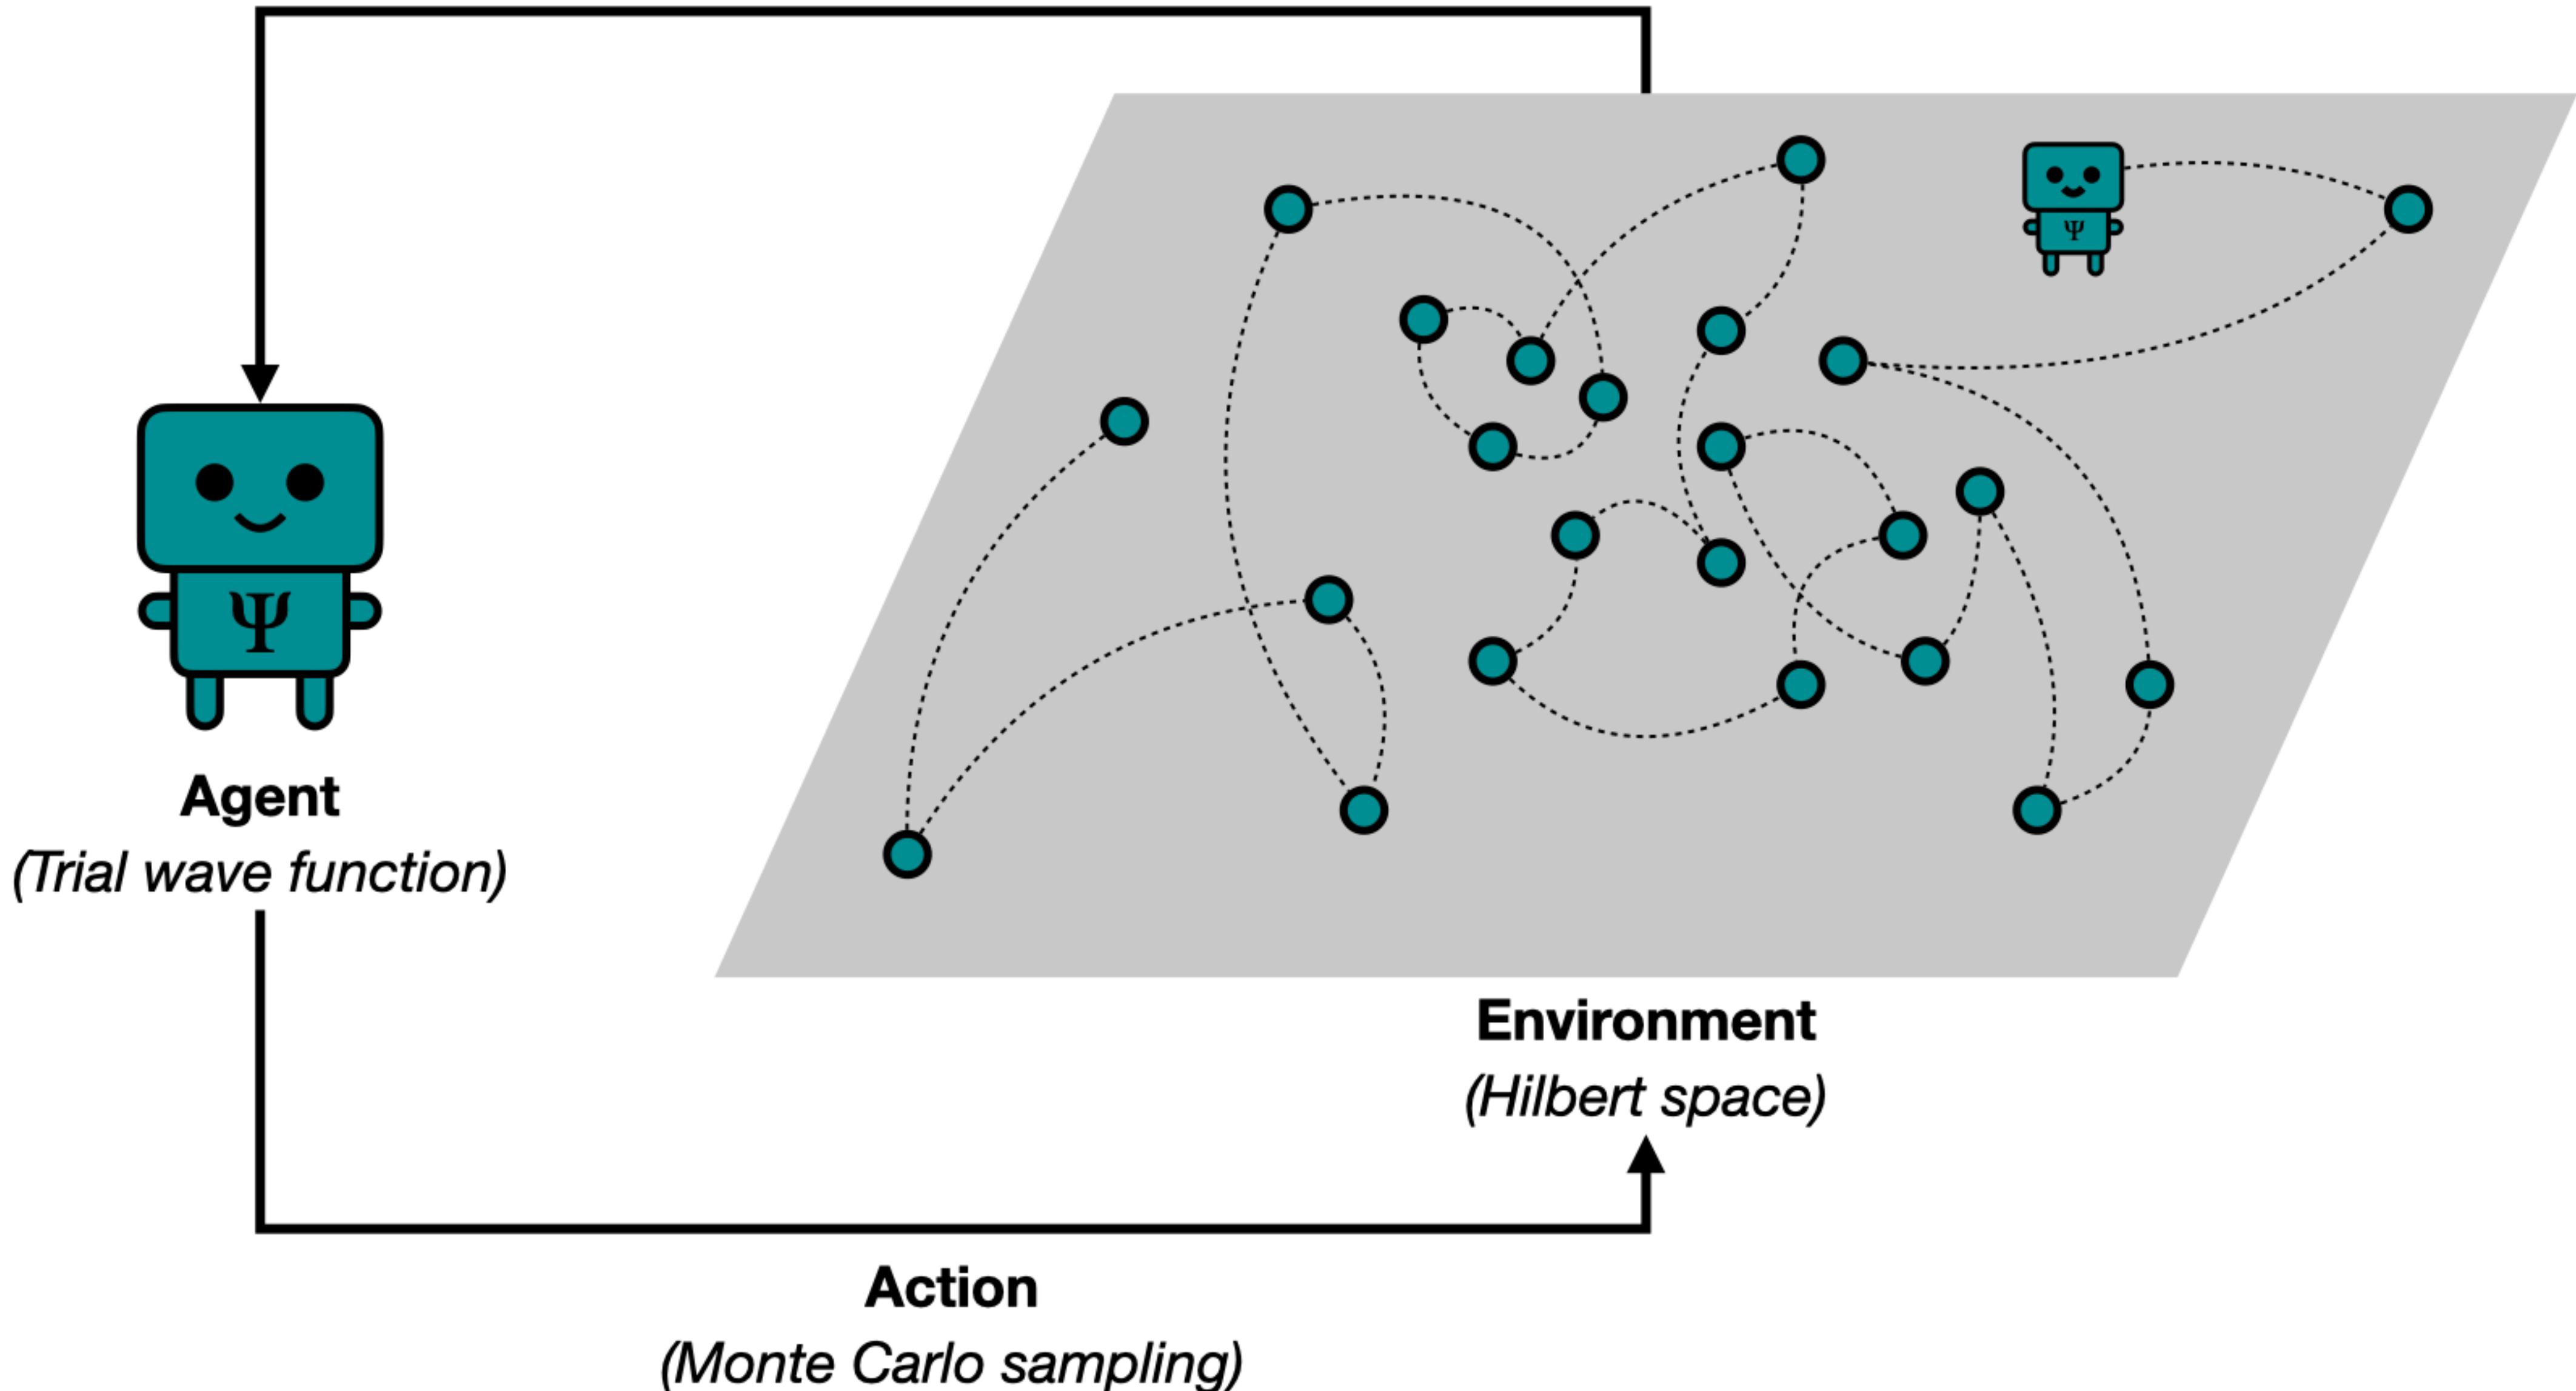
- Characterized by a data set that is dependent on the model itself
- Closest to how humans learn day-to-day (outside of the classroom)
- Cost function is recursive; it depends on all previous states of the model $\hat{\mathbf{y}}_i$ and data collected using that model $\mathbf{x}_i(\hat{\mathbf{y}}_i)$

$$C(\mathbf{x}_t(\hat{\mathbf{y}}_t(\mathbf{x}_{t-1}(\hat{\mathbf{y}}_{t-1}(\cdots(\mathbf{x}_0(\hat{\mathbf{y}}_0)\cdots))))))$$

- Commonly used to train robots in game-playing and navigation tasks

Compute cost or reward, update state

(Compute energy, update variational parameters)



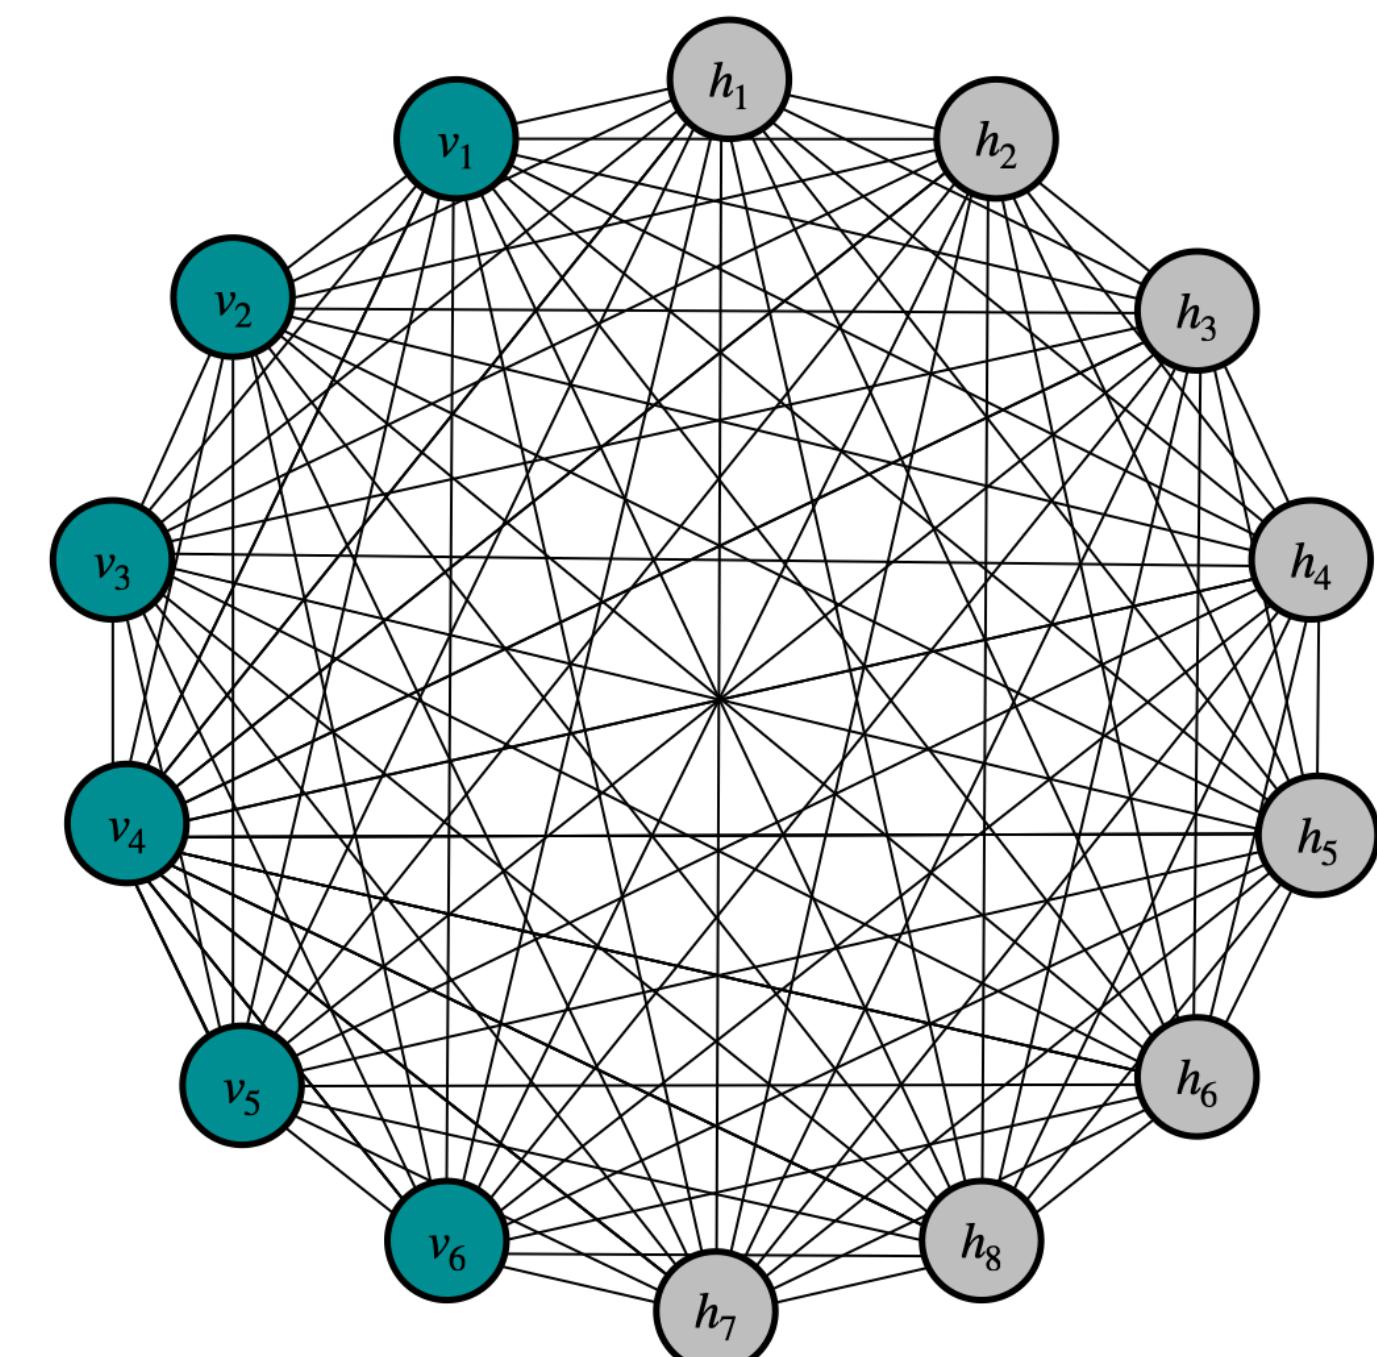
Artificial Neural Networks

Artificial Neural Networks

- Artificial neural networks (ANNs) are machine learning models inspired by the structure of the human brain, consisting of interconnected nodes (neurons) and nonlinear connections between them (synapses)
- The term is often used interchangeably with feedforward neural networks, the most common type of artificial neural network
- ANNs typically require a large amount of training data, only use when necessary
- Generally have low bias compared to other machine learning models, but there is still wide variation between different ANNs

Boltzmann Machines

- Generative stochastic artificial neural networks
- Typically used to learn the probability distribution of input data
- Two types of nodes that are fully connected:
 - ▶ **Visible nodes:** represents the inputs to the network
 - ▶ **Hidden nodes:** latent variables that capture higher-level features or patterns in the data
- General Boltzmann machines have limited applicability due to inefficient training

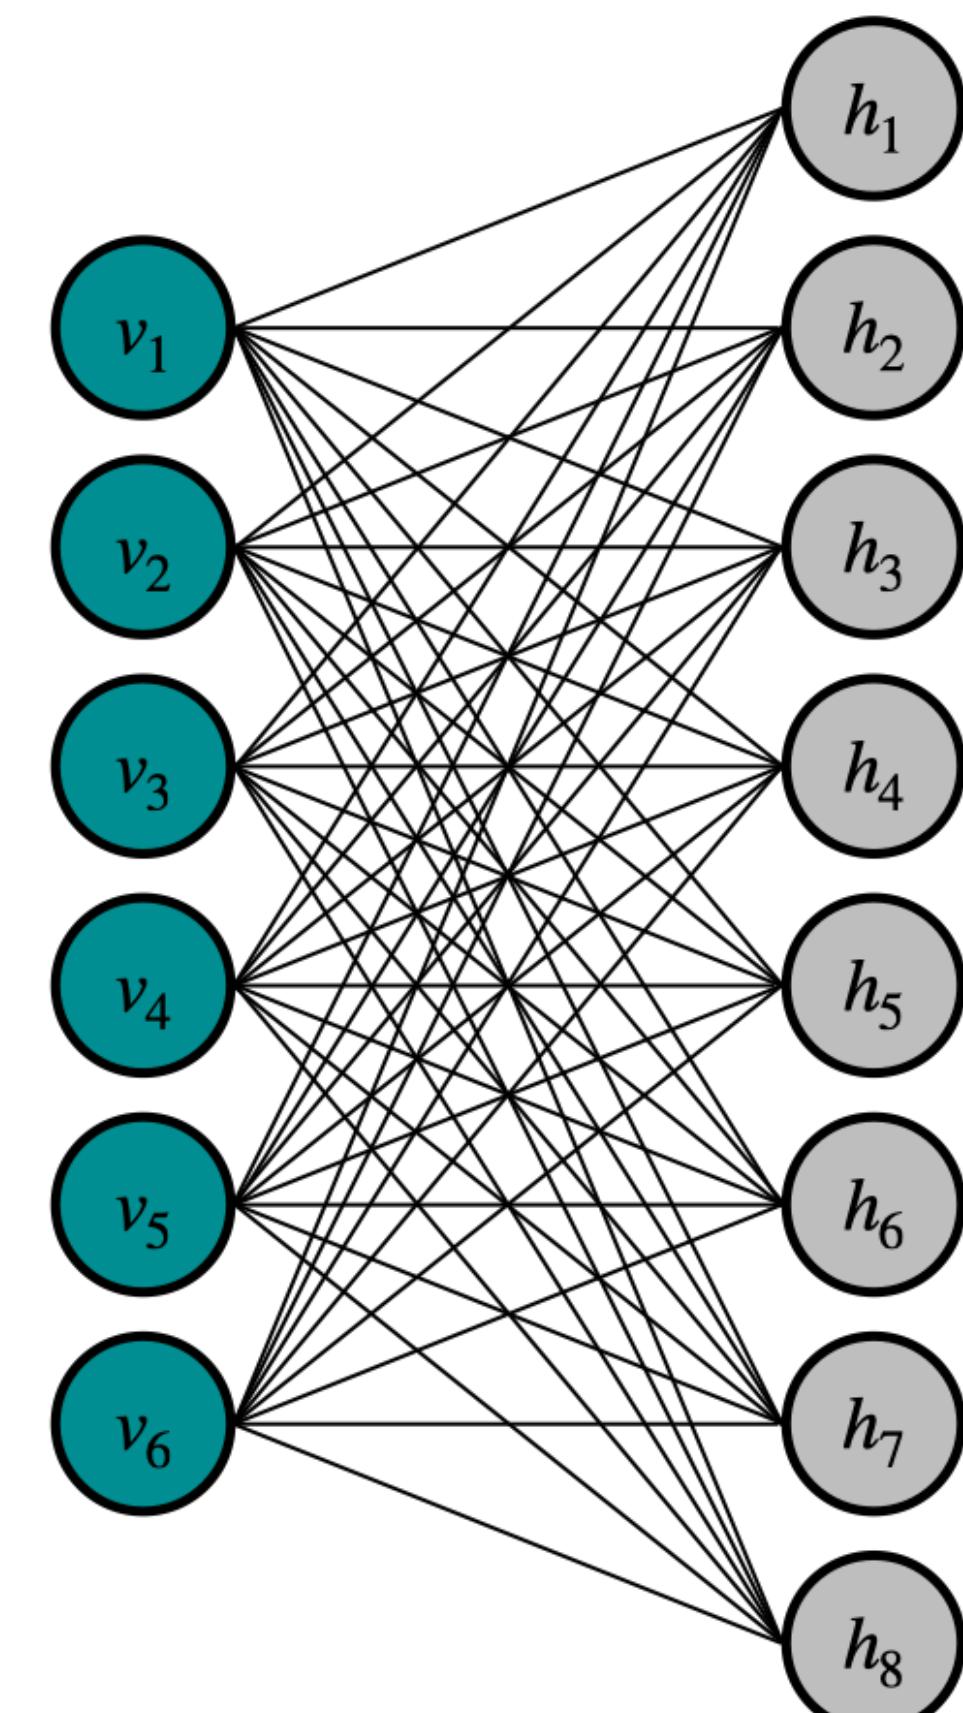


Restricted Boltzmann Machines

- A restricted Boltzmann machine (RBM) has no connections between nodes of the same type
- For continuous problems, the visible nodes are Gaussian and the hidden nodes are binary
- RBMs are constructed by defining an energy-like quantity

$$E(\mathbf{v}, \mathbf{h}) = \sum_{i=1}^V \frac{(v_i - a_i)^2}{2\sigma_i^2} - \sum_{j=1}^H b_j h_j - \sum_{i=1}^V \sum_{j=1}^H \frac{v_i}{\sigma_i^2} W_{ij} h_j$$

- Reparameterize the inverse variances as $1/\sigma_i^2 = \exp(s_i)$



Restricted Boltzmann Machines

- Taking the Boltzmann distribution yields a corresponding joint probability distribution

$$P(\mathbf{v}, \mathbf{h}) = \frac{1}{Z} e^{-E(\mathbf{v}, \mathbf{h})}$$

- Construct the marginal probability distribution of the visible nodes by summing over the possible values of the hidden nodes

$$P(\mathbf{v}) = \sum_{\mathbf{h}} P(\mathbf{v}, \mathbf{h})$$

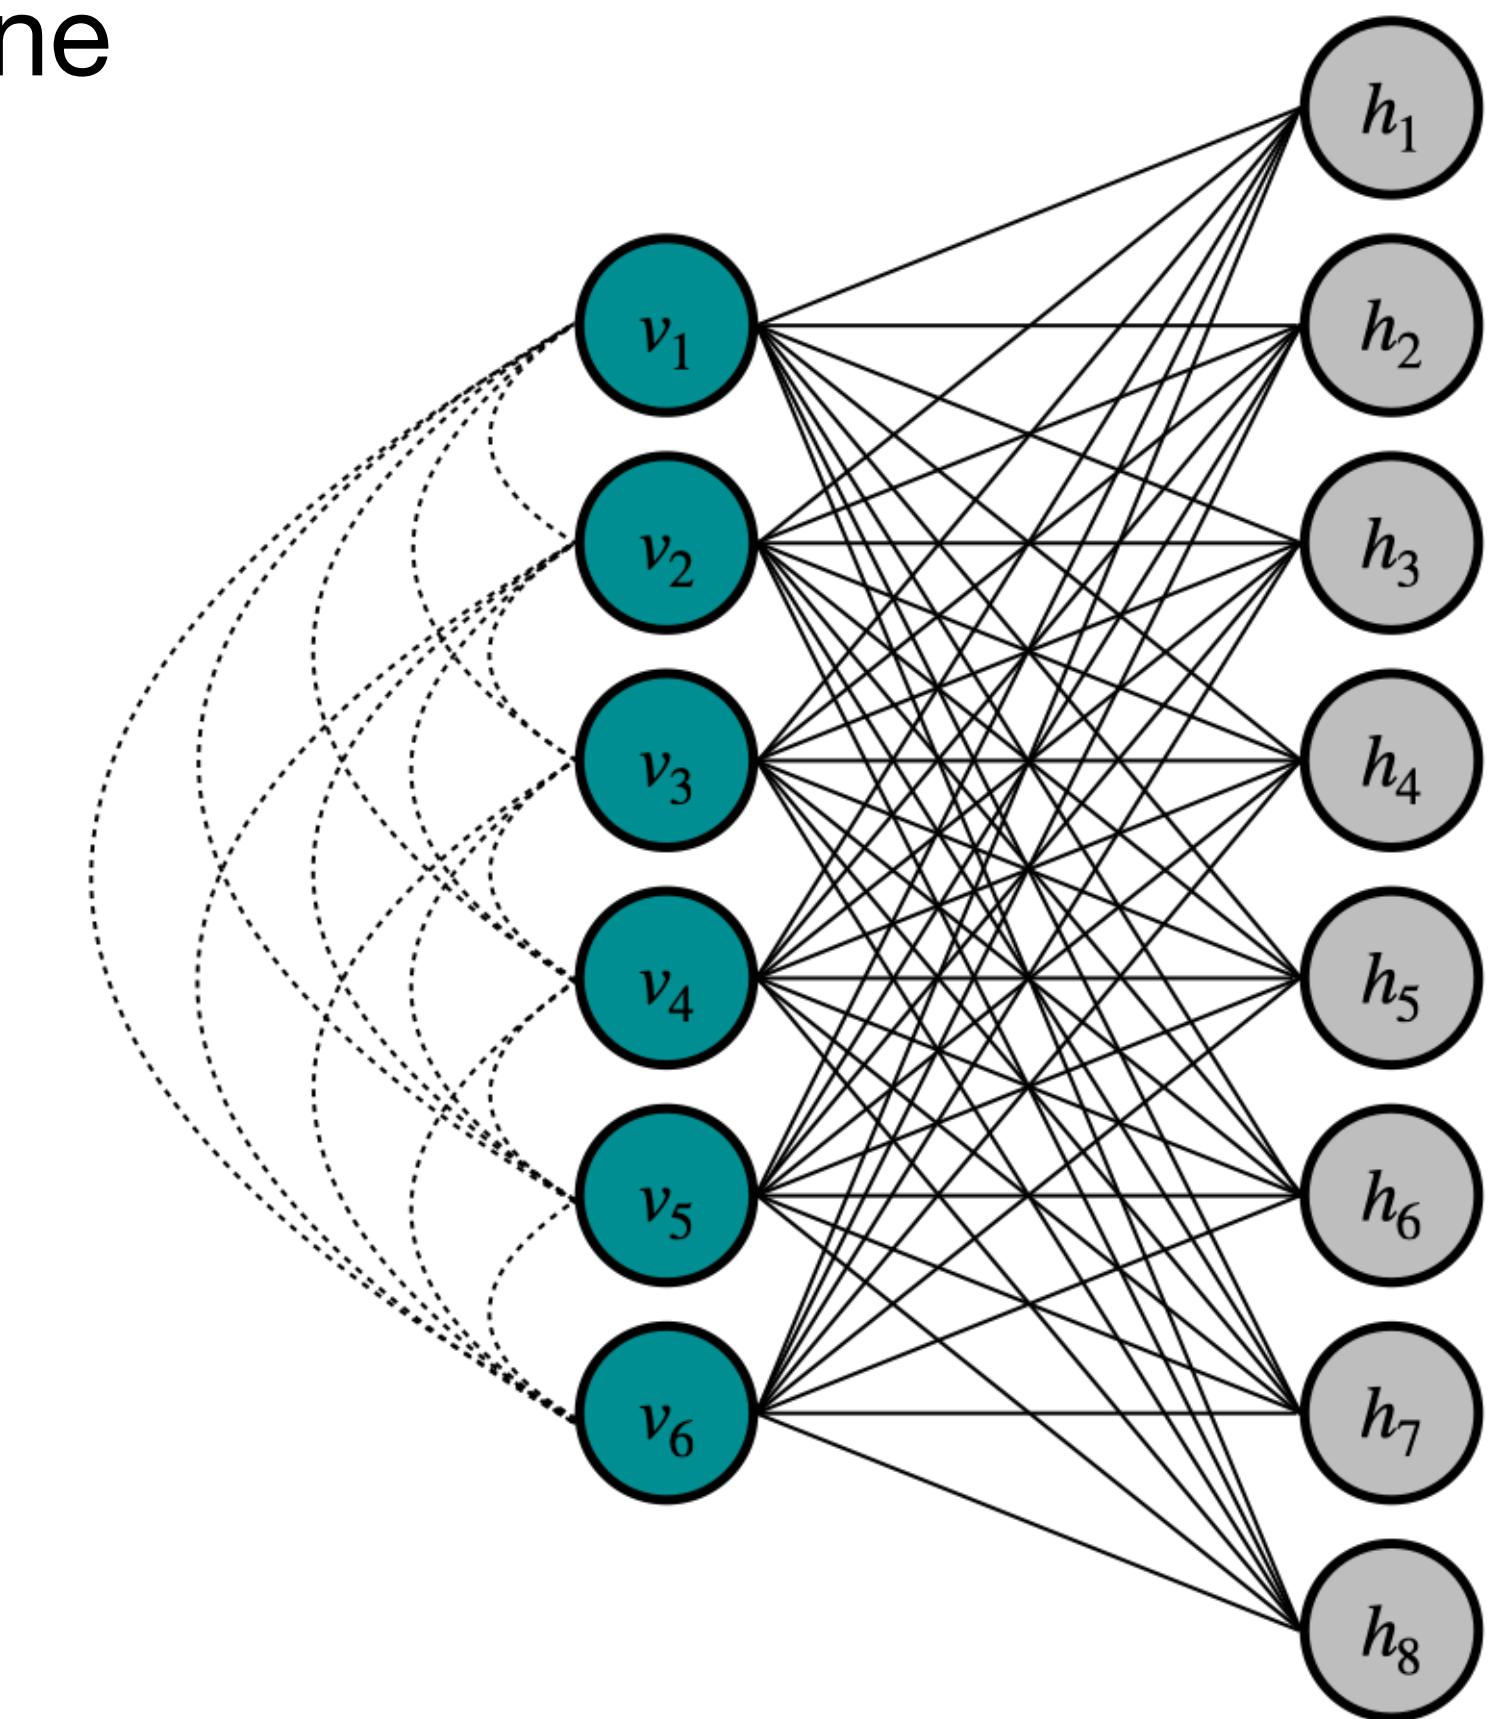
- Since the wave function is a probability amplitude, RBMs are a natural choice for a neural-network quantum state $\Psi_\theta(\mathbf{X}) = P(\mathbf{v}(\mathbf{X}))$

Multivariate Gaussian-Binary RBMs

- A multivariate Gaussian-binary restricted Boltzmann Machine (mGB-RBM) is halfway a standard Gaussian-Binary restricted Boltzmann Machine and a general Boltzmann machine
- Math becomes easier to write when connections between visible nodes are allowed
- Energy of a node configuration:

$$E(\mathbf{v}, \mathbf{h}) = \frac{1}{2}(\mathbf{v} - \mathbf{a})^T \Sigma^{-1} (\mathbf{v} - \mathbf{a}) - \mathbf{b}^T \mathbf{h} - \mathbf{v}^T \Sigma^{-1} W \mathbf{h},$$

where $\Sigma \equiv \text{Cov}(\mathbf{v}, \mathbf{v})$



Multivariate Gaussian-Binary RBMs

- The covariance matrix is reparameterized as

$$\Sigma_{ij}^{-1} = \begin{cases} \exp(S_{ij}), & \text{for } i = j, \\ S_{ij}, & \text{for } i \neq j. \end{cases}$$

- The log-likelihood becomes

$$\log P(\mathbf{v}) = -\frac{1}{2}(\mathbf{v} - \mathbf{a})^T \Sigma^{-1} (\mathbf{v} - \mathbf{a}) + \sum_{j=1}^H f(z_j(\mathbf{v})),$$

where $\mathbf{z}(\mathbf{v}) = \mathbf{b} + W^T \Sigma^{-1} \mathbf{v}$ and

$$f(x) \equiv \log(1 + \exp(x)) \quad (\text{softplus function})$$

Other Variations of Continuous RBMs

- If hidden binary nodes are allowed to take values of -1 and 1 (instead of 0 and 1), then

$$f(x) \equiv \log(\cosh(x))$$

(log-cosh function)

- If hidden nodes are continuous-valued between 0 and 1, then

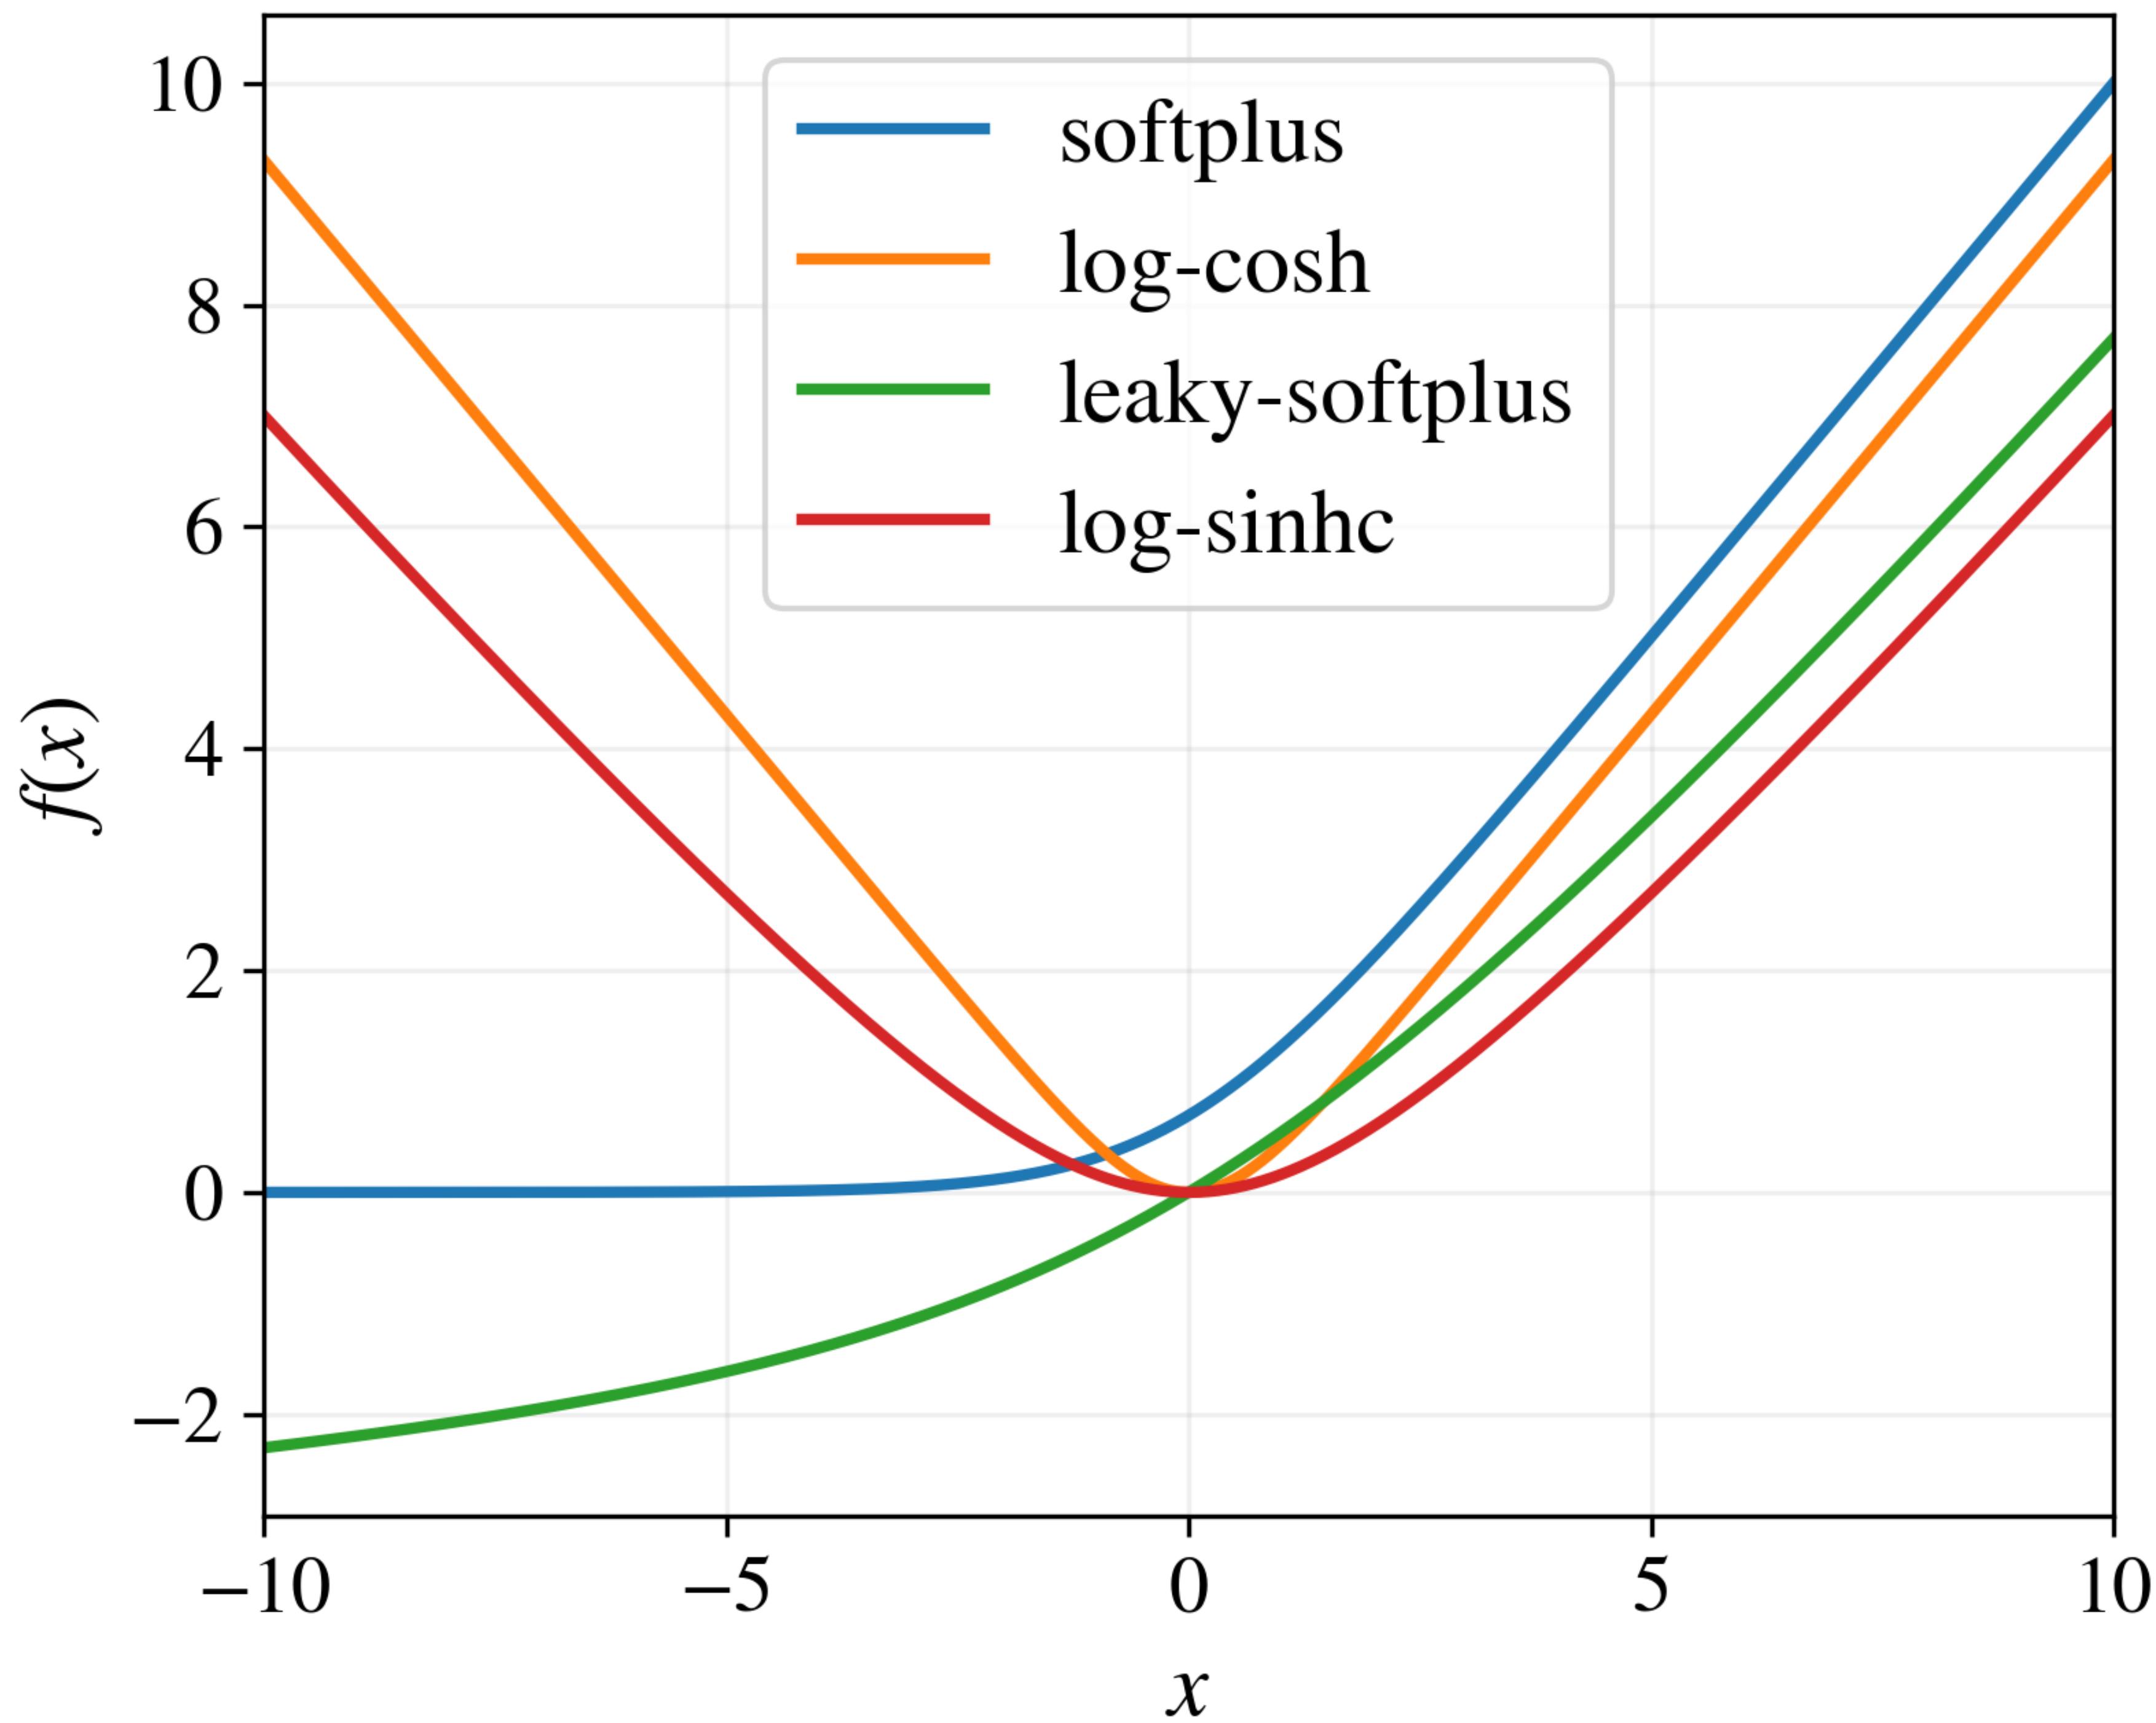
$$f(x) \equiv \log \frac{\exp(x) - 1}{x}$$

(leaky-softplus function)

- If hidden nodes are continuous-valued between -1 and 1, then

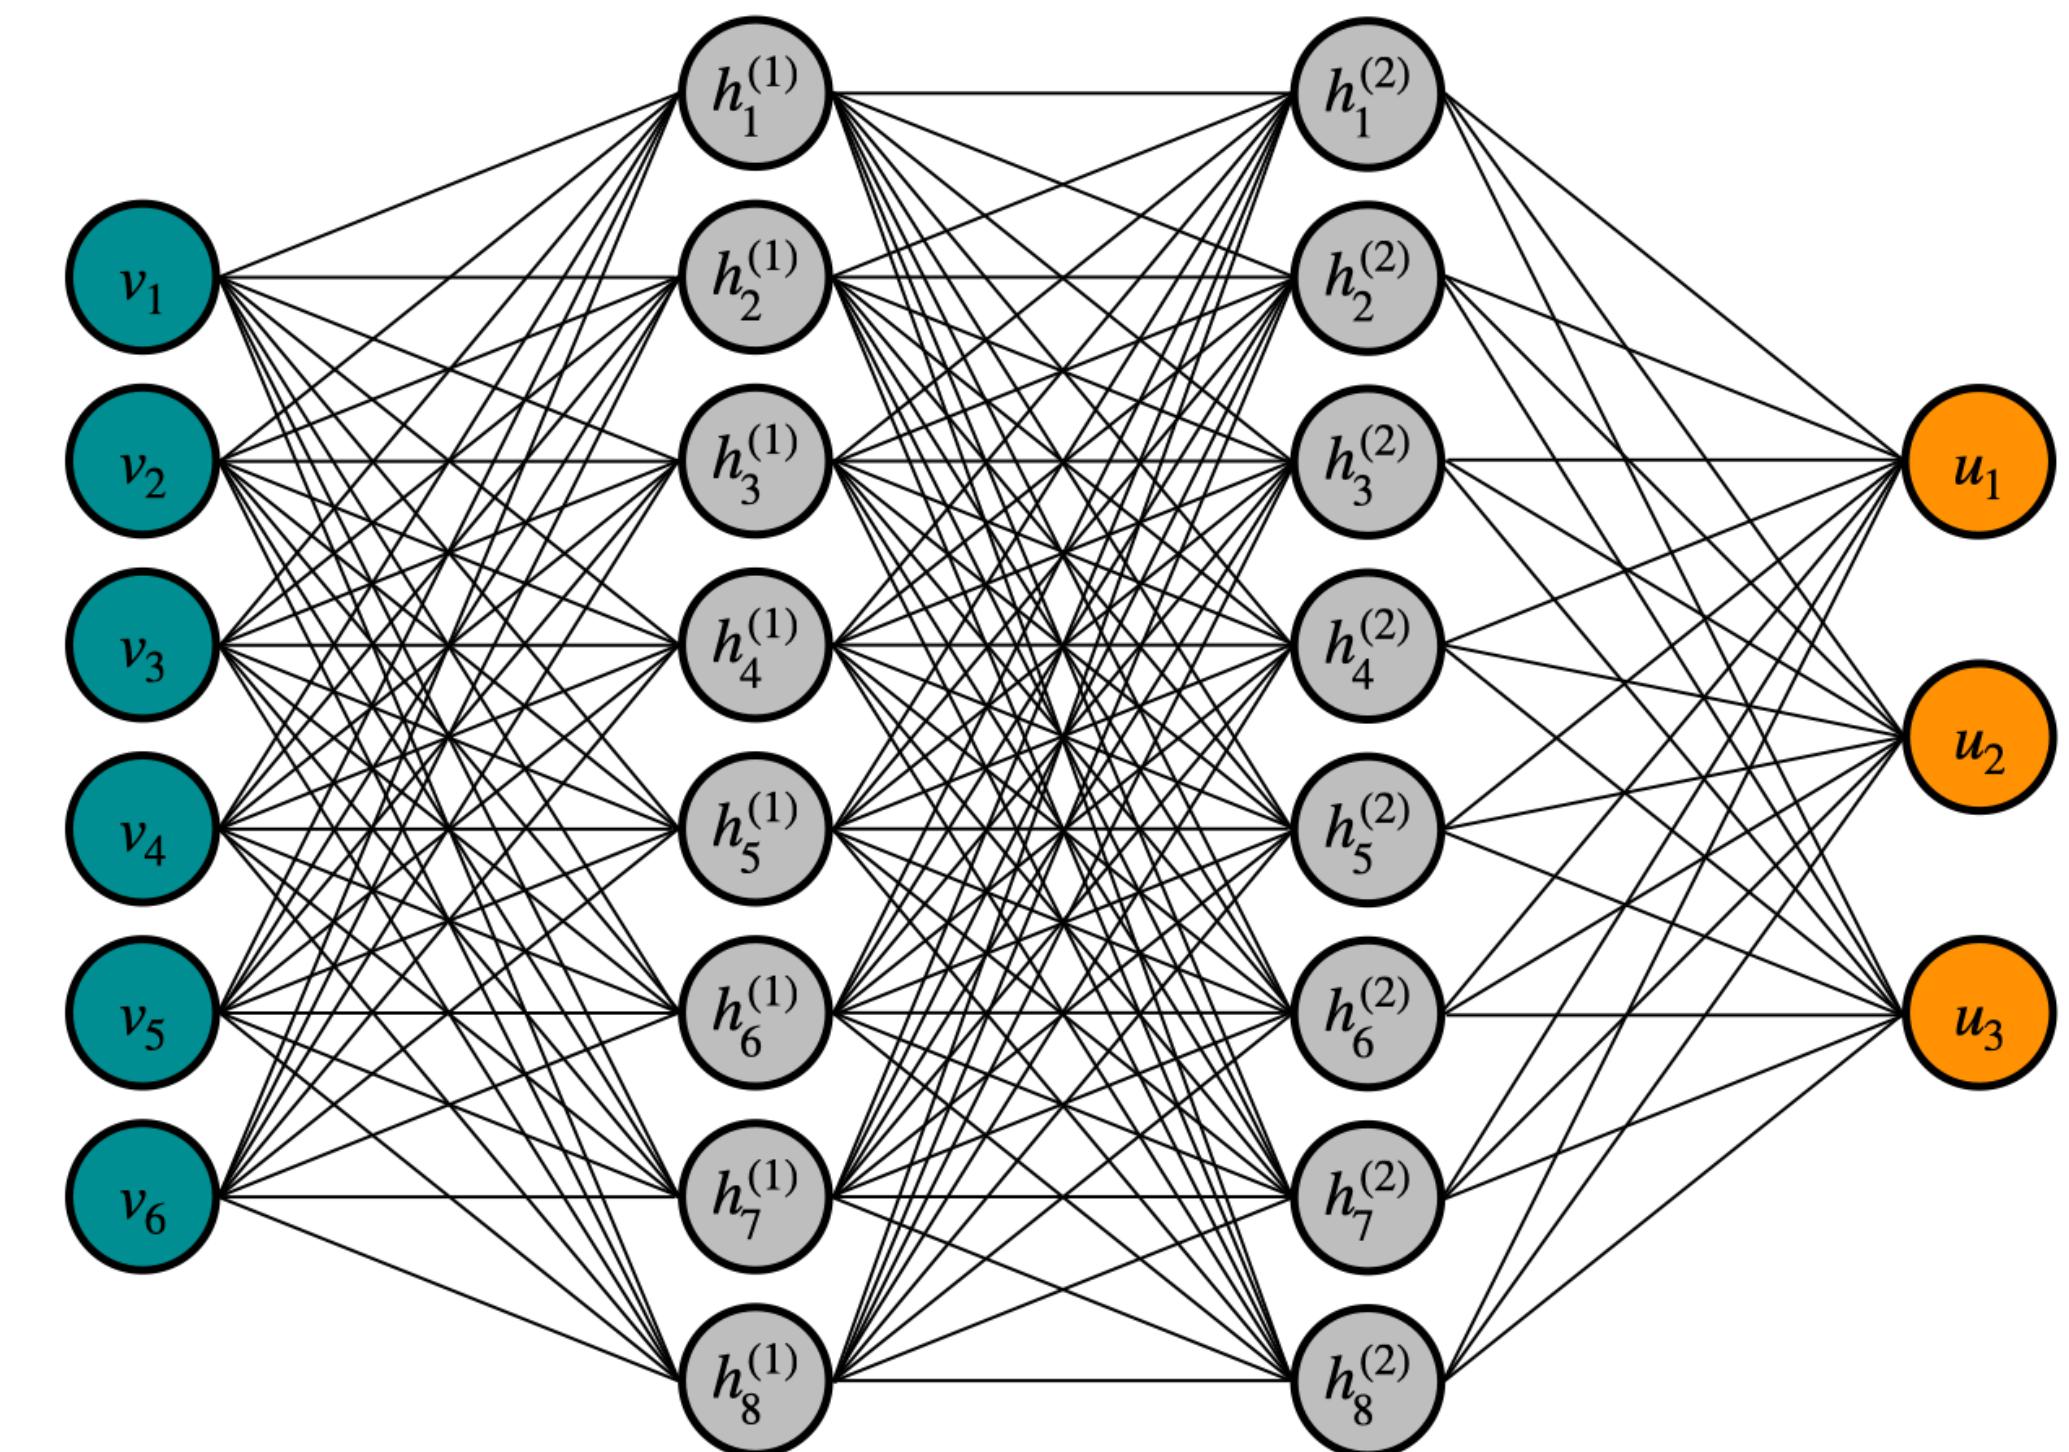
$$f(x) \equiv \log \frac{\sinh(x)}{x}$$

(log-sinhc function)



Feedforward Neural Networks

- Feedforward neural networks (FNNs) are the most widely used type of artificial neural network
- **Universal approximation theorem:** a shallow FNN can approximate any continuous function on a compact subset of Euclidean space, given there are a sufficient number of hidden nodes and the activation function satisfies certain conditions.
- Unlike RBMs, FNNs have explicit outputs, connections are directed, but trainable parameters are difficult to interpret



Feedforward Neural Networks

- Alternating compositions of affine transformations and simple, nonlinear transformations

$$\begin{aligned}\mathbf{h}^{(0)} &= \mathbf{v}, \\ \mathbf{h}^{(\ell)} &= \mathbf{f}_\ell \left(W^{(\ell)} \mathbf{h}^{(\ell-1)} + \mathbf{b}^{(\ell)} \right), \text{ for } \ell = 1, 2, \dots, L\end{aligned}$$

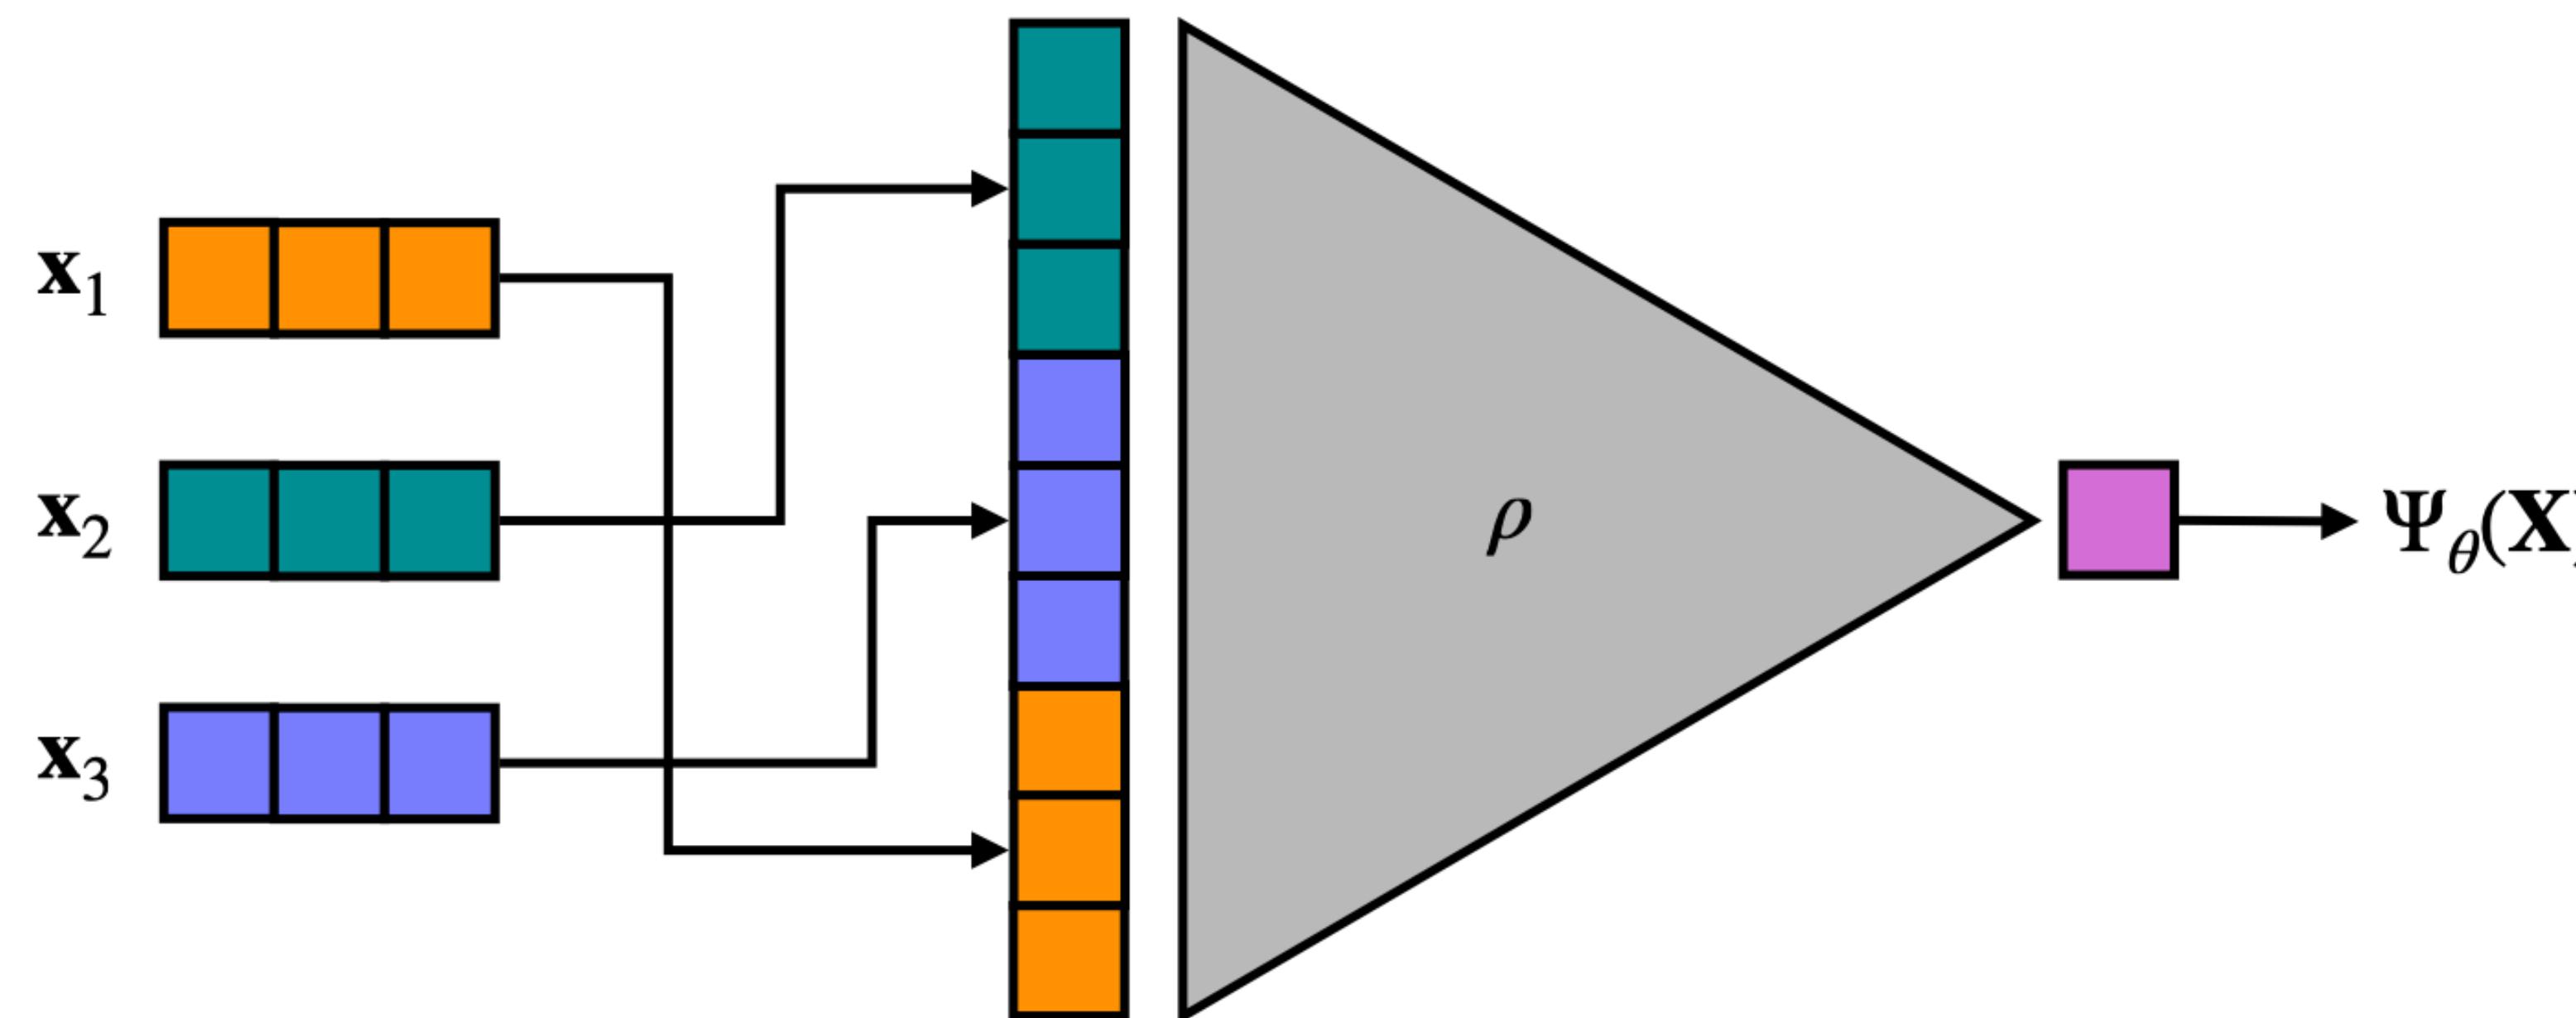
↑
Activation function

- Straightforward to write down the derivatives of the outputs w.r.t. the weights and biases using the chain rule, e.g.

$$\frac{\partial \mathbf{h}^{(L)}}{\partial W^{(\ell)}} = \frac{\partial \mathbf{h}^{(L)}}{\partial \mathbf{h}^{(L-1)}} \frac{\partial \mathbf{h}^{(L-1)}}{\partial \mathbf{h}^{(L-2)}} \cdots \frac{\partial \mathbf{h}^{(\ell+1)}}{\partial \mathbf{h}^{(\ell)}} \frac{\partial \mathbf{h}^{(\ell)}}{\partial W^{(\ell)}}$$

Permutation Symmetry

- Permutation symmetry can be enforced in two ways:
 - ▶ Data preprocessing: sort the inputs according to a fixed rule
 - ▶ Architecture design: set structure is ingrained into the network itself



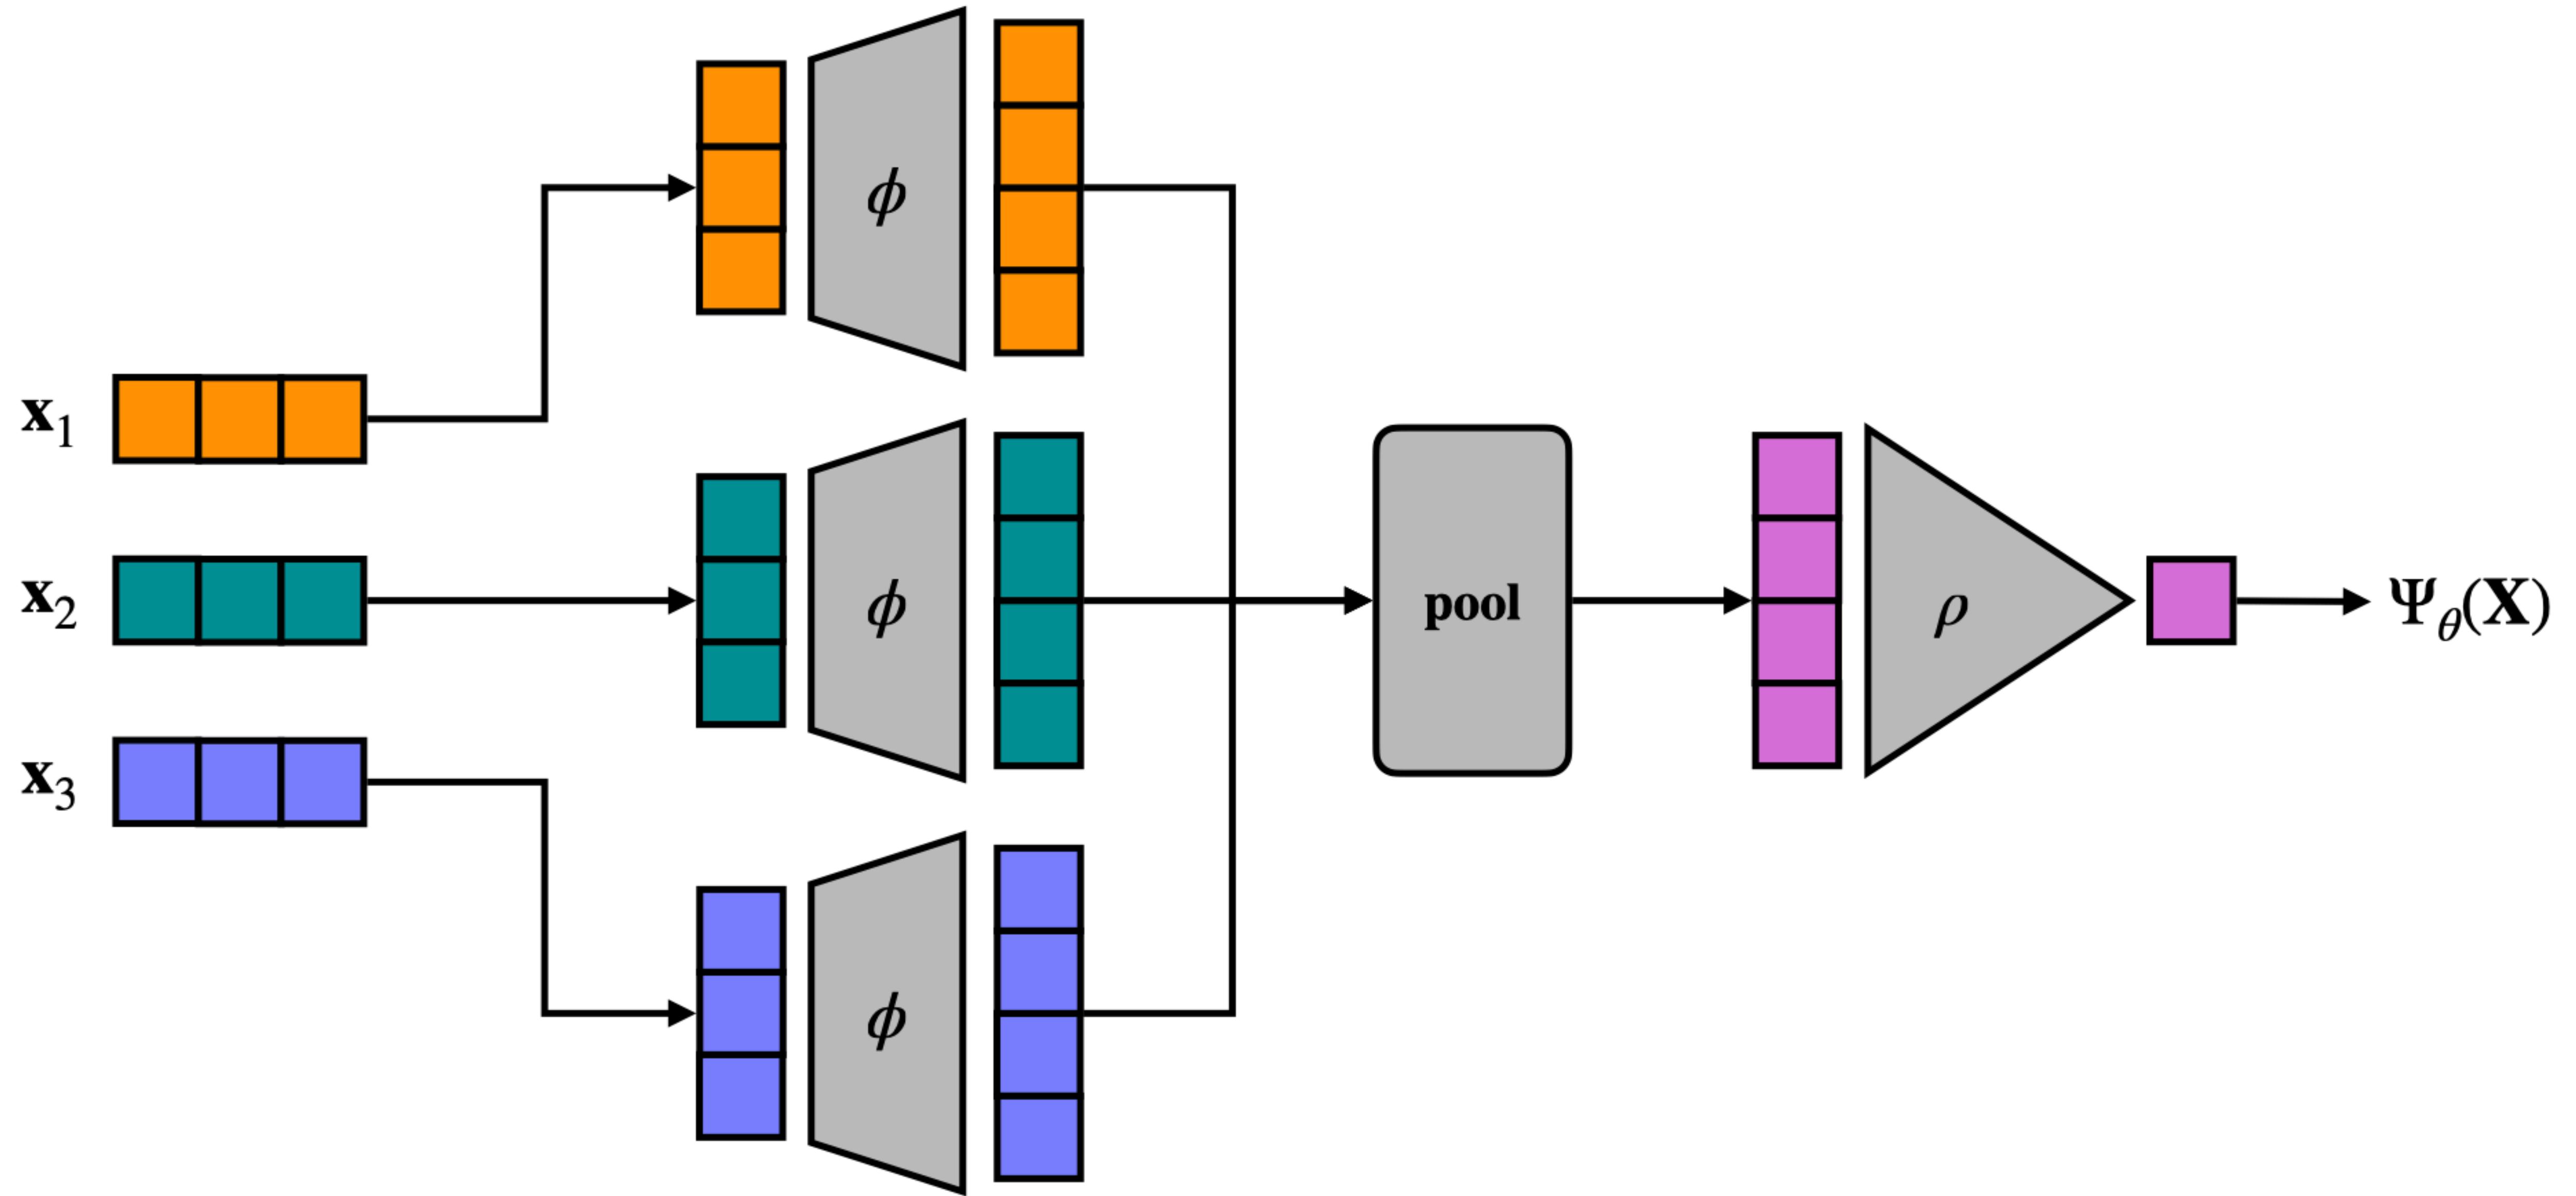
Deep Sets

- Introduced by Zaheer et al. in 2017 (arXiv:1703.06114)
- Suppose we want to learn on sets of N elements $\{\mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_N\}$
- Any permutation-invariant function can be decomposed as

$$f(\{\mathbf{v}_i\}) = \rho \left(\text{pool}(\{\phi(\mathbf{v}_i)\}) \right)$$



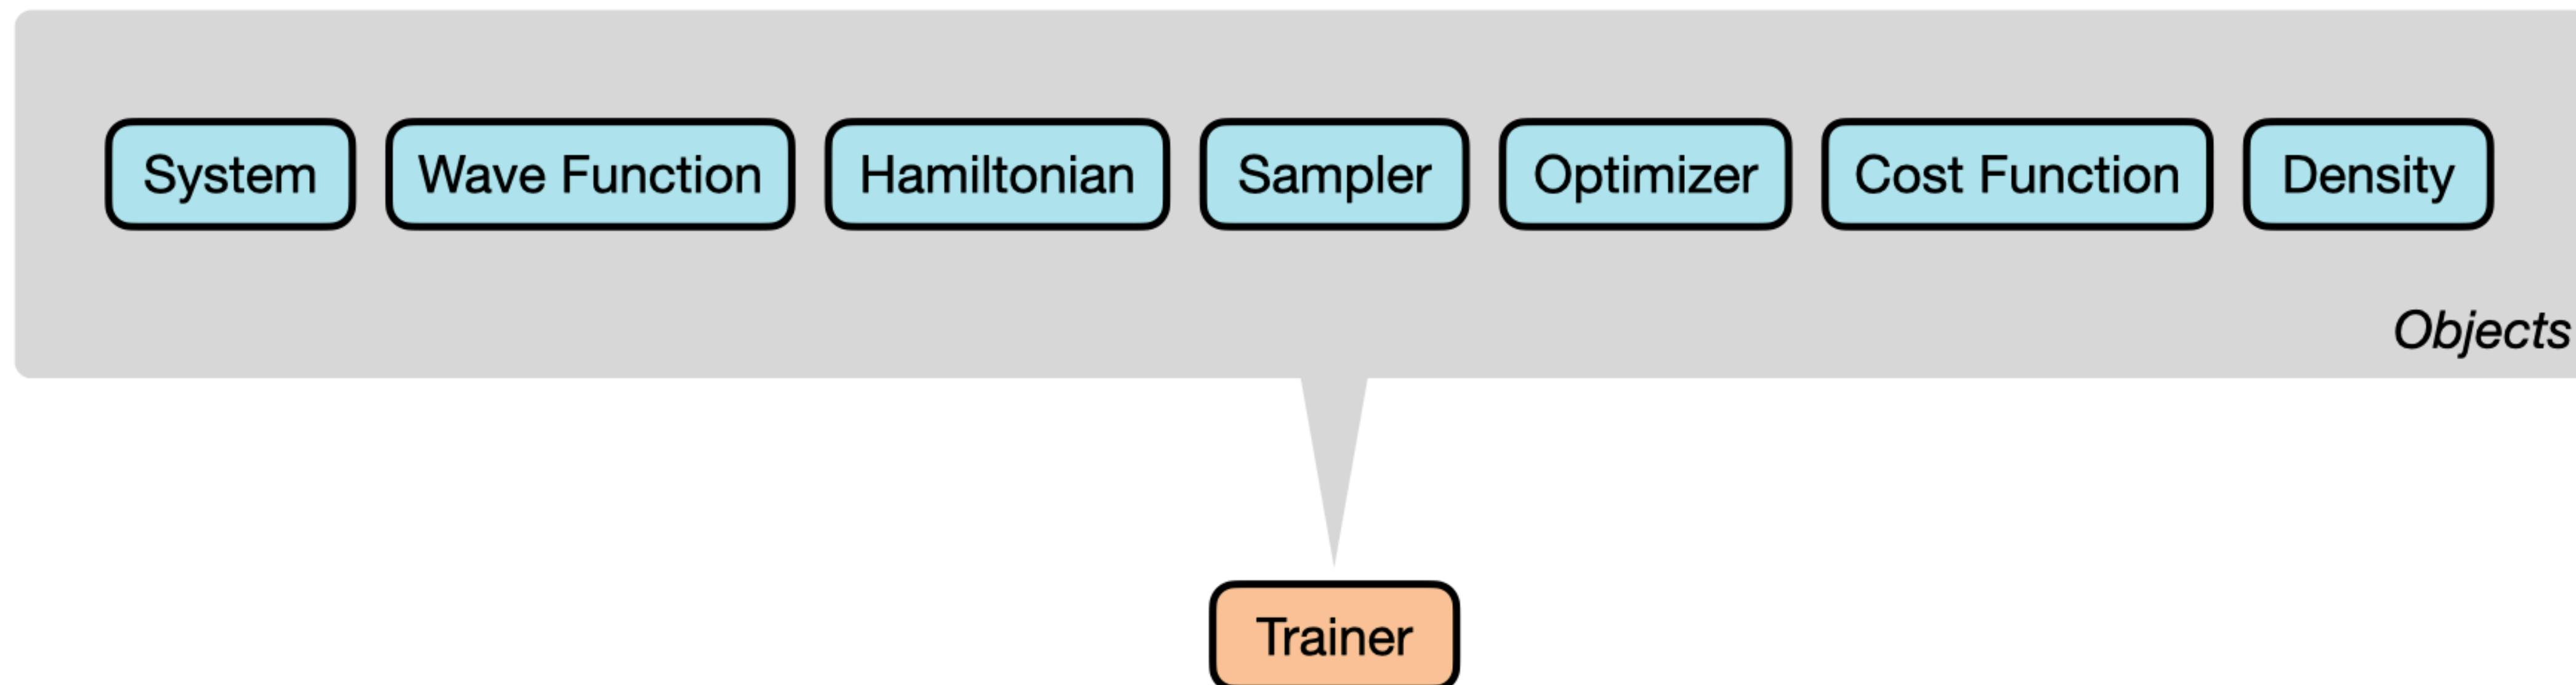
- If ρ and ϕ are feedforward neural networks, then $f(\{\mathbf{v}_i\})$ is a universal approximator for permutation-invariant functions



Implementations

NeuralAnsatz

- Object-oriented C++ software written from scratch on top of a lightweight, header-only linear algebra library called Eigen, parallelized with Open MPI
- Prioritizes modularity, reproducibility, low memory costs, and exact gradient calculations



NeuralAnsatz

- Can handle localized systems of bosons and fermions (periodic systems and nucleons are implemented but not yet validated)
- Available input formats: sorted and unsorted one- and two-body inputs
- Available wave functions: all variations of continuous RBMs, deep feedforward neural networks, deep sets (need validation), Slater determinants, Gaussian envelope, and product wrapper
- Local kinetic energy and gradients w.r.t. parameters computed exactly in terms of the trainable parameters
- For numerical stability, always use long double precision numbers, all computations are in terms of $\log \Psi_\theta$, and singularities in Hamiltonian are regularized.

NeuralAnsatz

- Brute force Metropolis sampler and Importance sampler both perturb all particles in one Monte Carlo step, and exchange spin to preserve total spin projection S_z
- Various optimizers are included, but simple stochastic gradient descent works best with stochastic reconfiguration
- Cost functions include the energy with stochastic reconfiguration (for reinforcement learning) and the overlap integral (for supervised learning)
- Density class computes one- or two-body distribution snapshots during training
- **Trainer can work for both supervised and reinforcement learning**

NeuralAnsatz

```
// main.cpp
mpi::initialize();
rng::initialize();

std::shared_ptr<Object> pObject = std::make_shared<Object>(...);

Trainer T(pSystem, pWaveFunc, pSampler, pOptimizer,
          pCostFunc, pDensity, filename);
T.train();
```

Initializes MPI and automatically finalizes MPI when program is finished

Initializes random number generator with a unique seed on each parallel process

Construct pointers to all required objects like this

Place arguments here

Trainer takes all pointers to required objects

Python/JAX Implementation

- Software initially developed by Alessandro Lovato for nuclei, and extended by Bryce Fore and myself for periodic systems
- JAX is an open-source software that combines the flexibility and ease-of-use of numpy with the efficiency of compiled languages like C++
 - ▶ Compute gradients with **grad**
 - ▶ Just-in-time compilation with **jit**
 - ▶ Vectorization with **vmap**
 - ▶ Parallelization with **pmap**
- Numerical differentiation provides far more freedom to construct more complex neural networks

Python/JAX Implementation

- Constructing deep feedforward neural networks with JAX is easy:

```
init_func, apply_func = stax.serial(*layers)
```

List of layer objects (e.g. stax.Dense, stax.Gelu, etc...)

- The “init” functions initialize the weight and biases of the network, given some input dimension
- The “apply” functions compute the output of the network given a set of parameters and input data

Python/JAX Implementation

- To create Deep Sets, just need to create two feedforward neural networks:

```
lat_init, lat_apply = stax.serial(stax.Dense(n), stax.Gelu, stax.Dense(n))
out_init, out_apply = stax.serial(stax.Dense(n), stax.Gelu, 1)
```

- Numpy-like syntax makes implementation of Deep Set easy:

```
phi = self.lat_apply(phi_params, X)
phi = jnp.sum(phi, axis=0)
logpsi = self.out_apply(rho_params, phi)
```

- Almost all artificial neural networks can be constructed by appropriately combining dense feedforward neural networks with the appropriate pooling and concatenation operations

Applications

Calogero-Sutherland Model

- Exactly-solvable model of strongly-interacting bosons in a one-dimensional harmonic oscillator trap
- Excellent basis to assess neural-network quantum states:
 - ▶ Learn how to handle the singularity without enforcing the cusp condition by hand
 - ▶ One-dimensional systems are easy to visualize
 - ▶ Benchmark with exact solutions for the ground state energy *and* the ground state wave function

Calogero-Sutherland Model

- Hamiltonian ($\hbar = m = 1$):

$$\hat{H} = \sum_{i=1}^N \left(-\frac{1}{2} \frac{\partial^2}{\partial x_i^2} + \frac{1}{2} \omega^2 x_i^2 \right) + \sum_{i < j} \frac{\beta(\beta-1)}{x_{ij}^2}$$

- Exact solutions:

$$\Psi_0(\mathbf{X}) = \exp \left(- \sum_{i=1}^N \frac{1}{2} \omega x_i^2 \right) \prod_{i < j} x_{ij}^\beta$$

$$E_0 = \frac{1}{2} N \omega + \frac{1}{2} \beta N(N-1) \omega$$

Throughout this work,
 $N = 6$ and $\beta = 2$

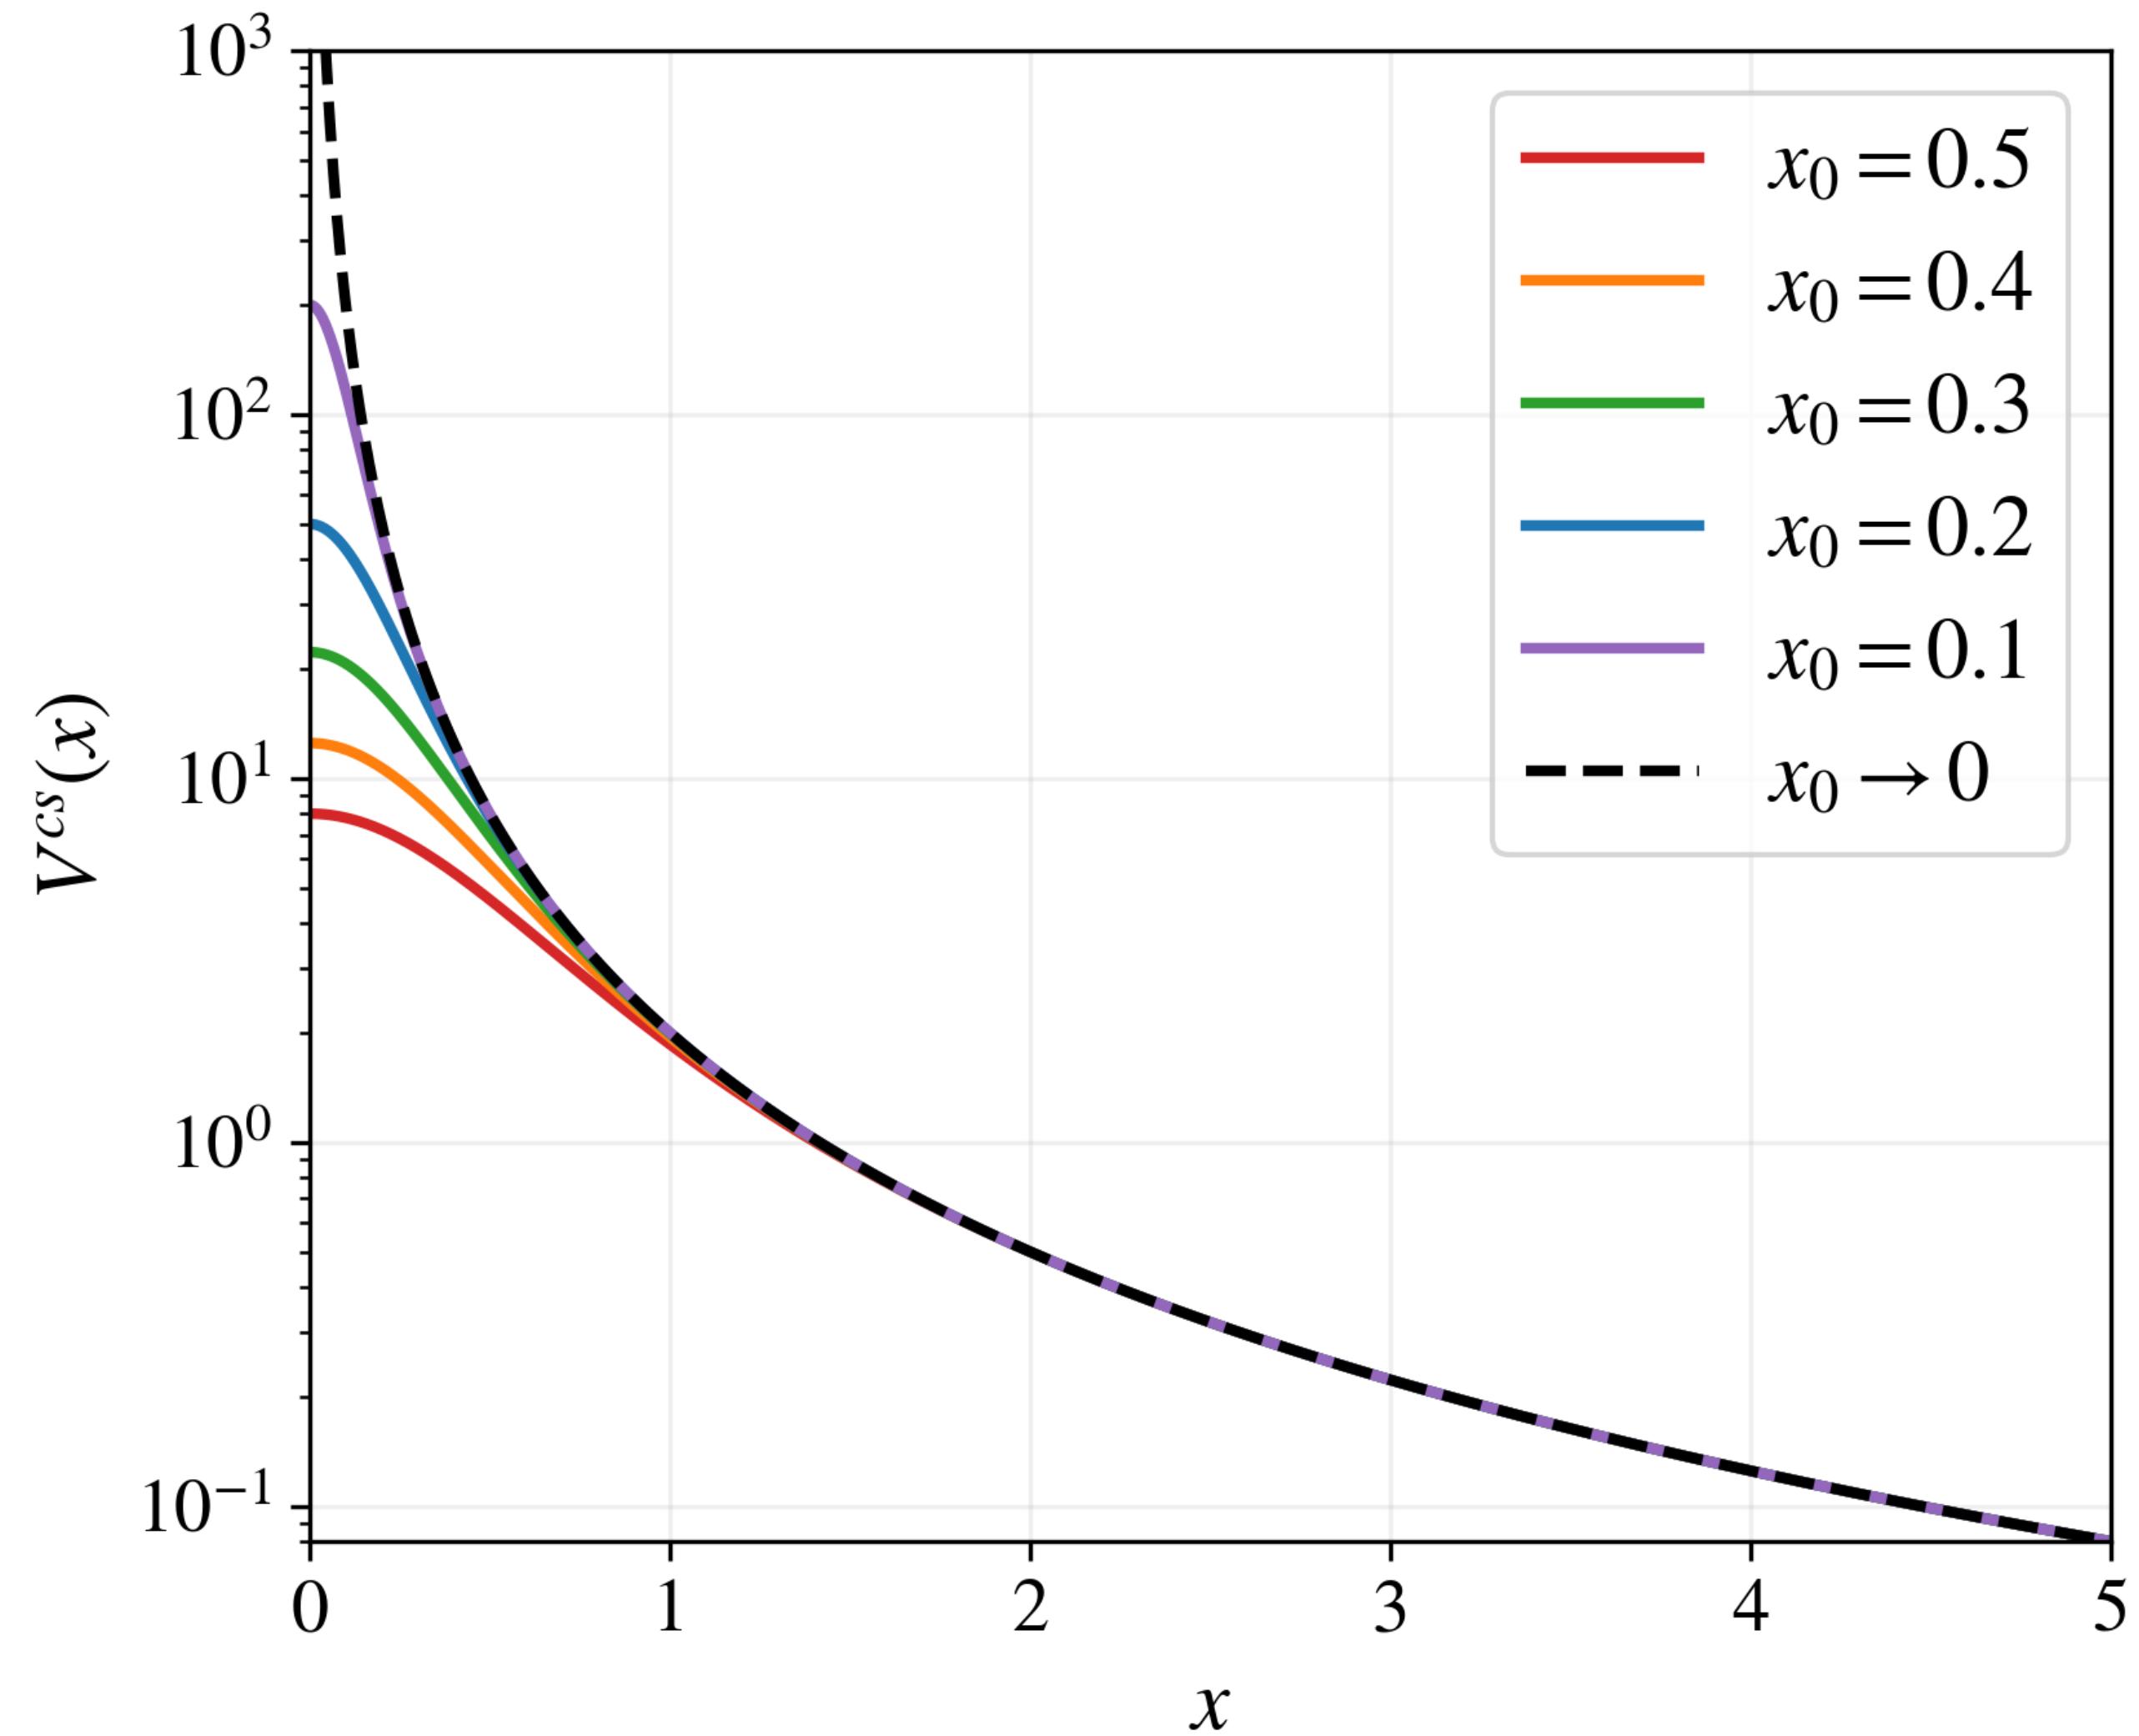
Calogero-Sutherland Model

- Regularize singularity:

$$\frac{1}{x^2} \mapsto \frac{\tanh^2(x/x_0)}{x^2}$$

- Transfer learning:

Quickly train networks on very soft interaction ($x_0 = 0.5$) before moving on to harder interactions ($x_0 = 0.01$)



Calogero-Sutherland Model

- Continuous RBMs:

$$\Psi_{\theta}^{rbm}(\mathbf{X}) = P(\mathbf{v}(\mathbf{X}))$$

- FNNs with a very wide Gaussian envelope ($\sigma = 16$):

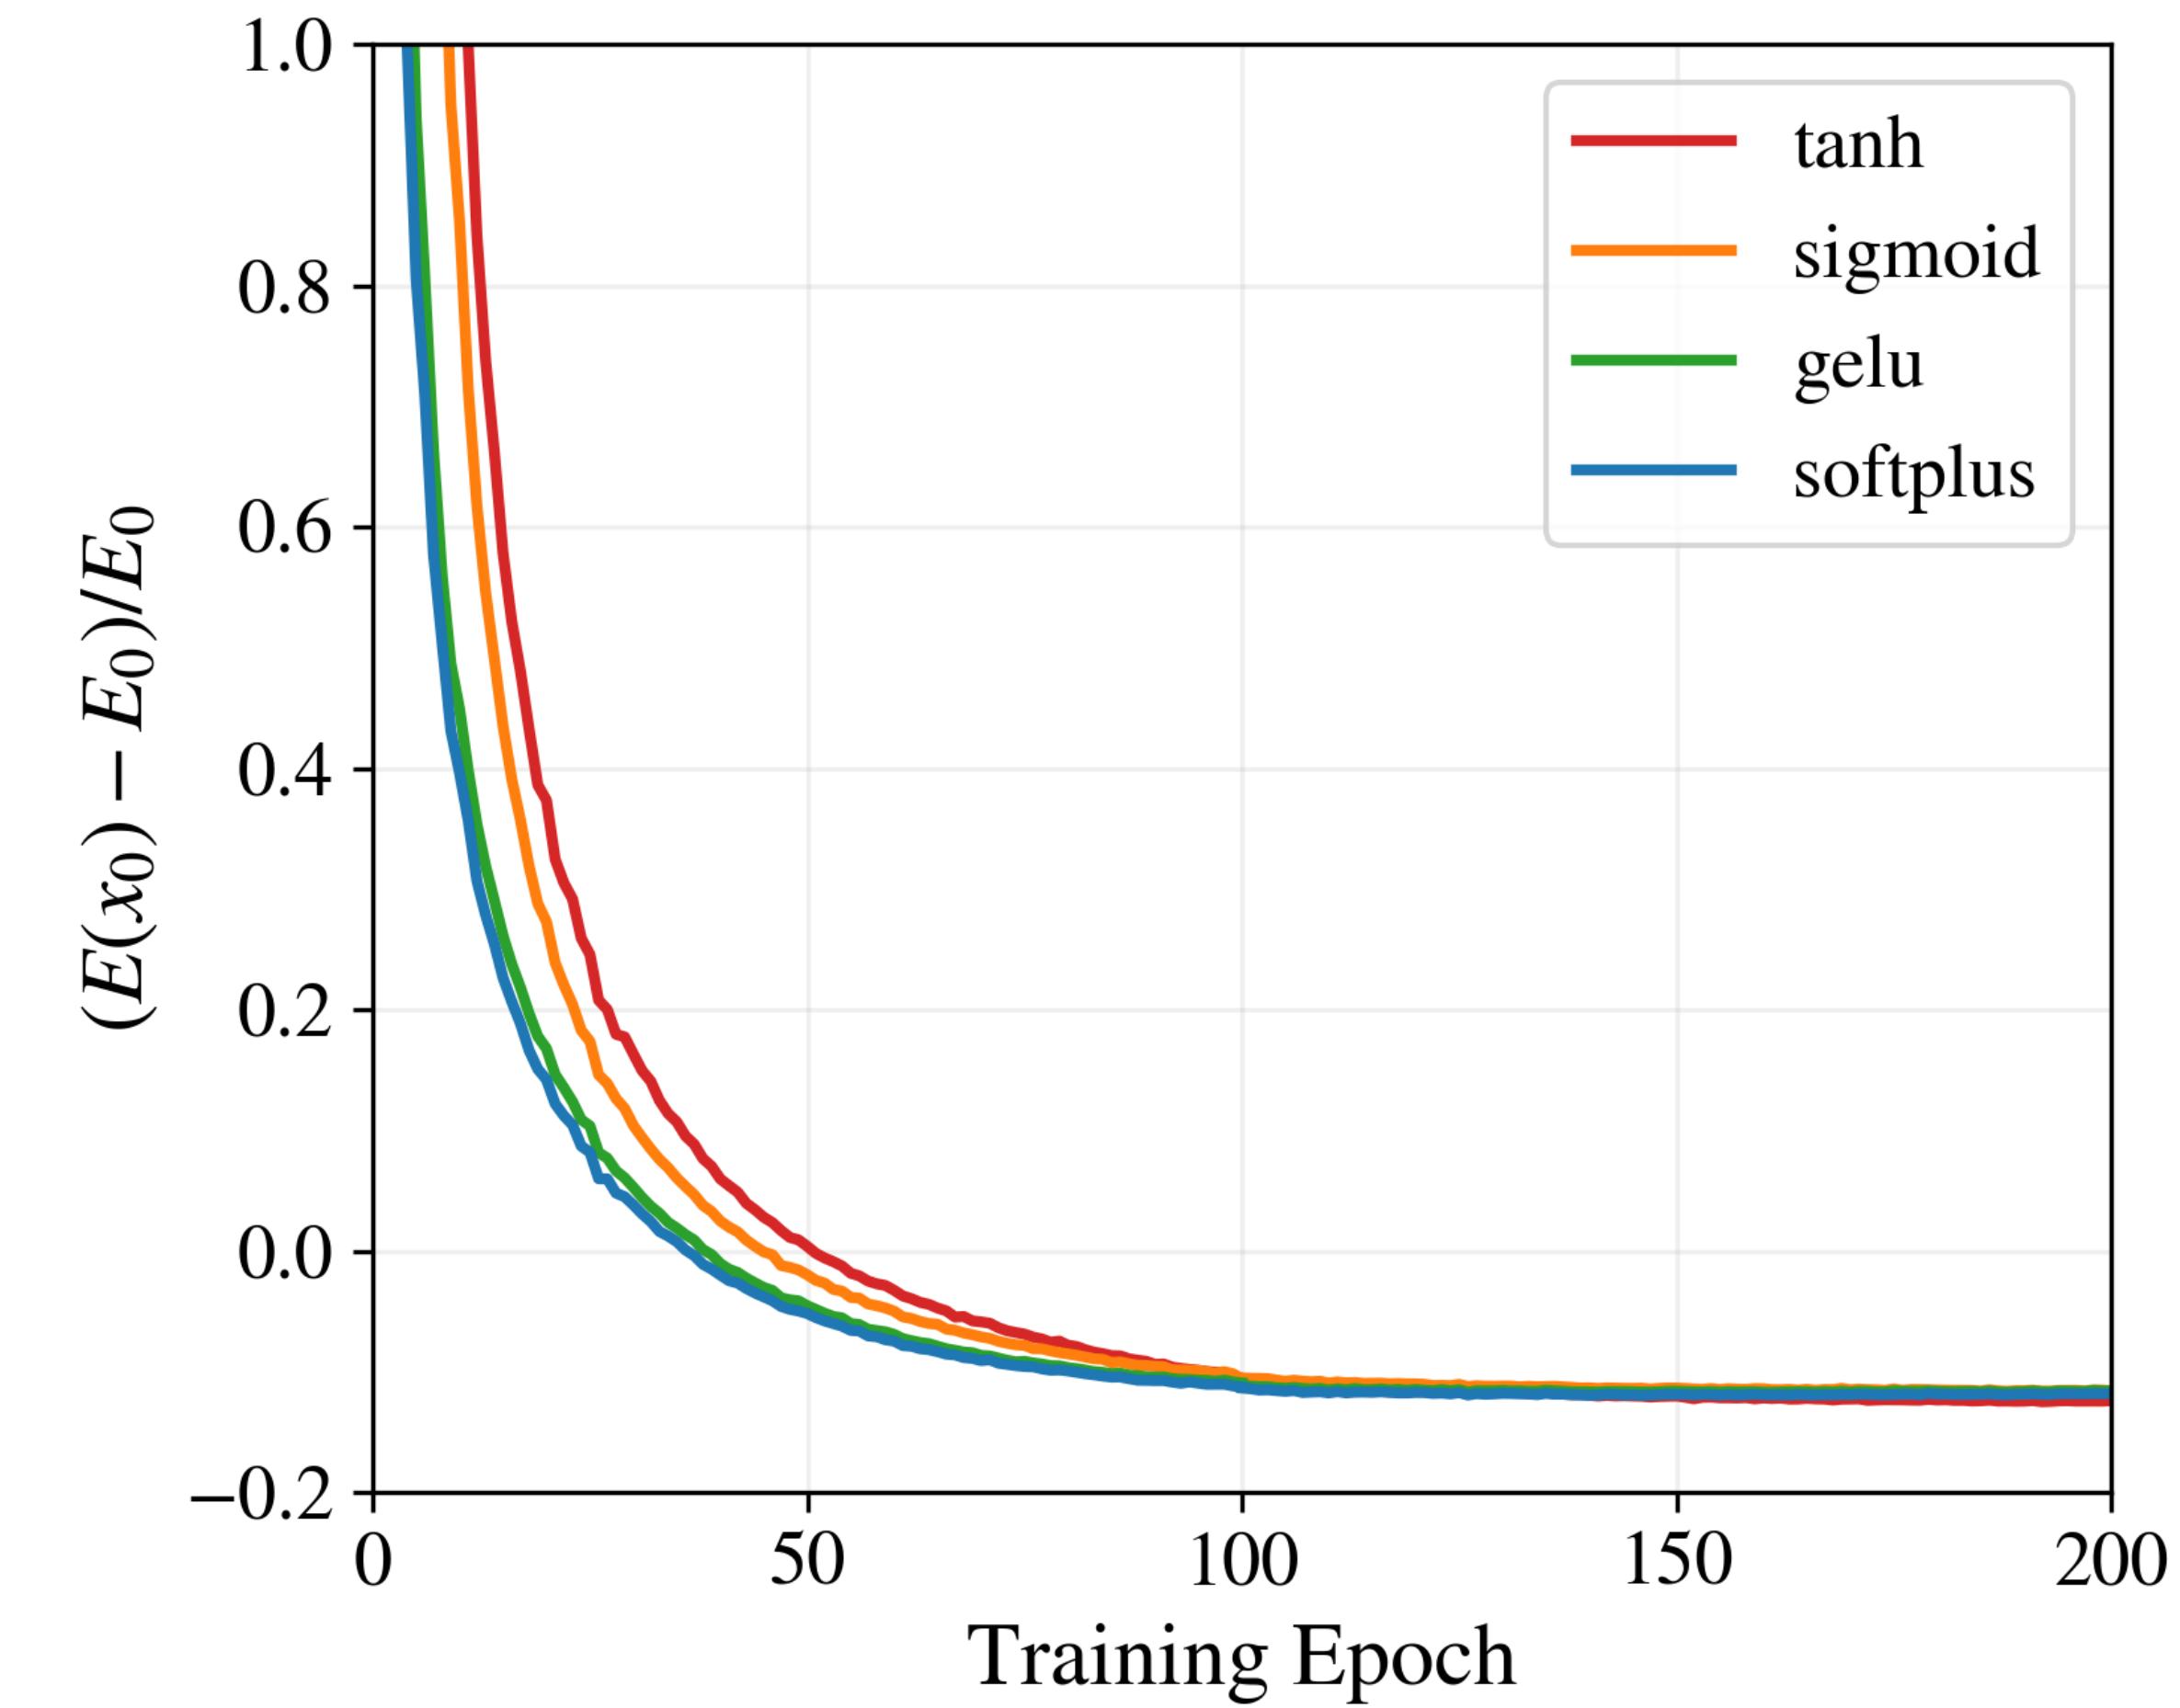
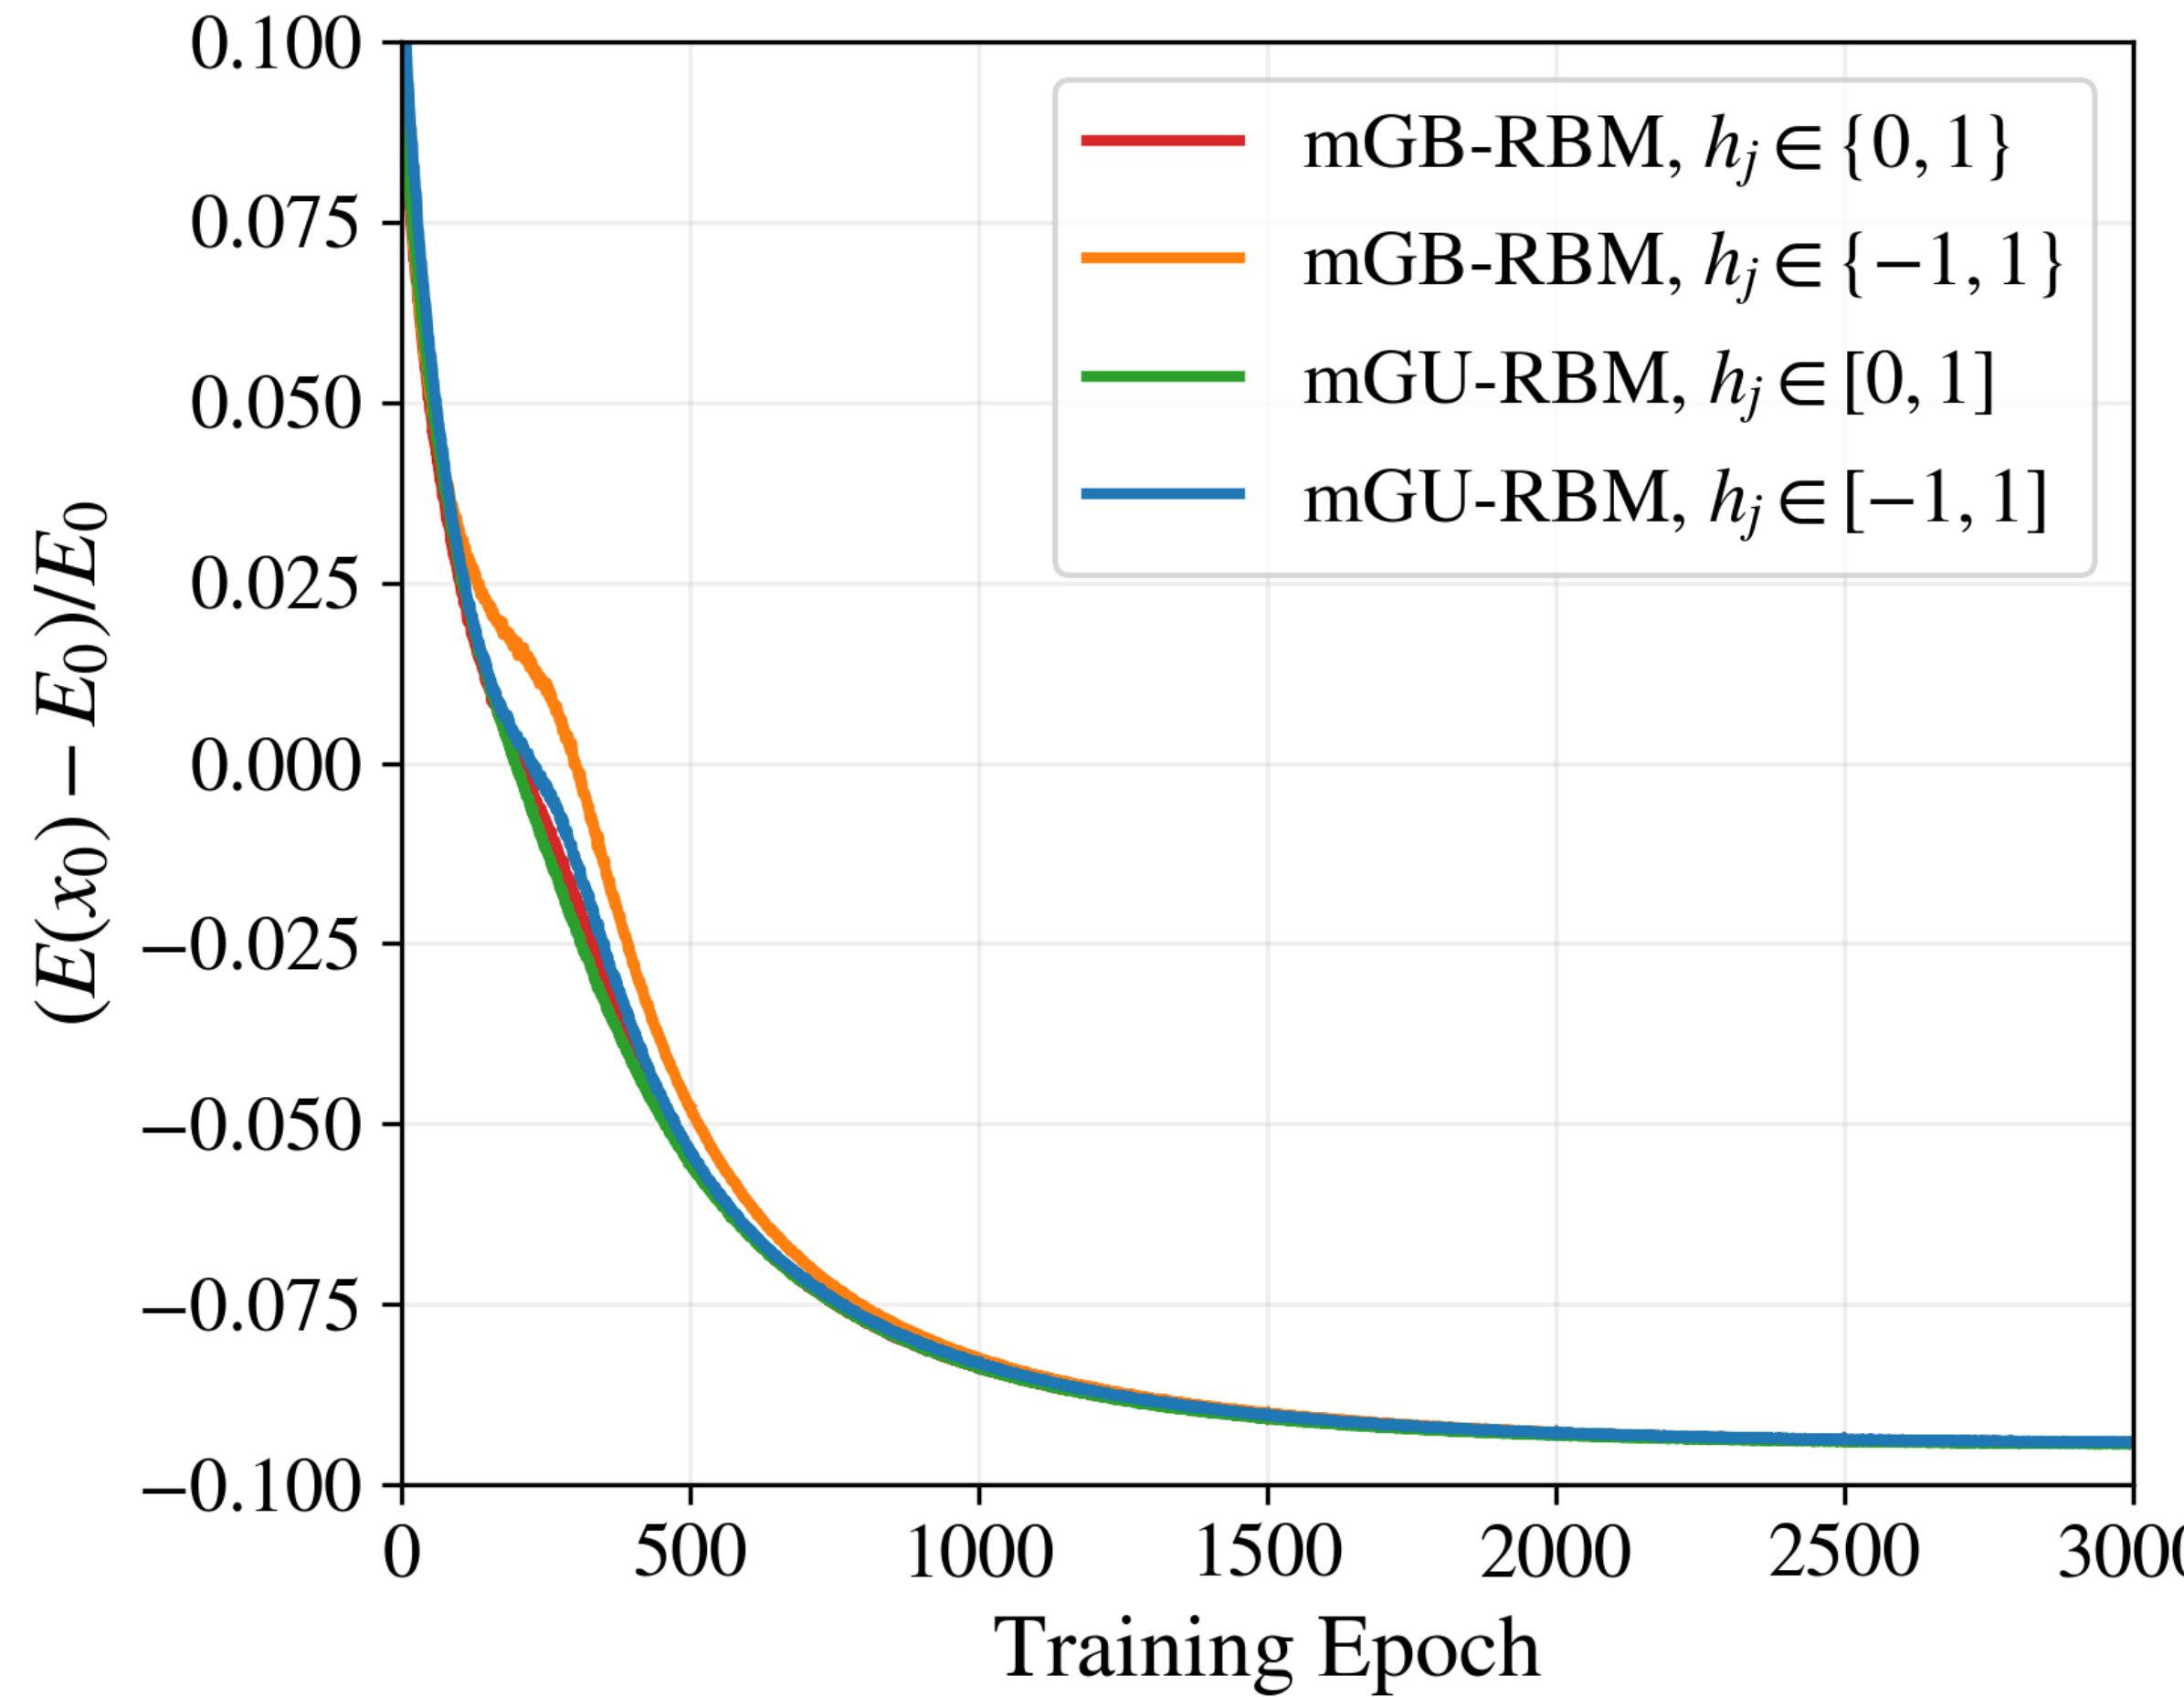
$$\Psi_{\theta}^{fnn}(\mathbf{X}) = \exp(\rho(\mathbf{v}(\mathbf{X}))) \exp\left(-\frac{\mathbf{X}^2}{2\sigma^2}\right)$$

(Alternatively can pretrain the FNN to a Gaussian using supervised learning)

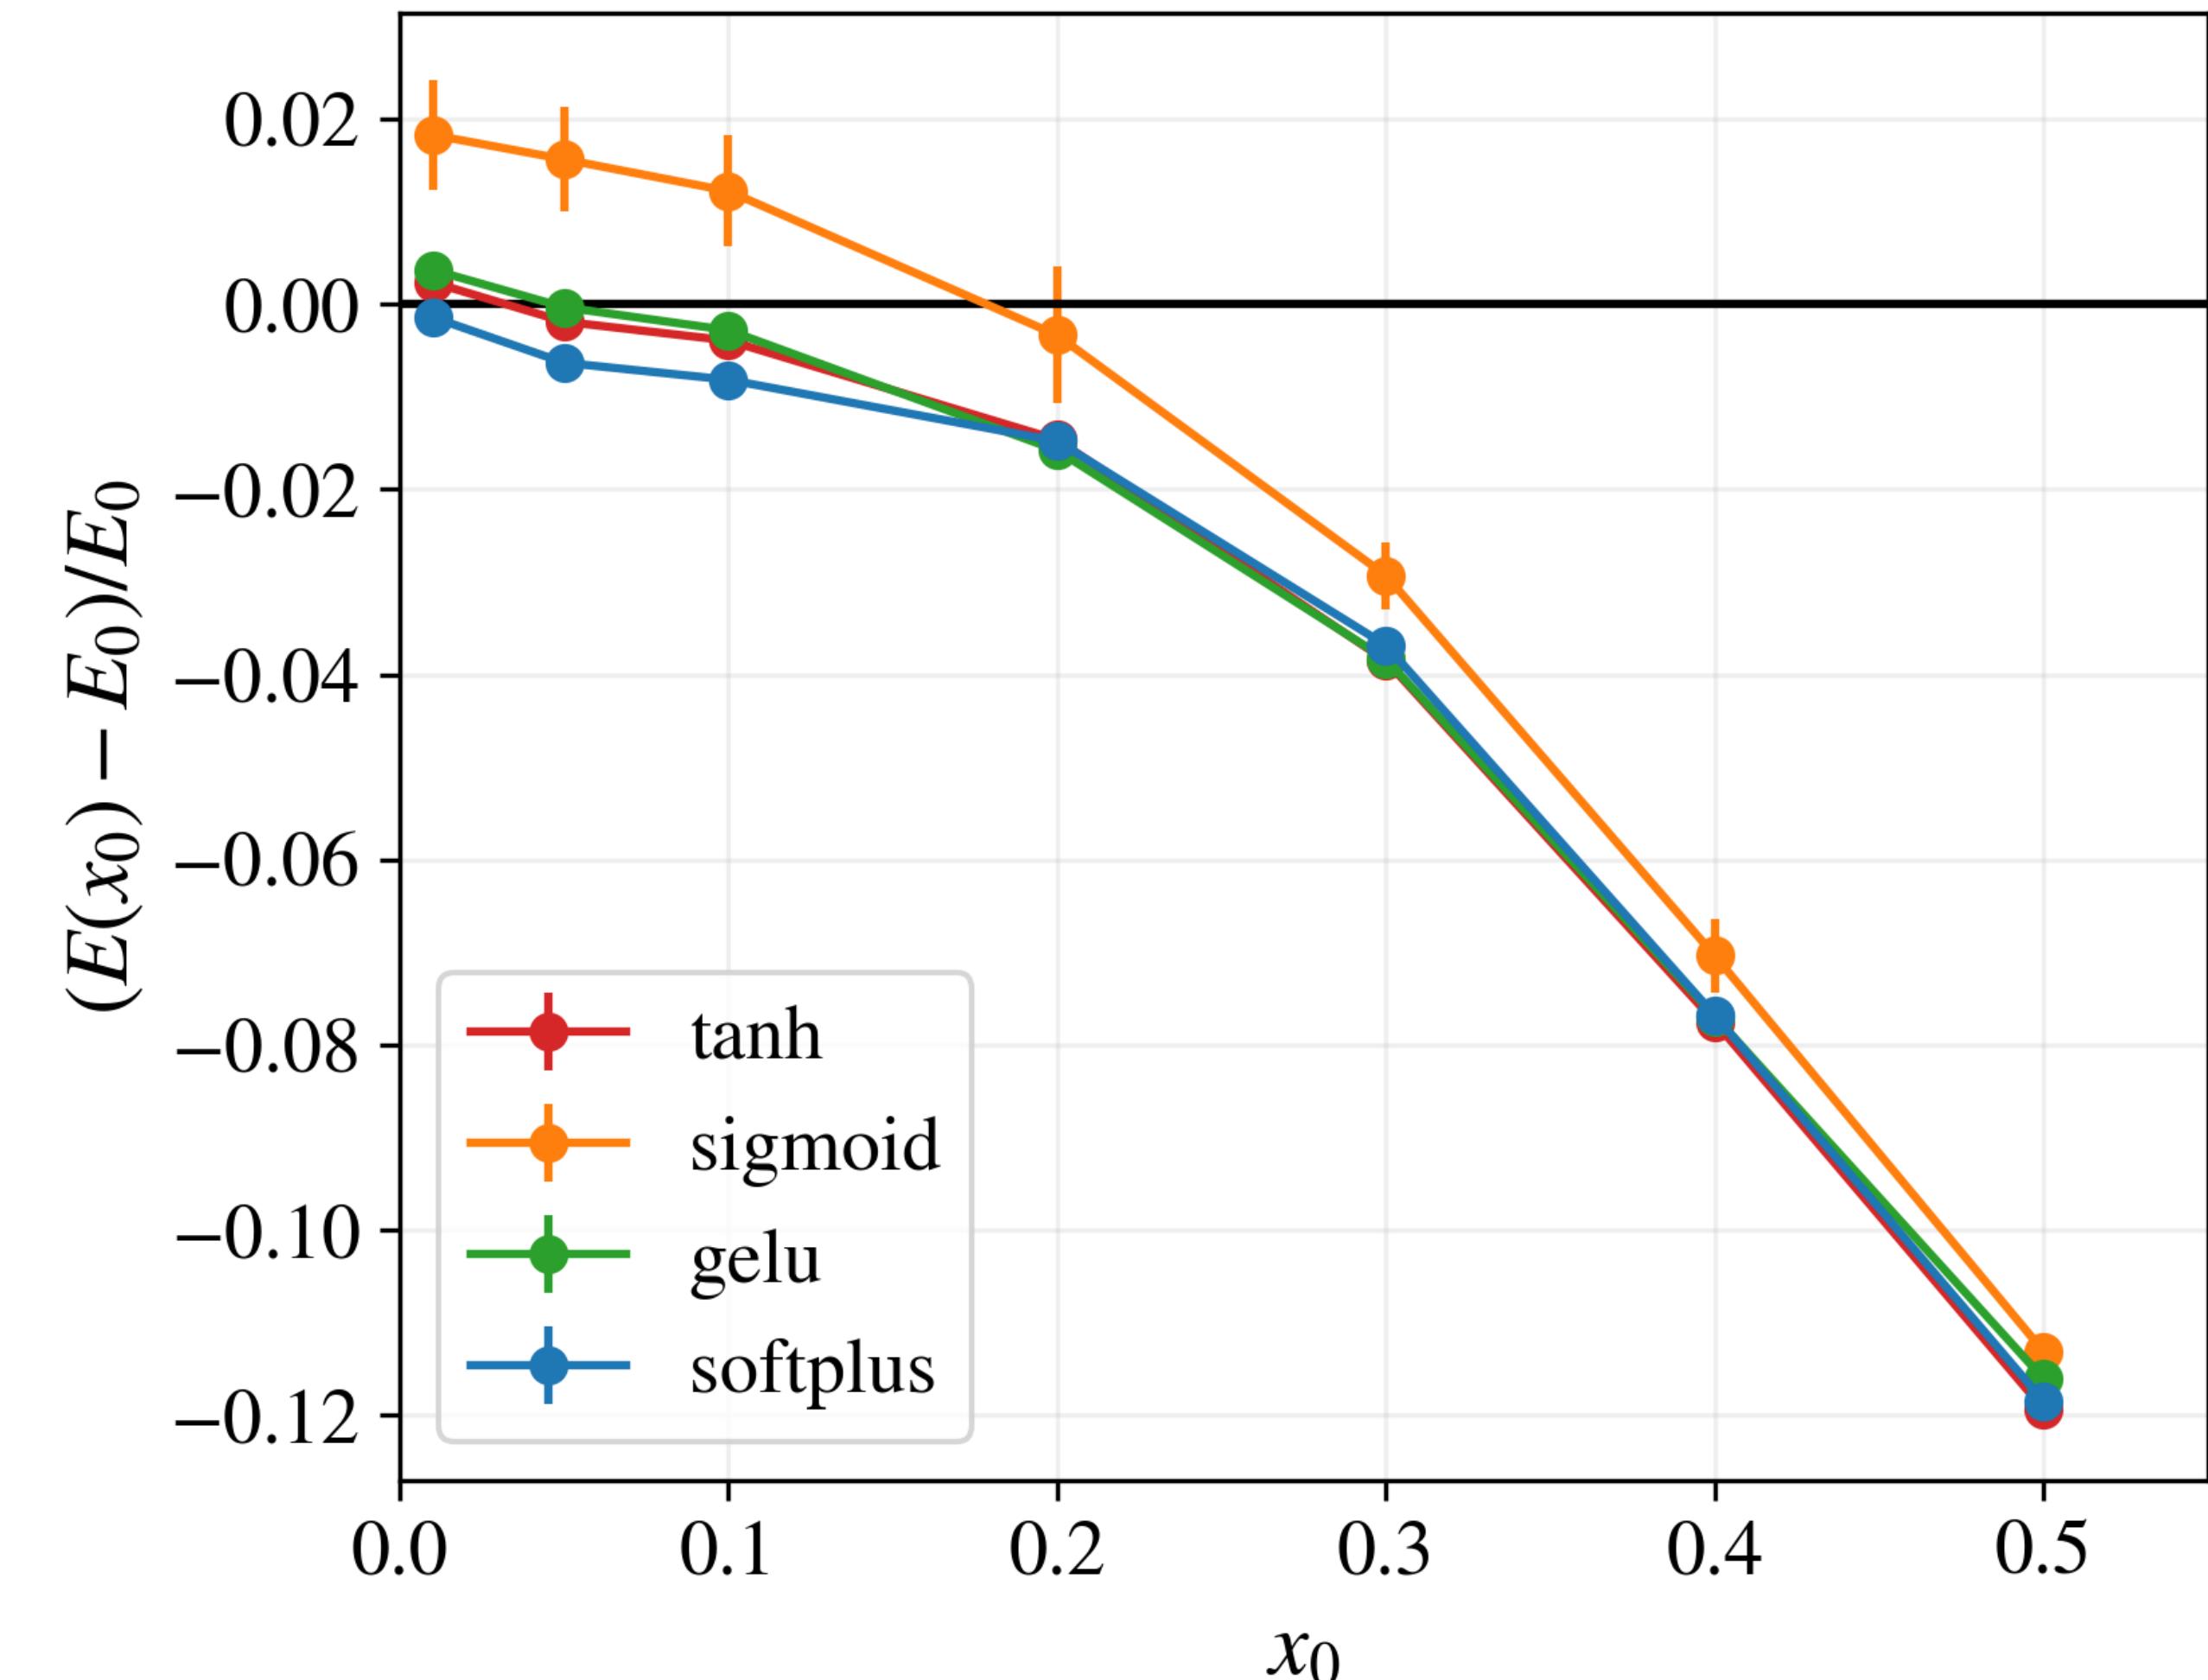
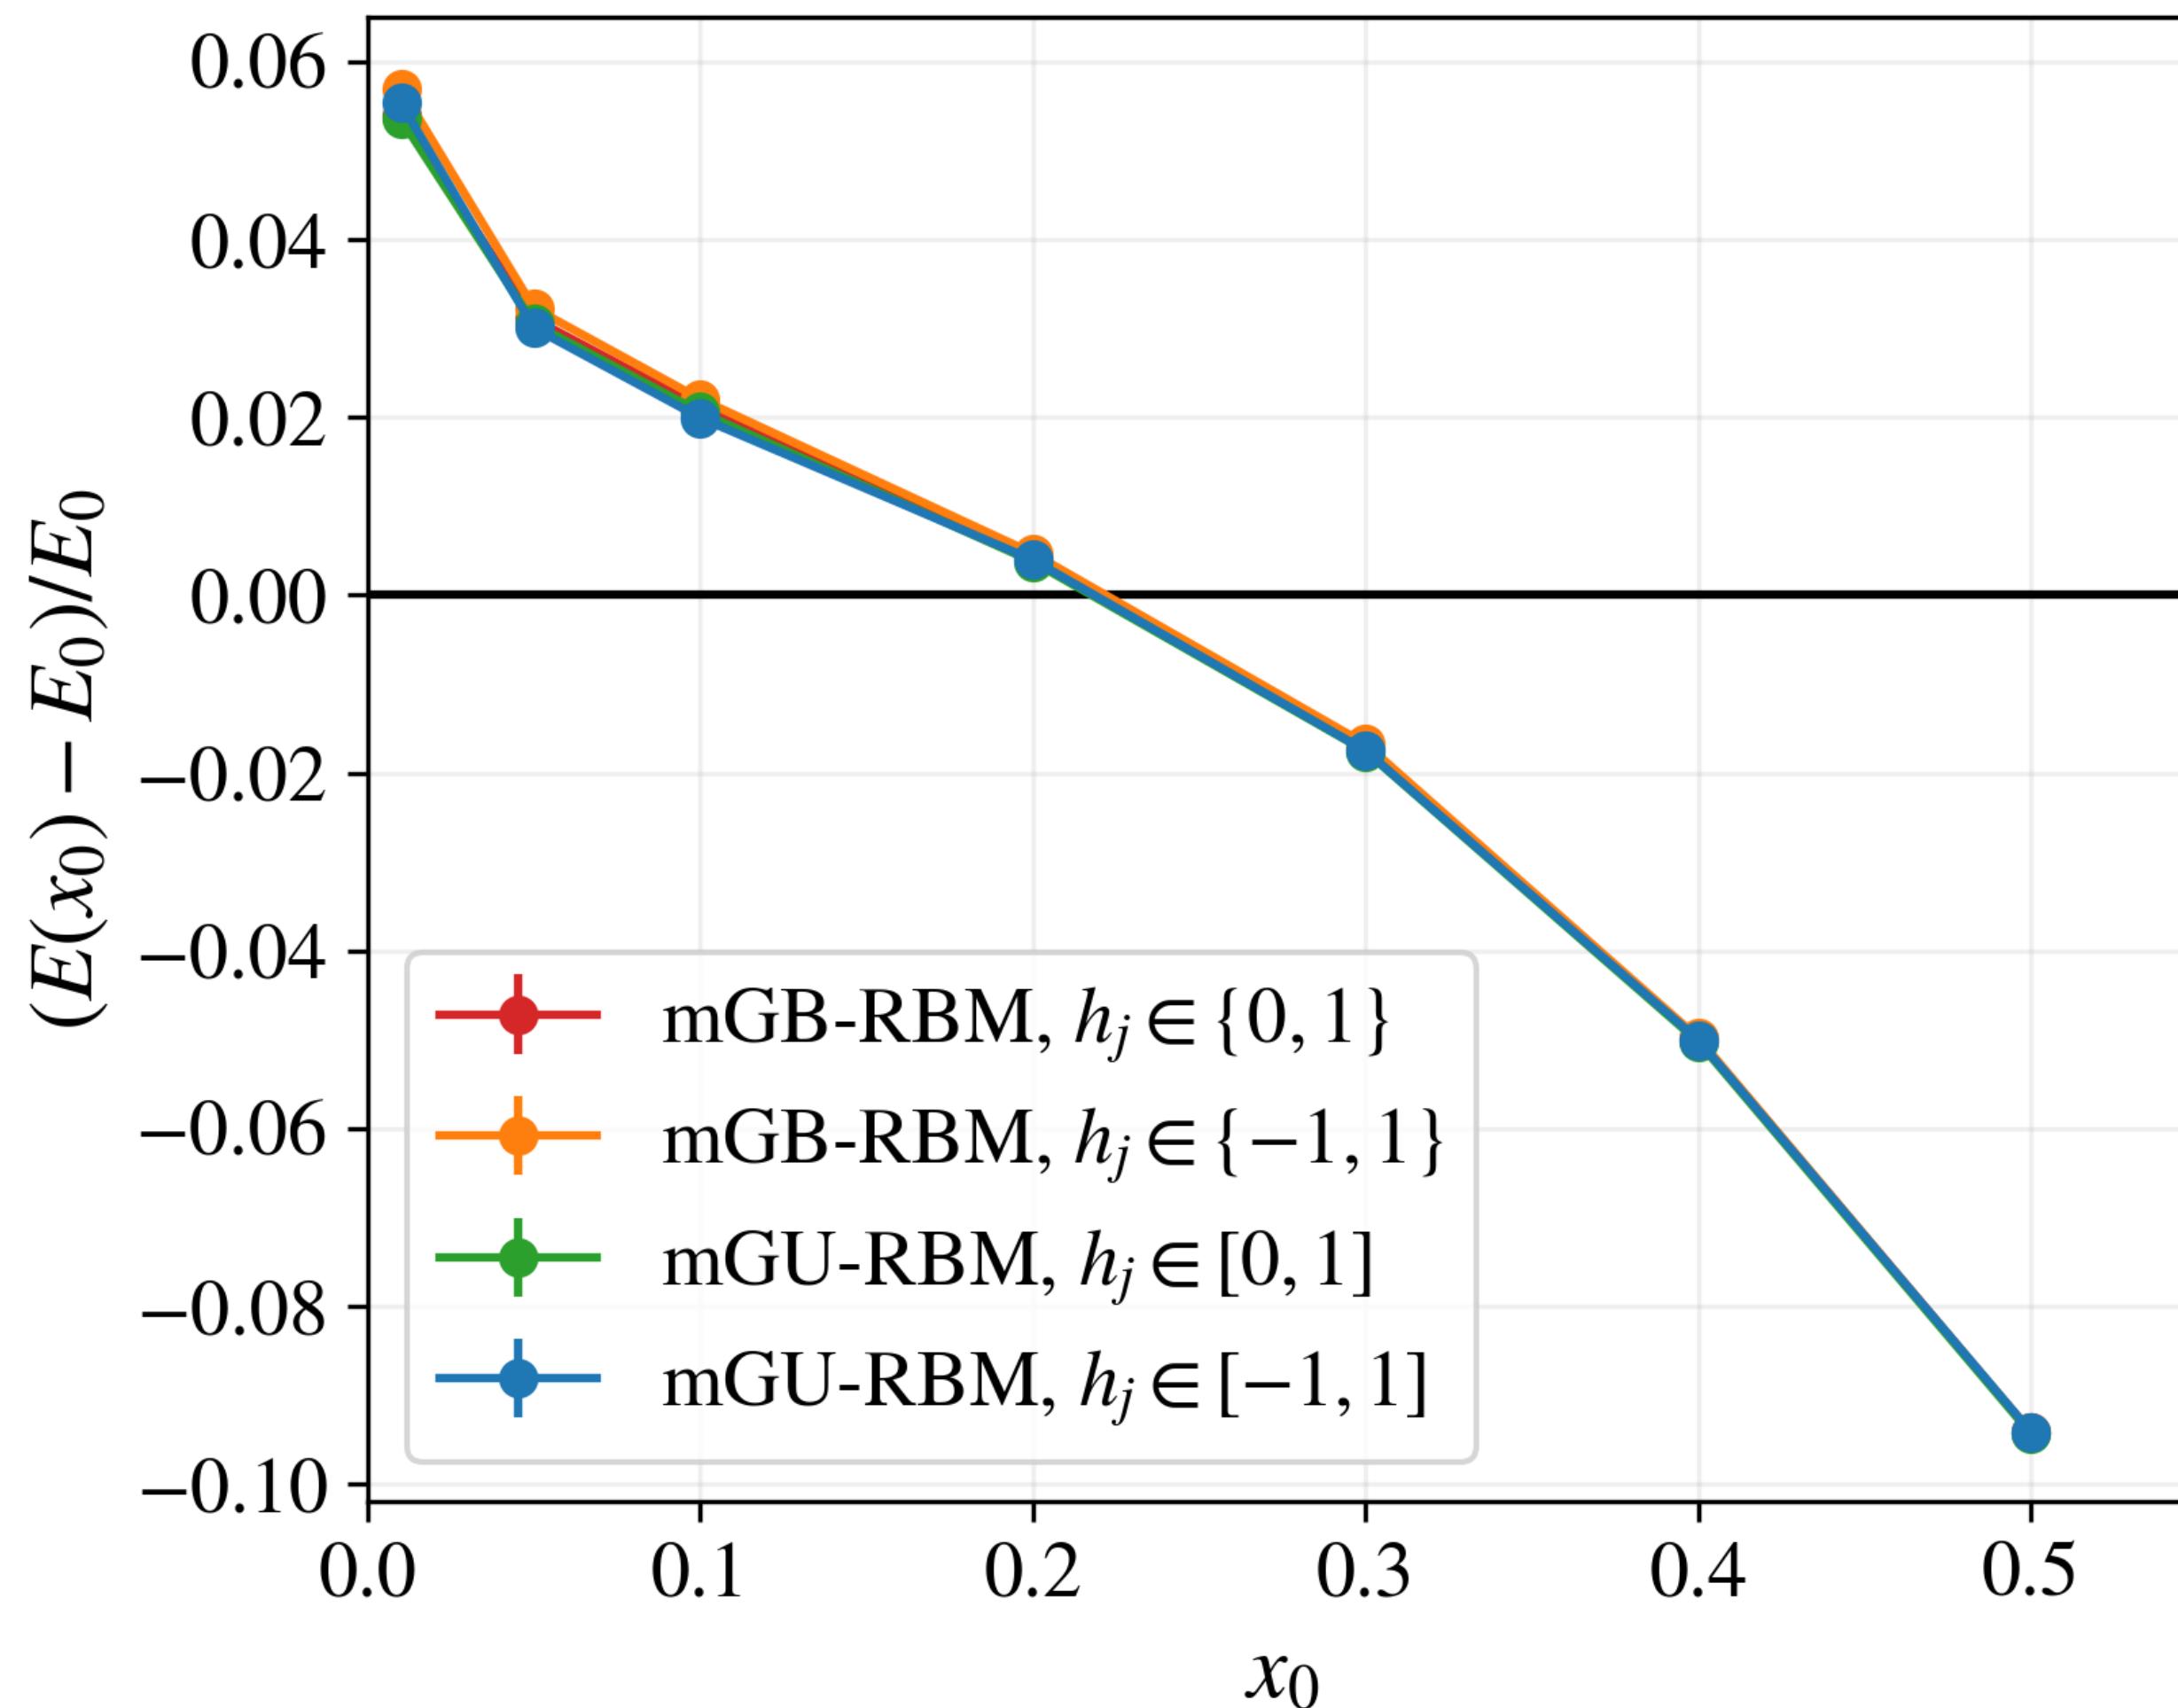
- $\mathbf{v}(\mathbf{X})$ denotes the visible nodes after sorting the input data

Calogero-Sutherland Model

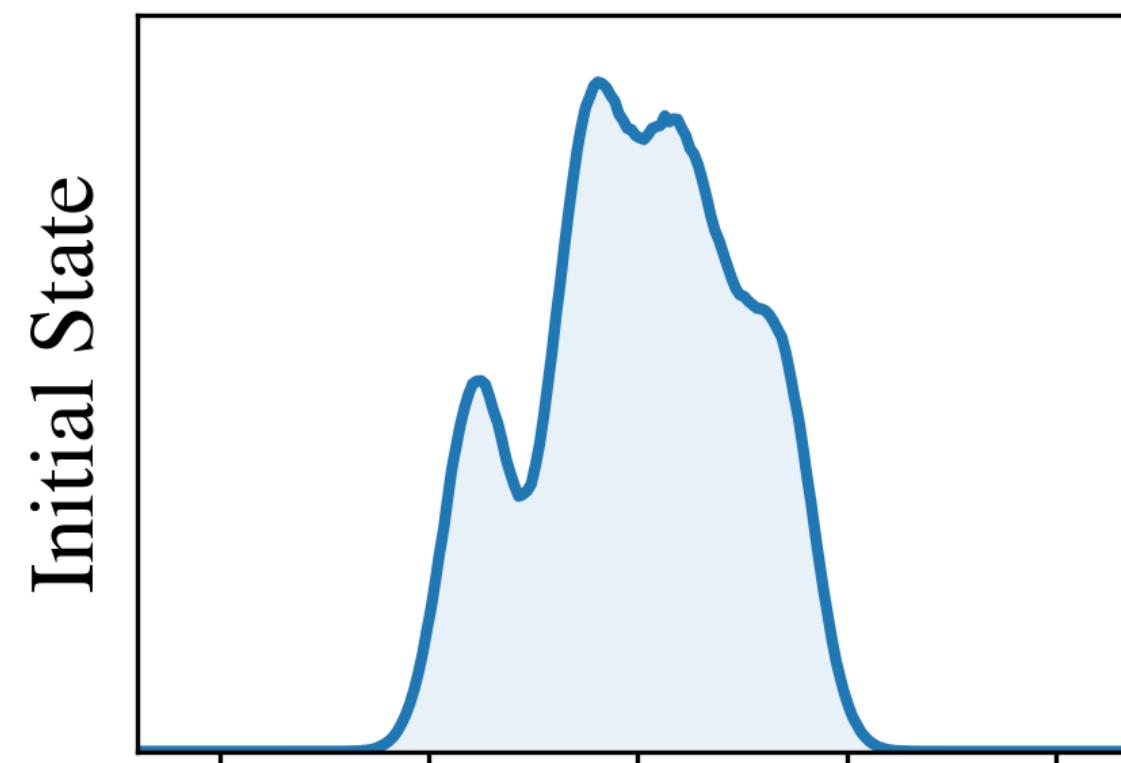
Initial pretraining stage with $x_0 = 0.5$



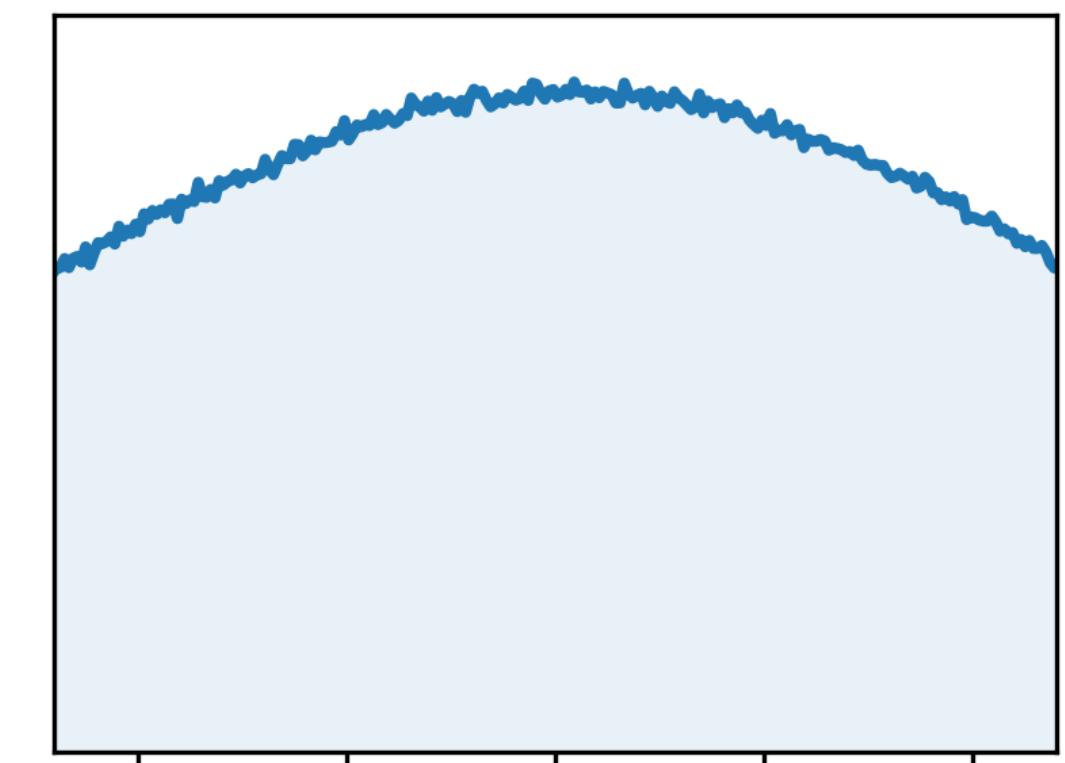
Calogero-Sutherland Model



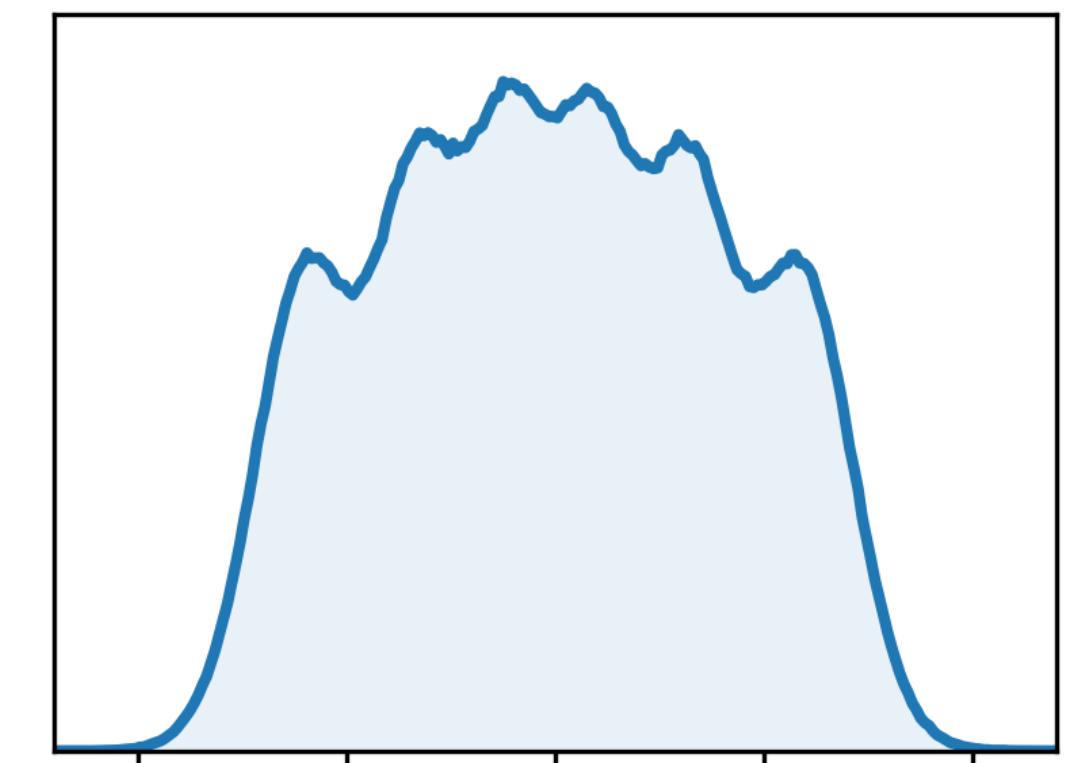
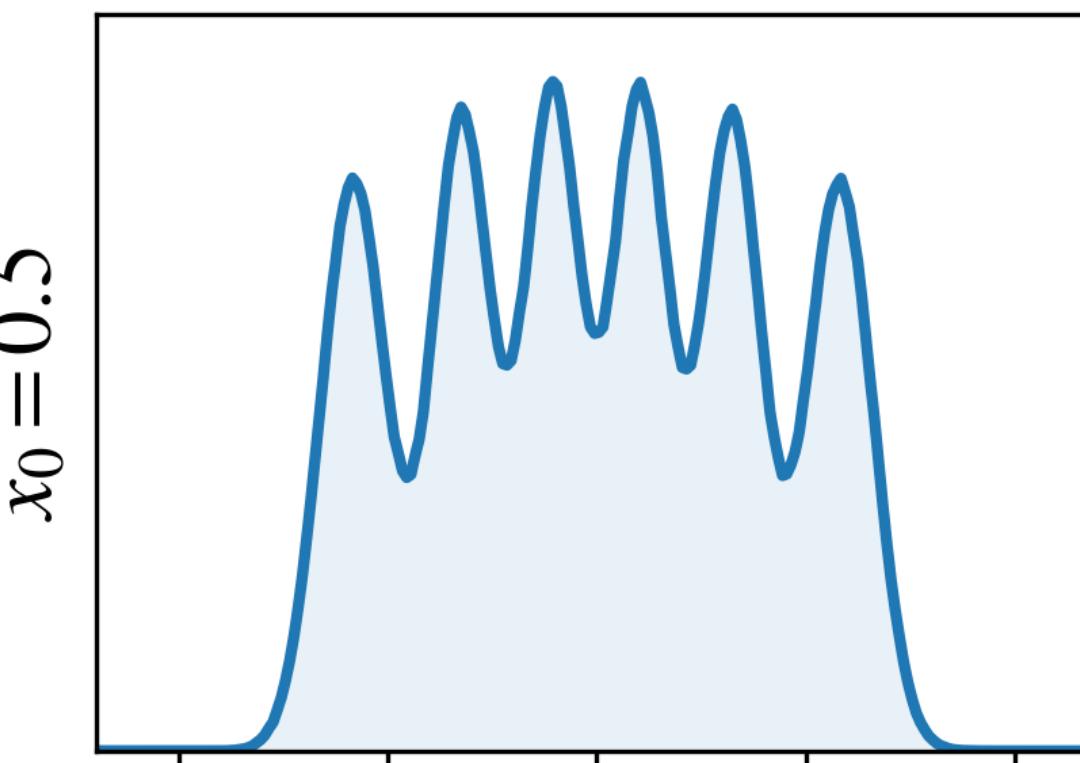
mGU-RBM



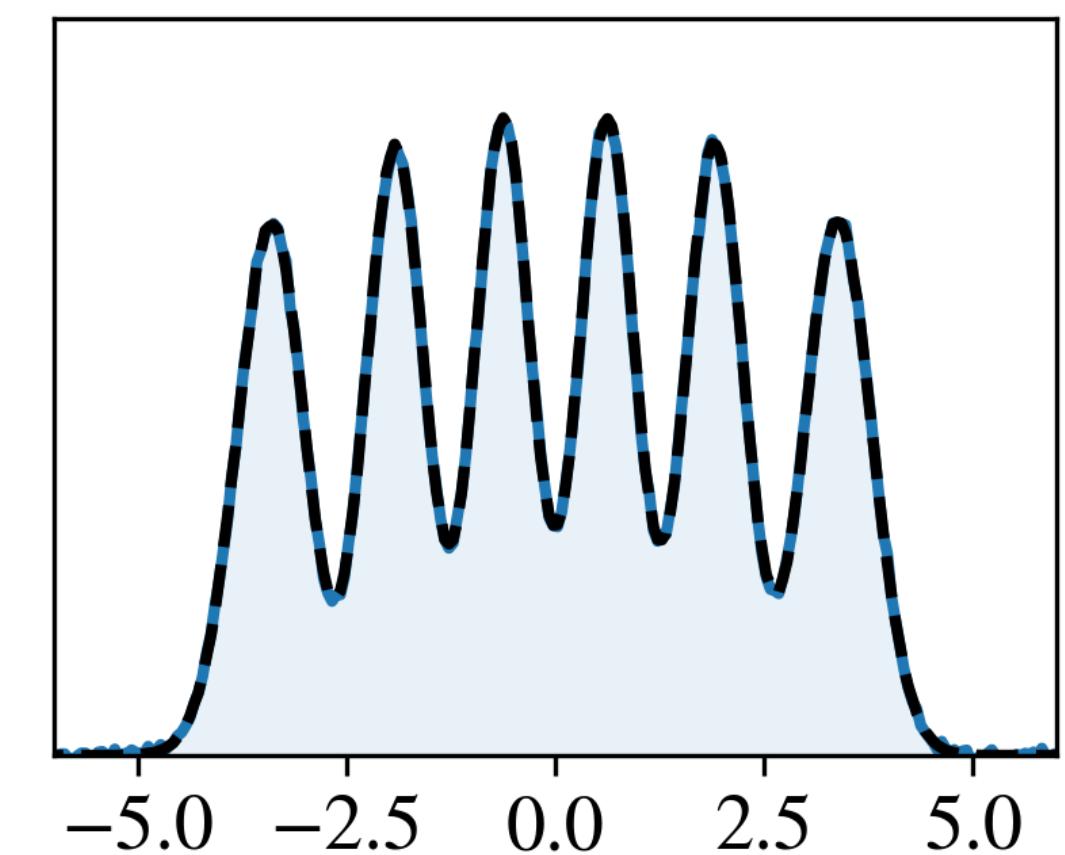
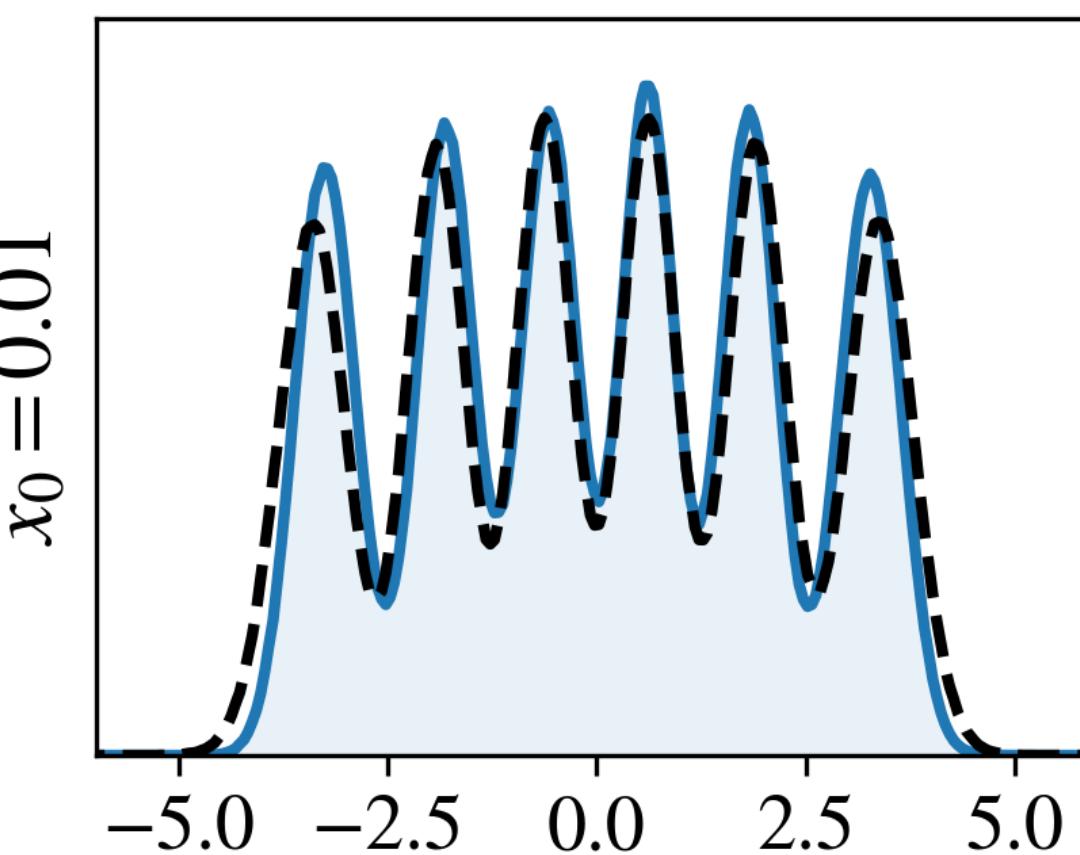
FNN



$x_0 = 0.5$



$x_0 = 0.01$



Calogero-Sutherland Model

- The different RBMs perform similarly, but the Gaussian-uniform variants perform slightly better at smaller values of x_0
- The assumption of Gaussian visible nodes biases results of RBMs
- FNNs provide more accurate results, but generally require more parameters to distinguish the individual peaks of the wave function (250 params vs. 70 params)
- Regularizing the singularity and using transfer learning is a controlled way of handling the singularity without enforcing the cusp condition

Dilute Neutron Matter

- Recently accepted for publication in Physical Review Research
- See our preprint titled “**Dilute neutron star matter from neural-network quantum states**” by B. Fore, J. Kim, G. Carleo, M. Hjorth-Jensen, and A. Lovato (arXiv:2212.04436, uploaded Dec 2022)
- We focus on lower densities $\rho \lesssim 0.08 \text{ fm}^{-3}$ relevant to our understanding of the inner crust at the boundary of outer core in neutron stars
- In this density region, neutron star matter consists of a Coulomb lattice of neutron-rich nuclei, relativistic electrons, and a superfluid neutron gas

Dilute Neutron Matter

- Leading order pionless-EFT Hamiltonian ($\hbar = 1$):

$$\hat{H} = - \sum_{i=1}^N \frac{\nabla_i^2}{2m} + \sum_{i < j} v_{ij} + \sum_{i < j < k} V_{ijk}$$

Non-relativistic kinetic energy



NN contact potential



3N contact potential

$$V_{ijk} = \frac{c_E}{f_\pi^4 \lambda_\chi} \frac{(\hbar c)^6}{\pi^3 R_3^6} \sum_{cyc} e^{-(r_{ij}^2 + r_{jk}^2)/R_3^2}$$

$$v = v^{EM} + v_{LO}^{CI}$$

$$v_{LO}^{CI} = C_{01}C_1(r_{ij})P_0^\sigma P_1^\tau + C_{10}C_0(r_{ij})P_1^\sigma P_0^\tau$$

Dilute Neutron Matter

- Simulate 14 neutrons in a periodic box of side length L
- Enforce periodicity by transforming the coordinates as

$$\mathbf{r}_i \mapsto \left(\cos\left(\frac{2\pi\mathbf{r}_i}{L}\right), \sin\left(\frac{2\pi\mathbf{r}_i}{L}\right) \right)$$

- Mitigate finite-size effects:
 - ▶ Sum over contributions from neighboring cells
 - ▶ Subtract 14 particle kinetic energy, add kinetic energy in thermodynamic limit

Dilute Neutron Matter

- We use an ansatz belonging to the “hidden fermion” family of states introduced by J. R. Moreno et al. (arXiv:2111.10420)
- Introduce M fictitious “hidden” nucleons that are constructed using Deep Sets over the “visible” nucleons

$$\Psi_{\theta}^{HN}(\mathbf{X}) = \det \begin{bmatrix} \phi_1(\mathbf{x}_1) & \cdots & \phi_1(\mathbf{x}_N) & \phi_1(\mathbf{y}_1) & \cdots & \phi_1(\mathbf{y}_M) \\ \phi_2(\mathbf{x}_1) & \cdots & \phi_2(\mathbf{x}_N) & \phi_2(\mathbf{y}_1) & \cdots & \phi_2(\mathbf{y}_M) \\ \vdots & \ddots & \vdots & \vdots & \ddots & \vdots \\ \phi_N(\mathbf{x}_1) & \cdots & \phi_N(\mathbf{x}_N) & \phi_N(\mathbf{y}_1) & \cdots & \phi_N(\mathbf{y}_M) \\ \chi_1(\mathbf{x}_1) & \cdots & \chi_1(\mathbf{x}_N) & \chi_1(\mathbf{y}_1) & \cdots & \chi_1(\mathbf{y}_M) \\ \chi_2(\mathbf{x}_1) & \cdots & \chi_2(\mathbf{x}_N) & \chi_2(\mathbf{y}_1) & \cdots & \chi_2(\mathbf{y}_M) \\ \vdots & \ddots & \vdots & \vdots & \ddots & \vdots \\ \chi_N(\mathbf{x}_1) & \cdots & \chi_N(\mathbf{x}_N) & \chi_N(\mathbf{y}_1) & \cdots & \chi_N(\mathbf{y}_M) \end{bmatrix}$$

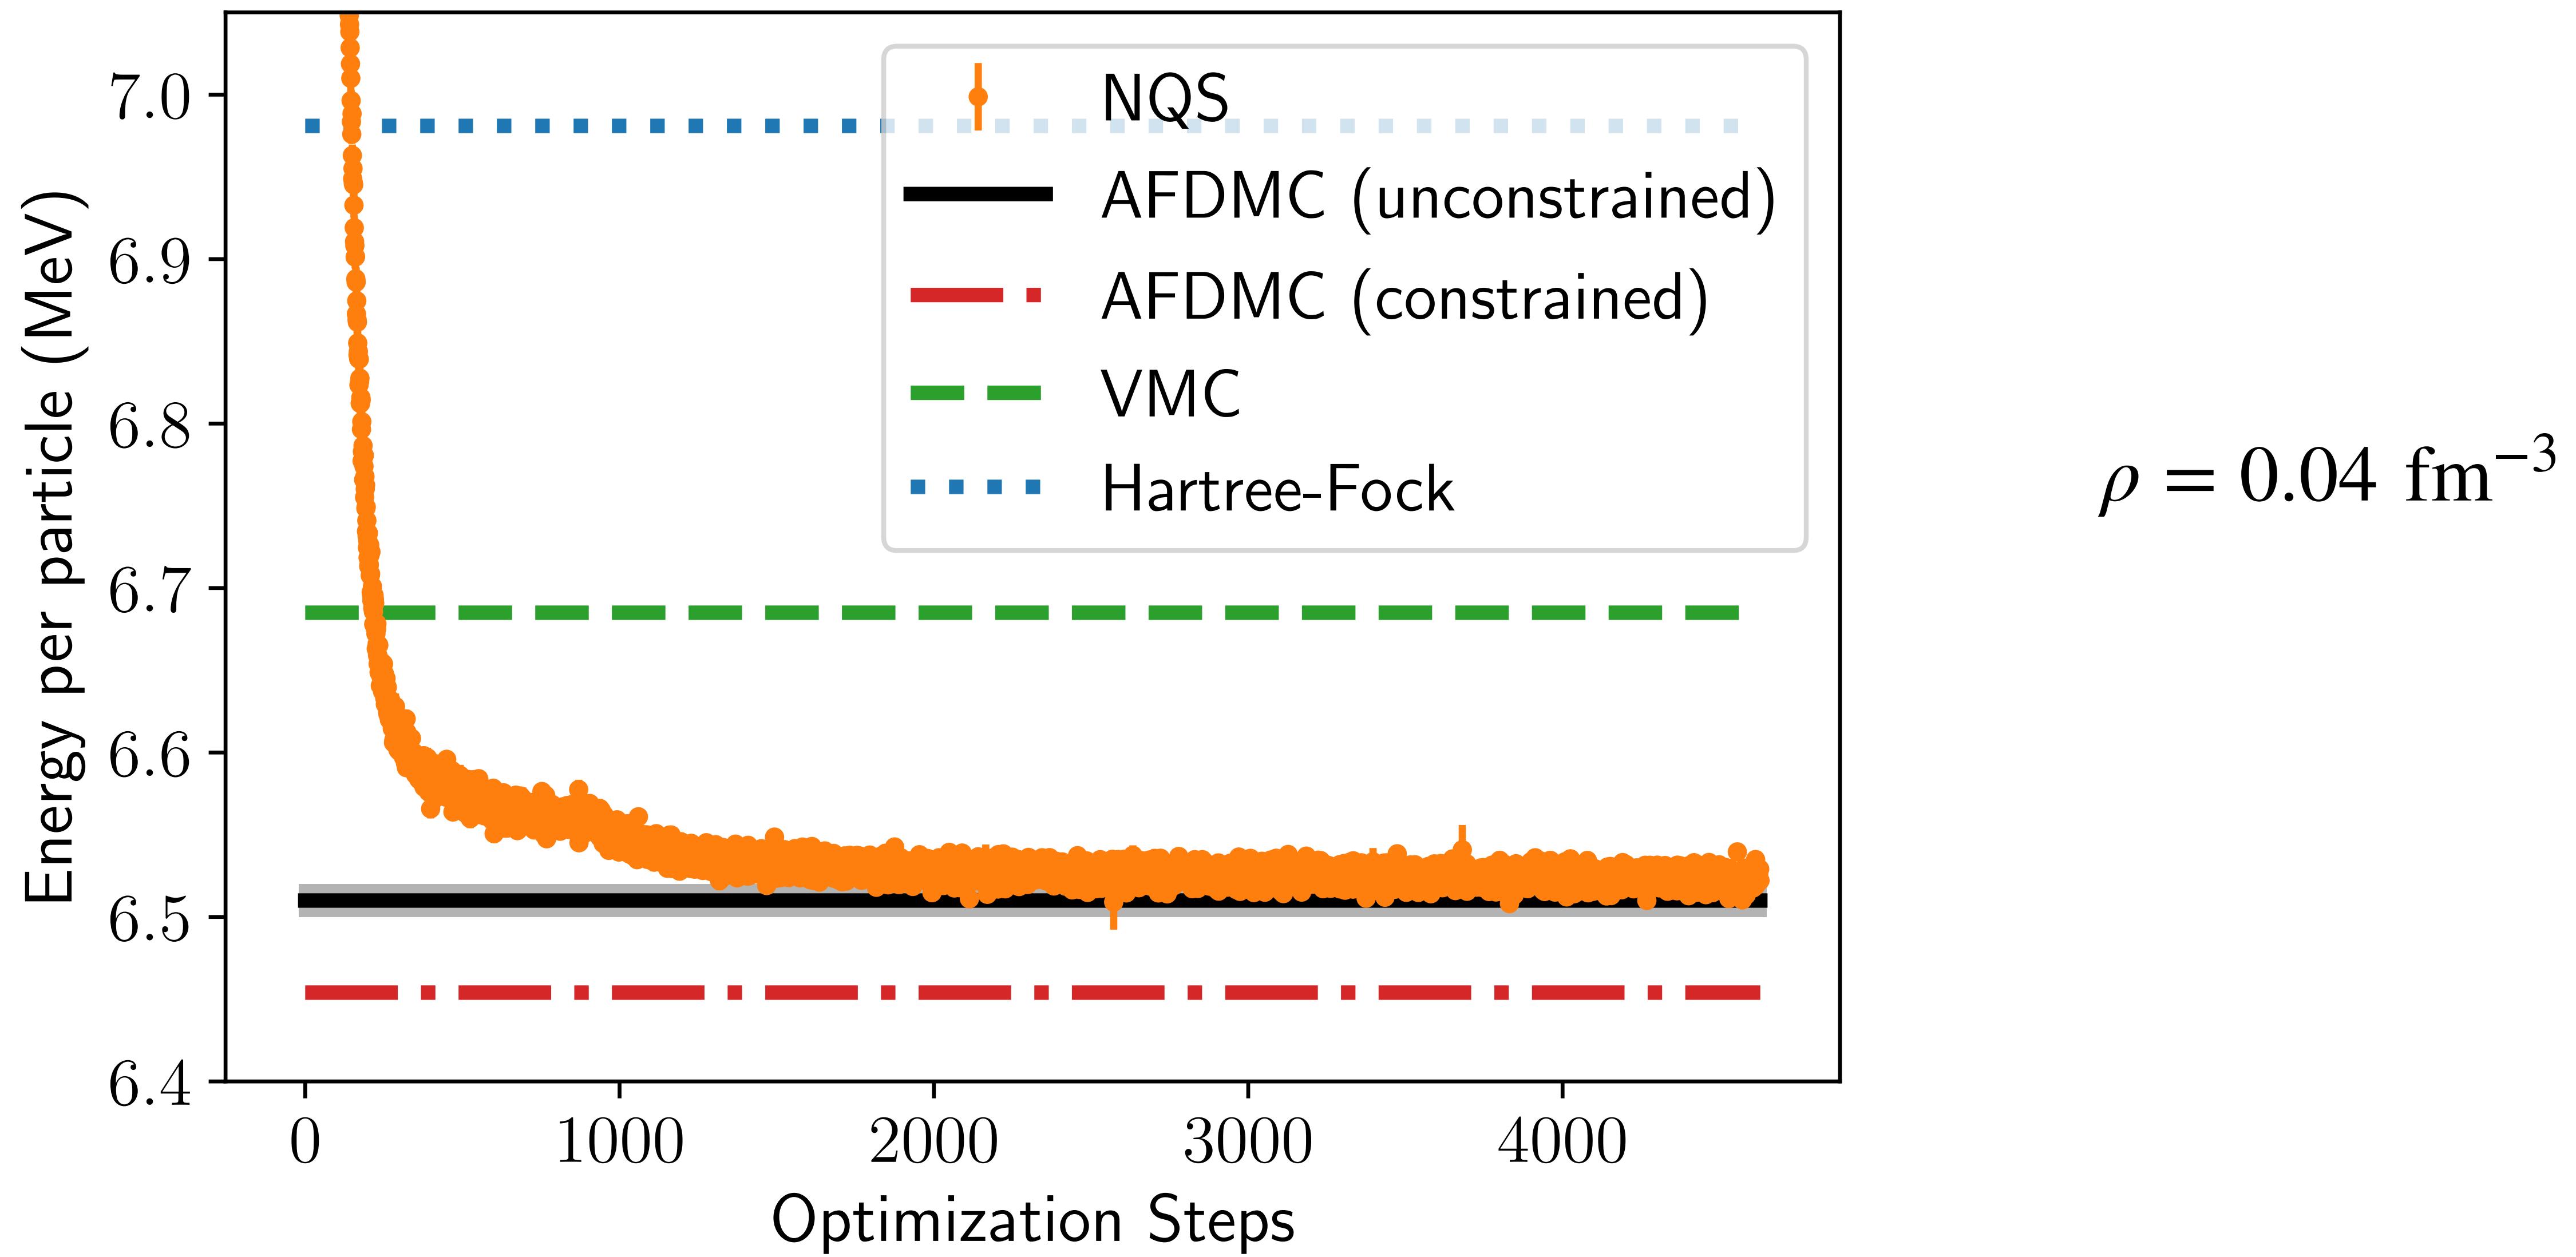
Visible s.p. states
Visible nucleons

Visible s.p. states
Hidden nucleons

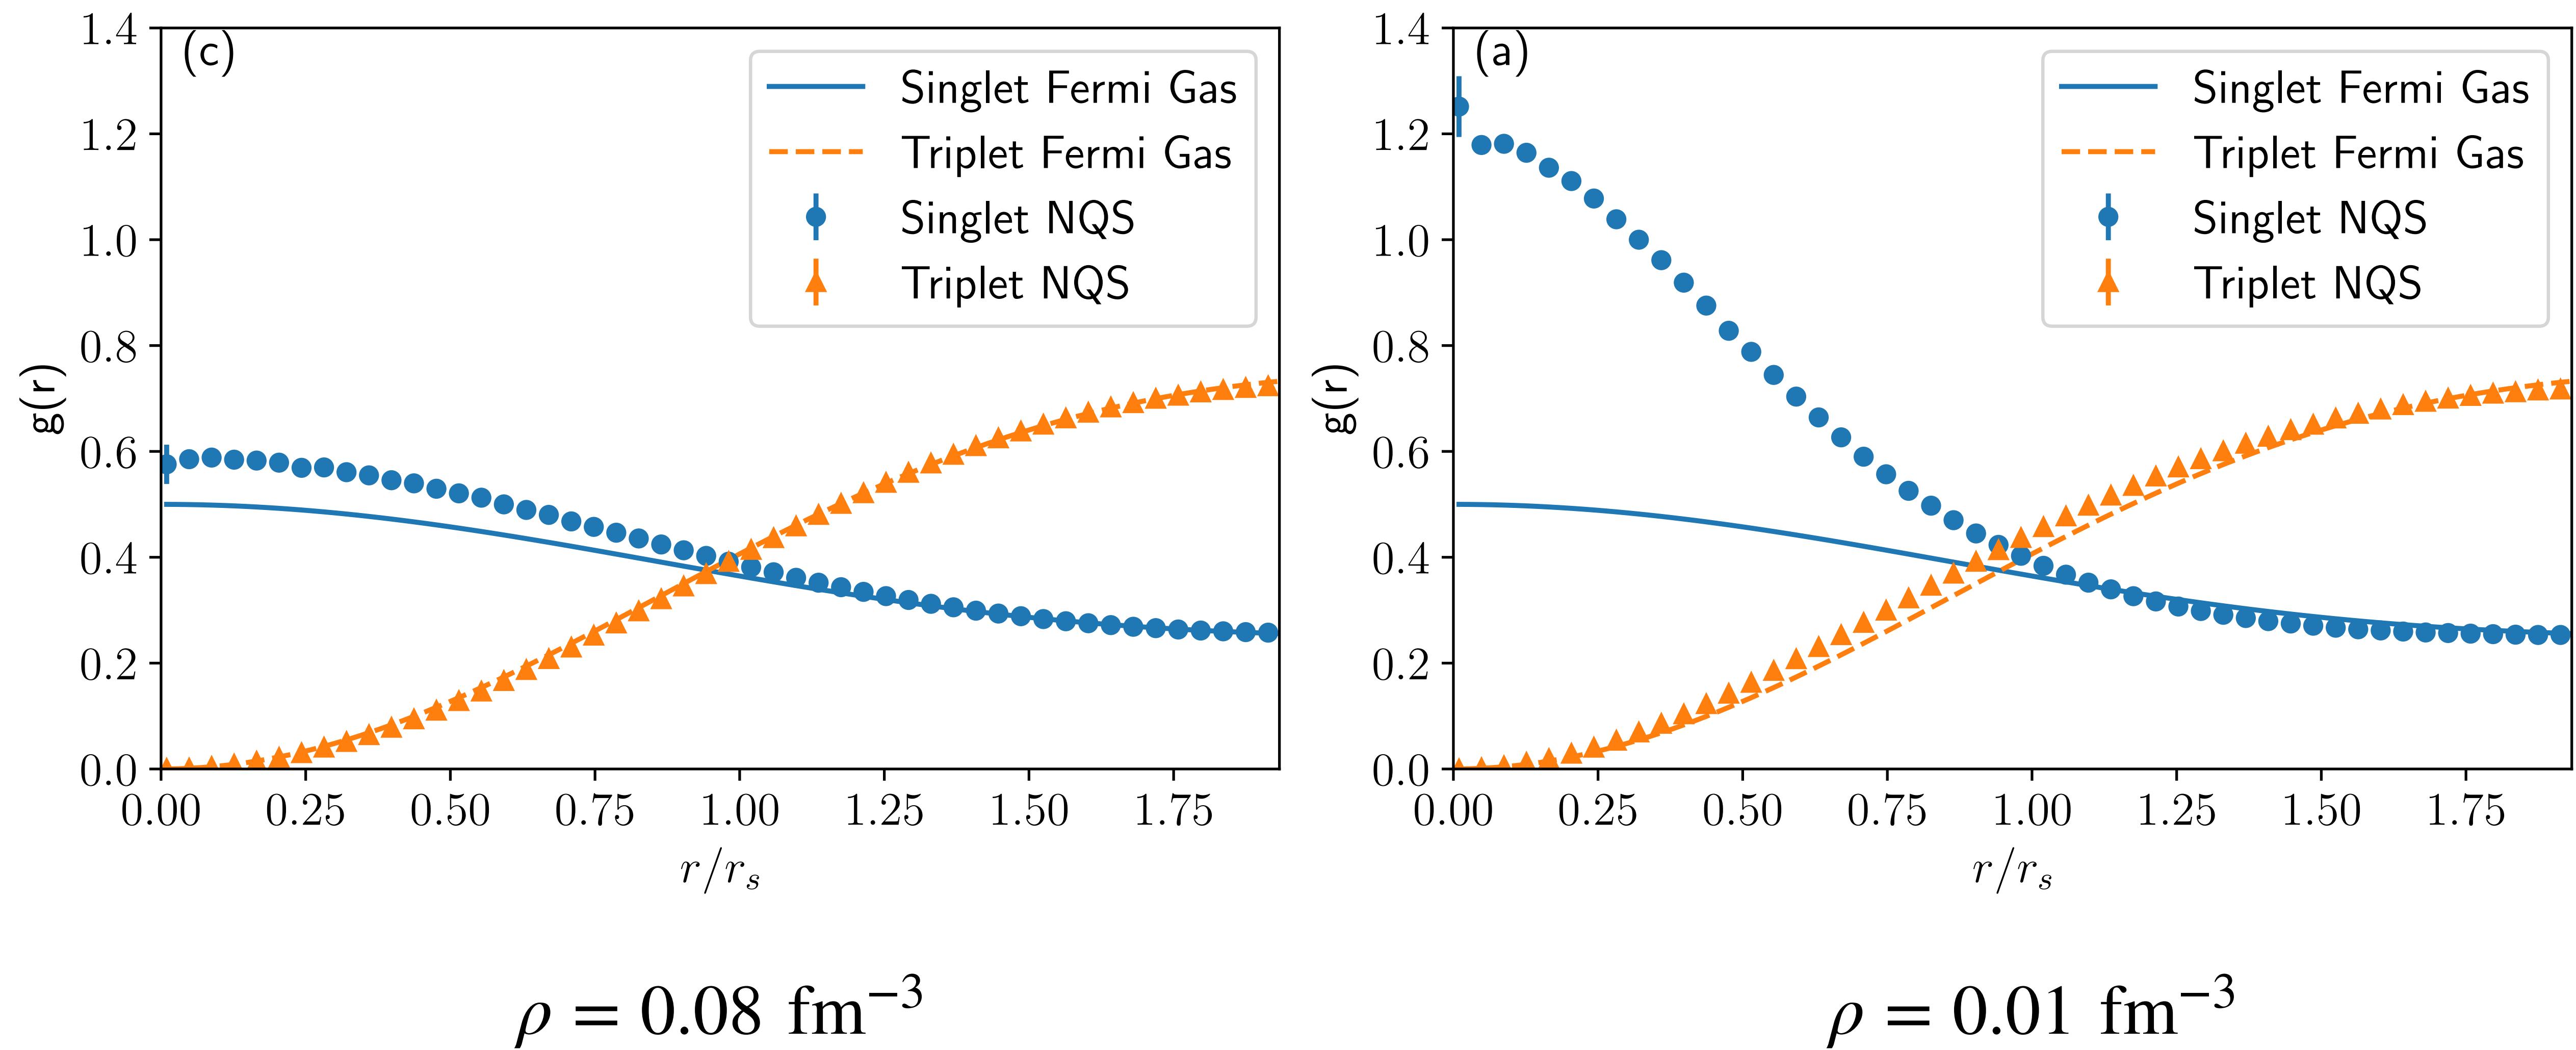
Hidden s.p. states
Visible nucleons

Hidden s.p. states
Hidden nucleons

Dilute Neutron Matter



Dilute Neutron Matter



Dilute Neutron Matter

- The hidden nucleon NQS agrees with unconstrained-path auxiliary field DMC
- Singlet pair distribution functions show evidence of pairing at lower densities
- A single ansatz can model both the normal and superfluid phases
- The formation of Cooper pairs is discovered by the NQS
- Future developments:
 - ▶ More realistic interactions (AV18+UIX and local Δ -full chiral-EFT)
 - ▶ Look for clustering in isospin-antisymmetric matter at low densities

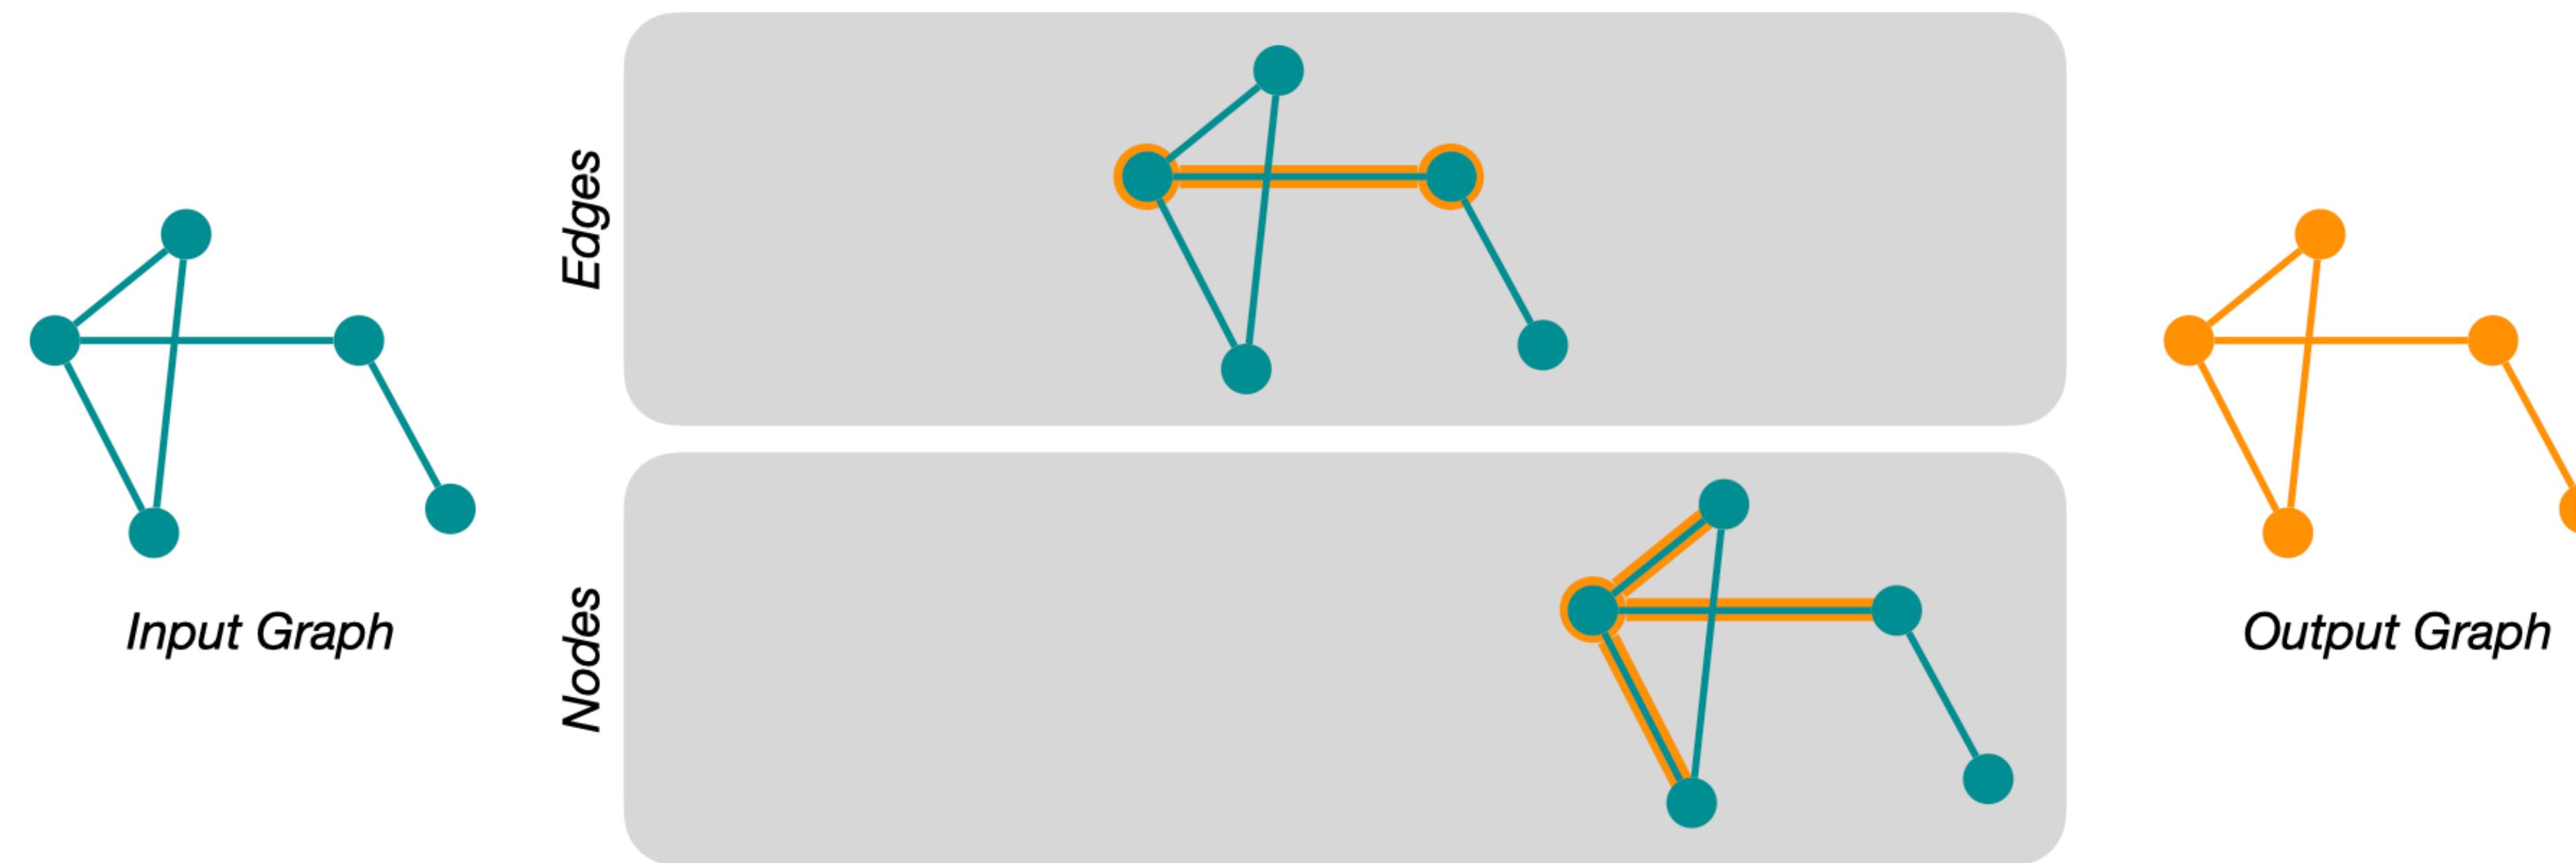
Homogeneous Electron Gas

- See our preprint “**Message-passing neural quantum states for the homogeneous electron gas**” by G. Pescia, J. Nys, J. Kim, A. Lovato, G. Carleo (arXiv:2305.07240)
- Model for the electronic structure of solids
- Electrons move freely within a static, positive background, interacting with neighboring electrons
- Long-range nature of Coulomb interaction requires Ewald summation
- Hamiltonian:

$$\hat{H} = -\frac{1}{2r_s^2} \sum_{i=1}^N \nabla_i^2 + \frac{1}{r_s} \sum_{i < j} \frac{1}{r_{ij}} + \text{const.}$$

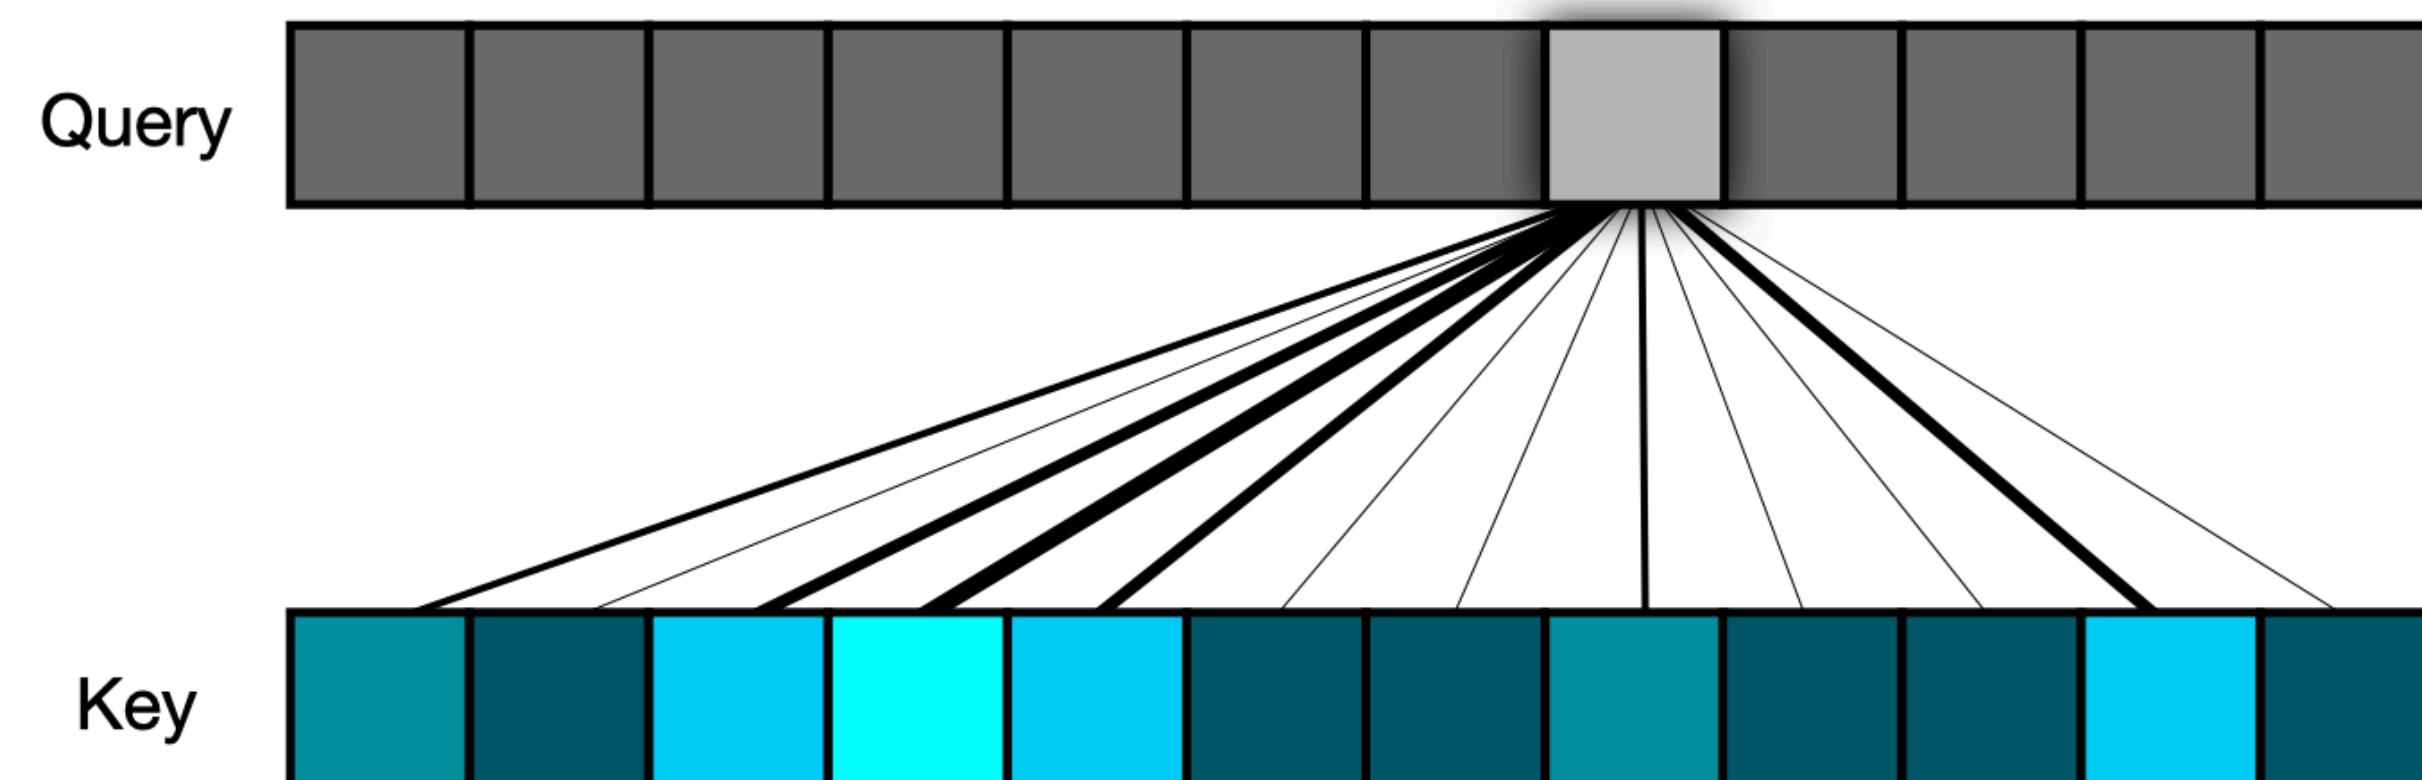
Homogeneous Electron Gas

- Employ a permutation-equivariant message-passing neural network, a type of graph neural network, that iteratively encodes backflow correlations into new positions (nodes) and interactions (edges) of the system (graph)
- To enforce translational invariance,



Homogeneous Electron Gas

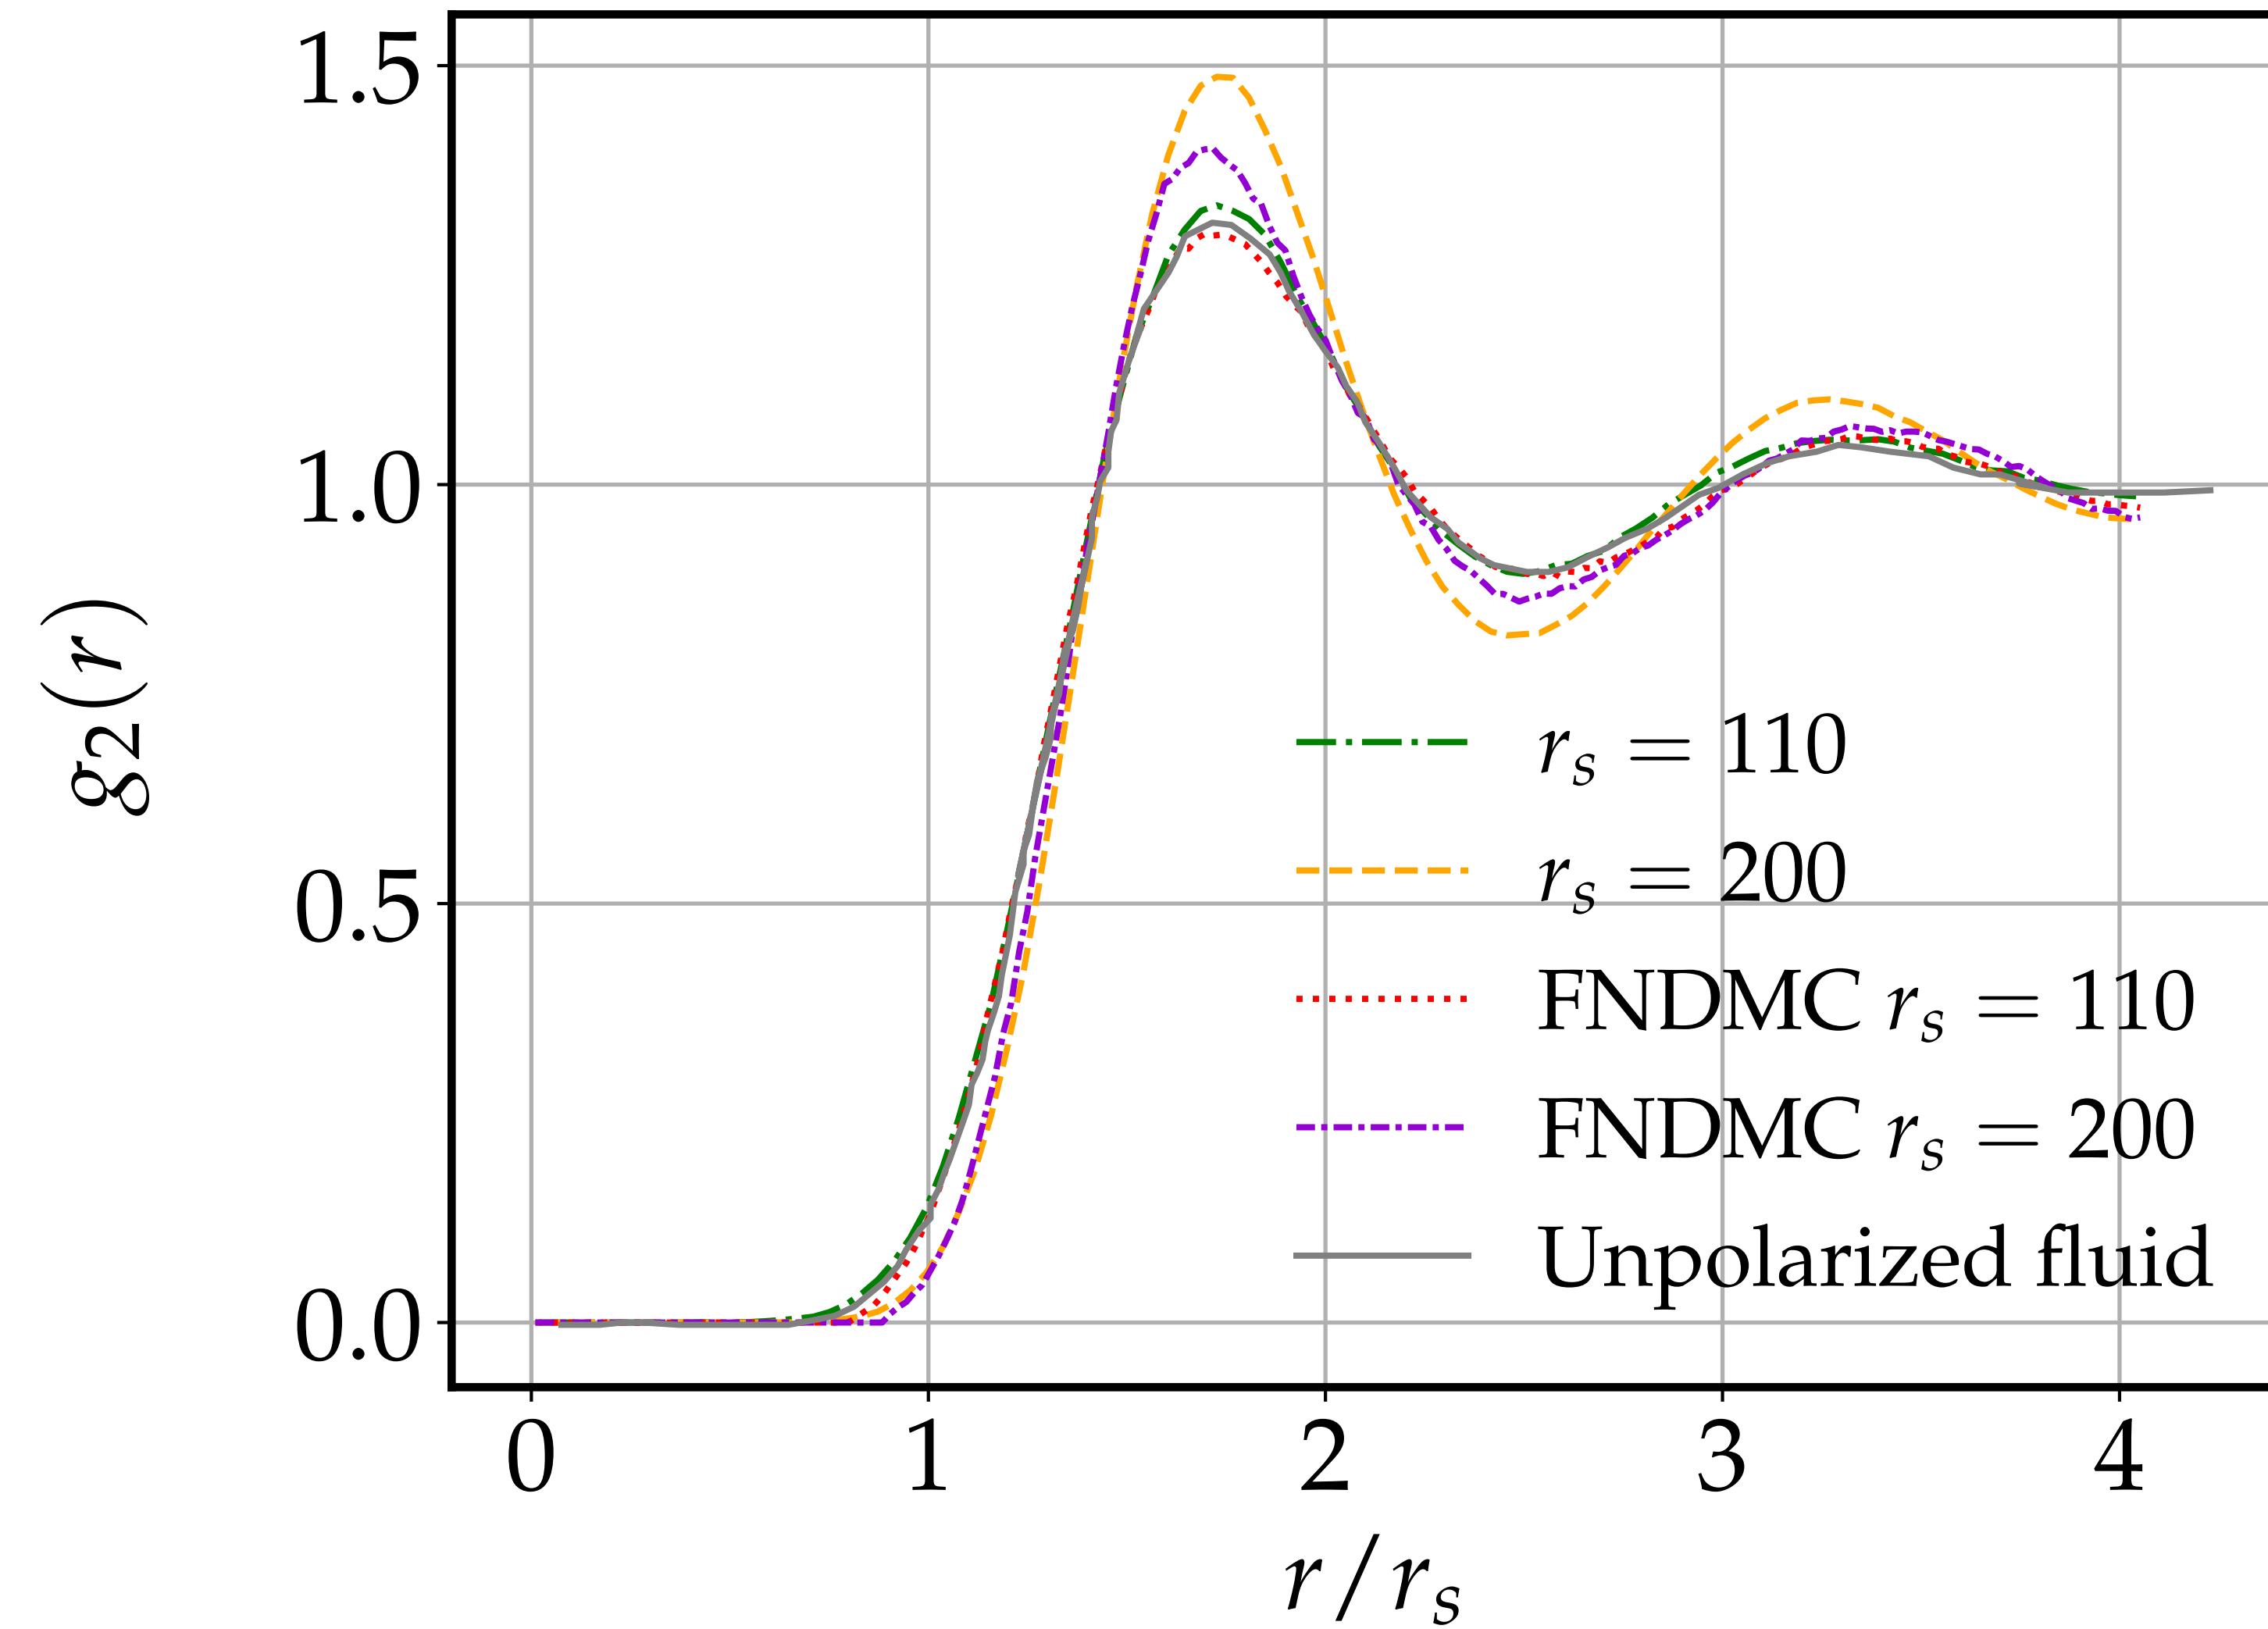
- Nodes and edges are updated using an attention mechanism
- Attention mechanisms are commonly used in natural language models to help neural networks focus on relevant data based on context clues
- Apply linear transformations to the edges, corresponding to queries and keys, then compute the overlap to determine the correlations between pairs of particles



Homogeneous Electron Gas

N	r_s	MP-NQS	LiNet [13]	FN-DMC	BF-DMC [14, 15]	FCI*/DCD** [12]
54	1	0.52973(1)	0.530019(1)	0.53094(2)	0.52989(4)	0.52973(3)*
	2	-0.014046(8)	-0.013840(1)	-0.01326(2)	-0.013966(2)	-0.01379**
	5	-0.079090(2)	-0.0788354(2)	-0.07867(1)	-0.079036(3)	-0.07837**
	10	-0.054448(1)	-0.0542785(1)	-0.054269(8)	-0.054443(2)	-0.05322**
	20	-0.0320524(5)	-0.0316886(1)	-0.031976(8)	-0.032047(2)	-0.03113**
	50	-0.0145015(1)	N/A	-0.01387(2)	-0.0144877(1)	-0.01281**
	100	-0.0076793(1)	N/A	-0.007674(3)	N/A	N/A

Homogeneous Electron Gas

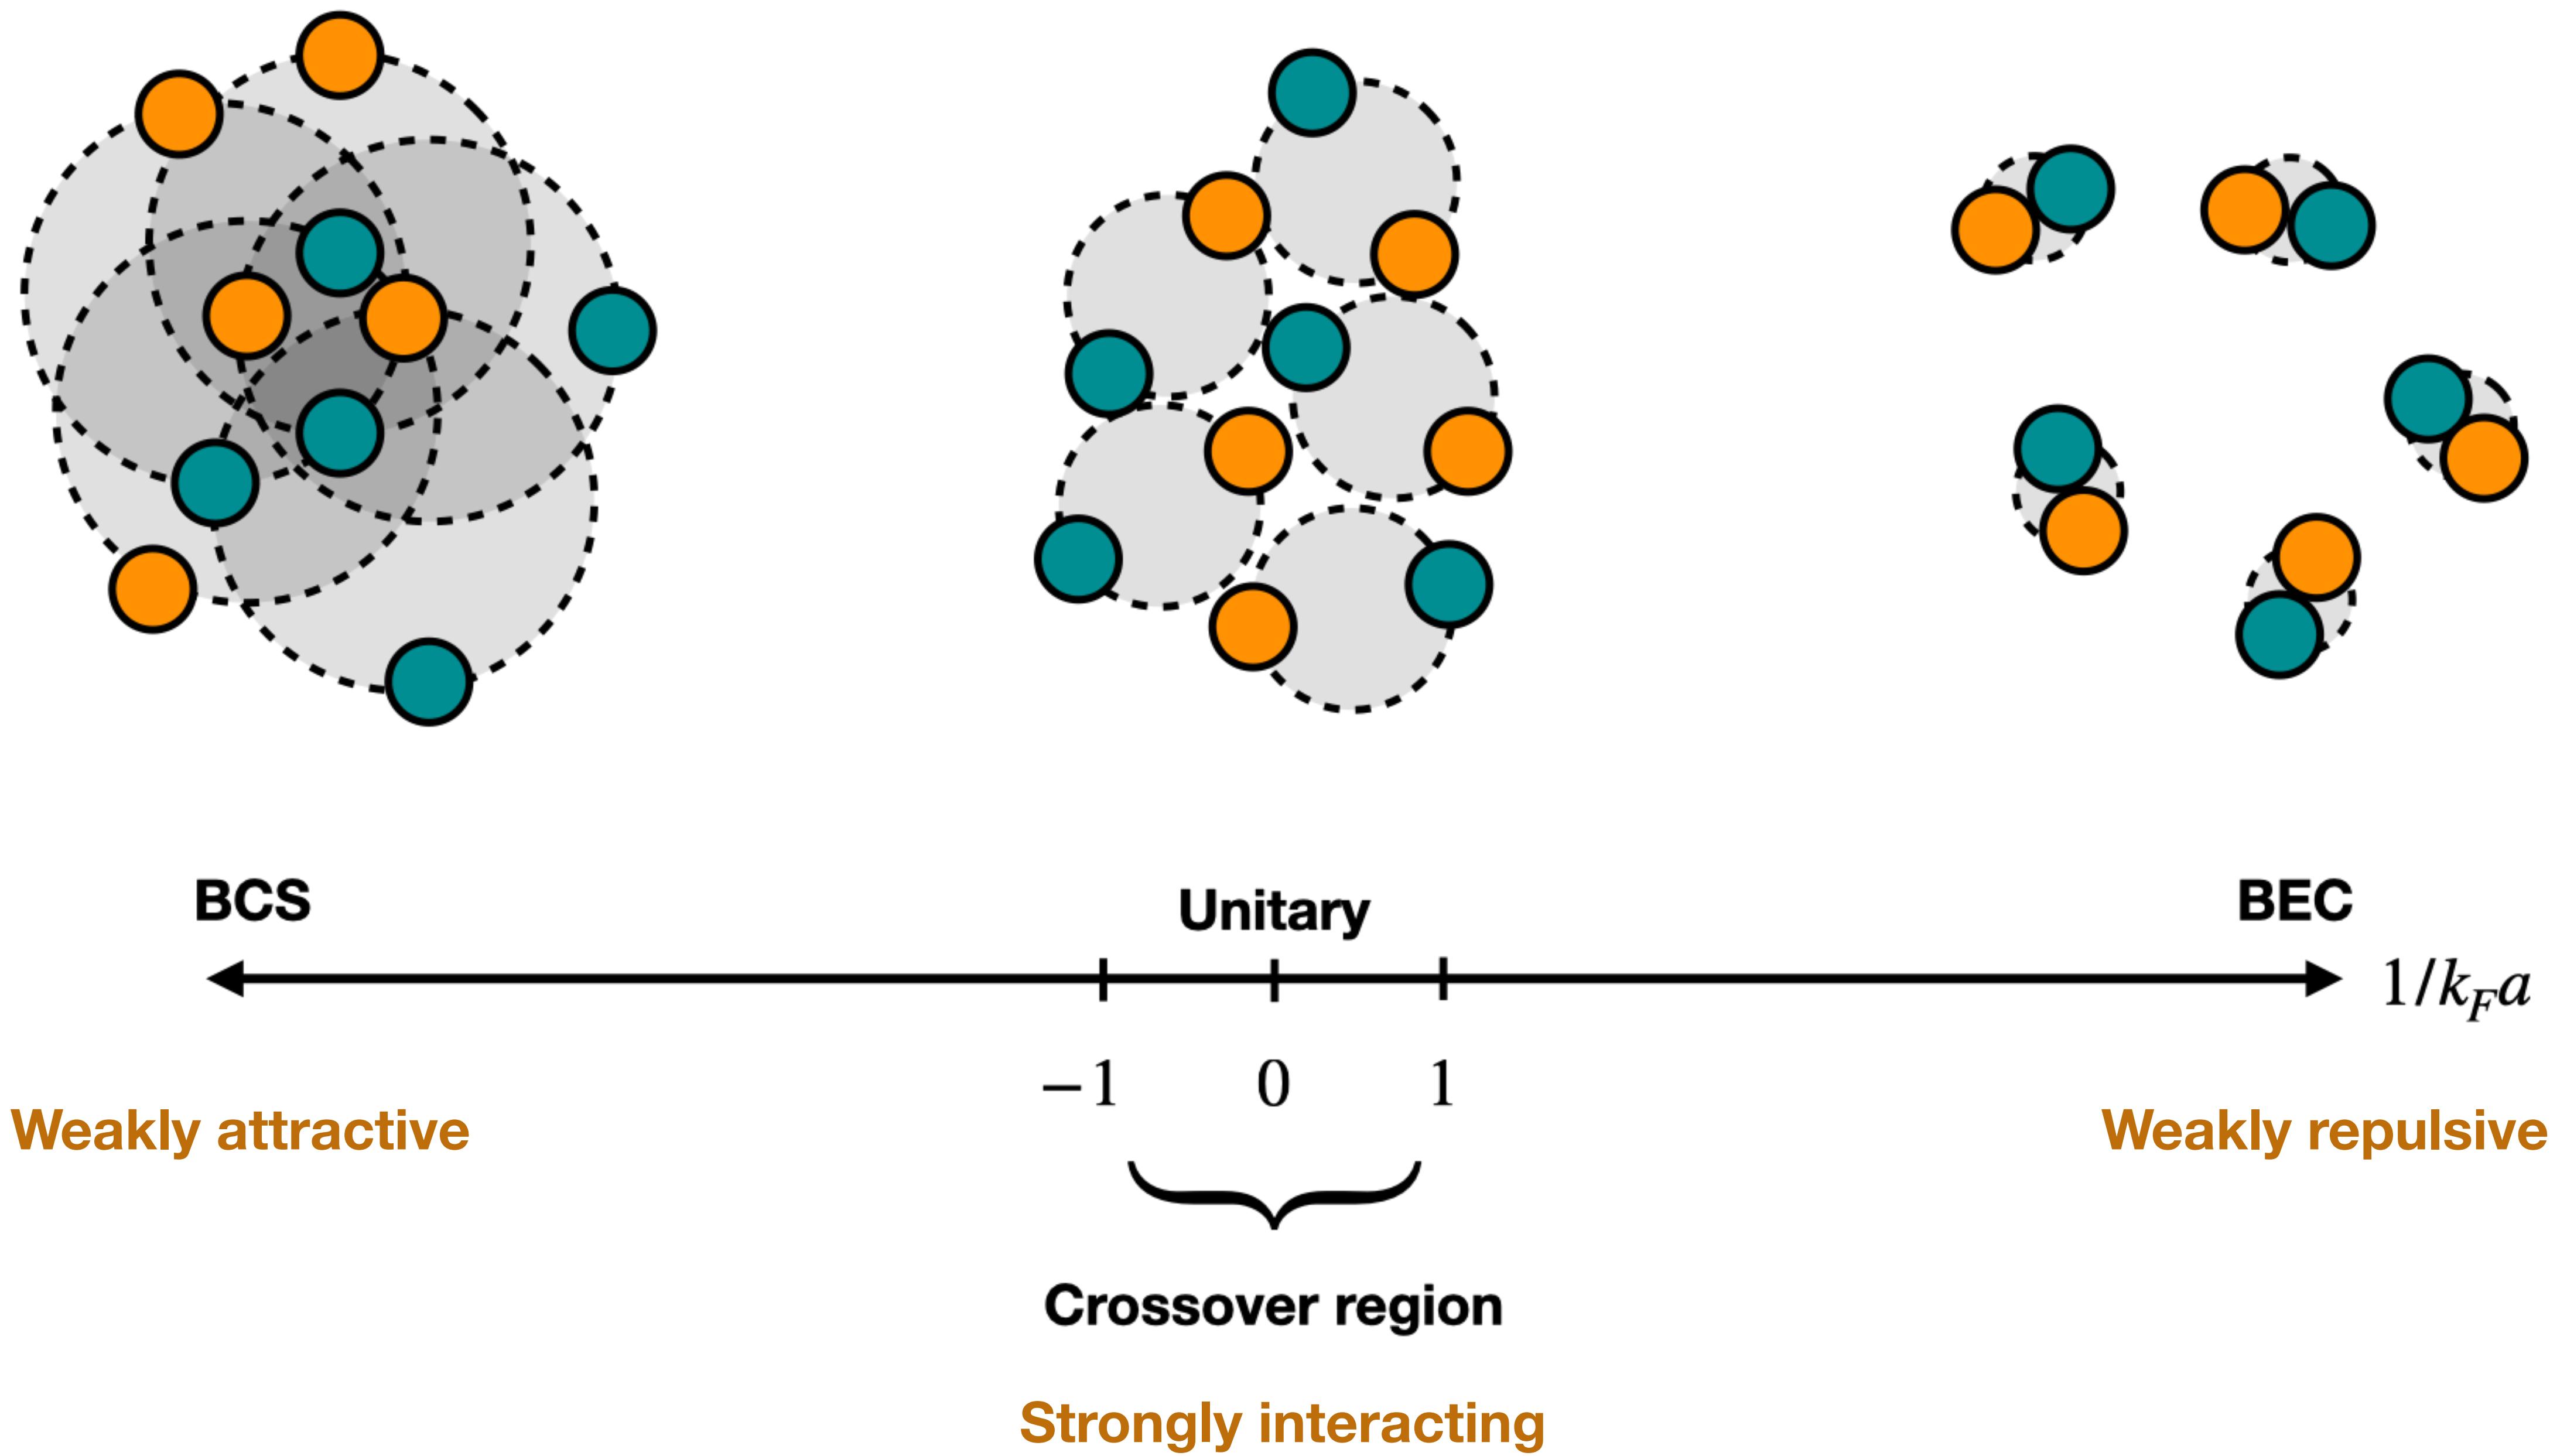


Homogeneous Electron Gas

- Message-passing neural quantum state efficiently encodes correlations with orders of magnitude fewer parameters than other neural network quantum states (e.g. FermiNet, WAPNet)
- Can handle system sizes and Wigner-Seitz radii as large as $N = 128$ and $r_s = 200$
- We observe more pronounced density fluctuations at $r_s = 200$ compared to DMC, suggesting signs of a liquid-crystal phase transition
- Since translational symmetry is preserved in our NQS, the observation of a true floating Wigner crystal can be possible, unlike other studies in which the NQS breaks translational symmetry

Ultra-cold Fermi Gases

- See our preprint “**Neural-network quantum states for ultra-cold Fermi gases**” by J. Kim, G. Pescia, B. Fore, J. Nys, G. Carleo, S. Gandolfi, M. Hjorth-Jensen, and A. Lovato (arXiv:2305.08831)
- Highly non-perturbative, short-range, attractive interaction between two-component fermions
- Behavior mainly governed by s -wave scattering length a and effective range r_e
- Universal in the unitary limit $a \rightarrow \infty$
- Relevant for understanding superfluidity in fermionic systems, dilute neutron star matter, and development of many-body methods

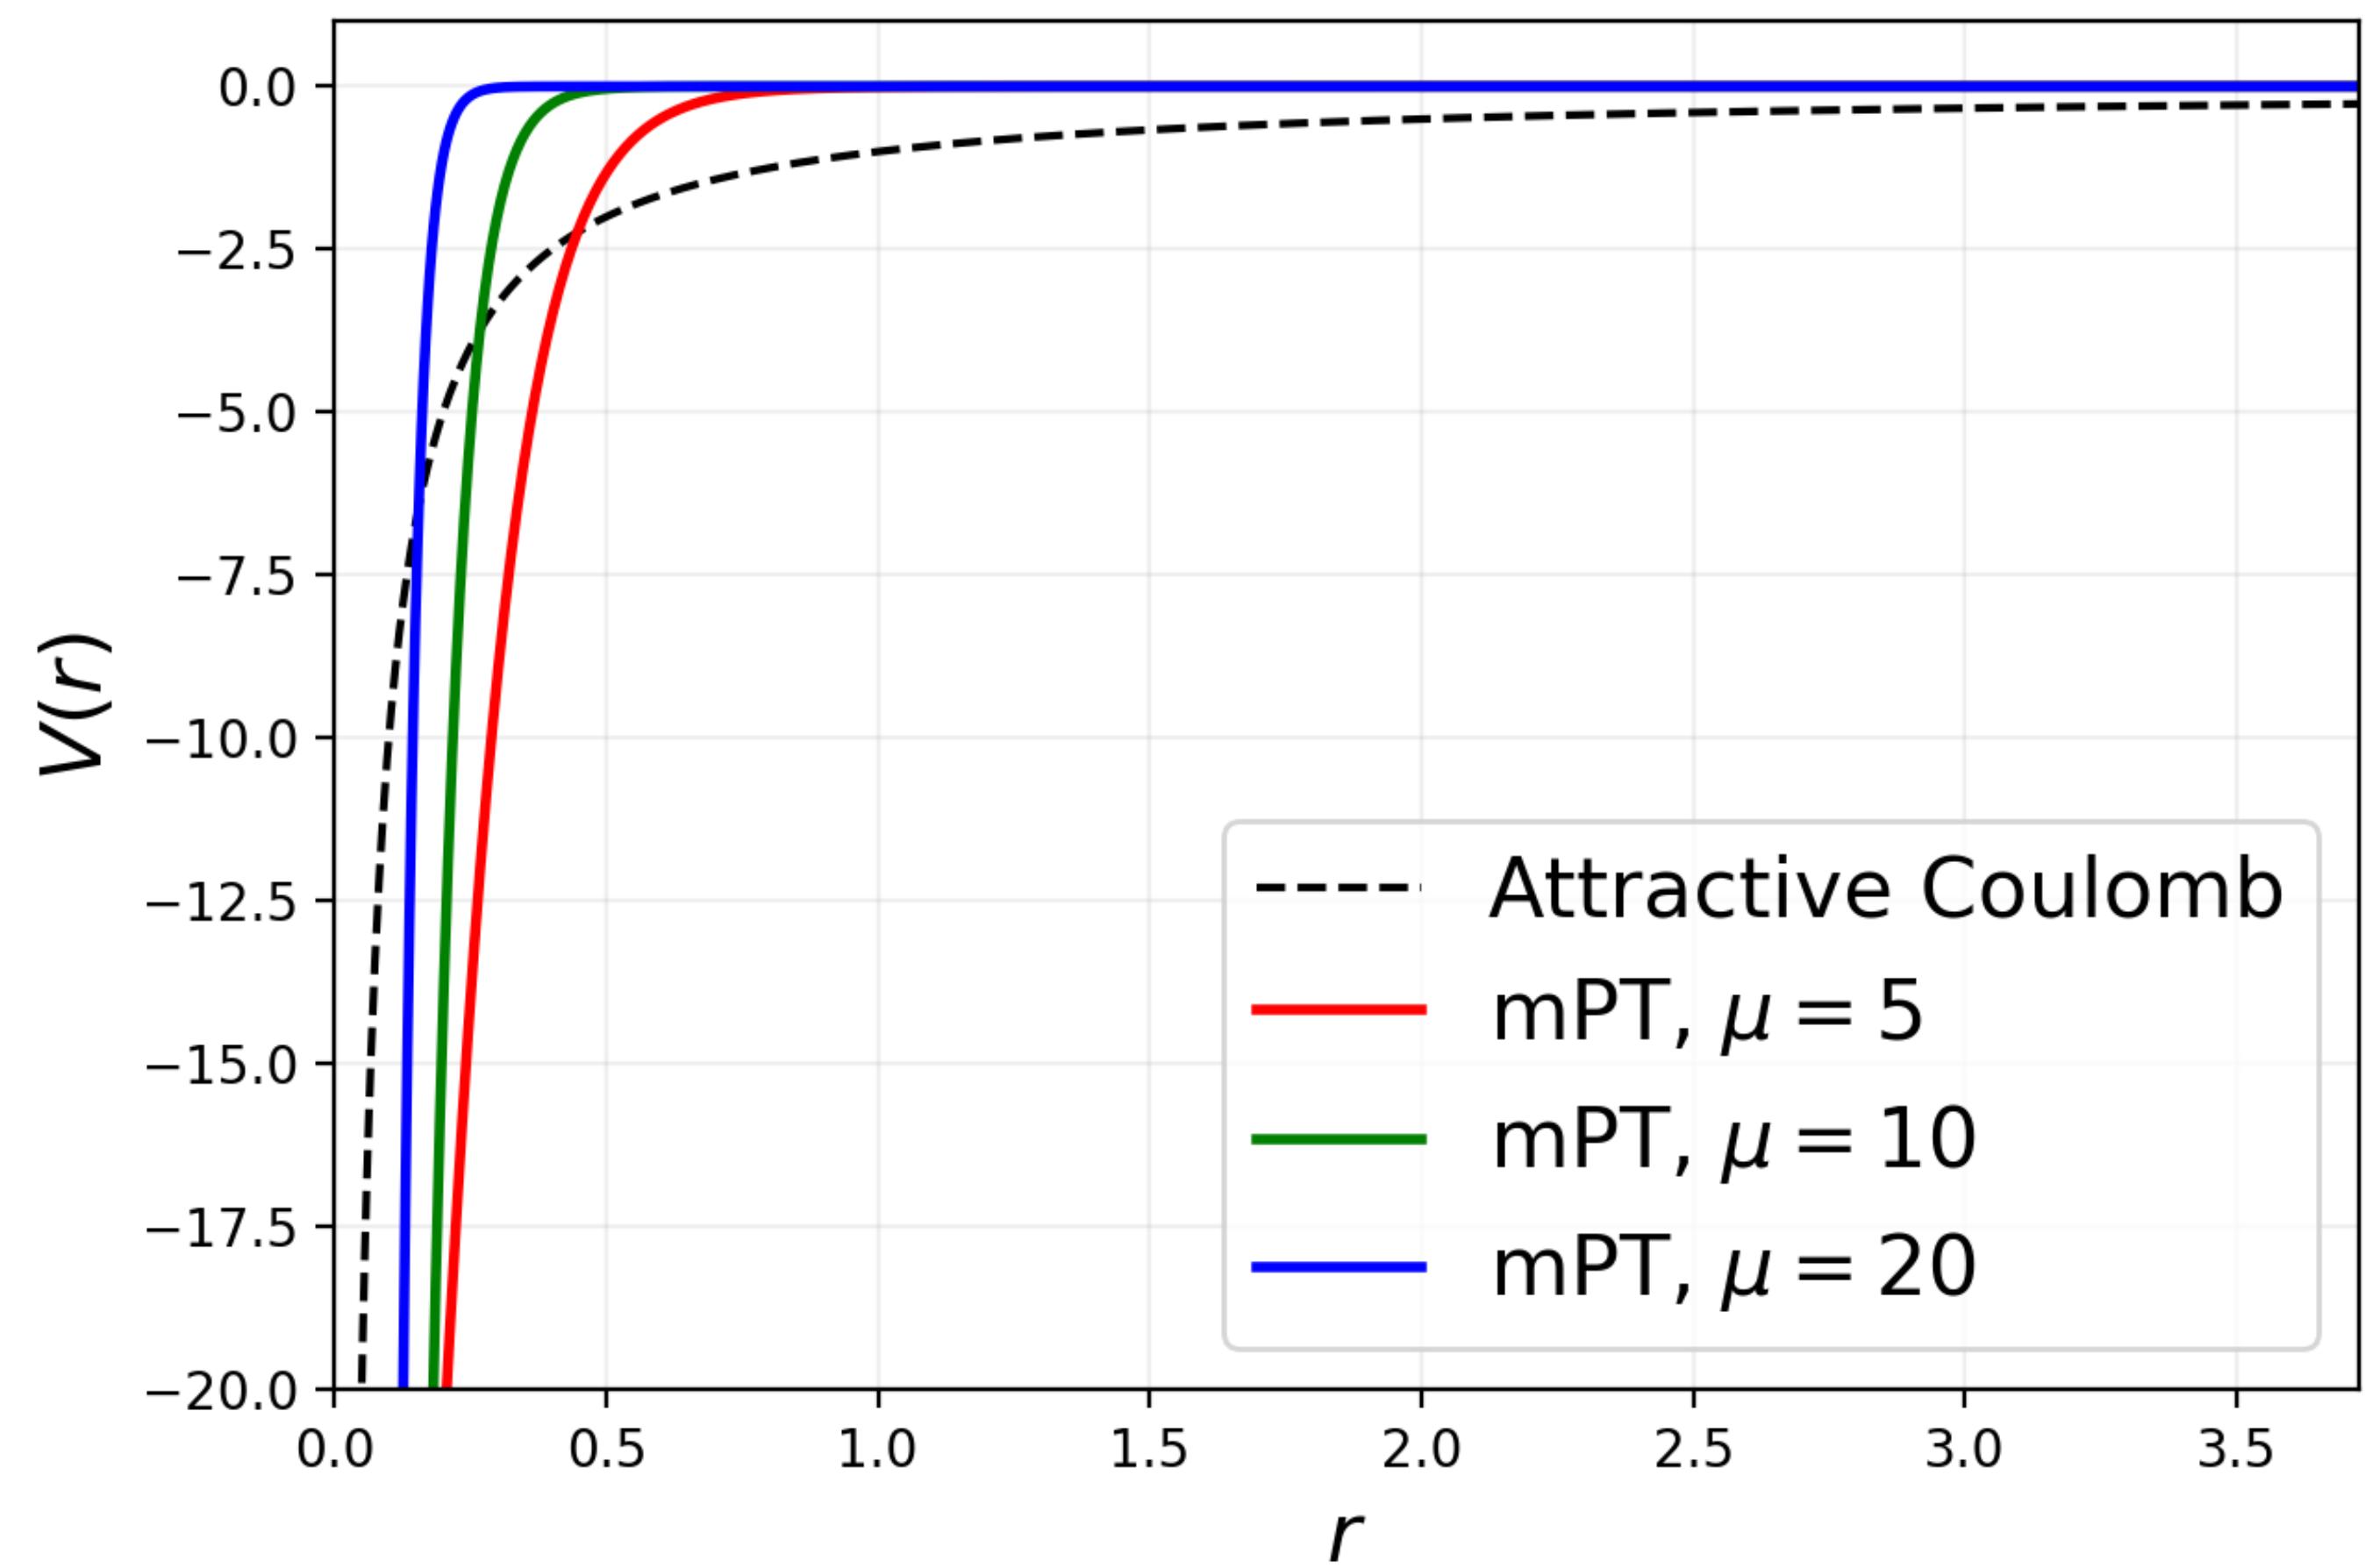


Ultra-cold Fermi Gases

- Regularized, short-range attraction between opposite-spin pairs (Pöschl-Teller)

$$V_{ij} = (s_i^z s_j^z - 1) v_0 \frac{\hbar^2}{2m} \frac{\mu^2}{\cosh^2(\mu r_{ij})}$$

- Effective range $\approx 2/\mu$
- Keep v_0 fixed, pretrain with small μ
- Provides exact solution of two-body problem
- Other interaction potentials give similar results near unitarity as long as effective range is the same



Pfaffian Wave Function

- Most general way to construct an antisymmetric wave function from a pairing orbital

$$\Phi(X) = \text{pf}[\phi(\mathbf{x}_i, \mathbf{x}_j)]$$

where ϕ must be antisymmetric.

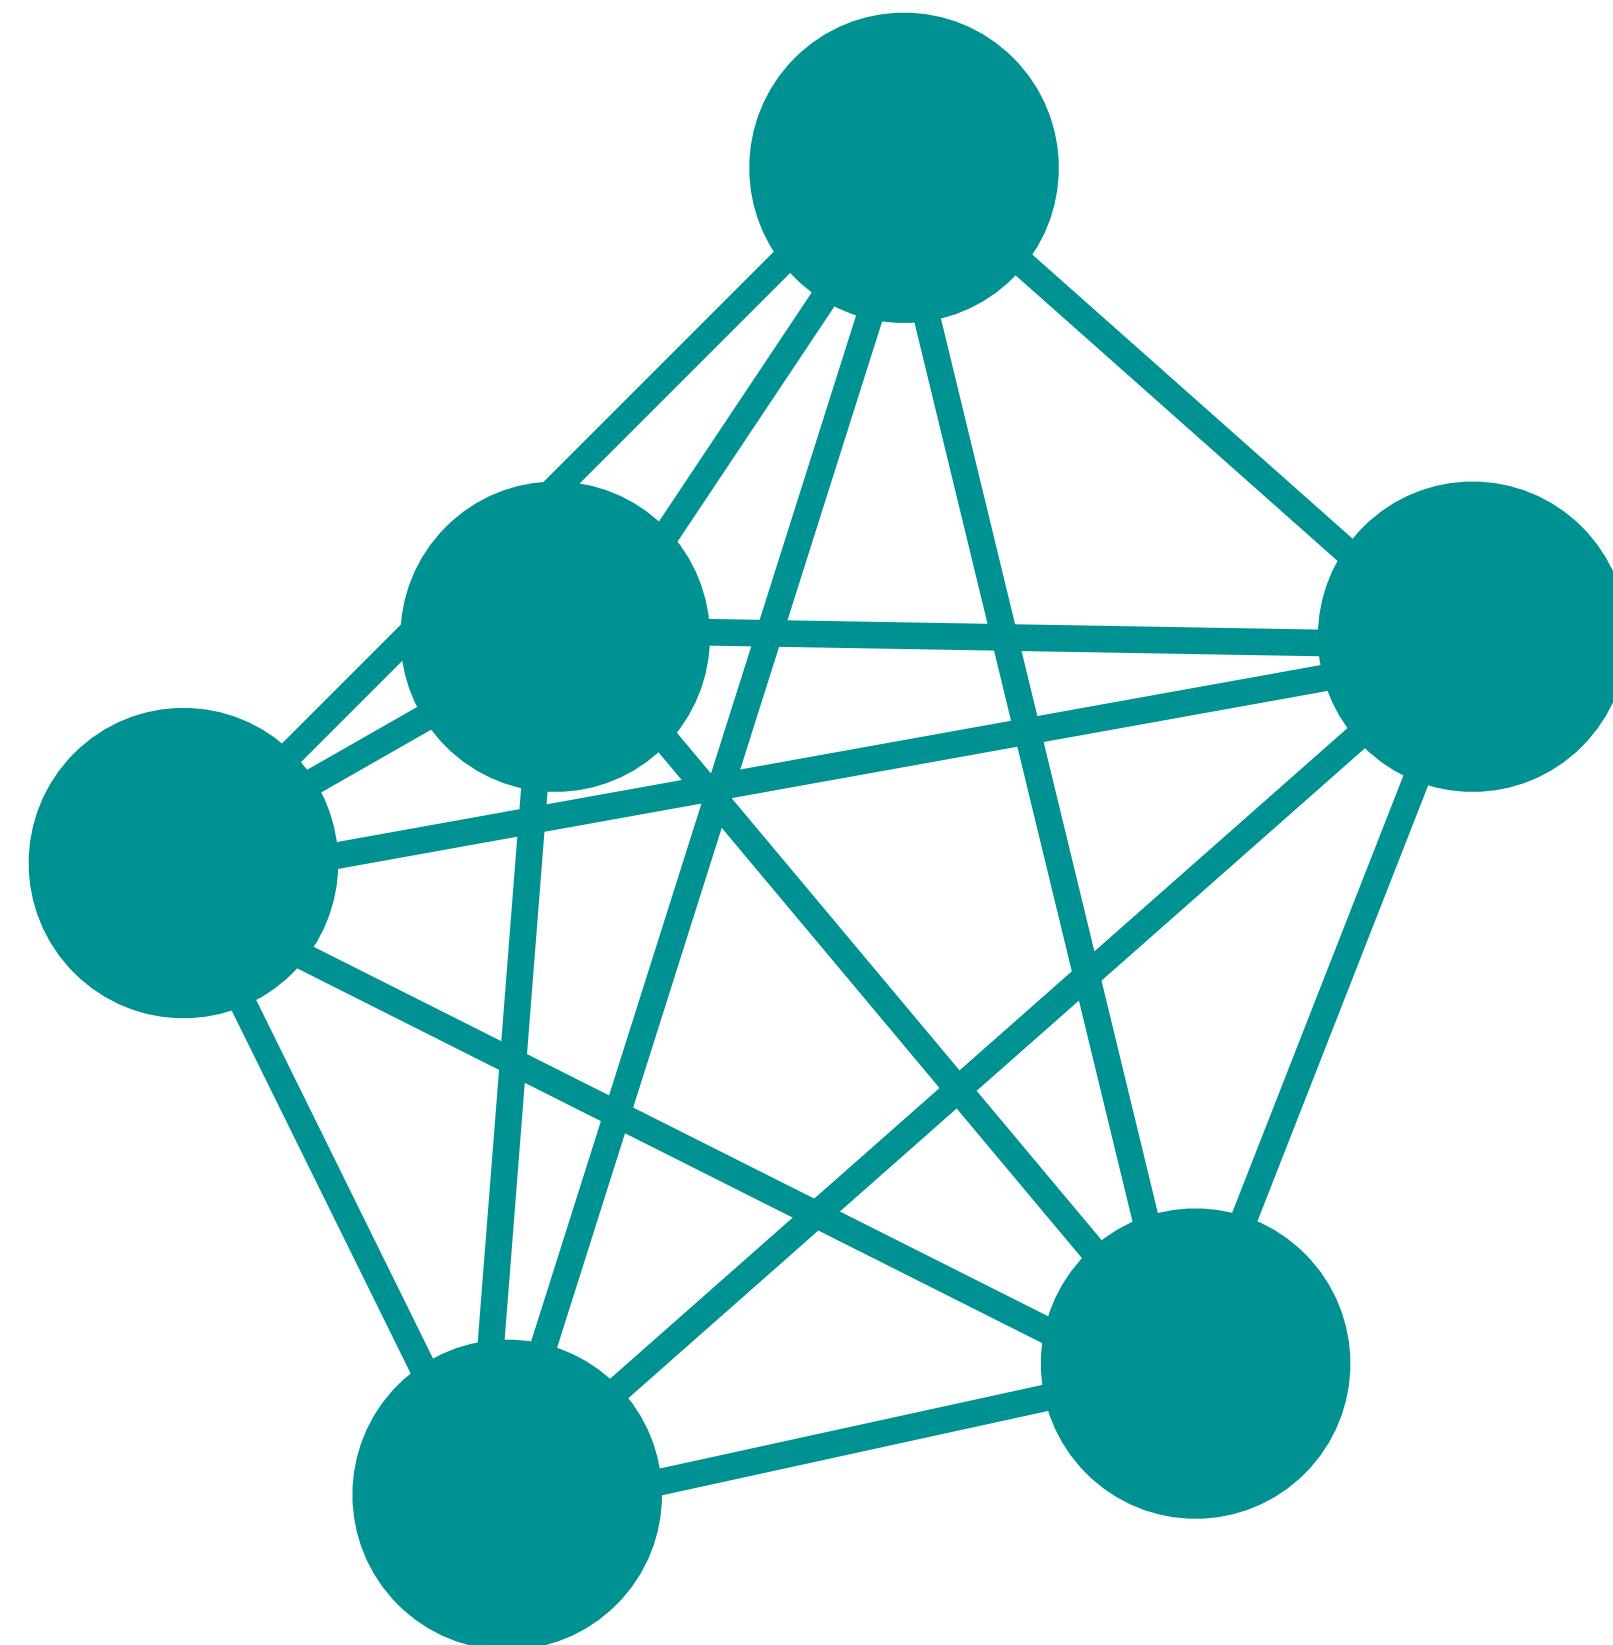
- For our NQS, we define

$$\phi(\mathbf{x}_i, \mathbf{x}_j) \equiv \nu(\mathbf{x}_i, \mathbf{x}_j) - \nu(\mathbf{x}_j, \mathbf{x}_i)$$

where ν is a neural network.

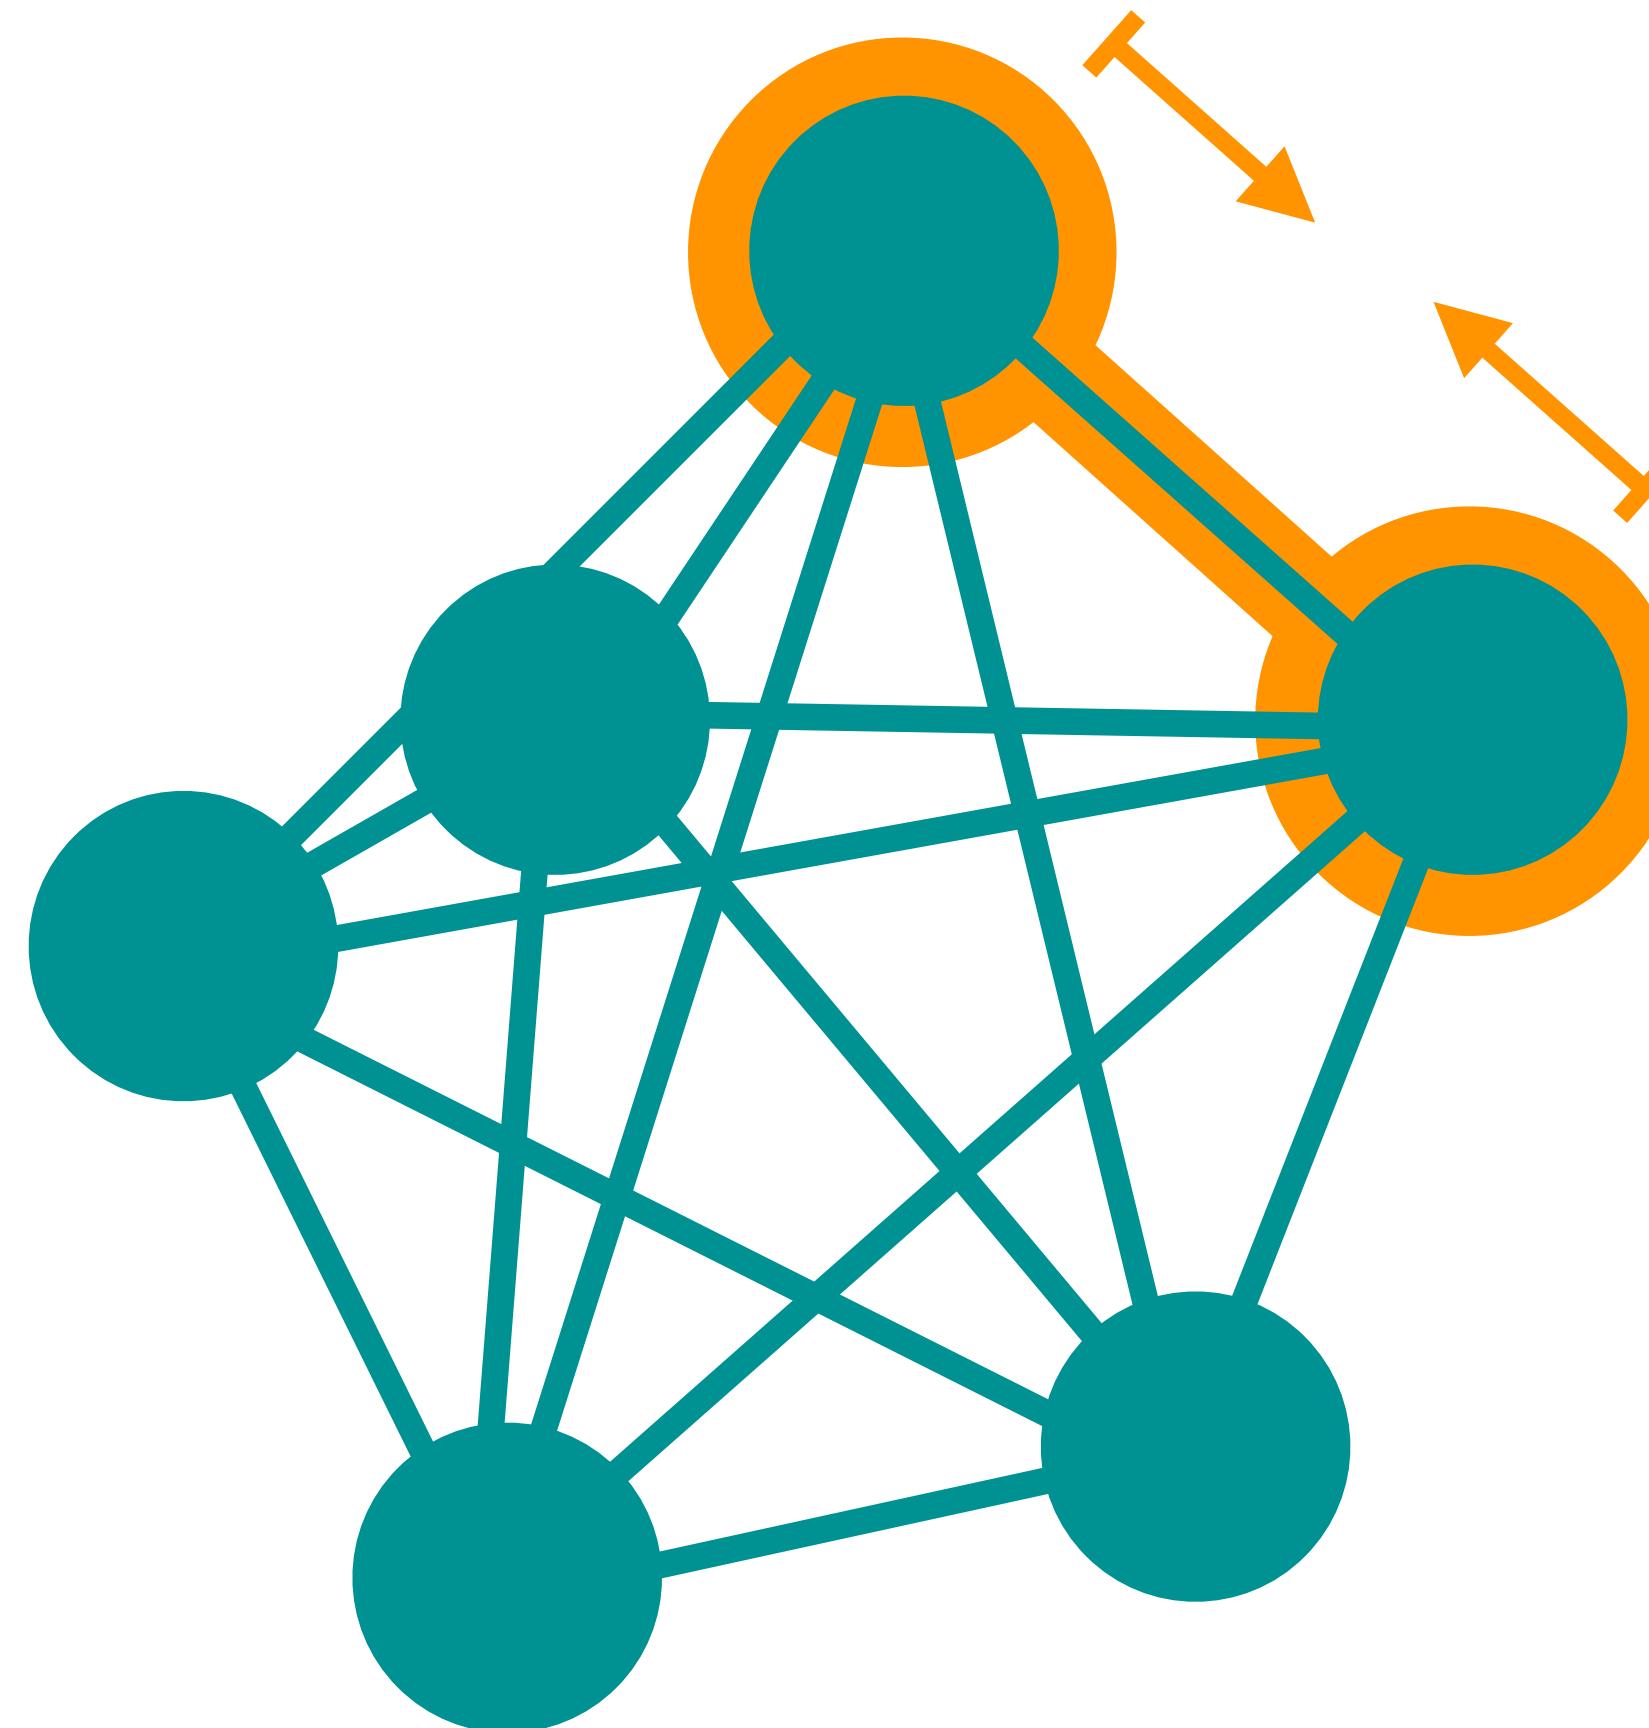
- Systematically improvable (universal approximation theorem)
- Naturally encodes singlet and triplet pairing because ν takes spins as input

Message-Passing Neural Network



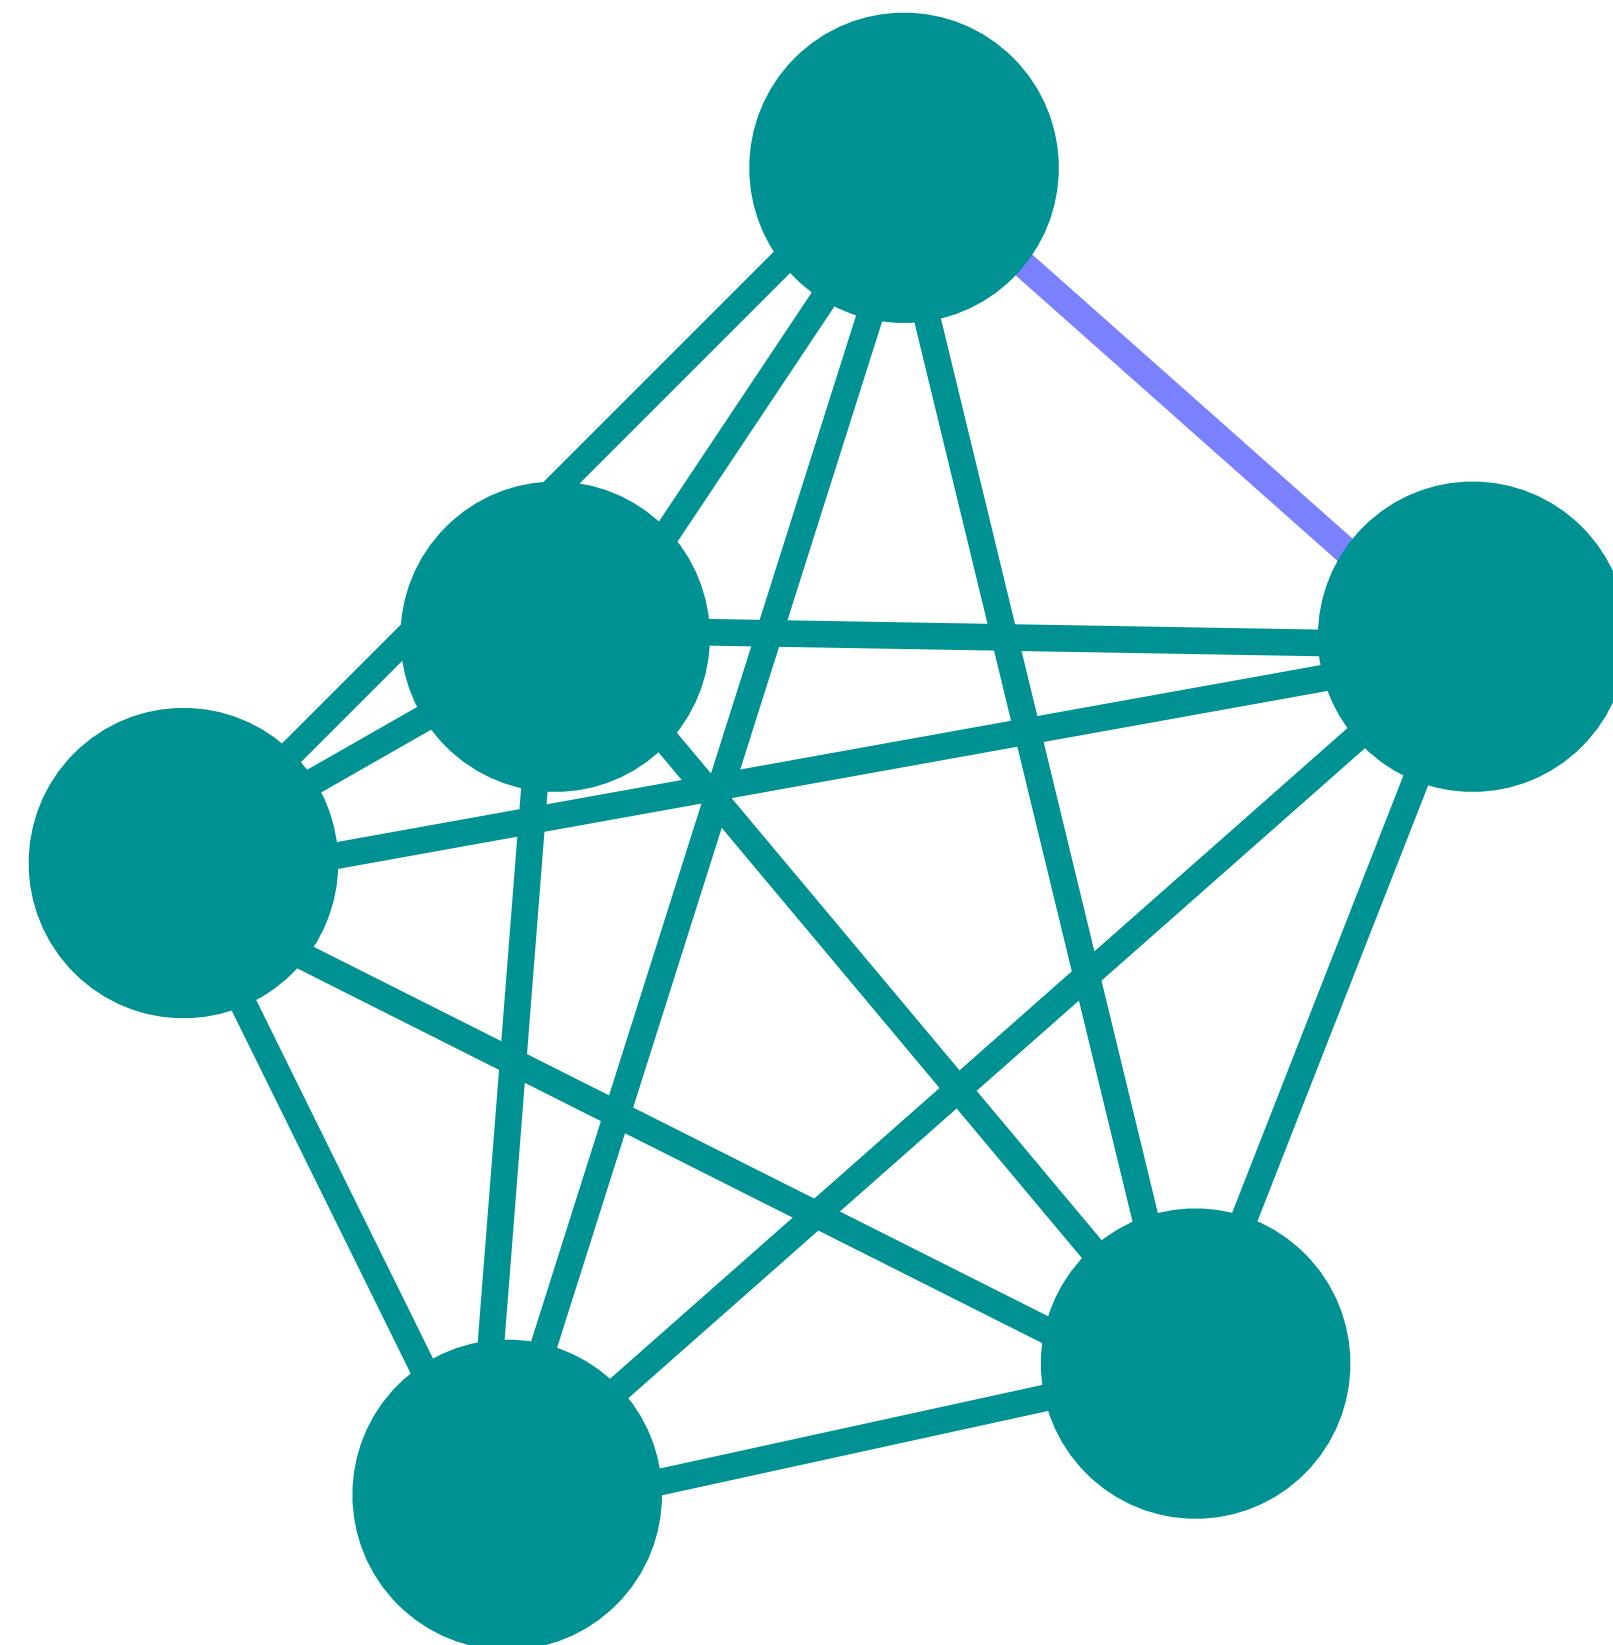
- Used in our study of the homogenous electron gas (Pescia et al. arXiv:2305.07240)
- Type of graph neural network
- Must be permutation equivariant to maintain antisymmetry
- Iteratively build correlations into new one-body and two-body features from original ones
- Nodes = one-body features
- Edges = two-body features

Message-Passing Neural Network



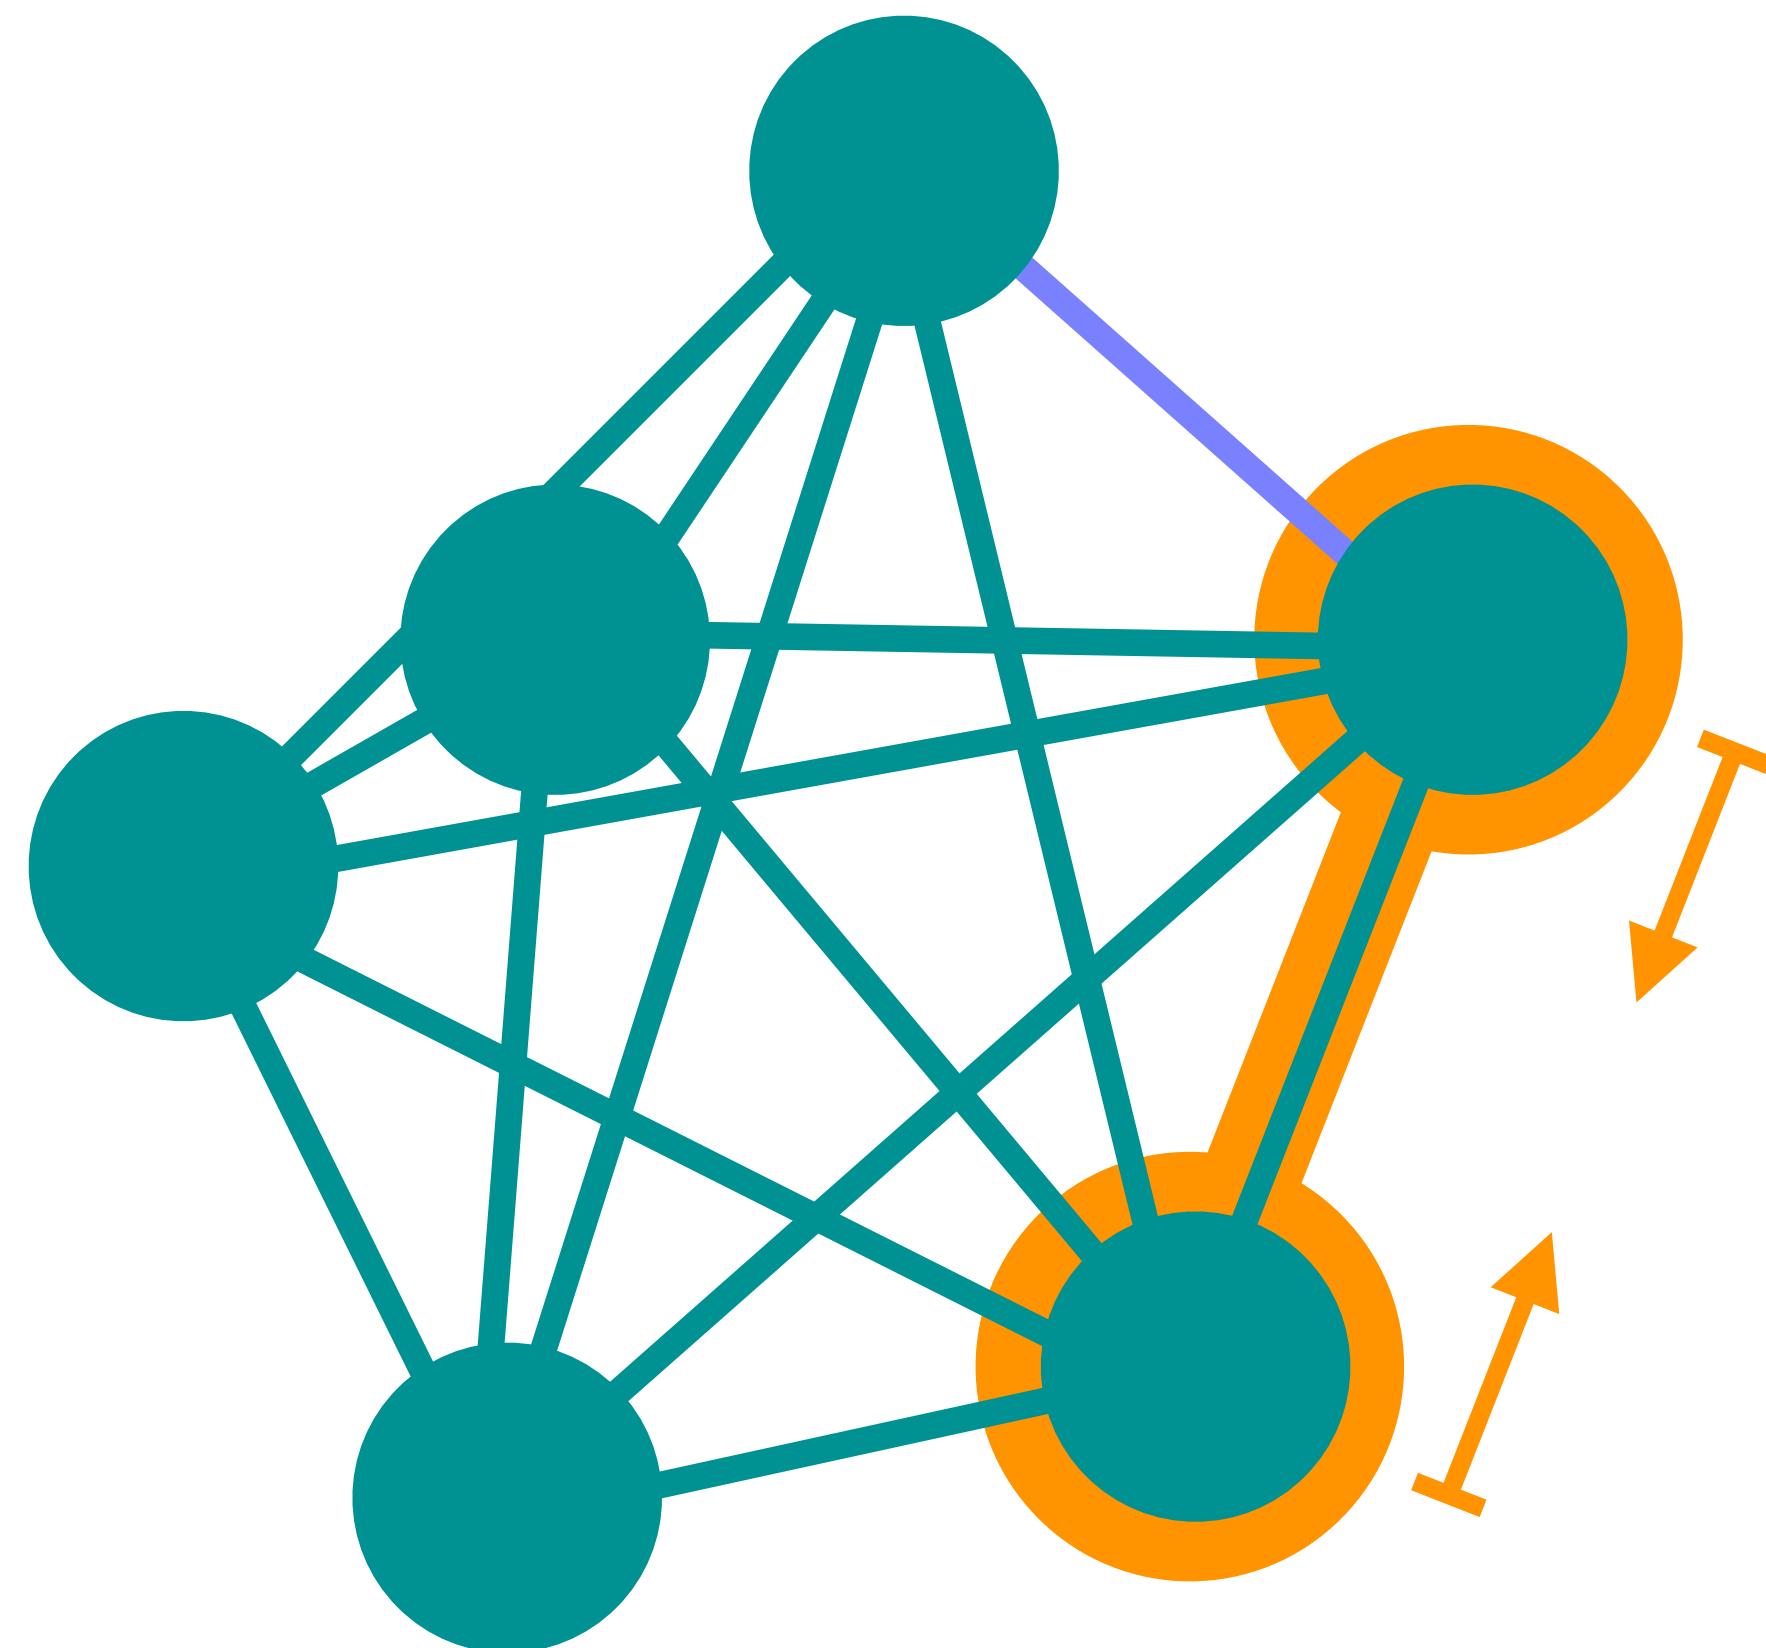
- Used in our study of the homogenous electron gas (Pescia et al. arXiv:2305.07240)
- Type of graph neural network
- Must be permutation equivariant to maintain antisymmetry
- Iteratively build correlations into new one-body and two-body features from original ones
- Nodes = one-body features
- Edges = two-body features

Message-Passing Neural Network



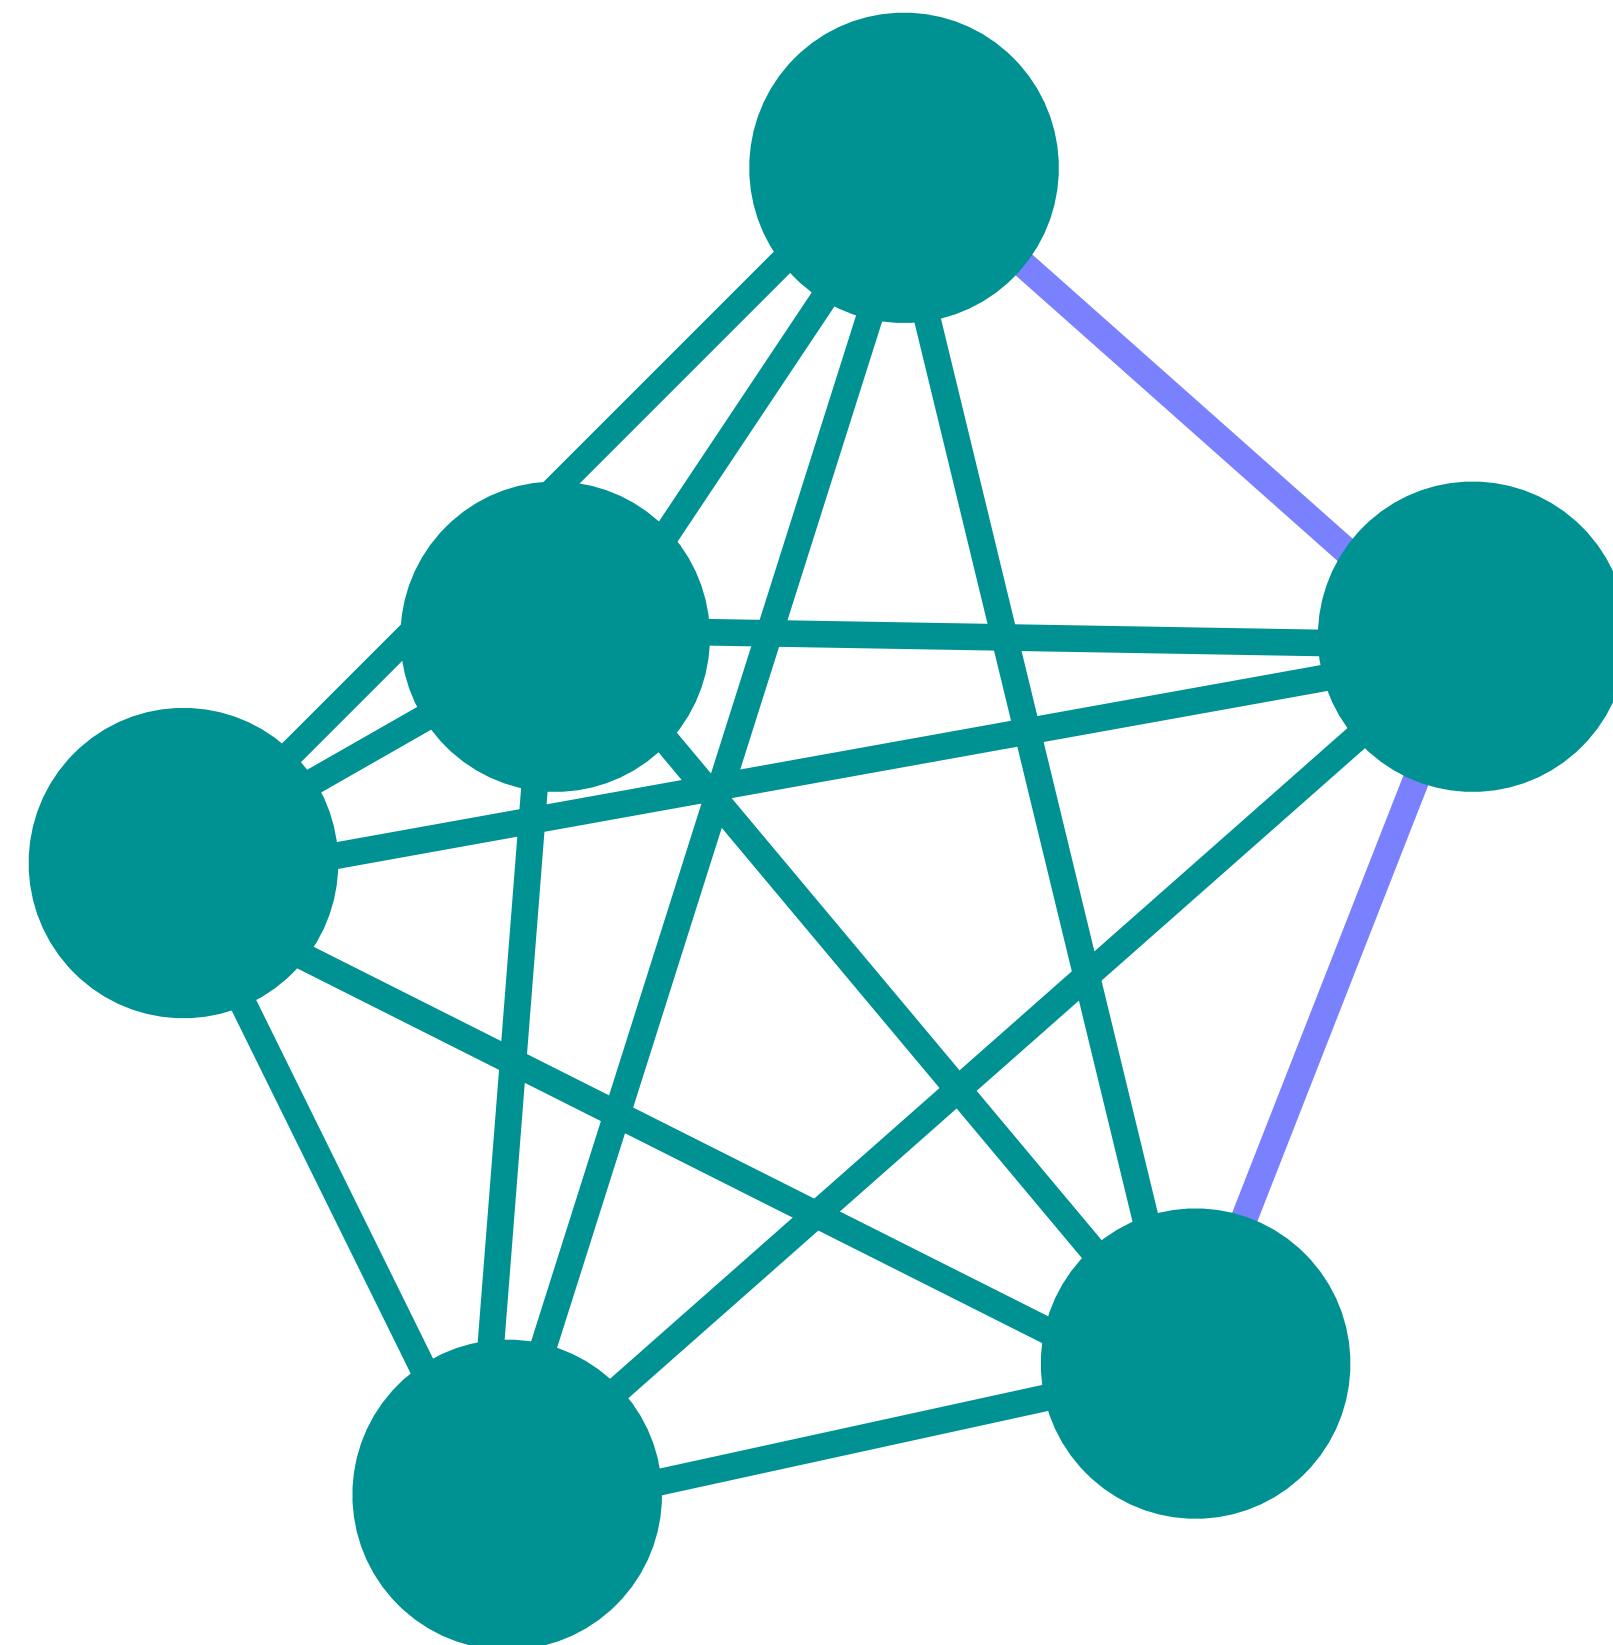
- Used in our study of the homogenous electron gas (Pescia et al. arXiv:2305.07240)
- Type of graph neural network
- Must be permutation equivariant to maintain antisymmetry
- Iteratively build correlations into new one-body and two-body features from original ones
- Nodes = one-body features
- Edges = two-body features

Message-Passing Neural Network



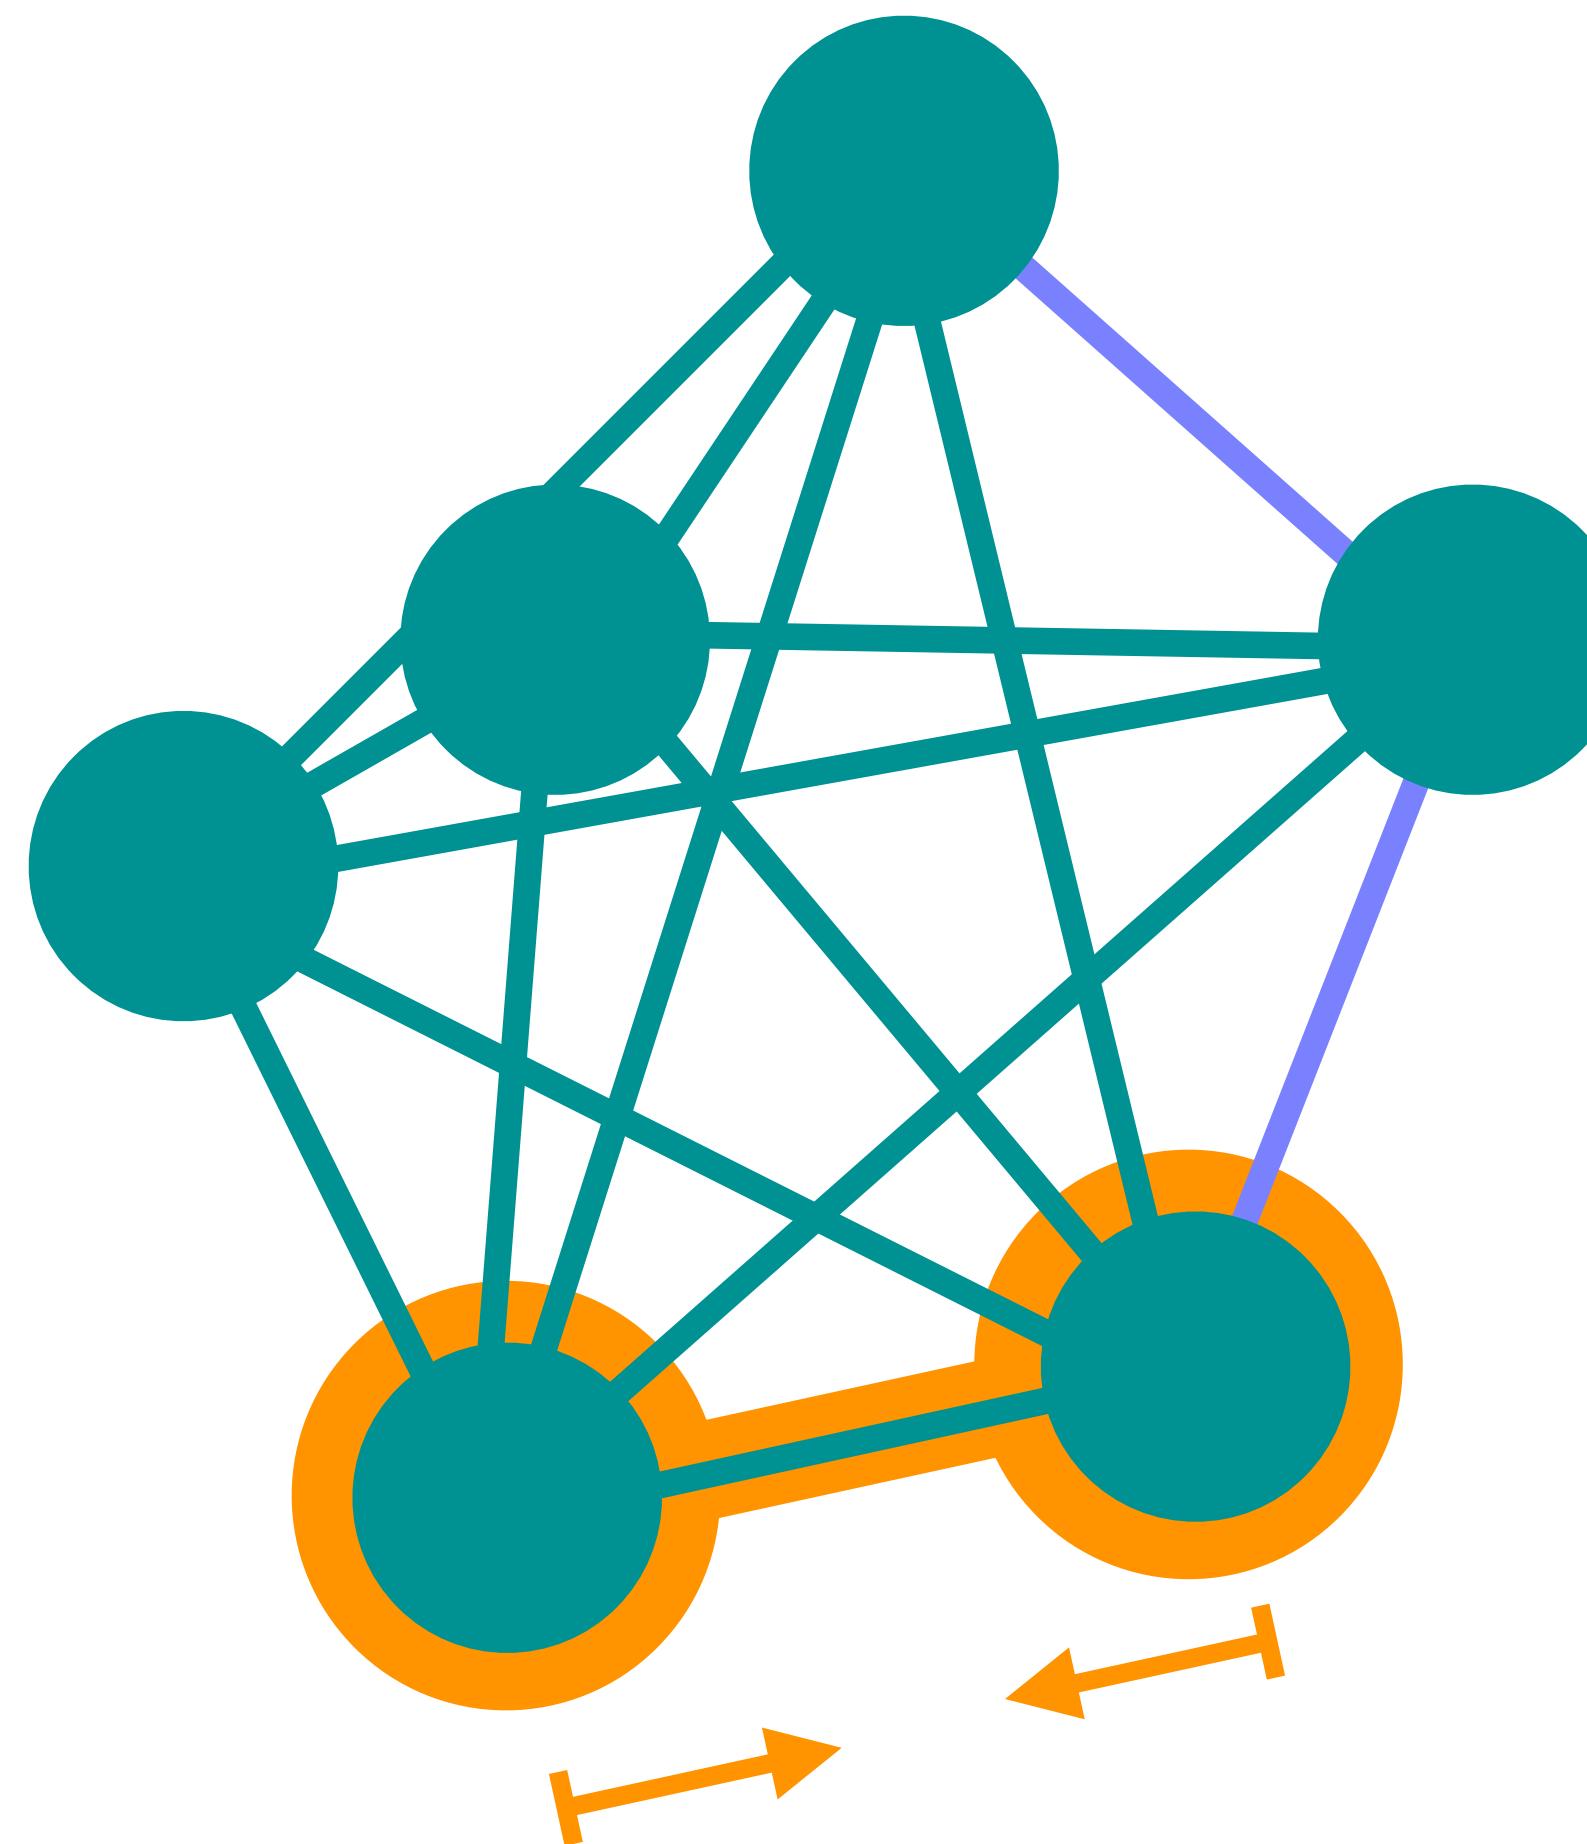
- Used in our study of the homogenous electron gas (Pescia et al. arXiv:2305.07240)
- Type of graph neural network
- Must be permutation equivariant to maintain antisymmetry
- Iteratively build correlations into new one-body and two-body features from original ones
- Nodes = one-body features
- Edges = two-body features

Message-Passing Neural Network



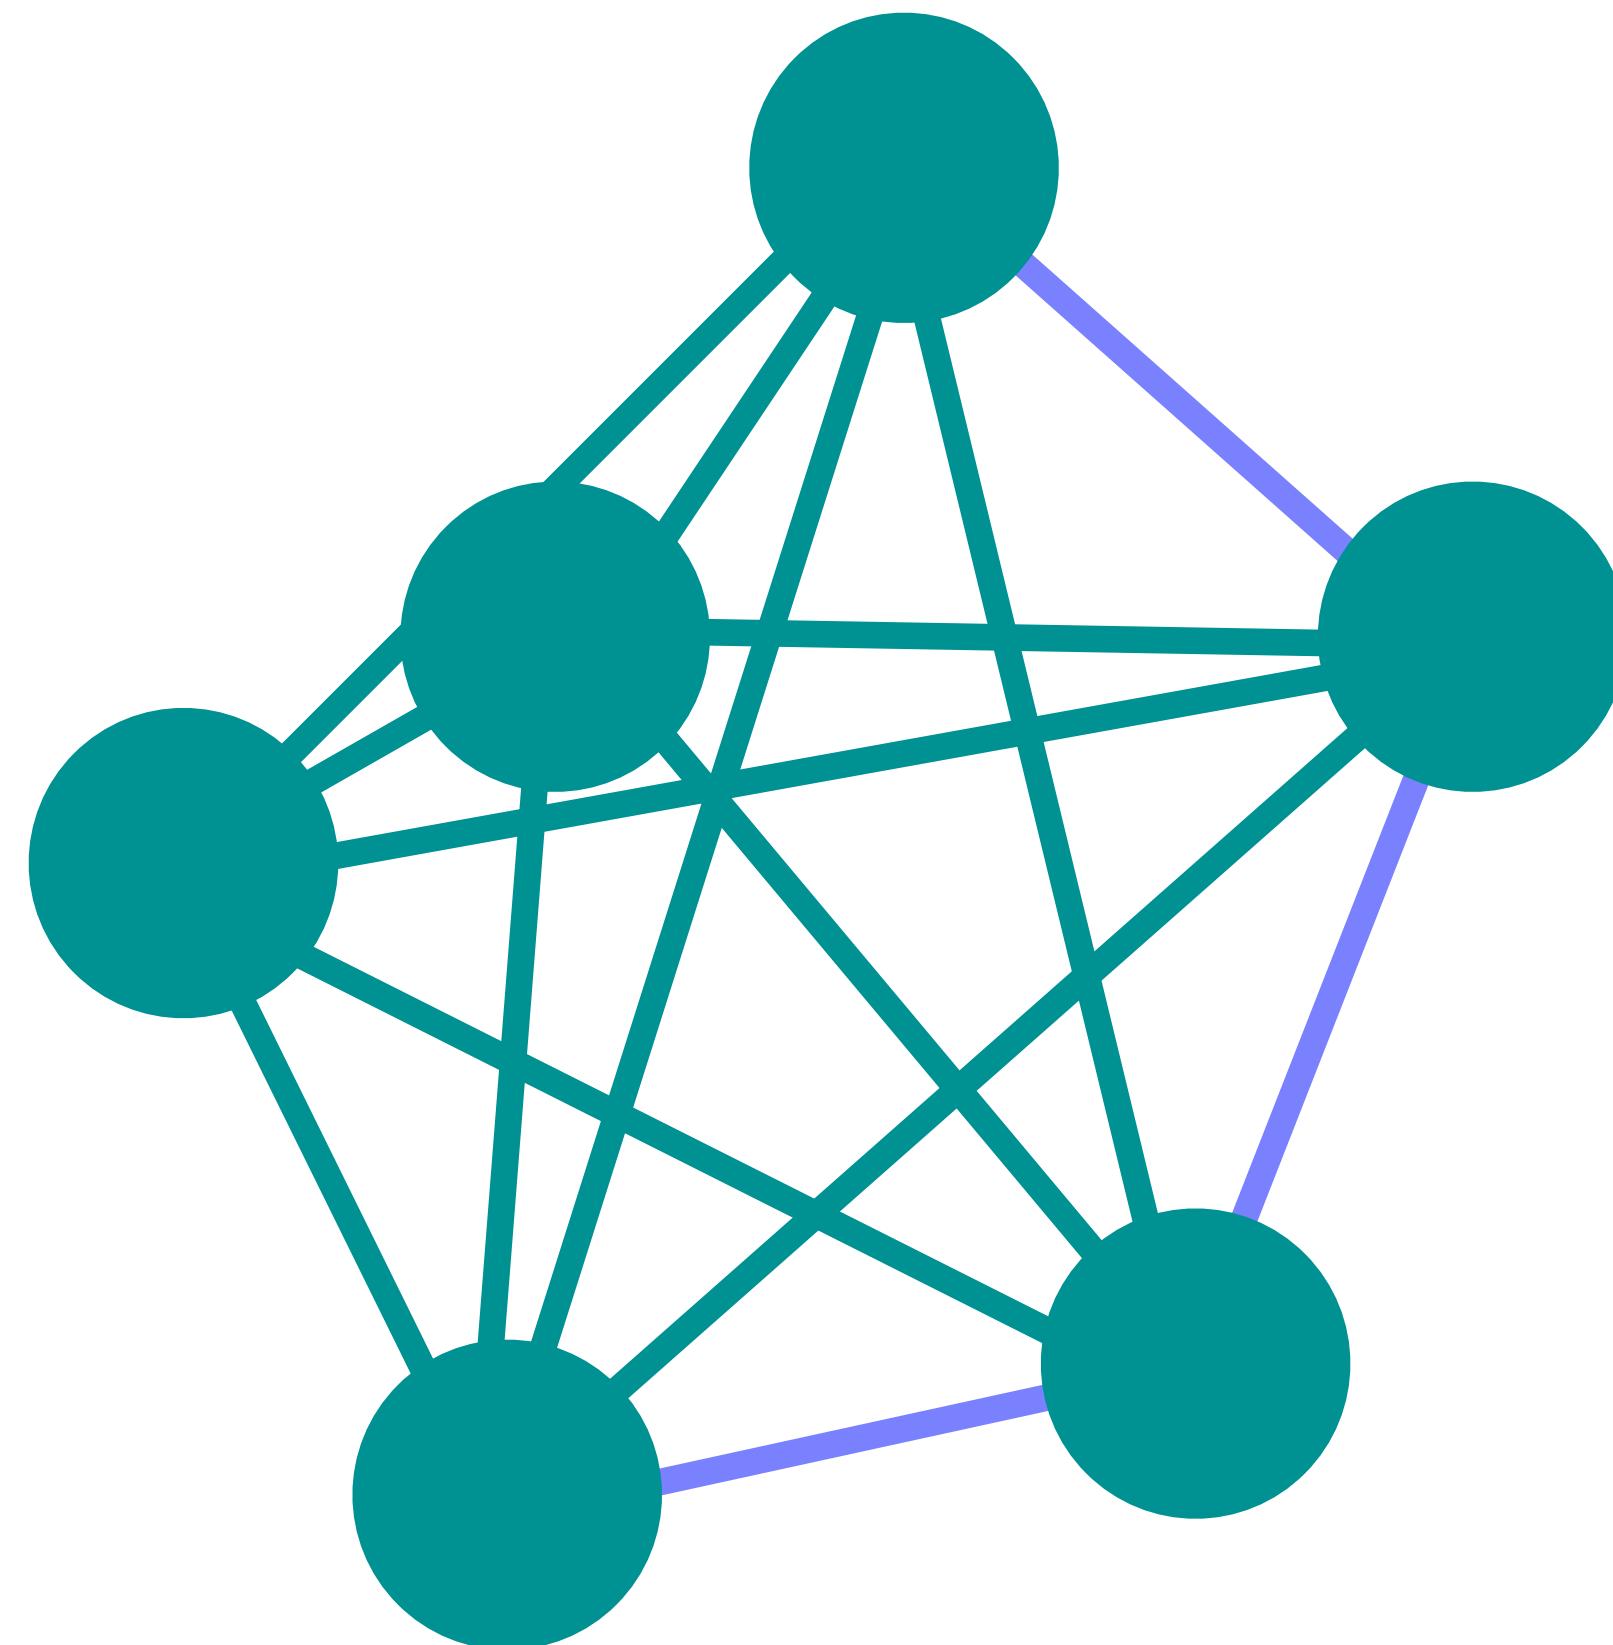
- Used in our study of the homogenous electron gas (Pescia et al. arXiv:2305.07240)
- Type of graph neural network
- Must be permutation equivariant to maintain antisymmetry
- Iteratively build correlations into new one-body and two-body features from original ones
- Nodes = one-body features
- Edges = two-body features

Message-Passing Neural Network



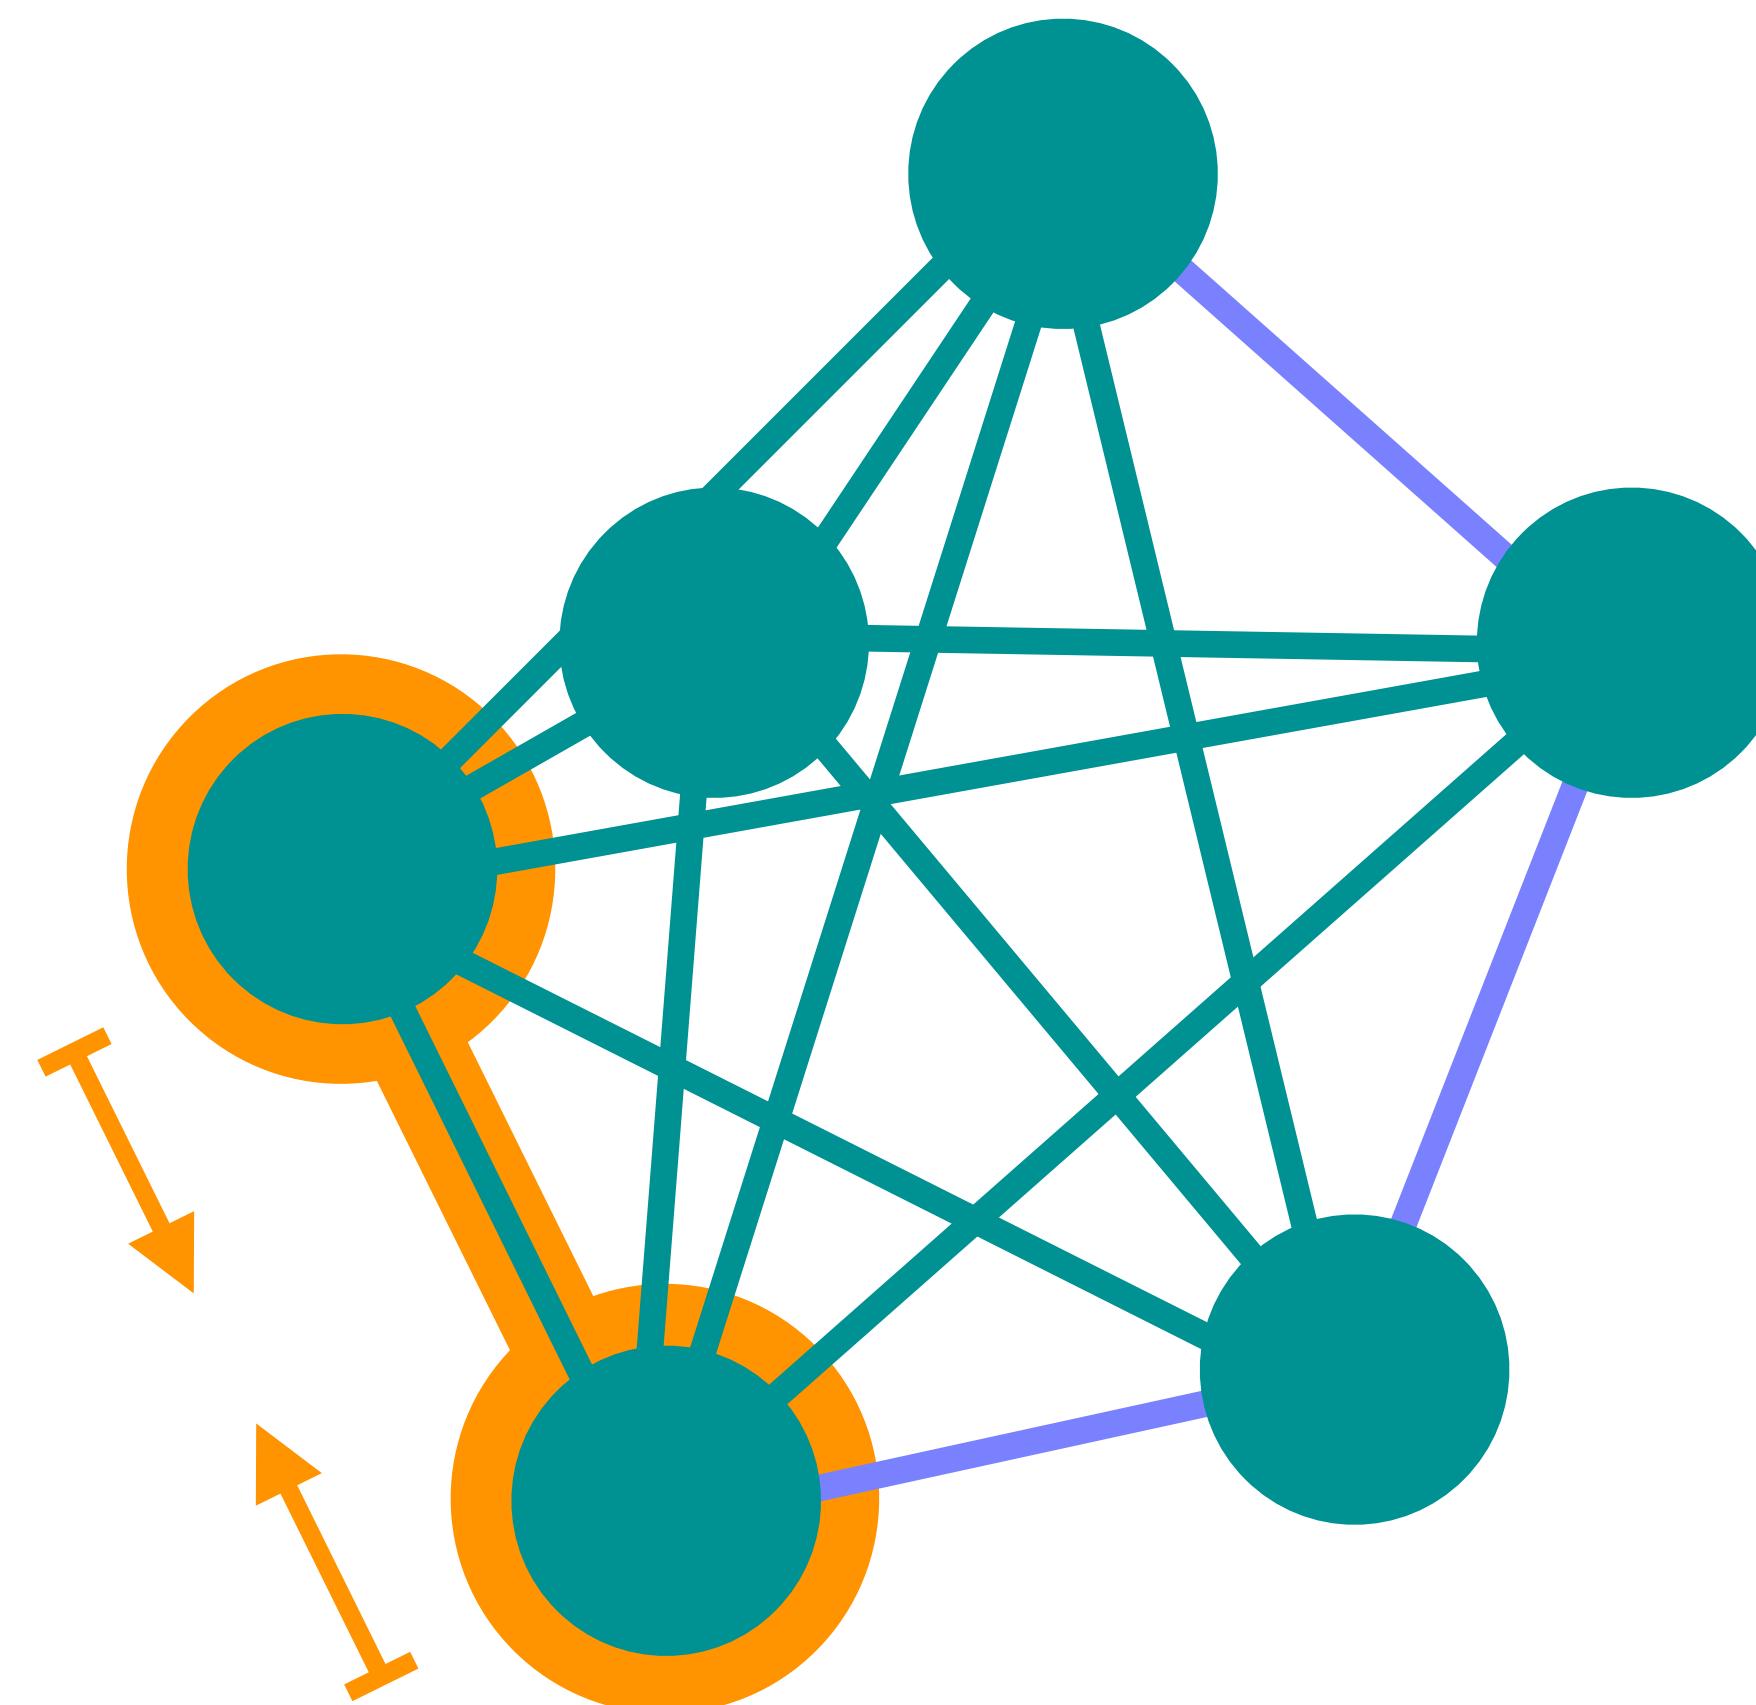
- Used in our study of the homogenous electron gas (Pescia et al. arXiv:2305.07240)
- Type of graph neural network
- Must be permutation equivariant to maintain antisymmetry
- Iteratively build correlations into new one-body and two-body features from original ones
- Nodes = one-body features
- Edges = two-body features

Message-Passing Neural Network



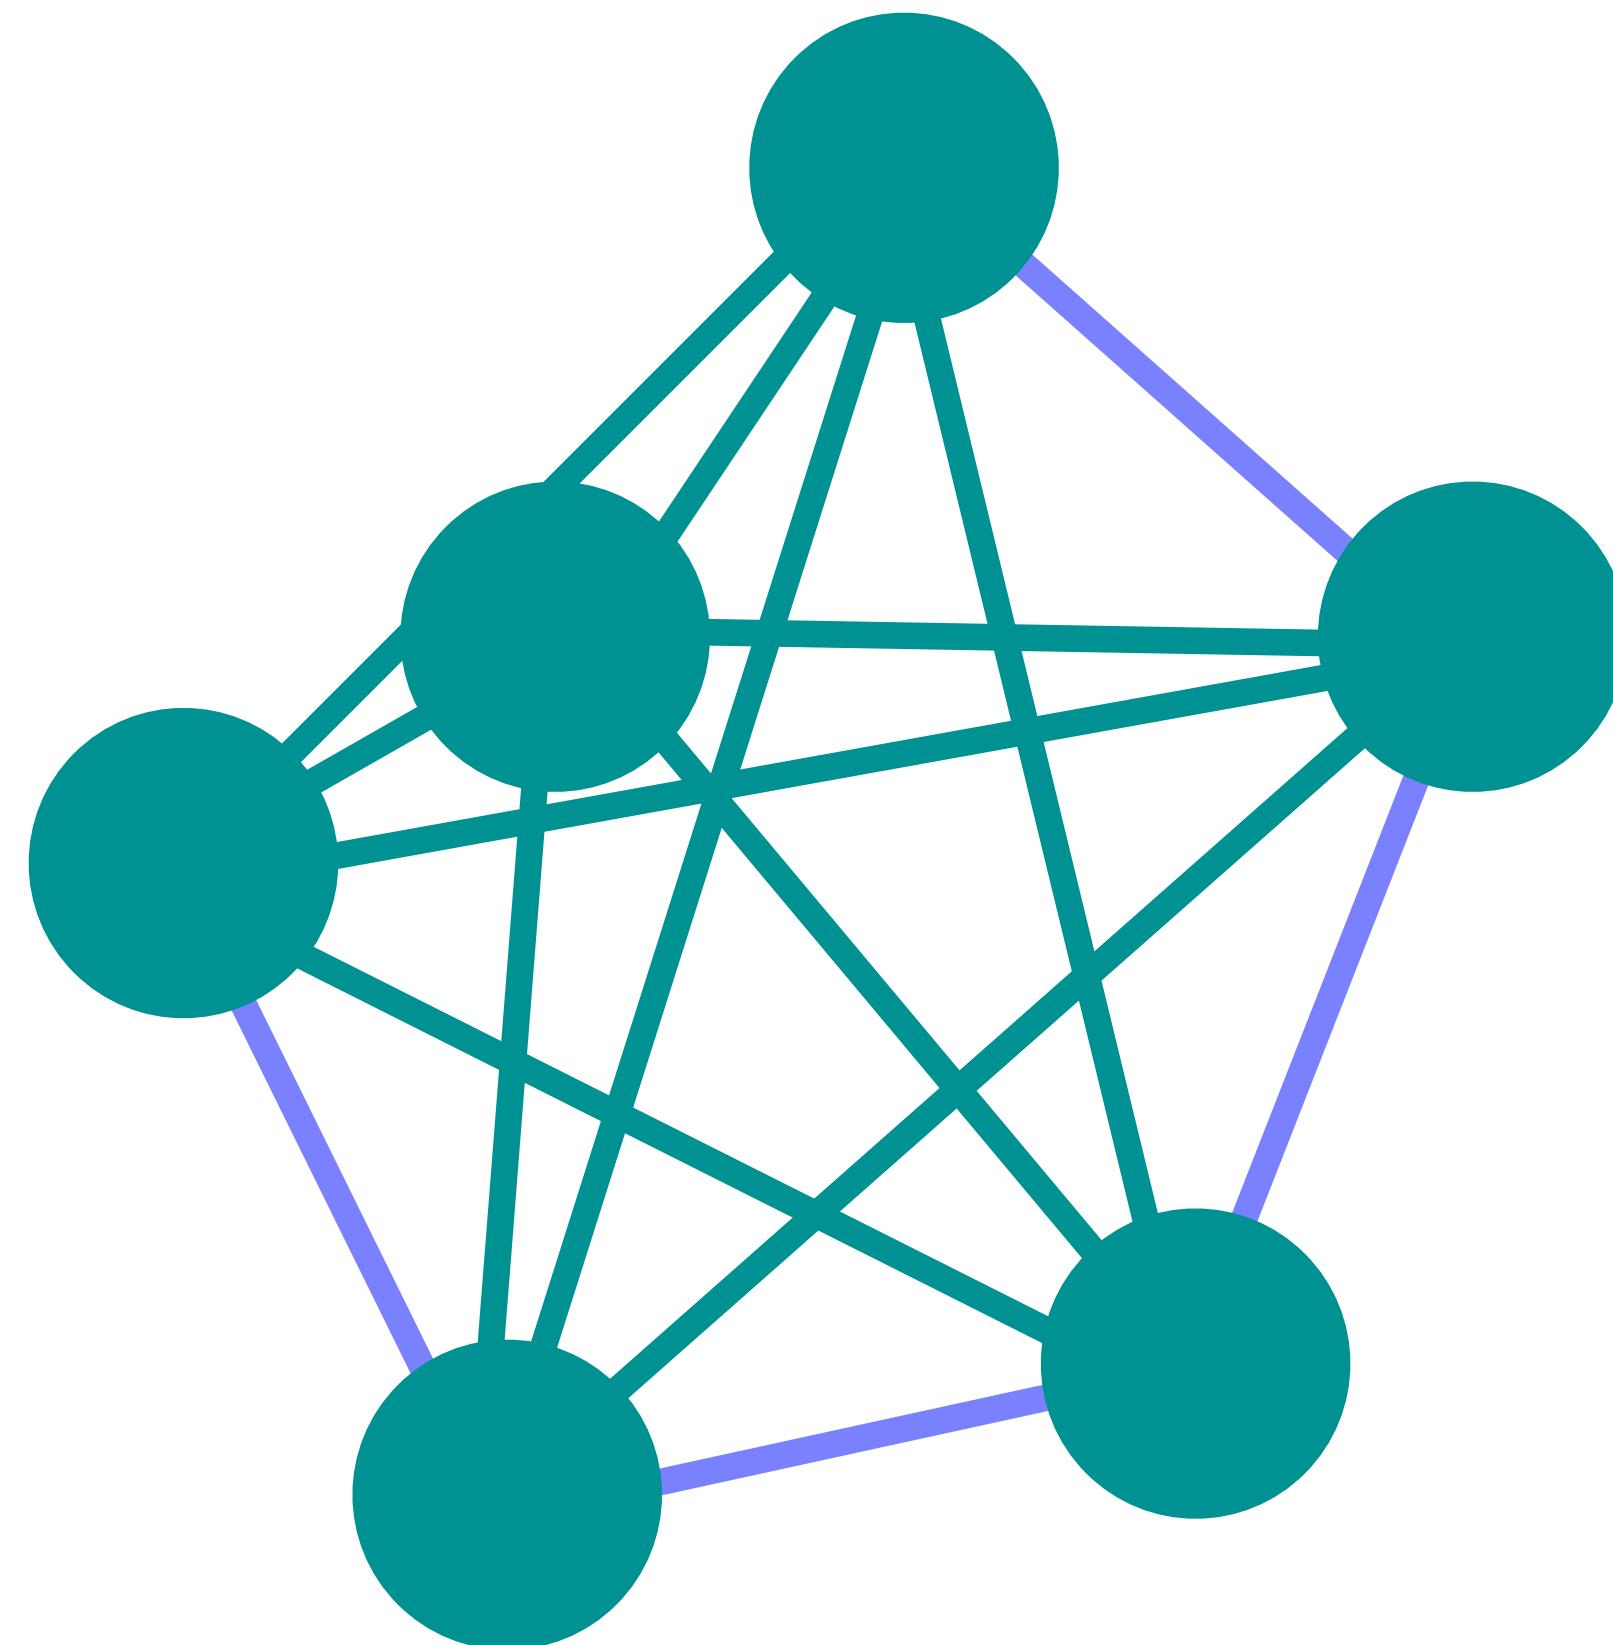
- Used in our study of the homogenous electron gas (Pescia et al. arXiv:2305.07240)
- Type of graph neural network
- Must be permutation equivariant to maintain antisymmetry
- Iteratively build correlations into new one-body and two-body features from original ones
- Nodes = one-body features
- Edges = two-body features

Message-Passing Neural Network



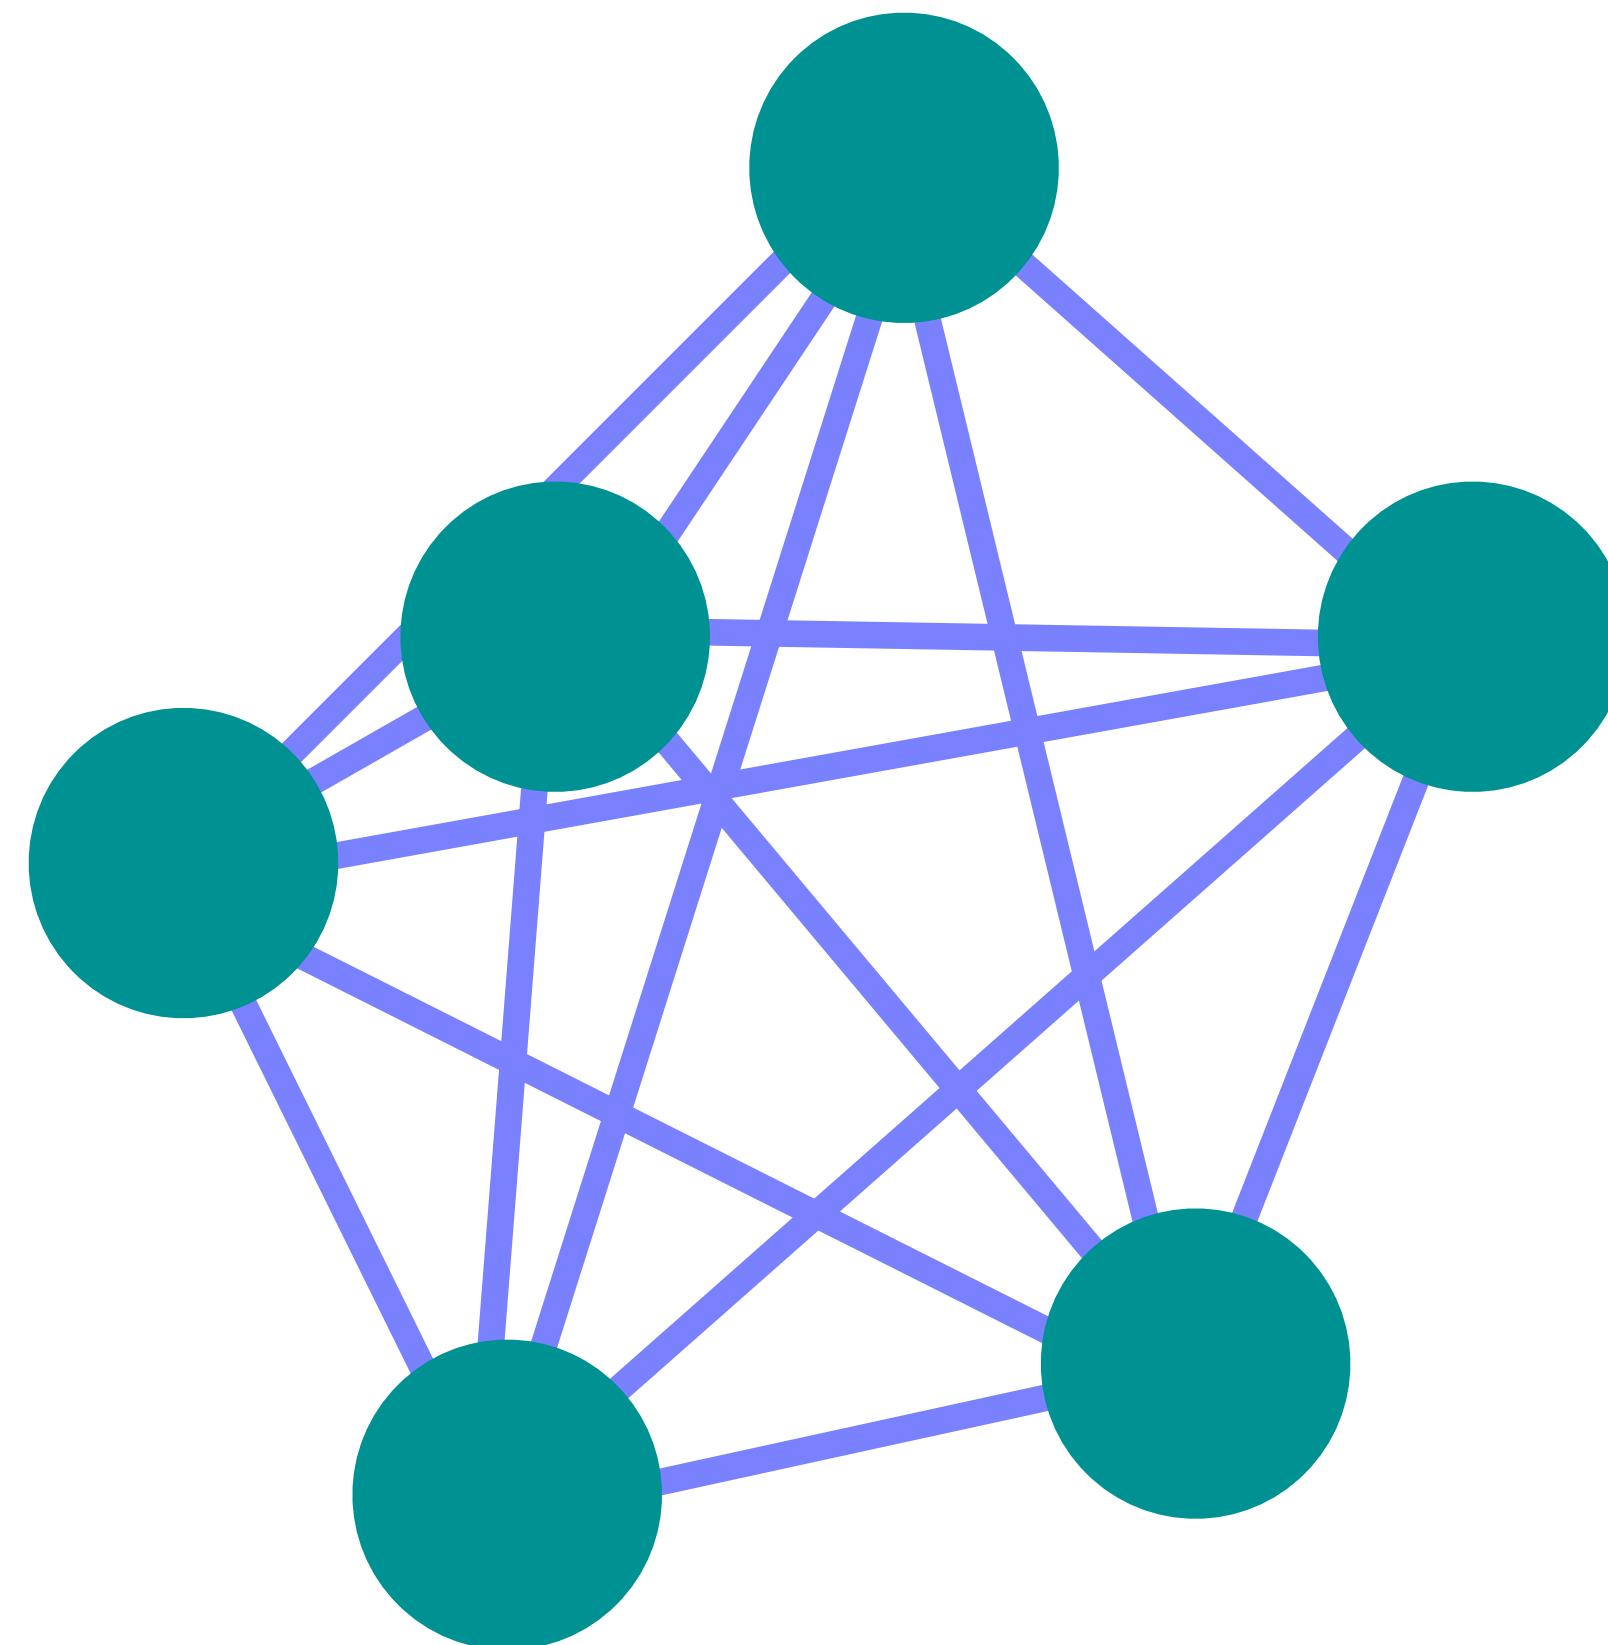
- Used in our study of the homogenous electron gas (Pescia et al. arXiv:2305.07240)
- Type of graph neural network
- Must be permutation equivariant to maintain antisymmetry
- Iteratively build correlations into new one-body and two-body features from original ones
- Nodes = one-body features
- Edges = two-body features

Message-Passing Neural Network



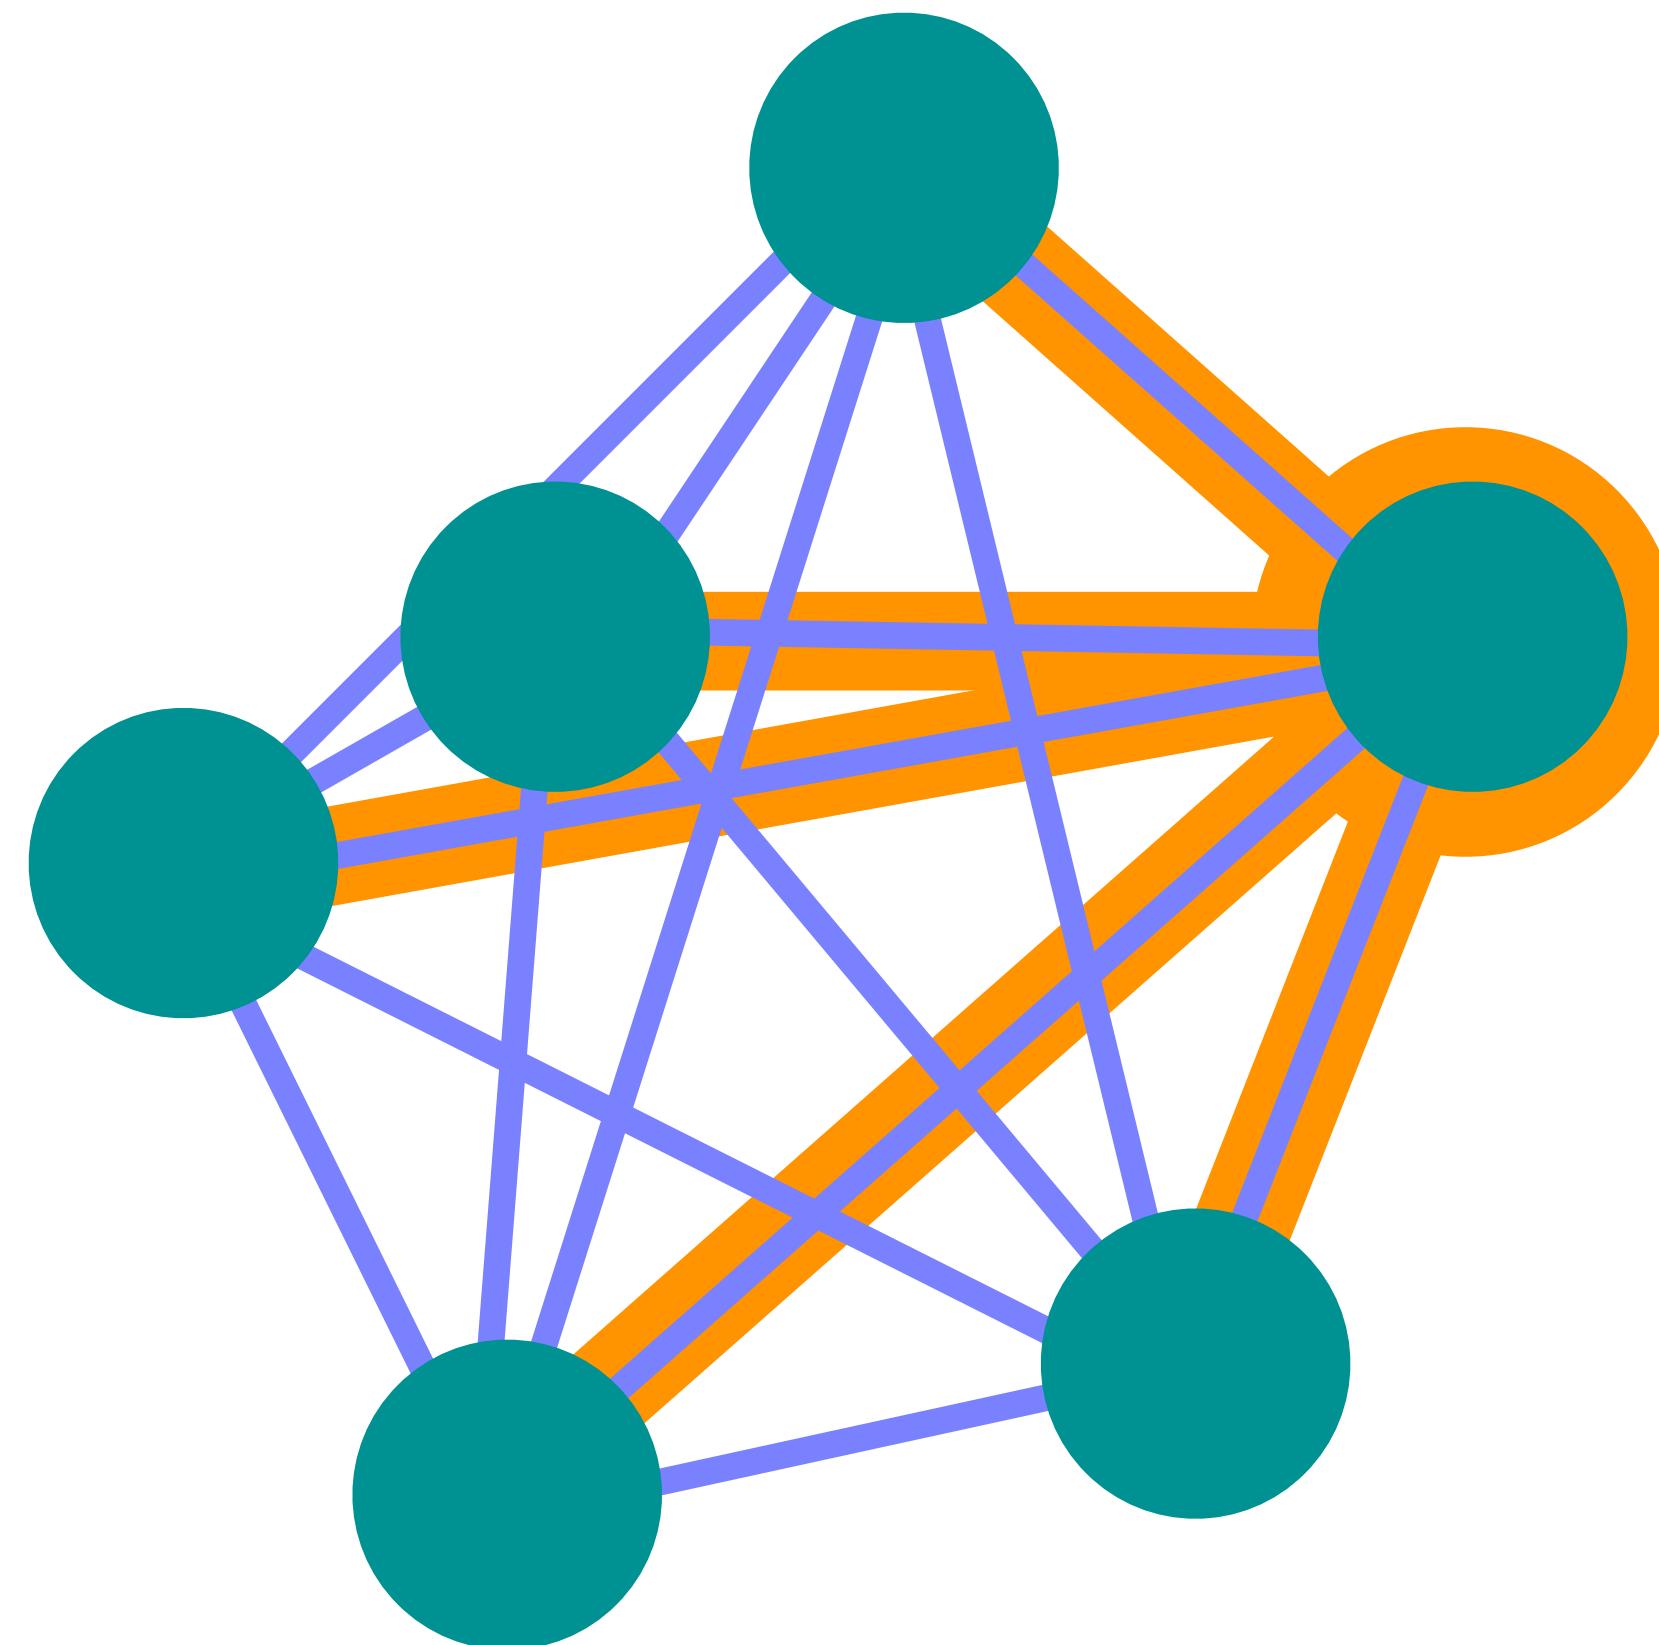
- Used in our study of the homogenous electron gas (Pescia et al. arXiv:2305.07240)
- Type of graph neural network
- Must be permutation equivariant to maintain antisymmetry
- Iteratively build correlations into new one-body and two-body features from original ones
- Nodes = one-body features
- Edges = two-body features

Message-Passing Neural Network



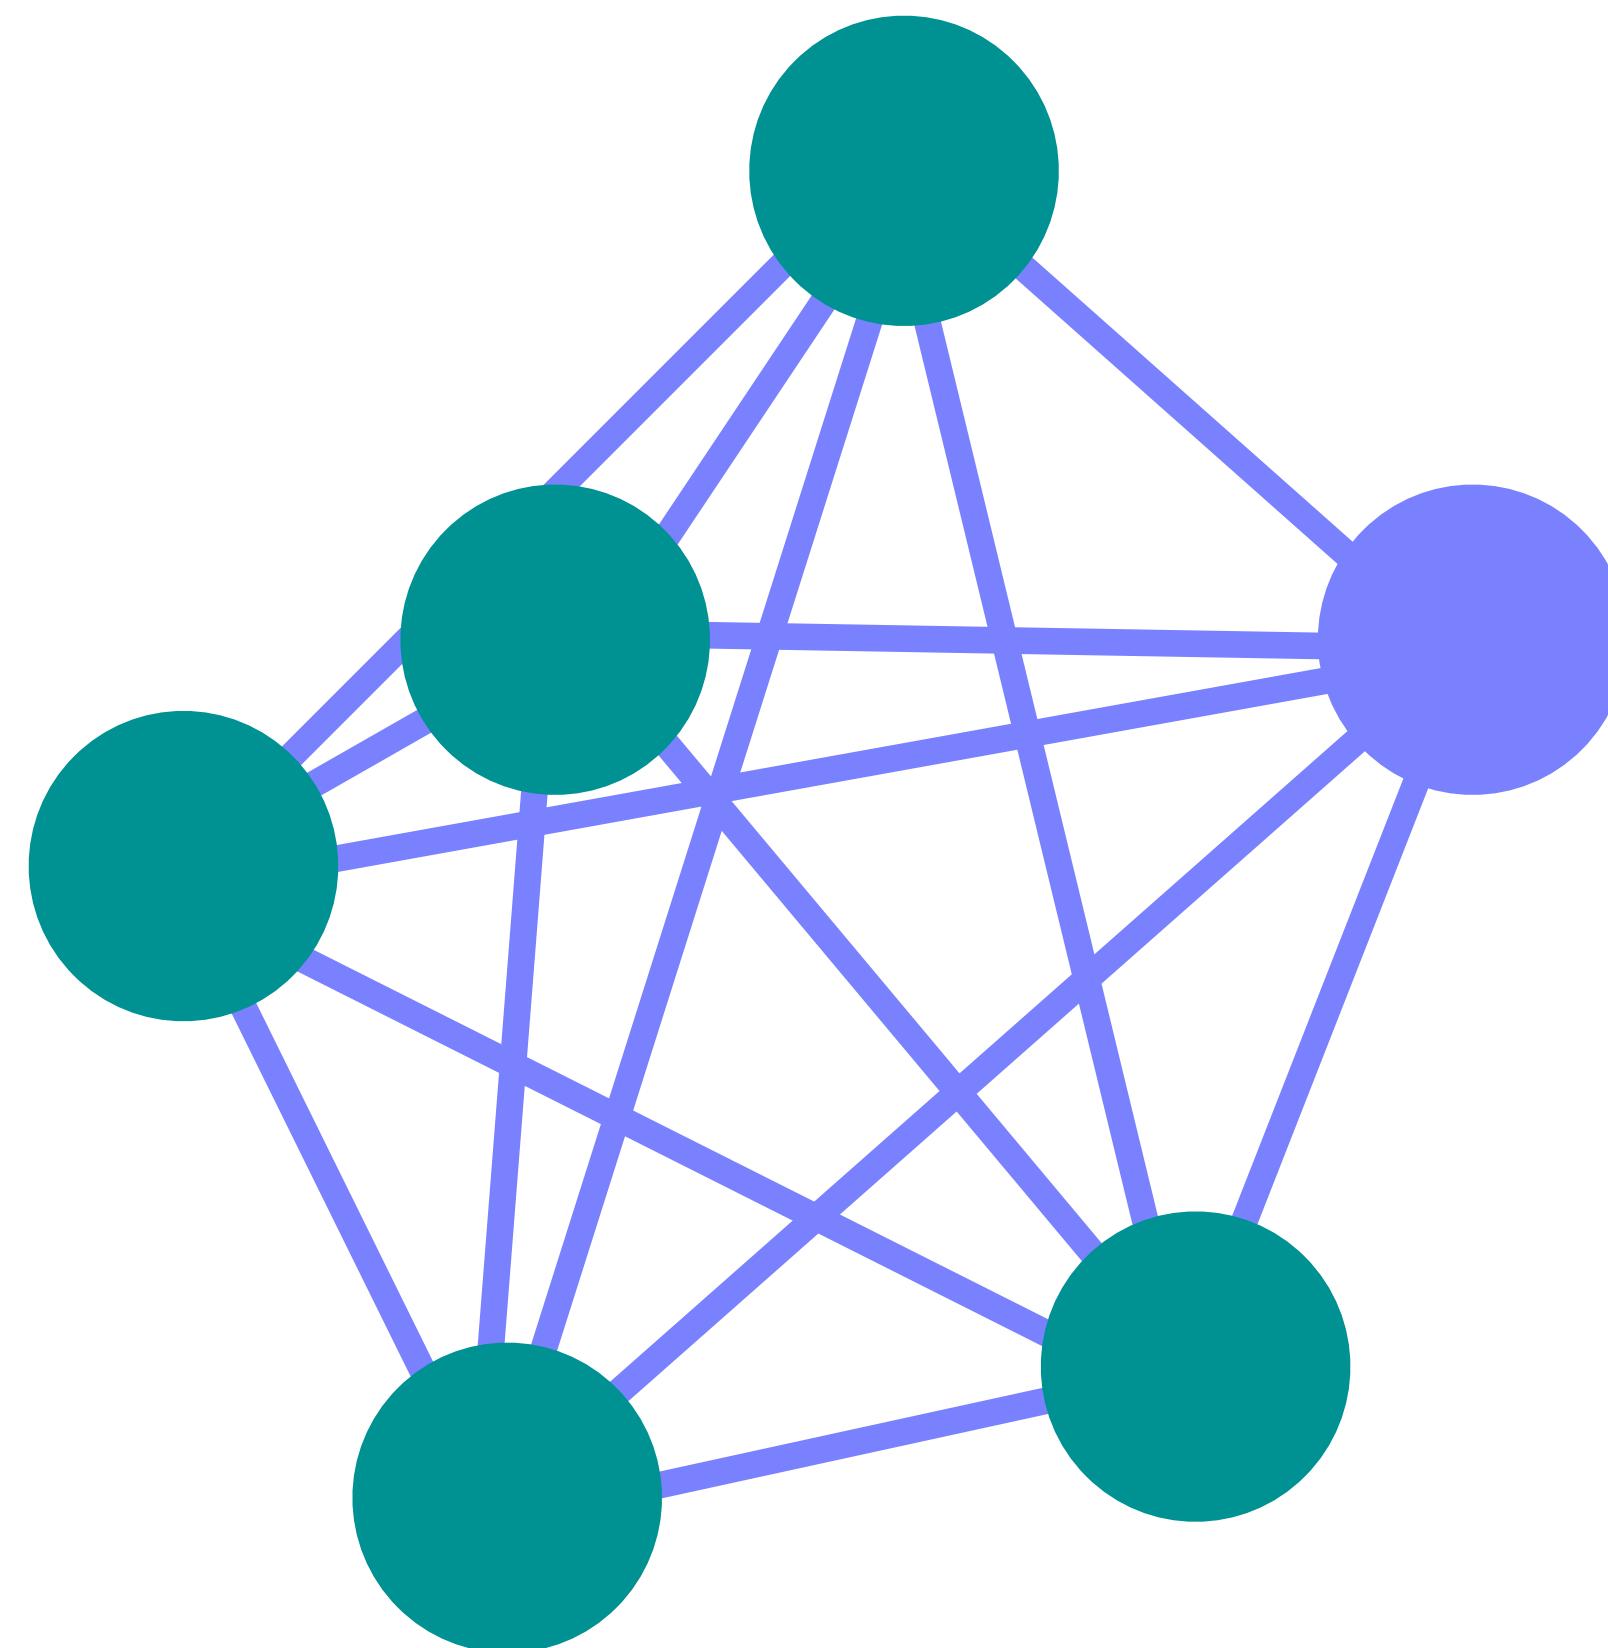
- Used in our study of the homogenous electron gas (Pescia et al. arXiv:2305.07240)
- Type of graph neural network
- Must be permutation equivariant to maintain antisymmetry
- Iteratively build correlations into new one-body and two-body features from original ones
- Nodes = one-body features
- Edges = two-body features

Message-Passing Neural Network



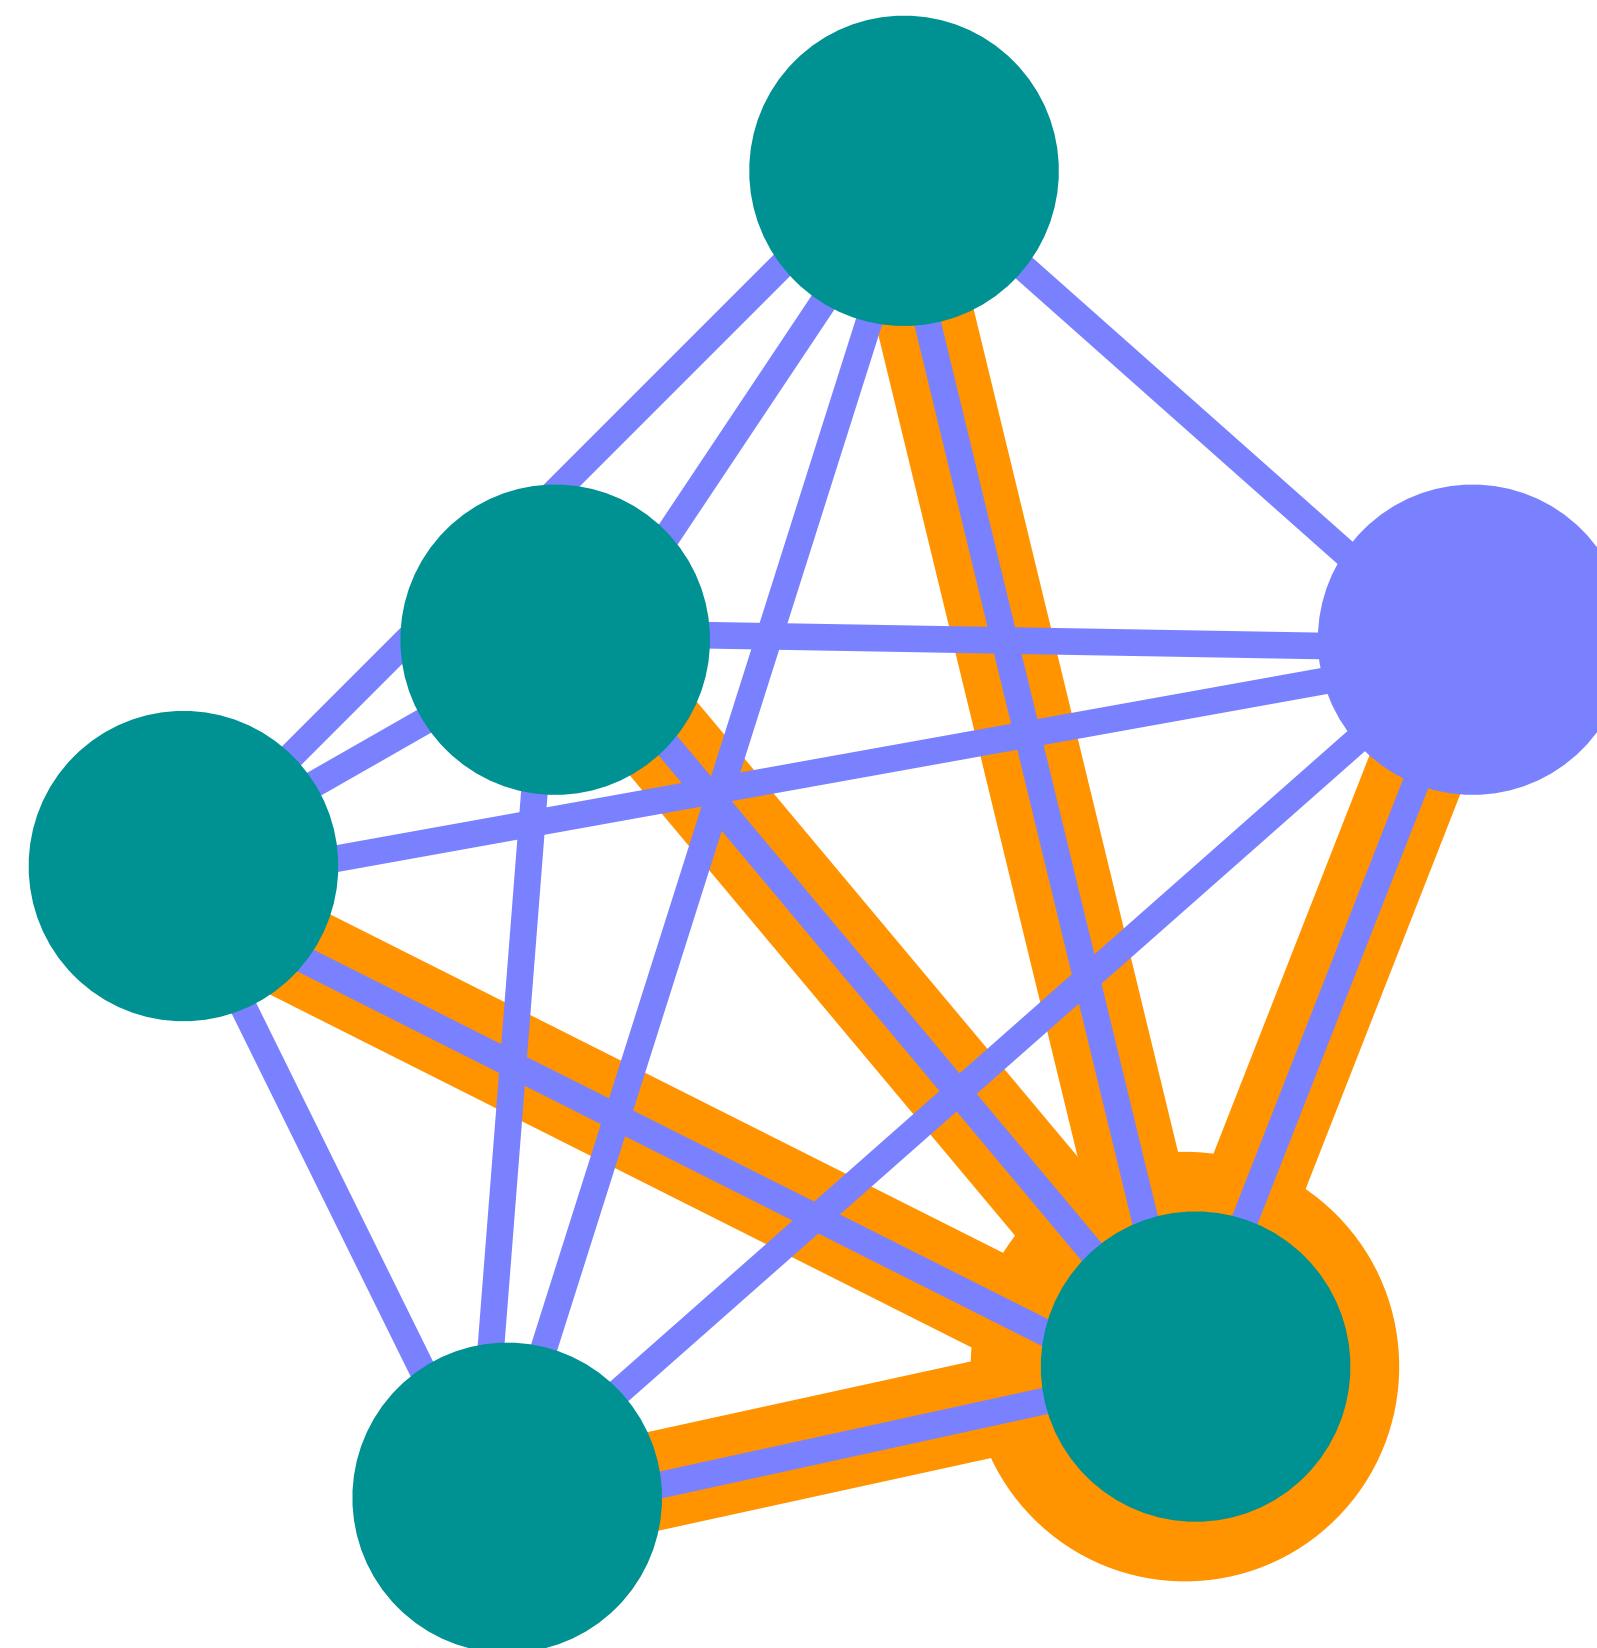
- Used in our study of the homogenous electron gas (Pescia et al. arXiv:2305.07240)
- Type of graph neural network
- Must be permutation equivariant to maintain antisymmetry
- Iteratively build correlations into new one-body and two-body features from original ones
- Nodes = one-body features
- Edges = two-body features

Message-Passing Neural Network



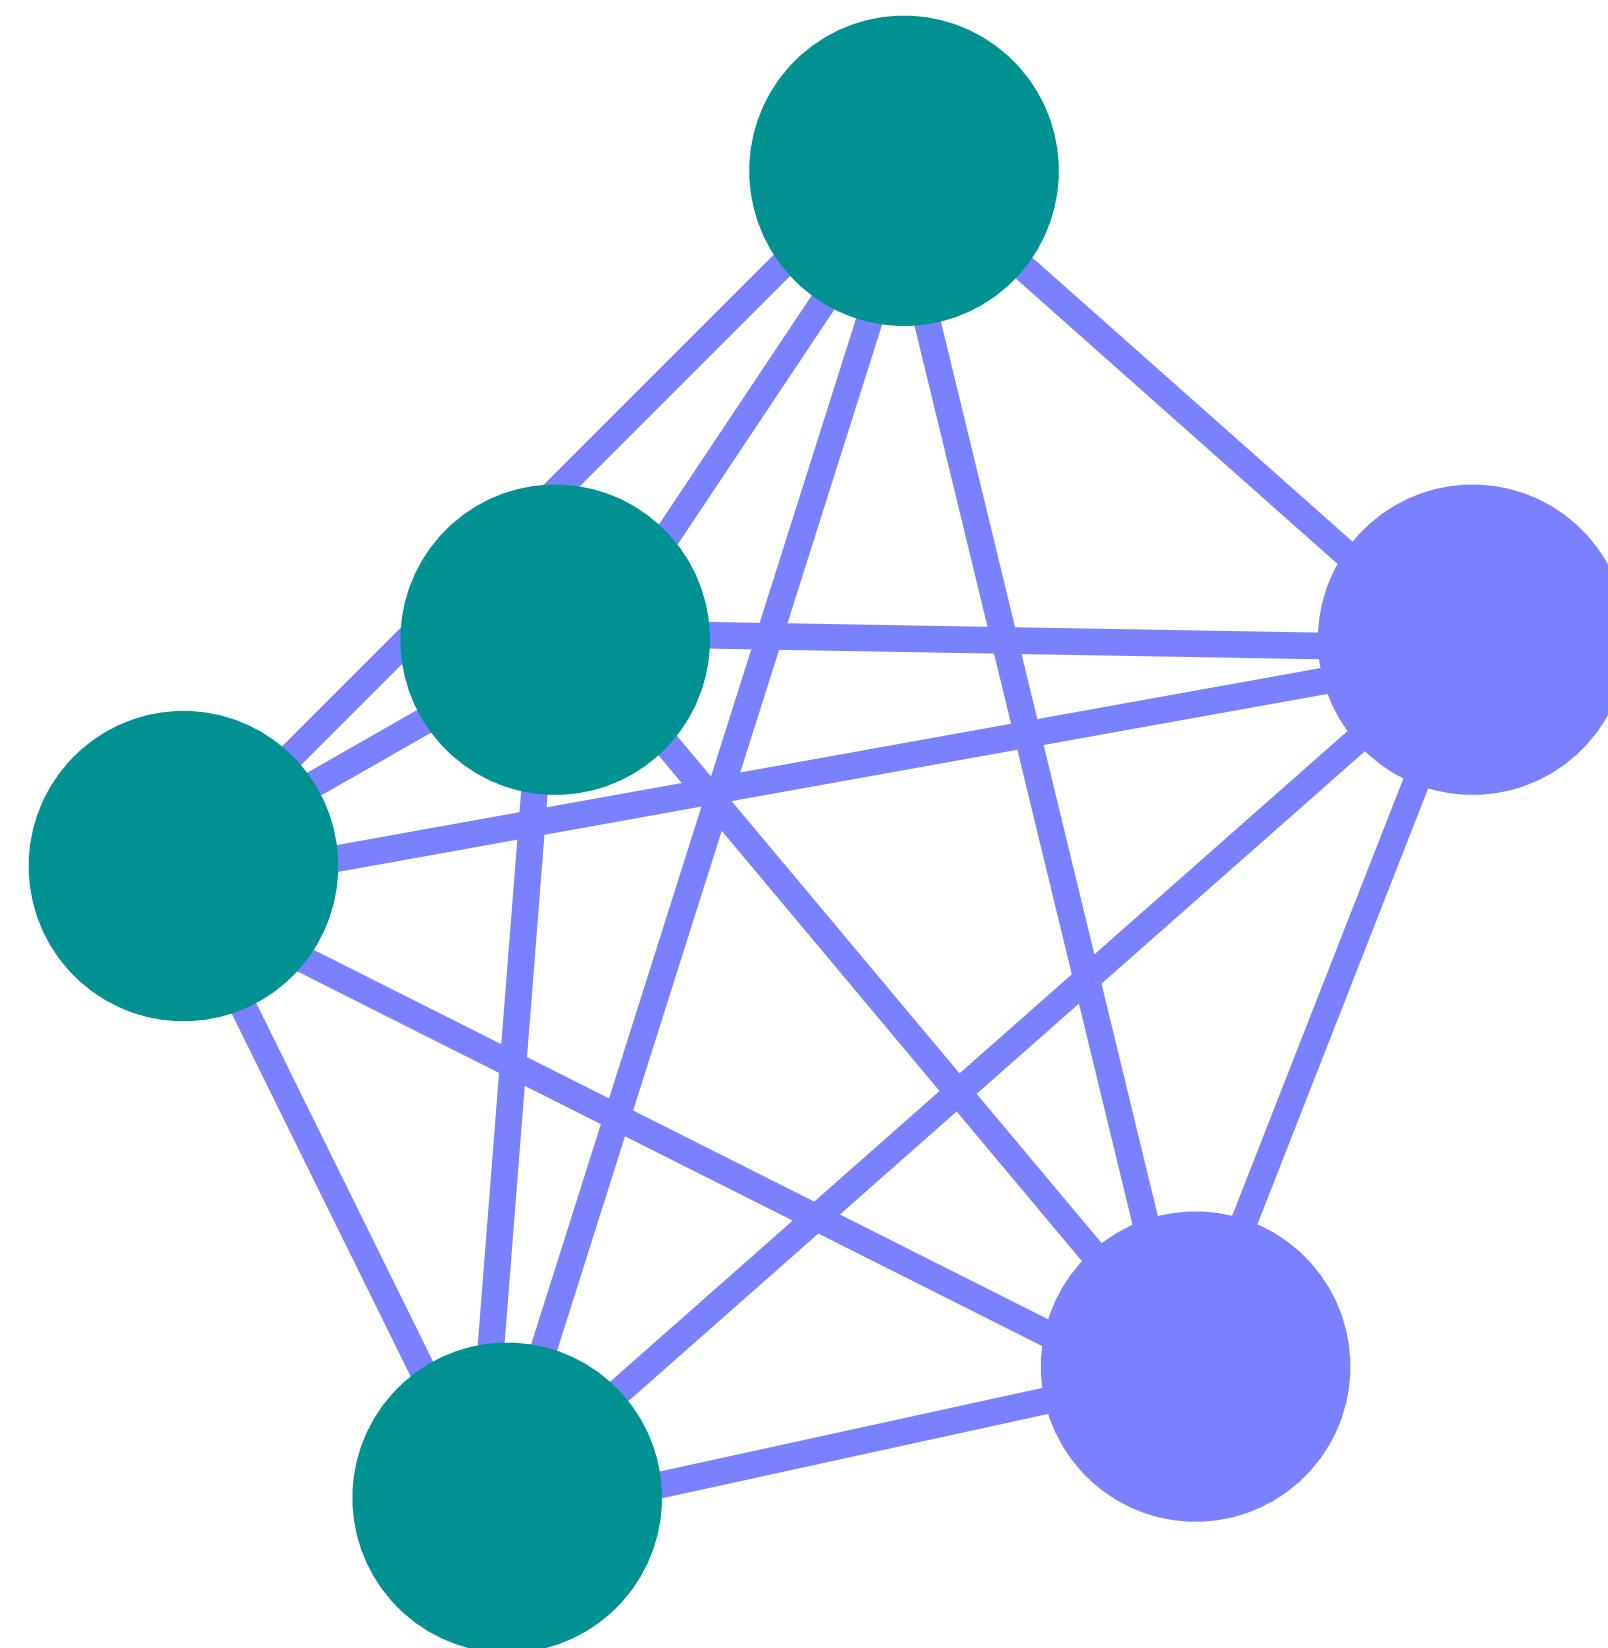
- Used in our study of the homogenous electron gas (Pescia et al. arXiv:2305.07240)
- Type of graph neural network
- Must be permutation equivariant to maintain antisymmetry
- Iteratively build correlations into new one-body and two-body features from original ones
- Nodes = one-body features
- Edges = two-body features

Message-Passing Neural Network



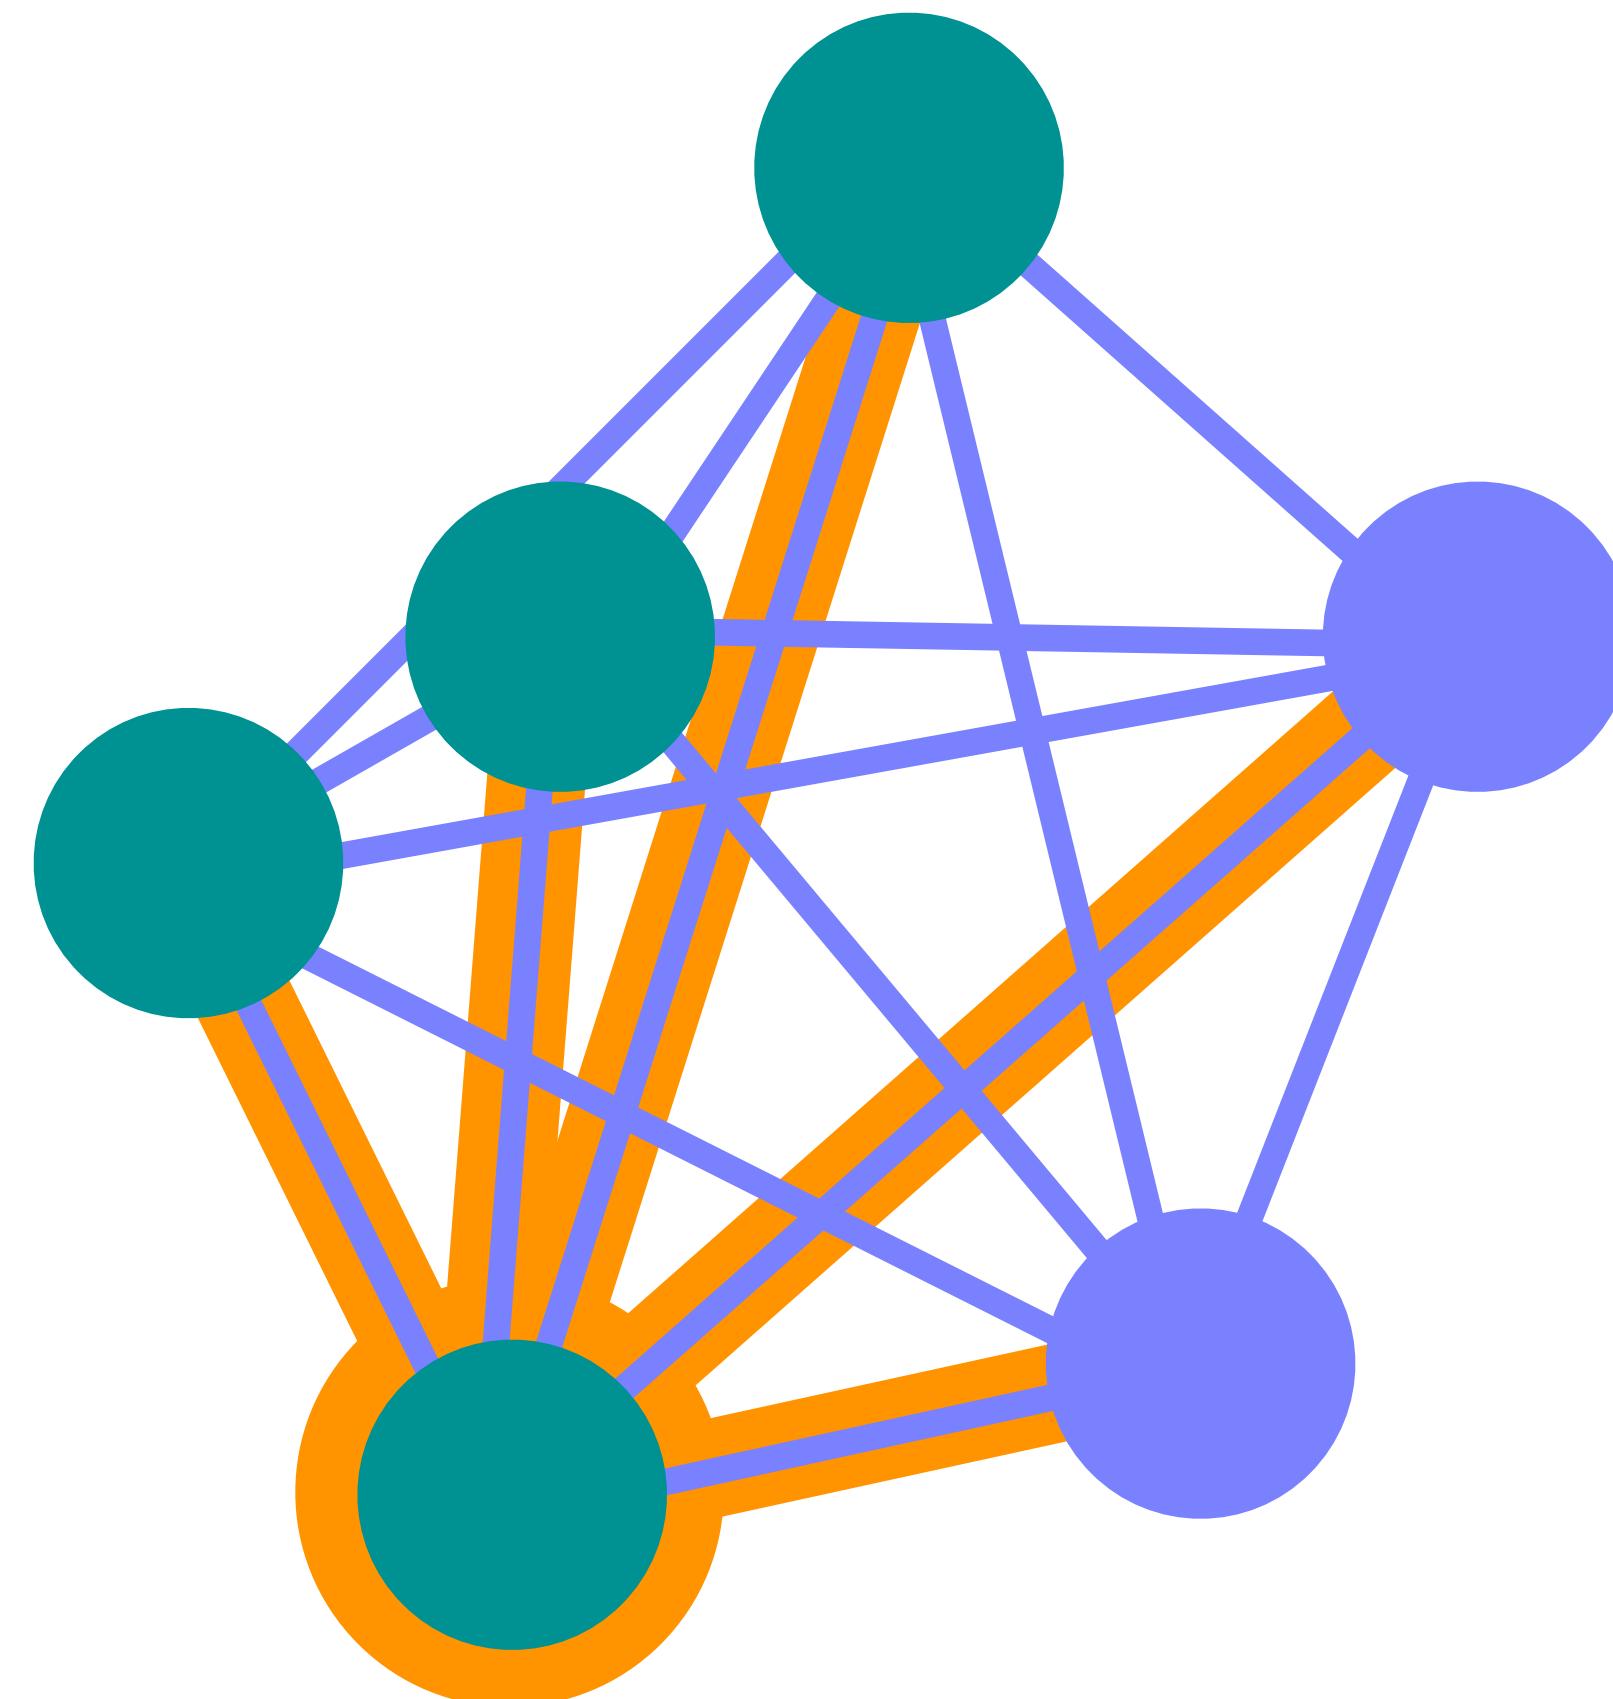
- Used in our study of the homogenous electron gas (Pescia et al. arXiv:2305.07240)
- Type of graph neural network
- Must be permutation equivariant to maintain antisymmetry
- Iteratively build correlations into new one-body and two-body features from original ones
- Nodes = one-body features
- Edges = two-body features

Message-Passing Neural Network



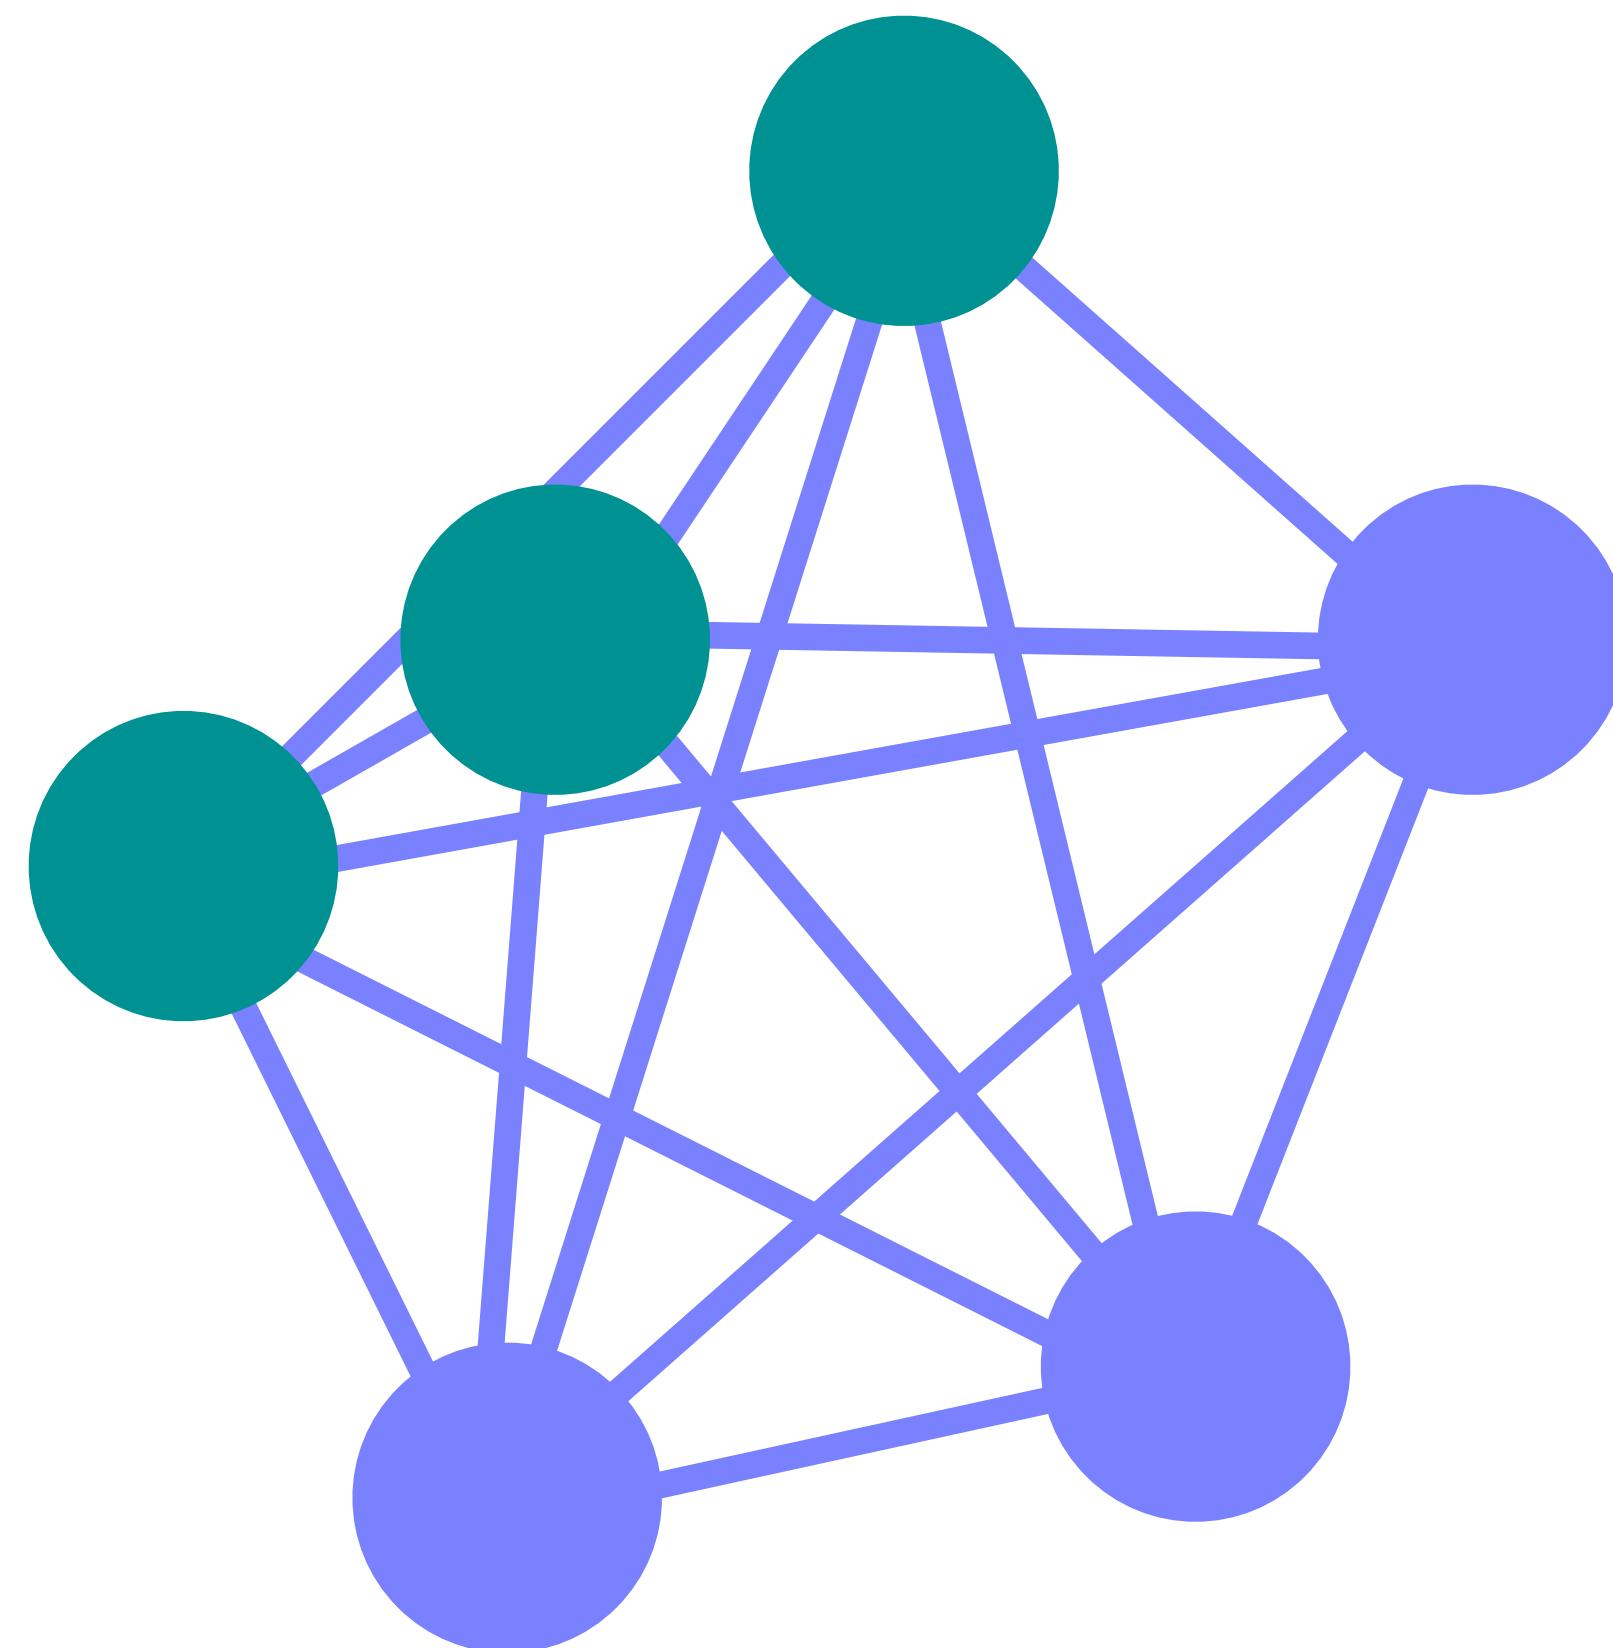
- Used in our study of the homogenous electron gas (Pescia et al. arXiv:2305.07240)
- Type of graph neural network
- Must be permutation equivariant to maintain antisymmetry
- Iteratively build correlations into new one-body and two-body features from original ones
- Nodes = one-body features
- Edges = two-body features

Message-Passing Neural Network



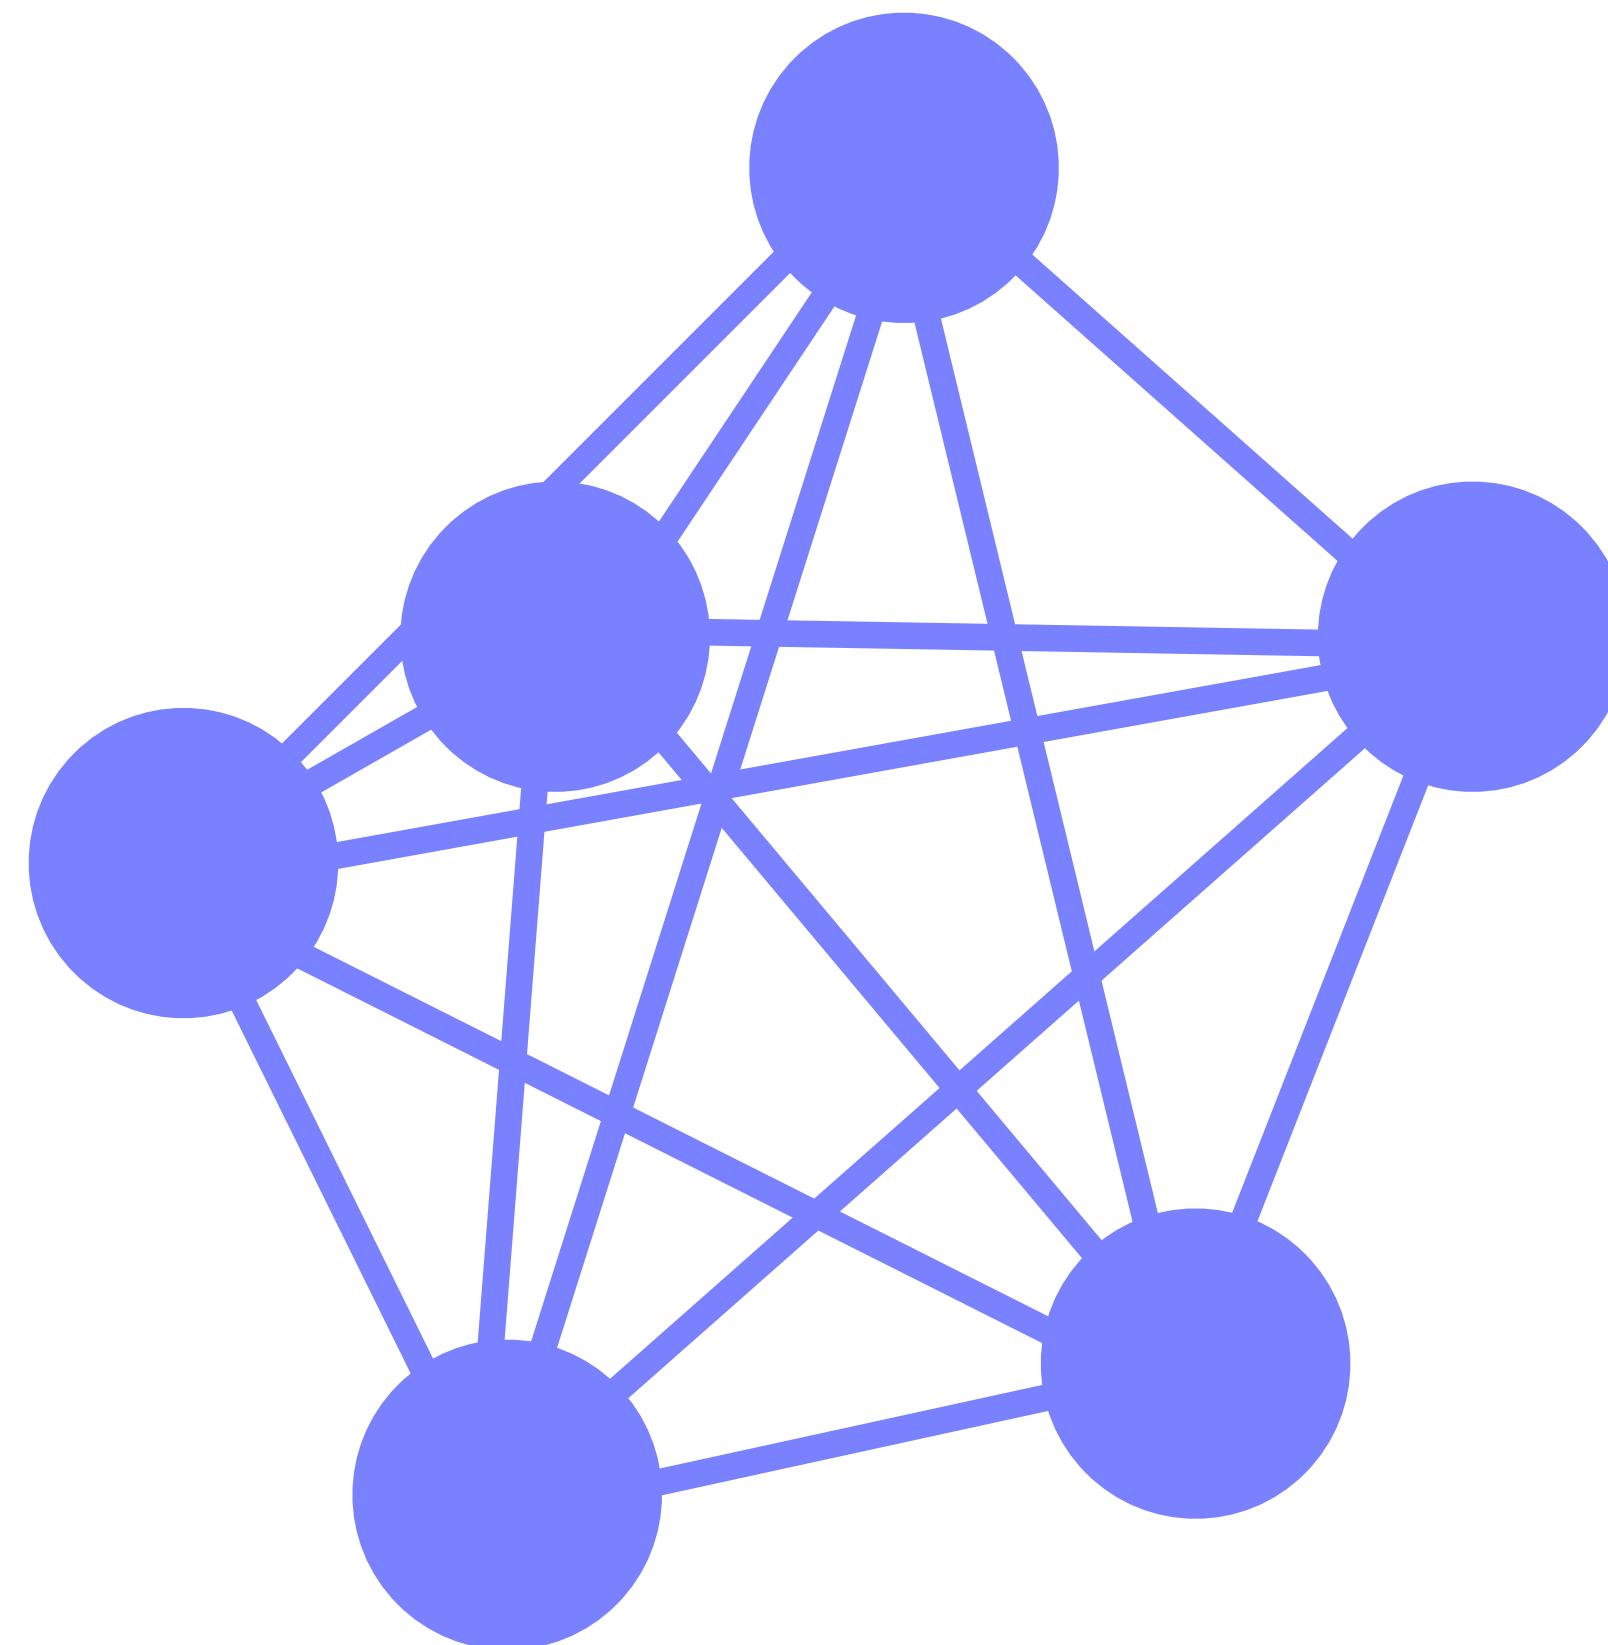
- Used in our study of the homogenous electron gas (Pescia et al. arXiv:2305.07240)
- Type of graph neural network
- Must be permutation equivariant to maintain antisymmetry
- Iteratively build correlations into new one-body and two-body features from original ones
- Nodes = one-body features
- Edges = two-body features

Message-Passing Neural Network

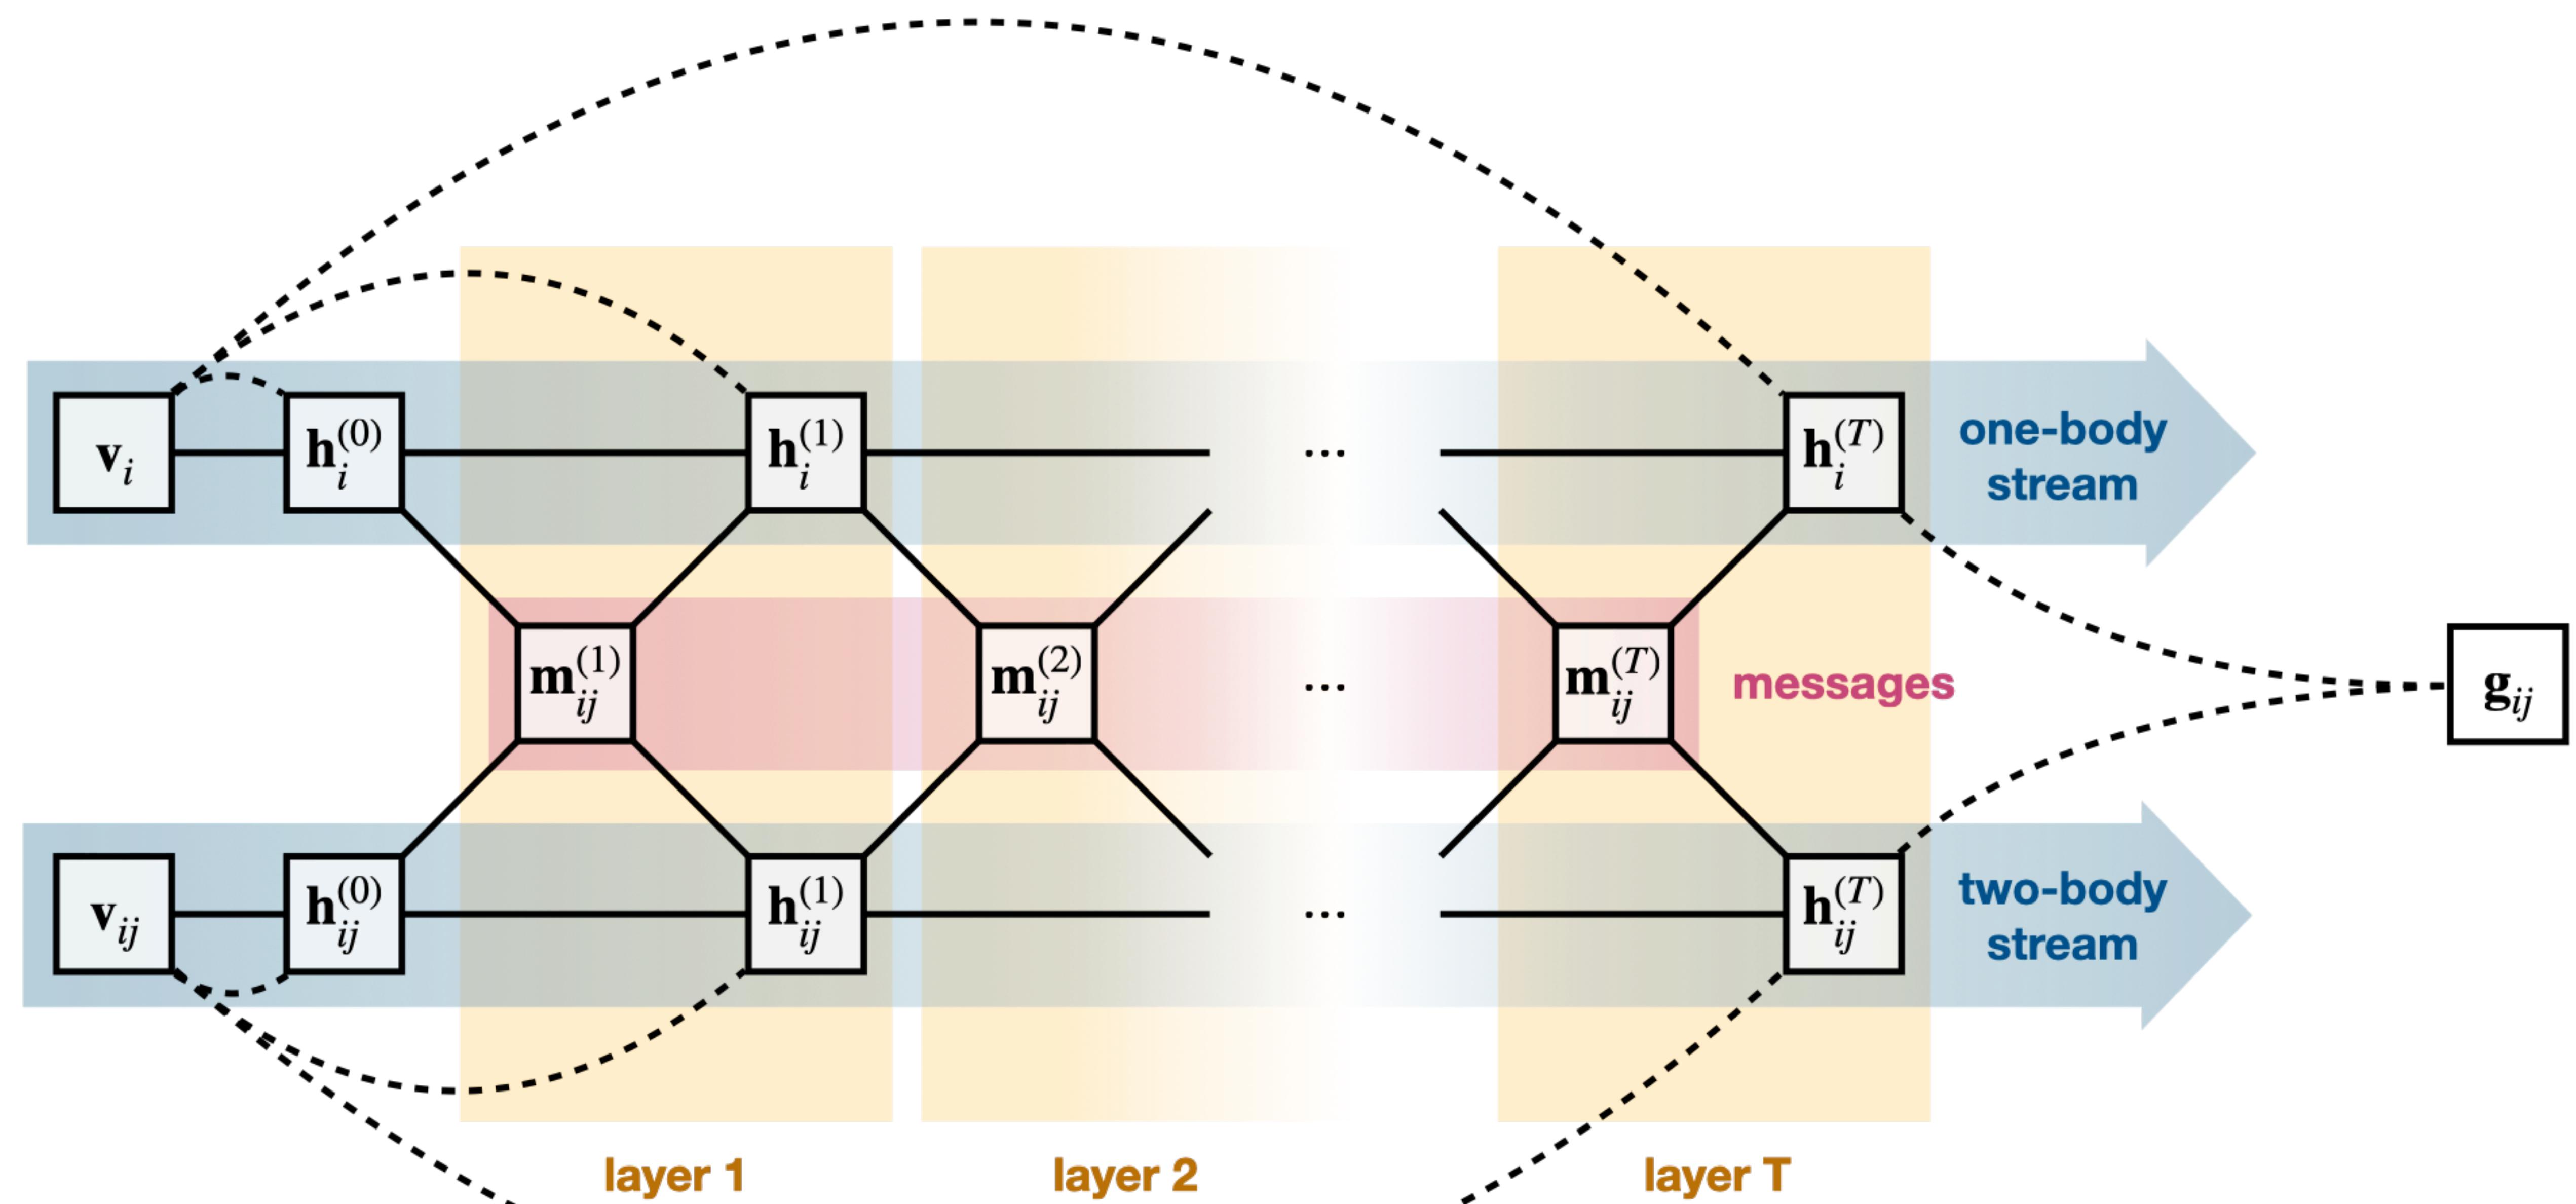


- Used in our study of the homogenous electron gas (Pescia et al. arXiv:2305.07240)
- Type of graph neural network
- Must be permutation equivariant to maintain antisymmetry
- Iteratively build correlations into new one-body and two-body features from original ones
- Nodes = one-body features
- Edges = two-body features

Message-Passing Neural Network



- Used in our study of the homogenous electron gas (Pescia et al. arXiv:2305.07240)
- Type of graph neural network
- Must be permutation equivariant to maintain antisymmetry
- Iteratively build correlations into new one-body and two-body features from original ones
- Nodes = one-body features
- Edges = two-body features



Pfaffian Wave Function with Backflow

- Use output of MPNN as input to pairing orbital instead of raw coordinates

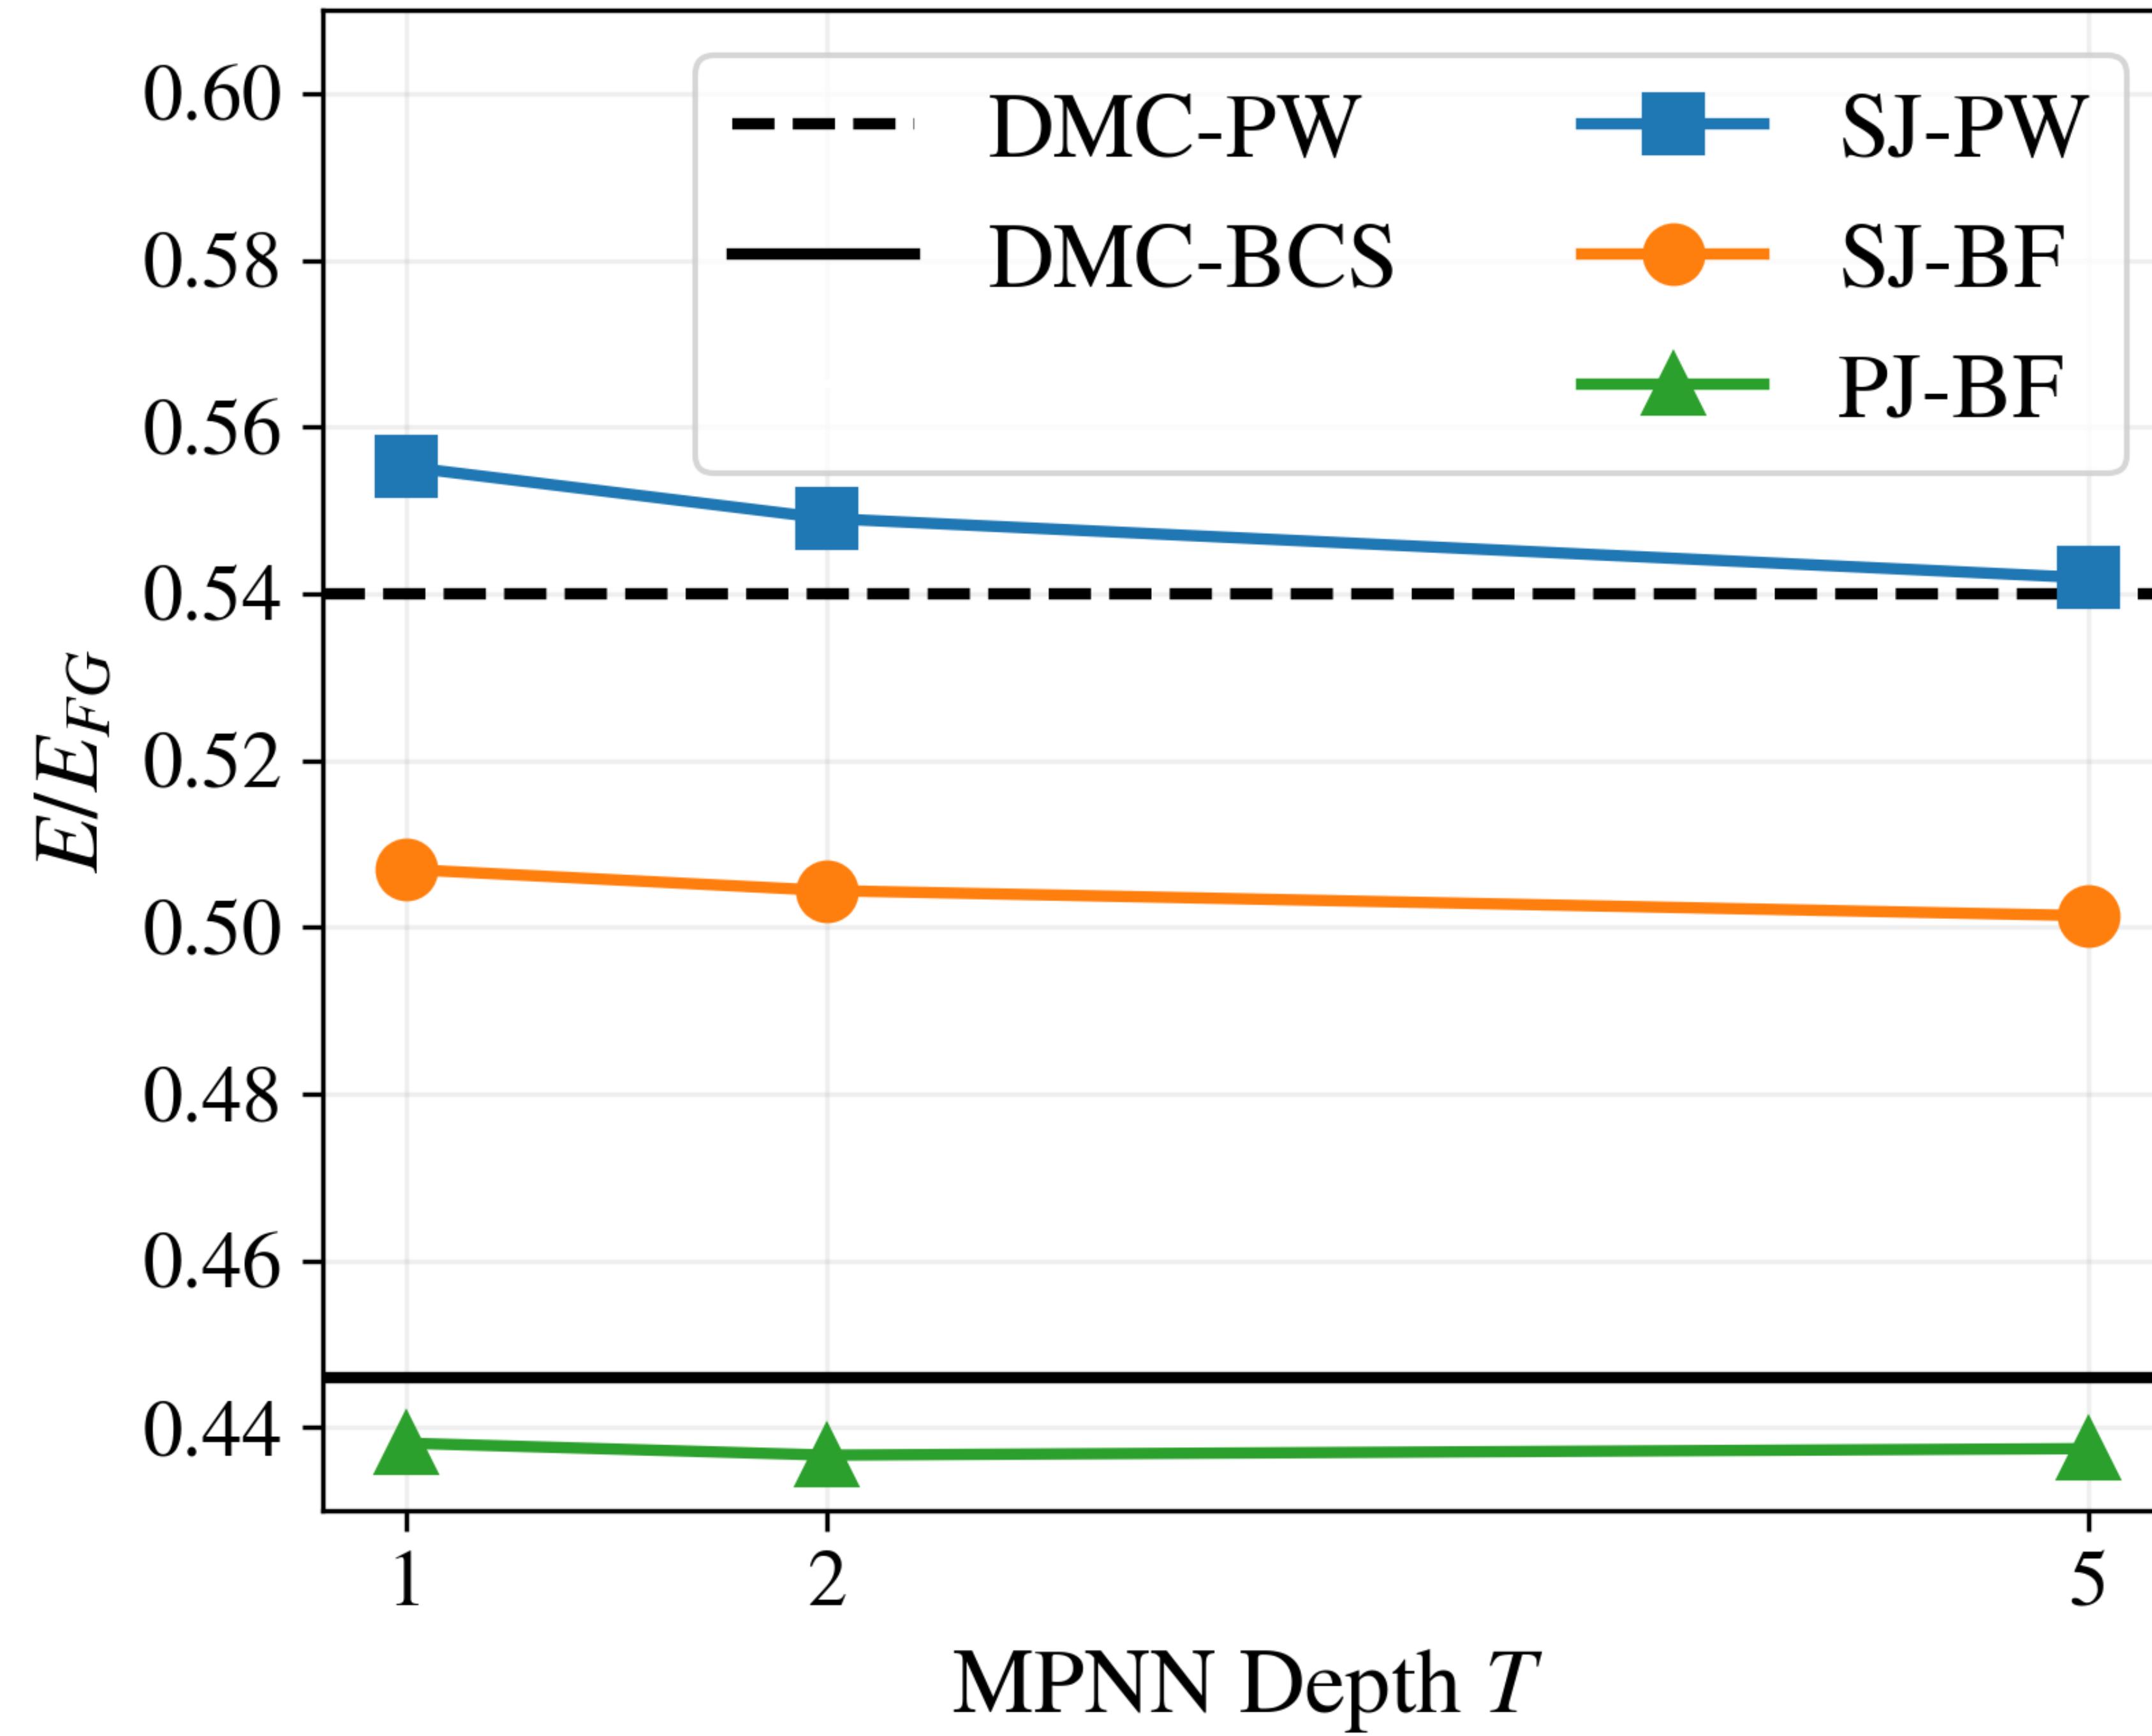
$$\Phi(X) = \text{pf}[\phi(\mathbf{g}_{ij})]$$

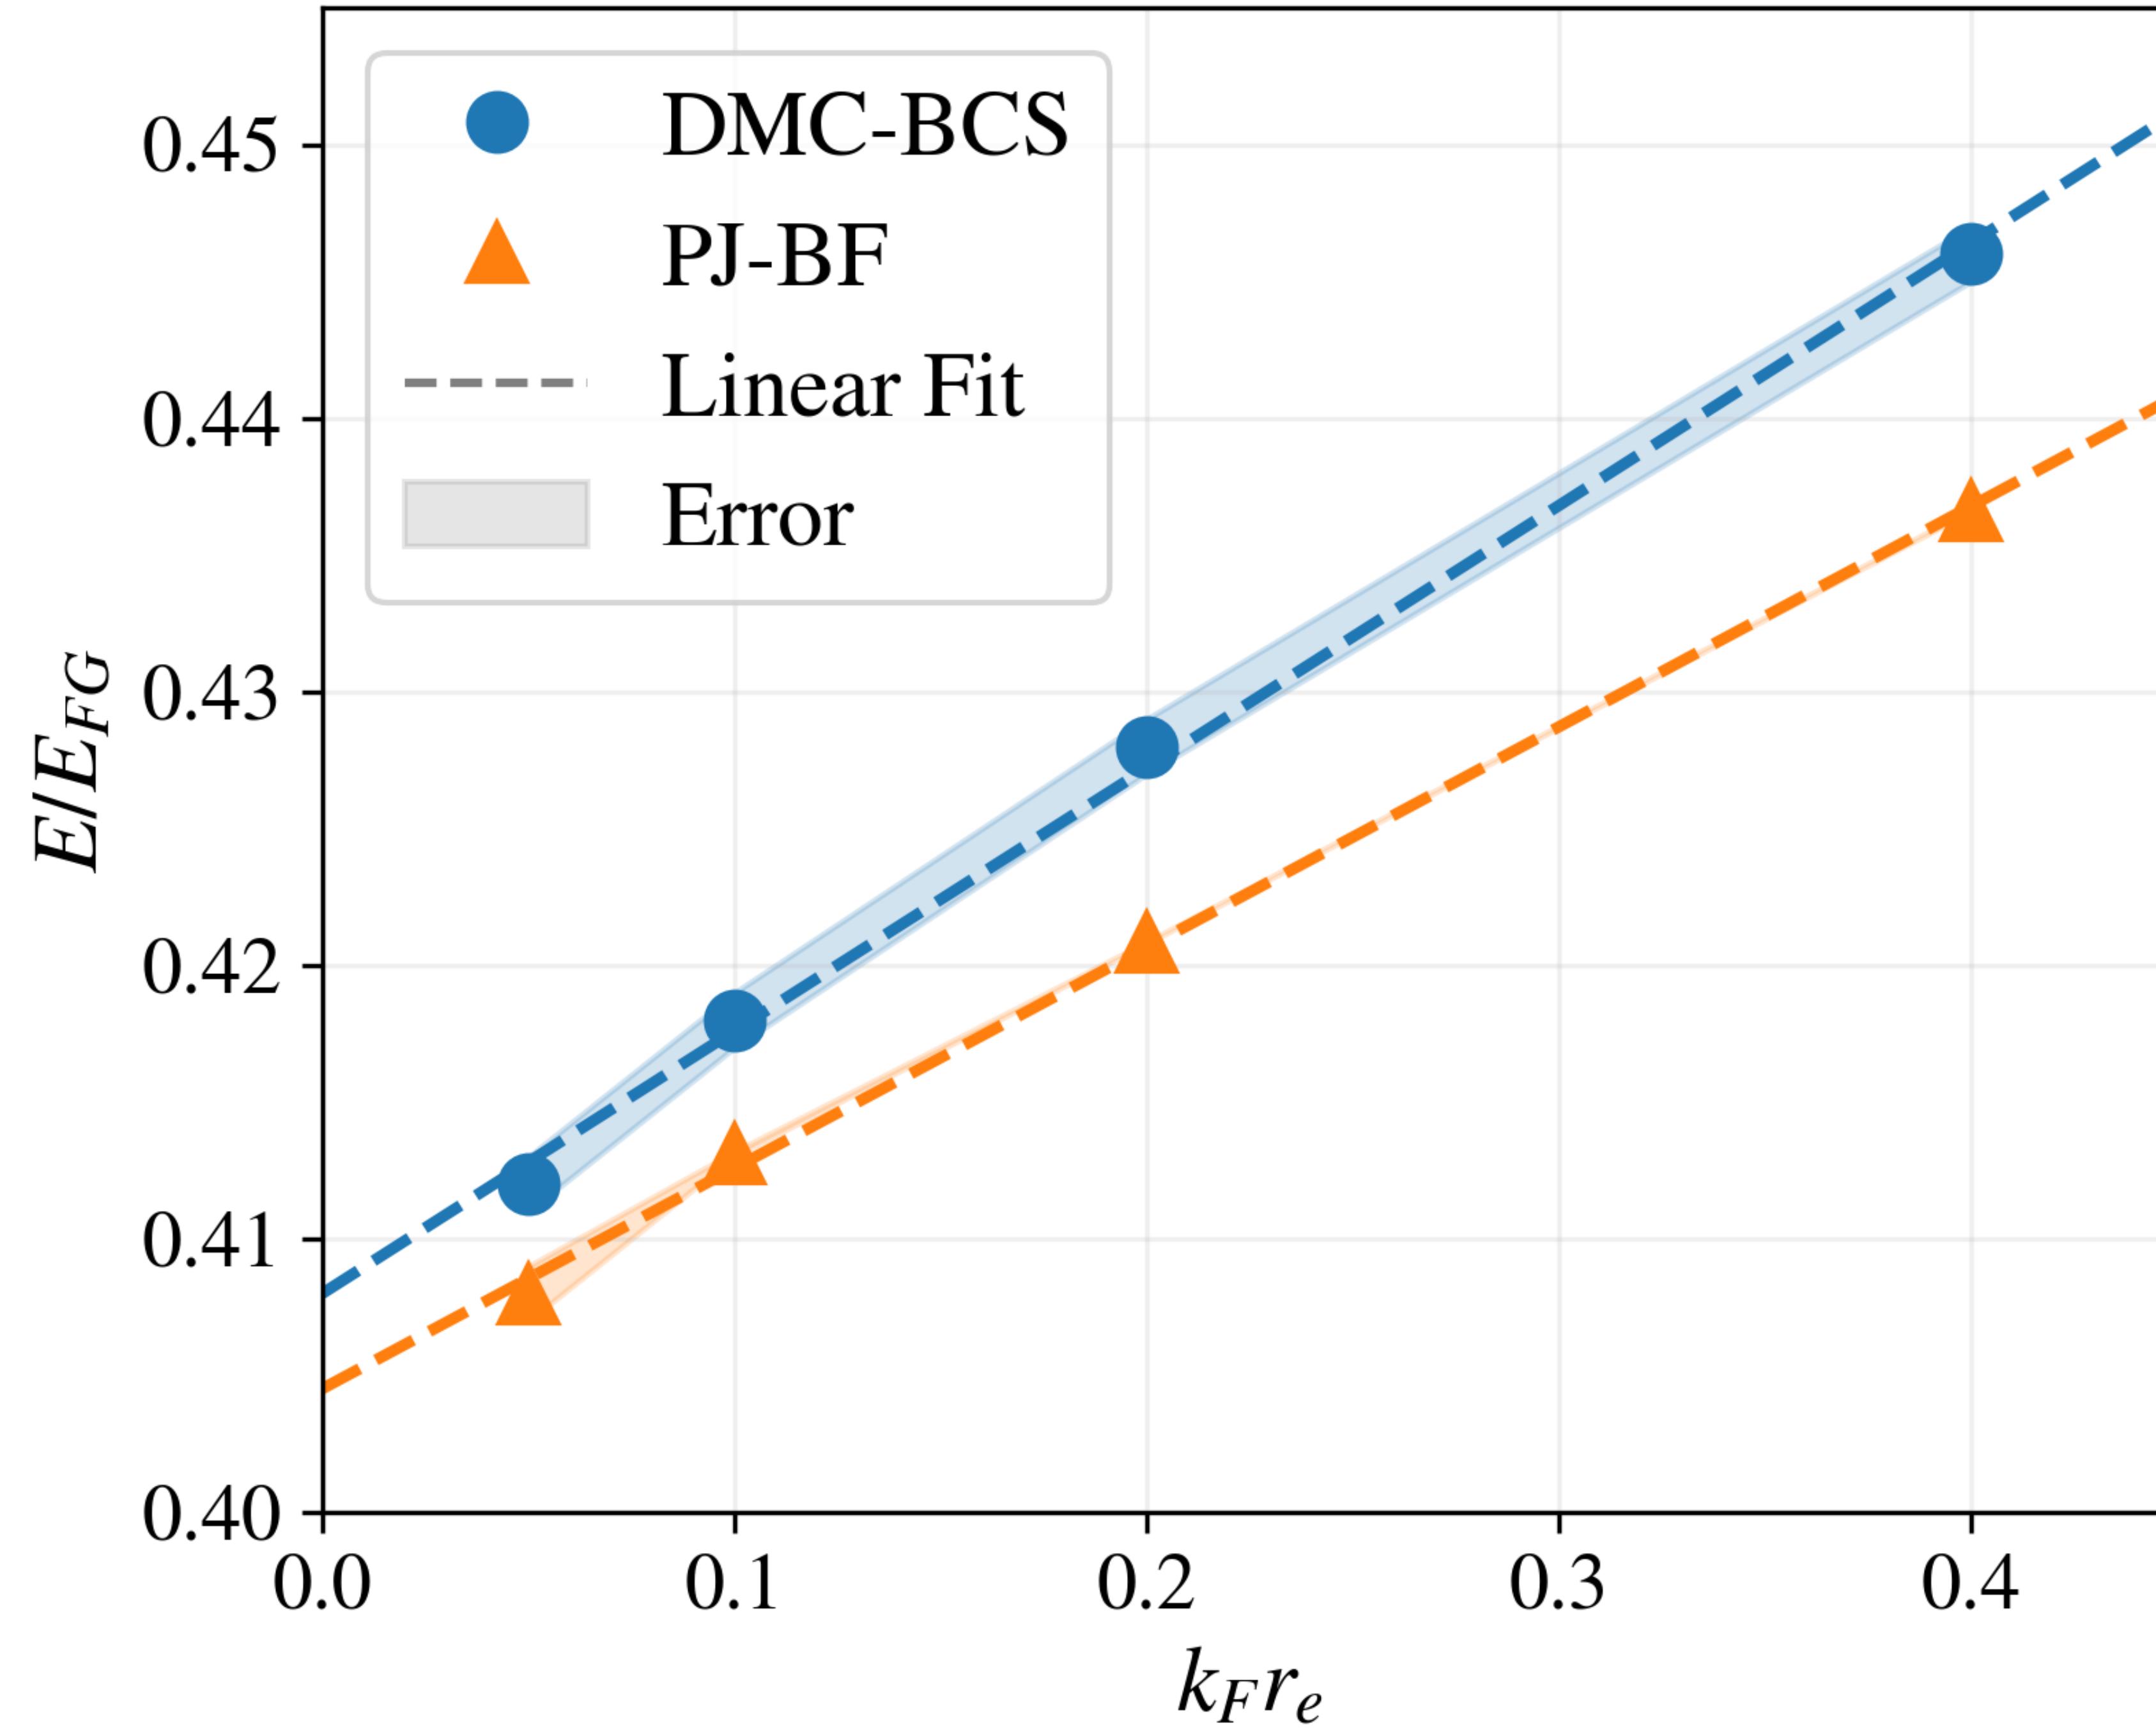
- Jastrow correlator based on a “Deep Set” (Zaheer et al. - arXiv:1703.06114)

$$J(X) = \rho \left(\sum_{i \neq j} \zeta(\mathbf{g}_{ij}) \right)$$

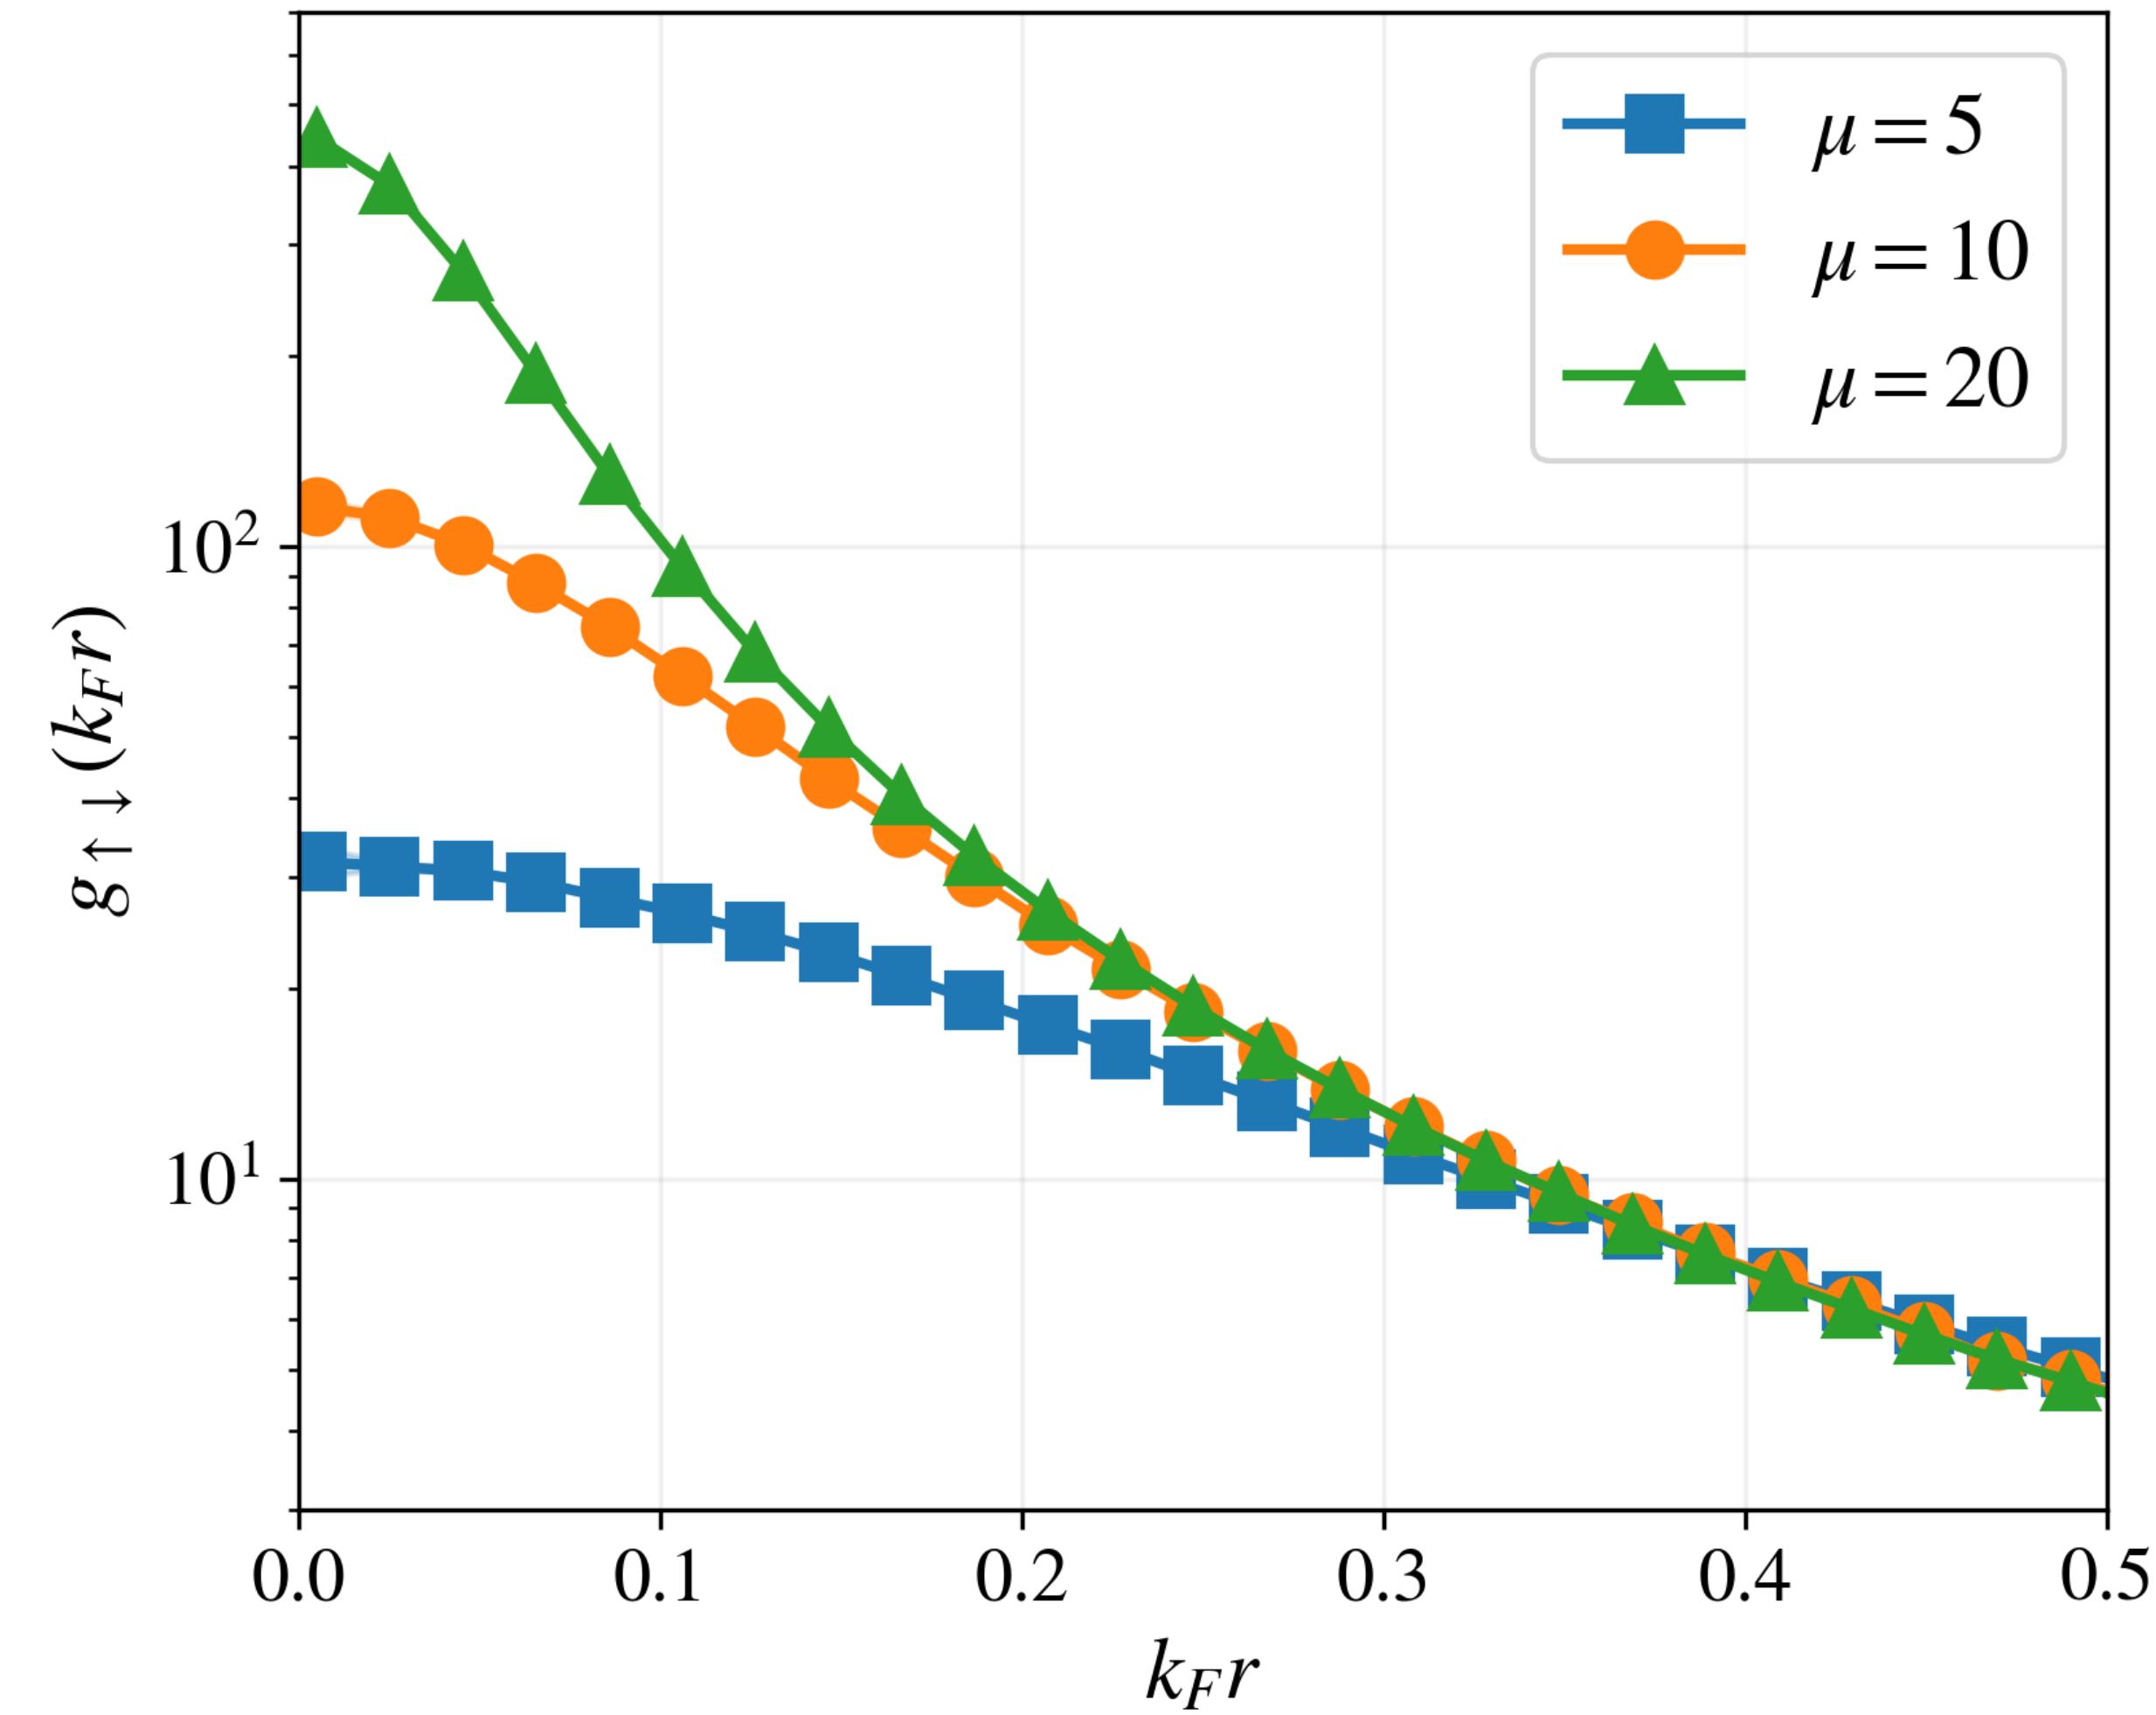
where ρ and ζ are neural networks.

- Most results shown use an MPNN with 2 iterations (~8500 parameters total)
- We also enforce periodicity, translational invariance, parity and time-reversal symmetry
- This is the first time neural backflow transformations have been applied to the Pfaffian :)

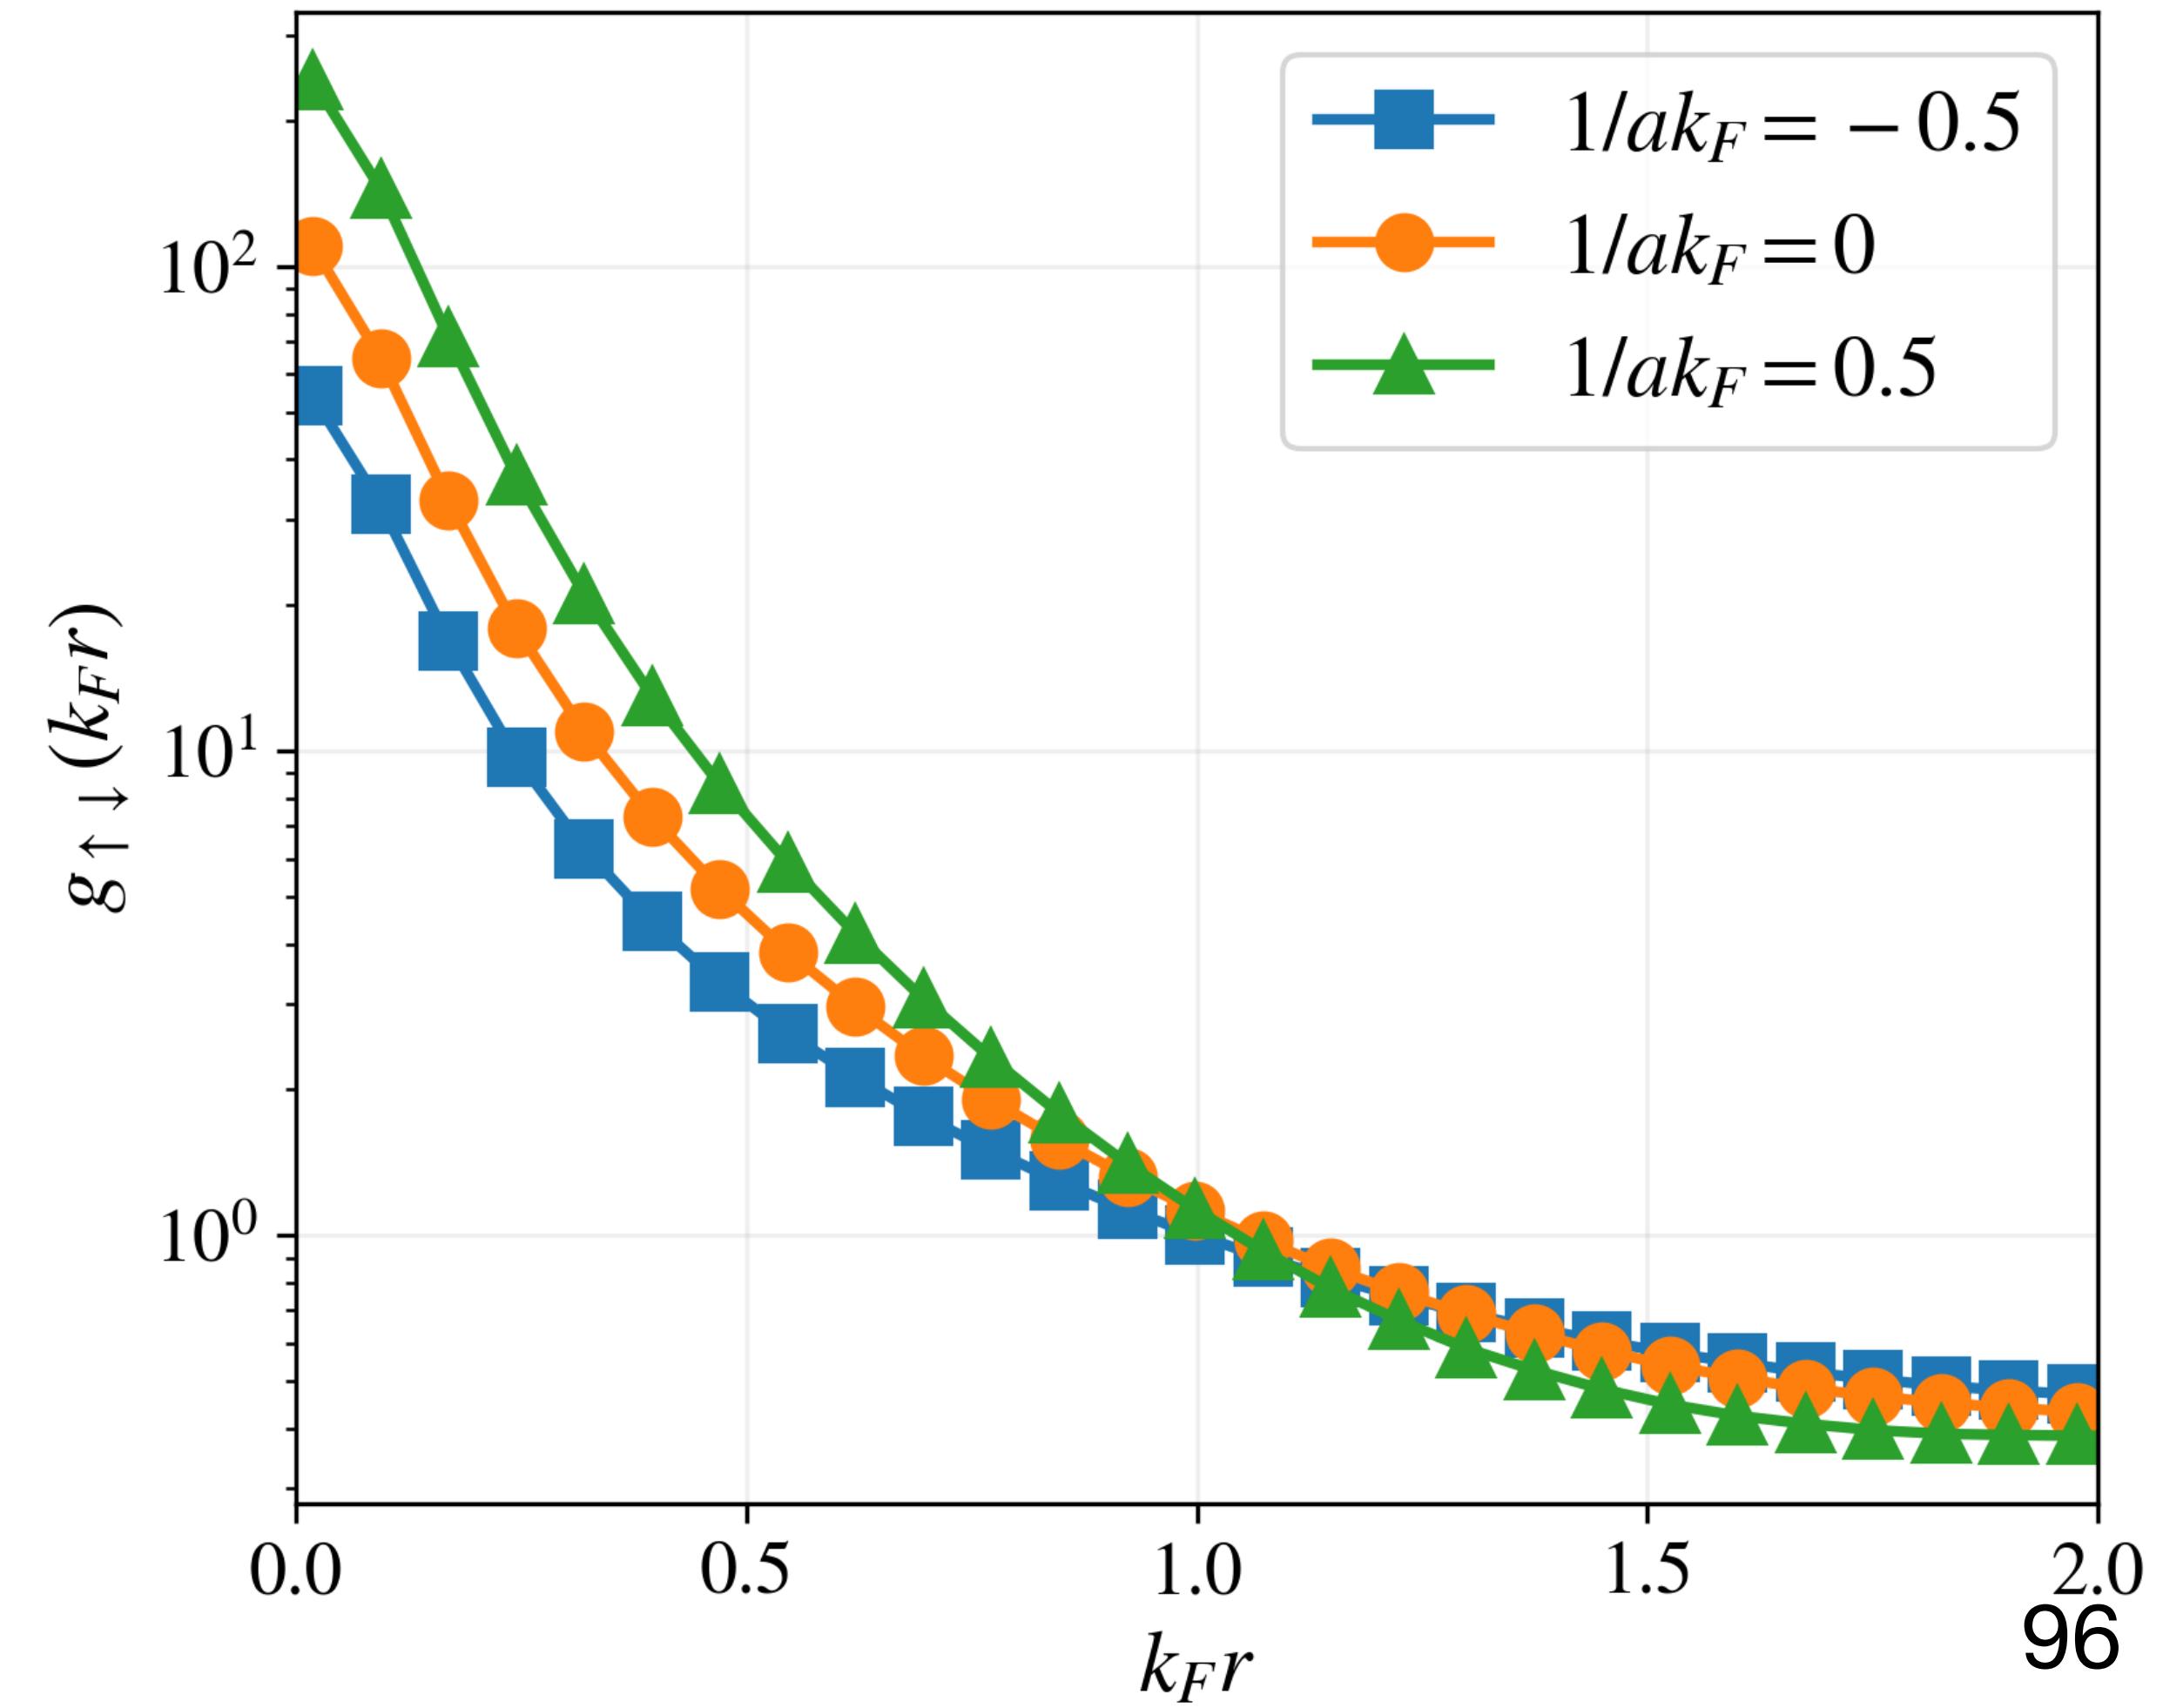


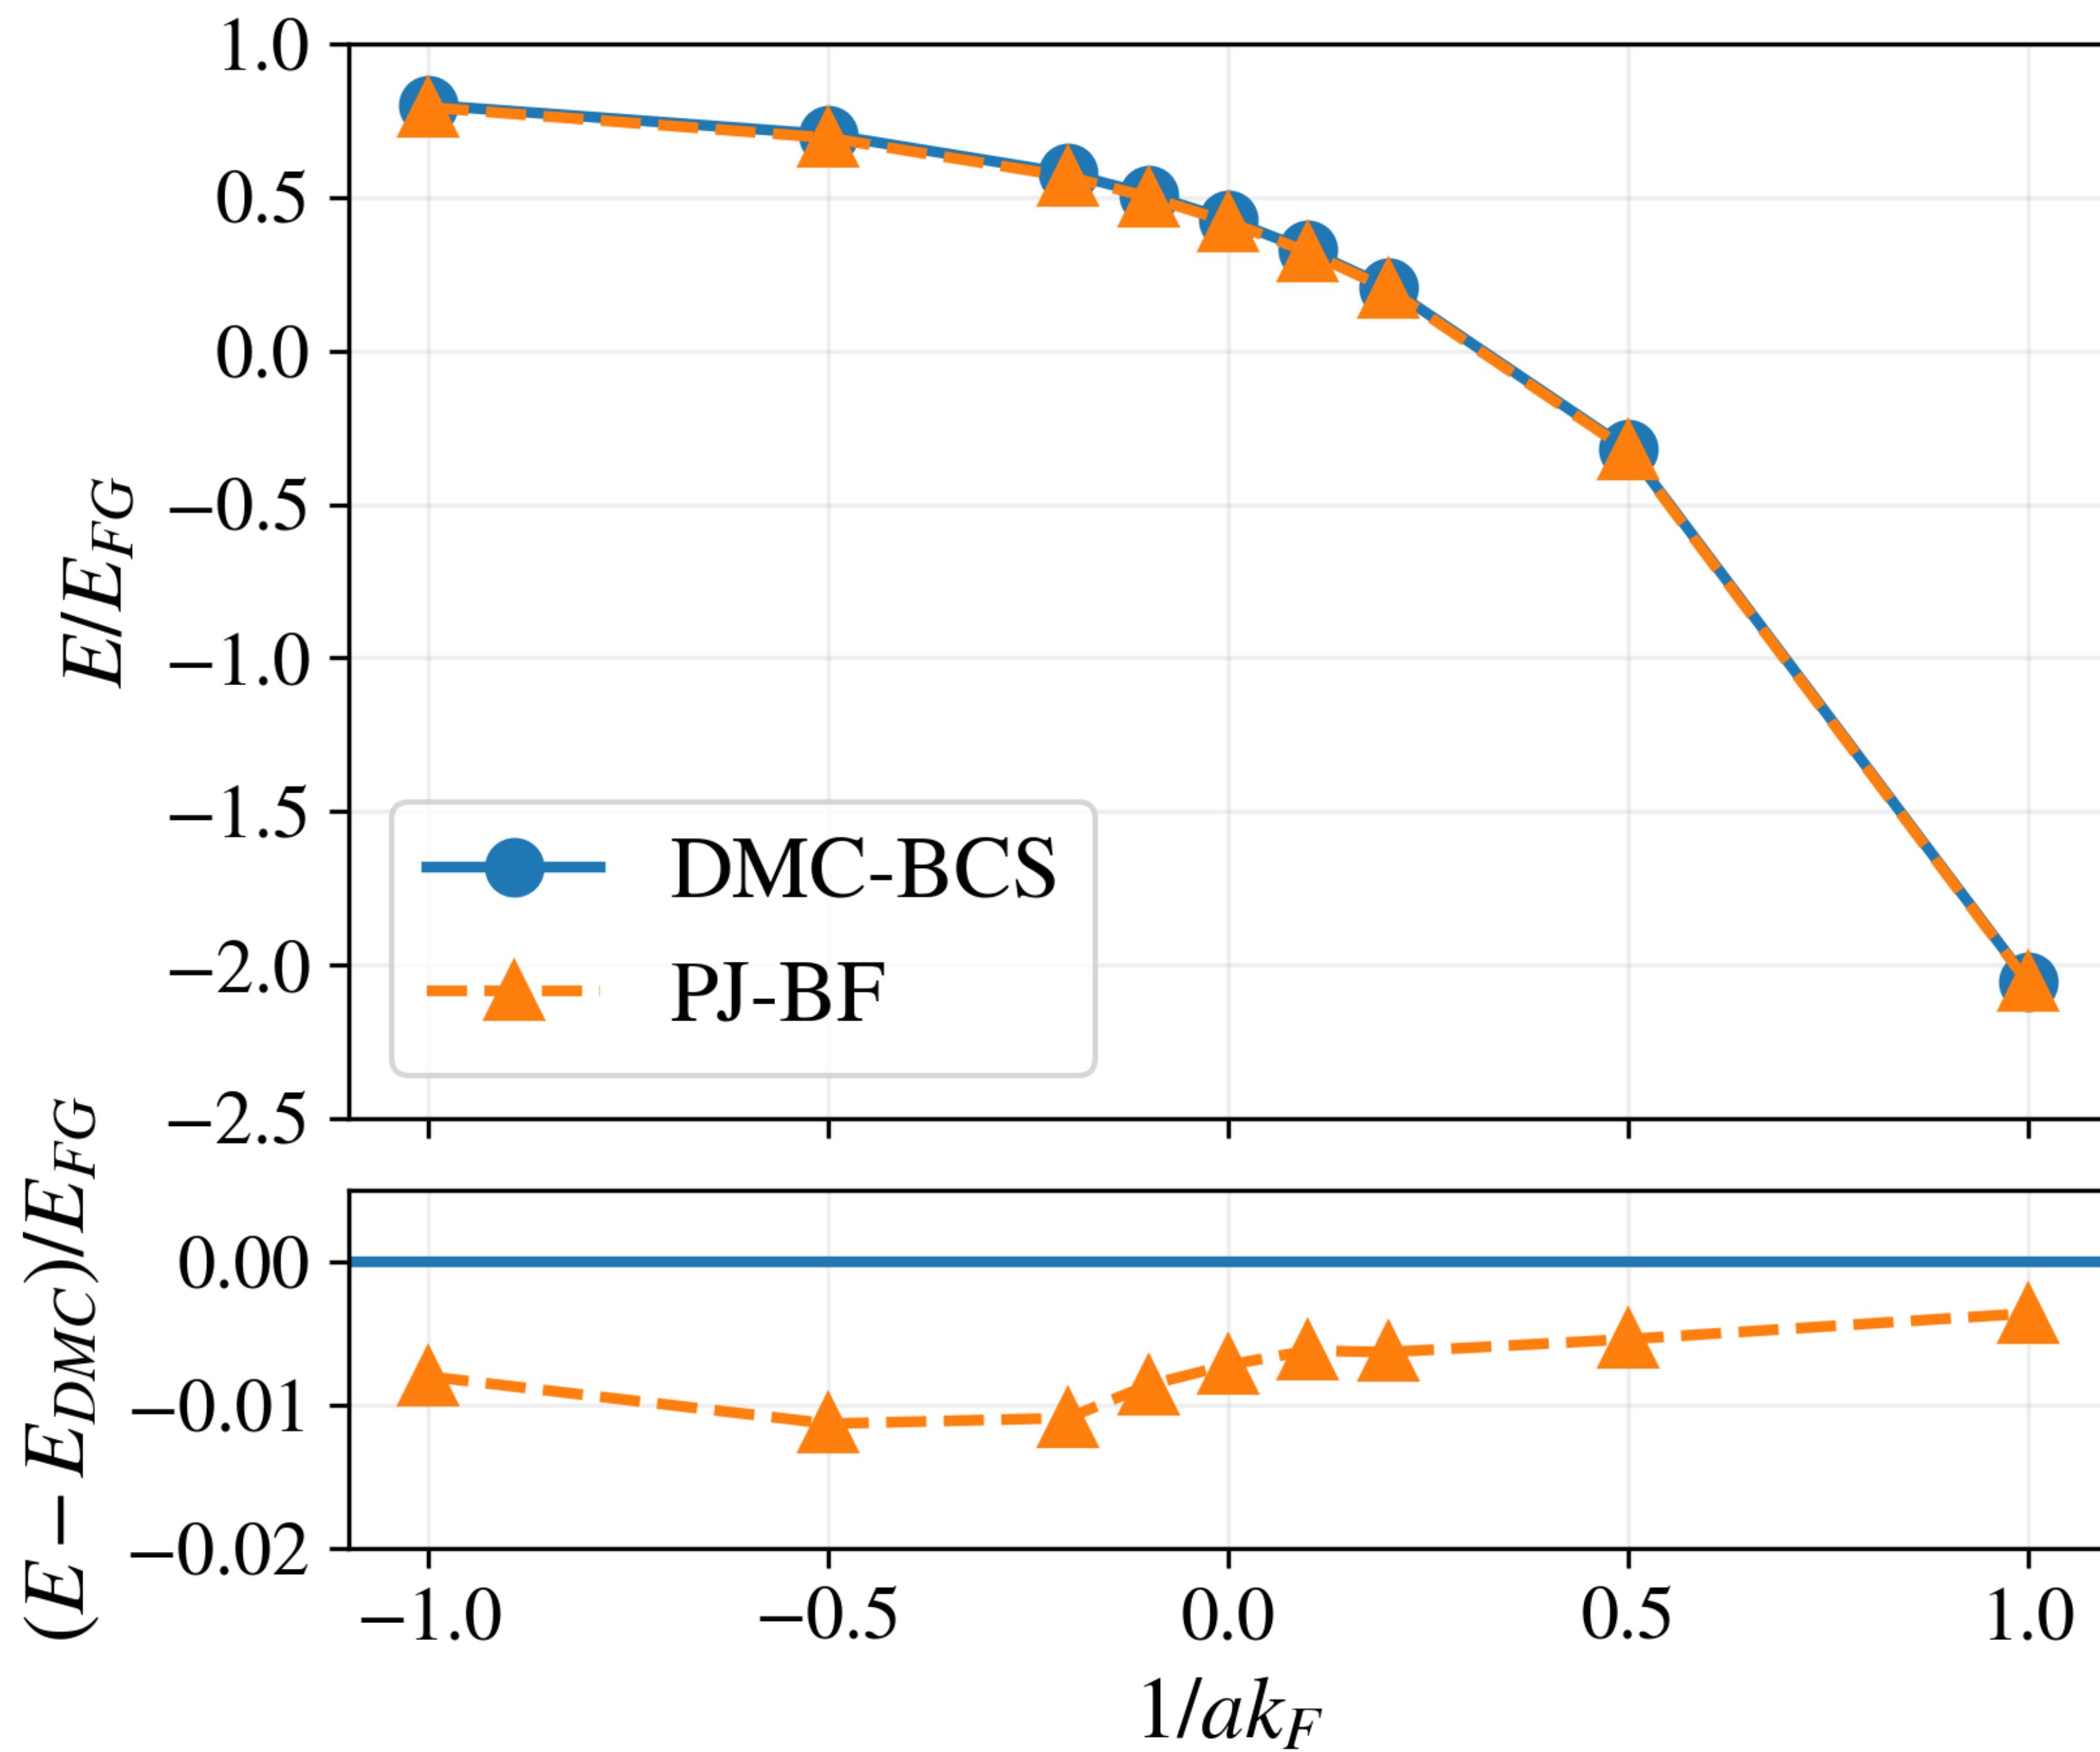


**Different effective ranges $r_e = 2/\mu$
at unitarity ($1/ak_F = 0$)**



**Different scattering lengths near unitarity
(fixed $r_e = 0.2$)**





Pairing Gap

- The pairing gap can be evaluated using

$$\Delta(N) = E(N) - \frac{1}{2}(E(N+1) + E(N-1))$$

for odd N .

- To accommodate odd- N cases in our Pfaffian NQS, we add one additional neural network to represent the unpaired single-particle spin-orbital

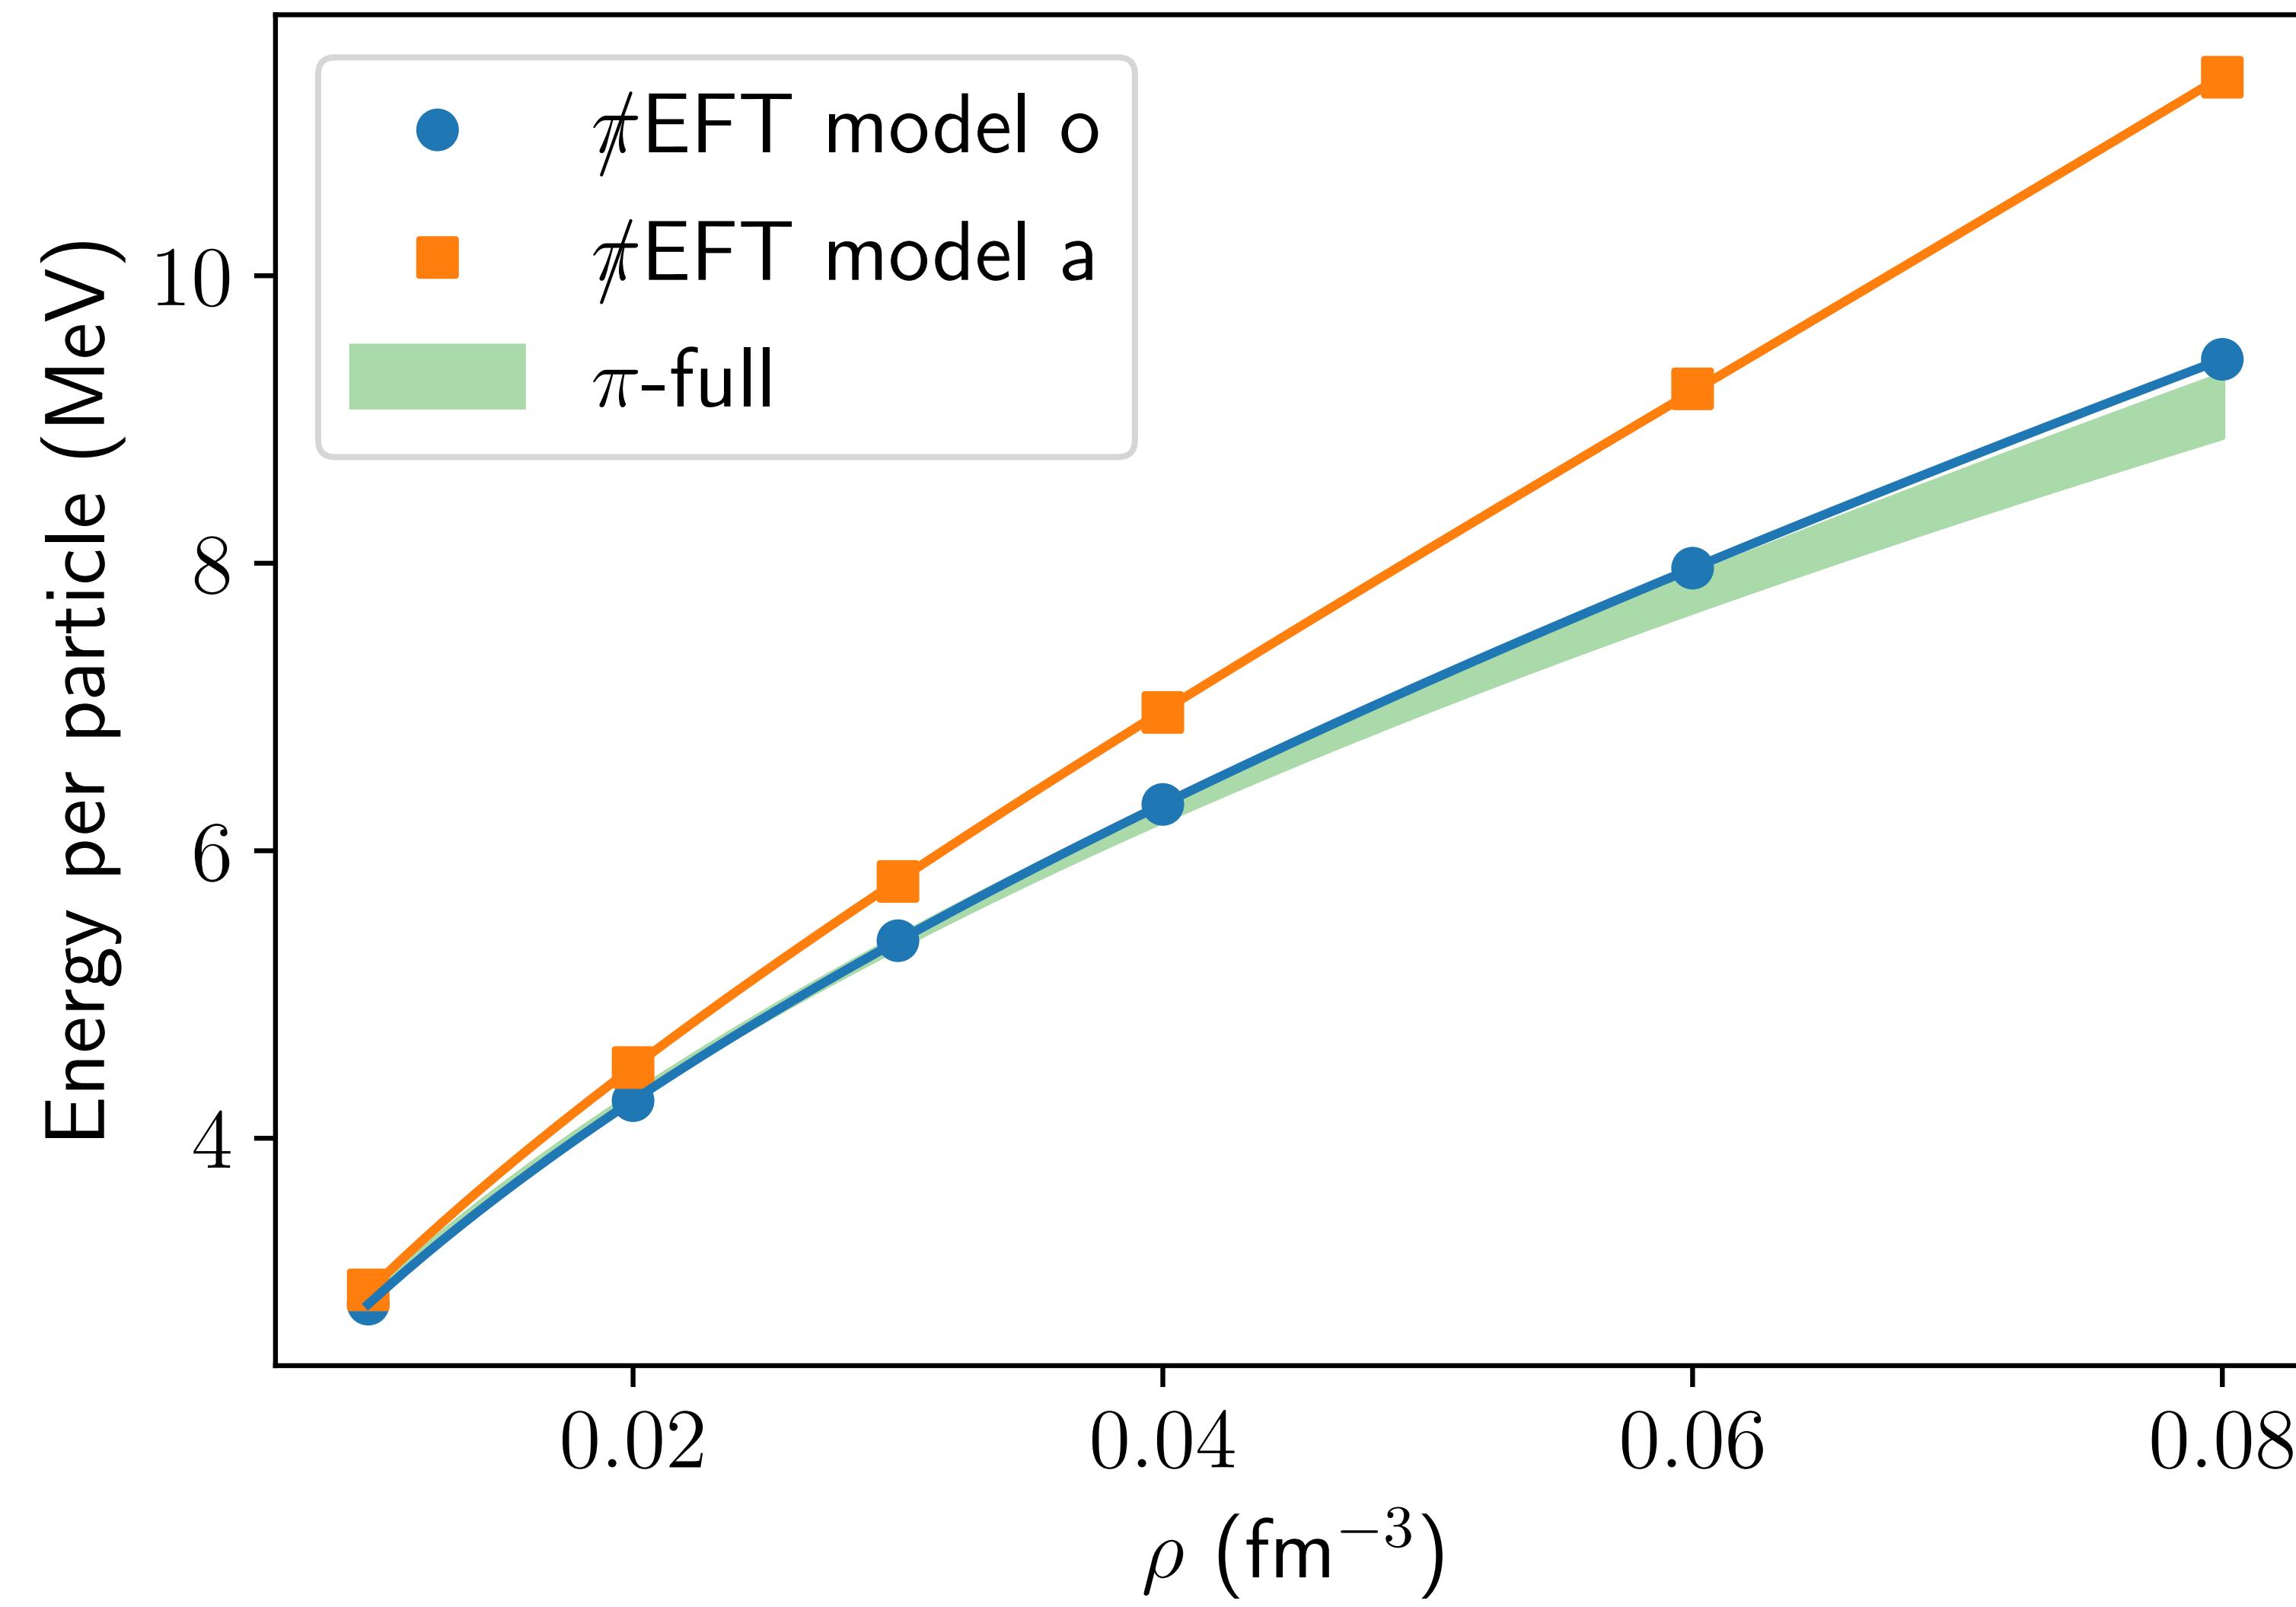
$\Delta(15)/E_{FG}$	DMC-BCS	PJ-BF
Non-translation invariant	0.9620	0.8618
Translation invariant	1.6475	1.5263

Ultra-cold Fermi Gases

- Our Pfaffian ansatz is very general — works for any Hamiltonian (even those that exchange spin!)
- Can obtain lower energies than state-of-the-art diffusion Monte Carlo methods
- We require far fewer parameters compared to other NQS applied to similar problems (~8500 vs millions)
- Future work:
 - Smaller r_e : better linear extrapolation to the $r_e \rightarrow 0$ limit
 - Larger N : As an initial test, we only used $N = 14$
 - Test if our Pfaffian NQS works just as well for partially-polarized systems

Thank you for your attention!

Dilute Neutron Matter



Number-Projected BCS Wave Function

$$\Phi(X) = \det [\phi(\mathbf{x}_{i\uparrow}, \mathbf{x}_{j\downarrow})] = \det \begin{bmatrix} \phi(\mathbf{x}_{1\uparrow}, \mathbf{x}_{1\downarrow}) & \phi(\mathbf{x}_{1\uparrow}, \mathbf{x}_{2\downarrow}) & \cdots & \phi(\mathbf{x}_{1\uparrow}, \mathbf{x}_{N/2\downarrow}) \\ \phi(\mathbf{x}_{2\uparrow}, \mathbf{x}_{1\downarrow}) & \phi(\mathbf{x}_{2\uparrow}, \mathbf{x}_{2\downarrow}) & \cdots & \phi(\mathbf{x}_{2\uparrow}, \mathbf{x}_{N/2\downarrow}) \\ \vdots & \vdots & \ddots & \vdots \\ \phi(\mathbf{x}_{N/2\uparrow}, \mathbf{x}_{1\downarrow}) & \phi(\mathbf{x}_{N/2\uparrow}, \mathbf{x}_{2\downarrow}) & \cdots & \phi(\mathbf{x}_{N/2\uparrow}, \mathbf{x}_{N/2\downarrow}) \end{bmatrix}$$

Periodic Inputs

- Separation vector

$$\mathbf{r}_{ij} = \mathbf{r}_i - \mathbf{r}_j \longmapsto \left(\cos(2\pi \mathbf{r}_{ij}/L), \sin(2\pi \mathbf{r}_{ij}/L) \right)$$

- Distance

$$\|\mathbf{r}_{ij}\| \longmapsto \|\sin(2\pi \mathbf{r}_{ij}/L)\|$$

- Positions are ignored to enforce translational invariance

Parity and Time-Reversal Symmetry

- We carry out the VMC calculations for the unpolarized gas using $\Psi^{PT}(R, S)$ given by

$$\Psi^P(R, S) = \Psi(R, S) + \Psi(-R, S)$$

$$\Psi^{PT}(R, S) = \Psi^P(R, S) + (-1)^{N/2} \Psi^P(R, -S)$$

where R and S are the set of all positions and spins, respectively.

MPNN Equations

- Iteratively builds correlations into new one- and two-body features from old ones
- Skip connections help stabilize training and avoid vanishing gradients
- Has been effective for the electron gas, despite having orders of magnitude fewer parameters compared to FermiNet

$$\left| \begin{array}{l}
 \mathbf{v}_i = (s_i^z) \\
 \mathbf{v}_{ij} = (r_{ij}, \mathbf{r}_{ij}, s_i^z \cdot s_j^z) \\
 \mathbf{h}_i^{(0)} = (\mathbf{v}_i, A\mathbf{v}_i) \\
 \mathbf{h}_{ij}^{(0)} = (\mathbf{v}_{ij}, B\mathbf{v}_{ij})
 \end{array} \right| \quad \begin{array}{l}
 \text{for } t = 1, \dots, T: \\
 \mathbf{m}_{ij}^{(t)} = \mathbf{M}_t \left(\mathbf{h}_i^{(t-1)}, \mathbf{h}_j^{(t-1)}, \mathbf{h}_{ij}^{(t-1)} \right) \\
 \mathbf{h}_i^{(t)} = \left(\mathbf{v}_i, \mathbf{F}_t \left(\mathbf{h}_i^{(t-1)}, \mathbf{m}_i^{(t)} \right) \right), \mathbf{m}_i^{(t)} = \text{Pool} \left(\{ \mathbf{m}_{ij}^{(t)} | j \neq i \} \right) \\
 \mathbf{h}_{ij}^{(t)} = \left(\mathbf{v}_{ij}, \mathbf{G}_t \left(\mathbf{h}_{ij}^{(t-1)}, \mathbf{m}_{ij}^{(t)} \right) \right)
 \end{array}$$

Determinants vs. Pfaffians

Defined for $n \times n$ matrices

$$\det(A) = \sum_{\sigma \in S_n} \operatorname{sgn}(\sigma) \prod_{i=1}^n a_{i,\sigma(i)}$$

$$\det(A^T) = \det(A)$$

$$\det(A) \det(B) = \det(AB)$$

$$\det(A) = \operatorname{pf}(A)^2$$

Defined for $2n \times 2n$ skew-symmetric matrices

$$\operatorname{pf}(A) = \frac{1}{2^n n!} \sum_{\sigma \in S_{2n}} \operatorname{sgn}(\sigma) \prod_{i=1}^n a_{\sigma(2i-1), \sigma(2i)}$$

$$\operatorname{pf}(A^T) = (-1)^n \operatorname{pf}(A)$$

$$\operatorname{pf}(A) \operatorname{pf}(B) = \exp\left(\frac{1}{2} \operatorname{tr} \log(A^T B)\right)$$